Conceptual alternatives: Competition in language and beyond[*]

Brian Buccola
Michigan State University

Manue Križ
University of Vienna

Emmanuel Chemla
LSCP (ENS, EHESS, CNRS)

August 13, 2020

Abstract

Things we can say, and the ways in which we can say them, compete with one another. And this has consequences: words we decide *not* to pronounce have critical effects on the messages we end up conveying. For instance, in saying *Chris is a good teacher*, we may convey that Chris is not an *amazing* teacher. How this happens is an unsolvable problem, unless a theory of *alternatives* indicates what counts, among all the things that have not been pronounced. It is sometimes assumed, explicitly or implicitly, that any word counts, as long as that word could have replaced one that was actually pronounced. We review arguments for going beyond this powerful idea. In doing so, we argue that the level of words is not the right level of analysis for alternatives. Instead, we capitalize on recent conceptual and associated methodological advances within the study of the so-called "language of thought" to reopen the problem from a new perspective. Specifically, we provide theoretical and experimental arguments that the relation between alternatives and words may be indirect, and that alternatives are not linguistic objects in the traditional sense. Rather, we propose that competition in language is better seen as primarily determined by general reasoning preferences, or thought preferences (preferences that may have forged the lexicons of modern languages in the first place, as argued elsewhere). We propose that such non-linguistic preferences can be measured and that these measures can be used to explain linguistic competition, non-linguistically, and more in depth.

1   The phenomenon associated with alternatives

When it comes to semantic interpretation in natural language, words that have *not* been pronounced matter. In fact, language even makes available specific expressions whose interpretation depends on unpronounced, alternative utterances, e.g. *only* and *even*, to mention just two. For instance, *Only Ann passed the exam* conveys not only that *Ann passed the exam* is true, but also that all alternatives of the form *x passed the exam* are false, where *Ann* is replaced by some relevant individual *x*: Bill, Chris, etc. (Horn, 1972; Rooth, 1985). Similarly, *Even Ann passed the exam* conveys that *Ann passed the exam* is not only true, but was less likely to be true than all relevant alternatives of the form *x passed the exam* (Karttunen and Peters, 1979).

Inferences based on what was said versus what could have been said but wasn't (i.e. what was unsaid; see Horn, 1992) arise even in the absence of particles such as *only* and *even*, and they are known in linguistics as *implicatures*. Perhaps the most widely studied kind are *scalar implicatures*, so called because they intuitively involve scales of some kind. For instance, the sentence *Ann or Bill passed the exam* is readily understood as conveying that Ann and Bill didn't both pass. A standard explanation of how this inference arises, pioneered by the philosopher H.P. Grice (Grice, 1975) and refined further by Horn, 1972; Gazdar, 1979; Gamut, 1991 (among others), goes as follows: *or* and *and* form a scale, ⟨*or*,*and*⟩; the speaker could have used the alternative utterance made up from the other member of that scale, *Ann and Bill passed the exam*, and presumably would've (and should've) done so, if this alternative were true; since the speaker didn't use this alternative, it must be because it's false. Thus, it's true that Ann or Bill passed, but false that they both did.

The same disjunctive sentence also typically conveys that the speaker of the sentence is uncertain which of the two passed, and this ignorance inference can be explained along similar lines: the speaker could have used one of the alternative sentences *Ann passed the exam* or *Bill passed the exam* if she had known which of them was true; since she didn't use either one, it must be because she wasn't certain which one of them was true.

One role that alternatives play, then, is that they trigger inferences, either plain (sometimes called exhaustive) inferences or ignorance inferences. As the flip side of this, alternatives also help us understand why certain sentences are judged to be odd. For example, the sentence *Italians come from Italy or France* is odd because, even if it is strictly speaking true, it may irresistibly trigger a false suggestion that an Italian may not come from Italy, but rather from France. Thus, alternatives are responsible for routinely derived inferences, and these inferences are so robustly

anchored in the interpretive process of language that alternatives may not be ignored, even when they result in plain oddness (Magri, 2009).

2 A theory of competition and constraints on alternatives

2.1 The symmetry problem and a recipe for alternatives

A complete theory of scalar implicature involves two components. First, one needs to describe a system predicting which inferences are generated by the competition between a given sentence and a given set of alternatives. This has been the subject of much research (Horn, 1972; Gazdar, 1979; Gamut, 1991; Chierchia et al., 2012). But a predictive theory also needs to specify, second, what the set of alternatives is for any given sentence (see Katzir, 2007 for a seminal proposal, which we'll rely on heavily). Indeed, the phenomenon we are interested in is not restricted to the ⟨*or*, *and*⟩ alternation, but arises similarly with other scales, such as ⟨*some*, *all*⟩, ⟨*possible*, *certain*⟩, ⟨*allowed*, *required*⟩, and ⟨*warm*, *hot*⟩. It's arguably for this reason that each of the sentences in [1] typically licenses an exhaustivity implicature, viz. that the corresponding stronger alternative is false (Horn, 1992).

[1]   a.   Dev ate some of the cookies.   (⇝ not all)
    b.   Emily read *Hamlet* or *Macbeth*.   (⇝ not both)
    c.   It's possible that it's raining.   (⇝ not certain)
    d.   You're allowed to stay home.   (⇝ not required to)
    e.   The soup is warm.   (⇝ not hot)

We can now highlight one of the most pressing and immediate challenges for a theory of alternatives, known as *the symmetry problem* (Fox, 2007; Katzir, 2007; Fox and Katzir, 2011), which arises if we do not constrain what scales are made of. Abstractly first, if α typically implicates 'not β', then a theory of alternatives needs to explain why β is the alternative to α, rather than *not β* (or α *and not β*). If the latter but not the former were an alternative, then α would actually implicate β (= not[not β]); and if both β and *not β* were alternatives, then α would implicate speaker uncertainty about β.[1] In either case, the implicature 'not β' would remain unexplained. The symmetry problem is so called because, without constraints on what counts as an alternative, every potential alternative β has a symmetric partner, *not β*, whose existence preempts the derivation of the scalar implicature actually observed.

---

1 In principle, if both β and *not β* are alternatives of α, then one could imagine one of them, say *not β*, being 'ignored' during implicature calculation, e.g. because context makes it less salient. In such a case, α would implicate 'not β'. However, in practice it has been argued that symmetry cannot be 'broken' in this way (by ignoring one of two symmetric alternatives) (Fox and Katzir, 2011).

To illustrate, an unconstrained theory fails to capture the fact that *some* is routinely understood as 'some but not all' rather than as 'all' (or as ignorance between 'some but not all' and 'all'). A better theory would explain why *all* is an alternative to *some*, while *not all* or *some but not all* aren't. An immediately obvious difference between *all* and *some but not all* is that the former, like *some* itself, is a single word, whereas the latter is a string of multiple words. That is, *some* and *all* are, in some intuitive sense, equally complex, while *some but not all* is more complex than *some*. Building on this observation, the following recipe for generating alternatives is often explicitly or implicitly endorsed; it states that alternatives are derived on just a word-by-word basis.

[2] **Recipe for alternatives.** Take the string of words, and replace one word with another word.

2.2 Successes

This theory scores (at least) two immediate successes.

Success 1: Solving the symmetry problem. Given the string *Dev ate some of the cookies*, we can replace *some* with *all* to yield the alternative *Dev ate all of the cookies* (whose negation correctly delivers the observed implicature), but we can't replace *some* with *some but not all* (since the latter is not a single word) to yield the alternative *Dev ate some but not all of the cookies* (whose negation would deliver the unobserved inference that Dev ate all of the cookies). Thus, we correctly predict that *some* implicates 'not all', rather than 'all'. Similarly, we correctly predict that *Ann or Bill passed the exam* has as an alternative *Ann and Bill passed the exam* (by replacing *or* with *and*), but not *Ann or Bill but not both passed the exam*, hence that *or* implicates 'not both', rather than 'both'. And similarly for all other scalar expressions.

Success 2: Explaining lexical constraints. Since the recipe makes reference to words, i.e. to single lexical items of the language, we predict that scalar inferences are constrained by the lexicon of the language under consideration. For instance, the English sentence *Fred is Gale's brother* doesn't imply one way or the other about whether (the speaker believes that, or is uncertain whether) Fred is older or younger than Gale. By contrast, Japanese has single-word sibling terms that semantically encode seniority (and, like English, gender), e.g. *ani* 'older brother' and *otooto* 'younger brother', as well as non-specific sibling terms that don't encode seniority, *kyoodai* 'brother'. As a result, the use of the non-specific term *kyoodai*, as in *Kochira wa Takashi-kun no kyoodai no Michio-kun desu* 'This is Takashi's brother Michio', implicates that the speaker was not in a position to use either of the more specific

terms, *ani* or *otooto*, and as such implicates speaker uncertainty about whether Michio is Takashi's older or younger brother (Matsumoto, 1995).

2.3 Challenges

Under a common view, alternatives are determined by the recipe in [2] and the lexicon of the language under consideration. Despite initial successes, however, careful attention reveals a number of challenges for such a view. We discuss five such challenges. The first two relate to the lexicon, from which alternatives are taken to be drawn. The second two relate to cases beyond the lexicon, where the alternative needed to explain an inference is not actually part of the language. The last is a more general, conceptual challenge. Taken together, the challenges motivate our proposal to move from a theory based on linguistic objects to one based more on conceptual objects.

Challenge 1: Hierarchy of logical/content words.   Consider the sentence in [3a], where *cards* refers to standard playing cards that can only be red (diamonds and hearts) or black (clubs and spades). This sentence can naturally be understood as conveying that not all of Mary's cards are red. On first glance, the recipe in [2] appears to explain this inference: [3b], but not [3c], is an alternative; thus, the sentence implicates the negation of [3b], rather than the negation of [3c].

[3]   a.   Some of Mary's cards are red.
     b.   All of Mary's cards are red.
     c.   Some but not all of Mary's cards are red.
     d.   Some of Mary's cards are black.

However, the recipe in [2] also allows us to replace *red* with *black*, yielding the alternative [3d], whose negation means that none of Mary's cards are black. If [3d], and not [3b], were the alternative of [3a], then [3a] would actually implicate that all of Mary's cards are red; and if both [3b] and [3d] were alternatives, then [3a] would implicate speaker ignorance about whether some of Mary's cards are black. Empirically, however, while [3a] can convey either that not all of Mary's cards are red, or that the speaker is ignorant about whether that is true, it simply cannot convey that all of Mary's cards are red. In theoretical terms, this means that [3b] can be an active alternative without [3d] being active (which yields the implicature that not all of Mary's cards are red), or they can both be active (which yields ignorance), but it cannot be the case that [3d] is active without [3b] (since that would implicate that all of Mary's cards are red). Thus, there is an observable asymmetry between [3b] and [3d] that the recipe in [2] is not sensitive to.

One possible (and traditional) solution is to revise the recipe in [2] so that only words in the same *scale* may replace each other. If we furthermore assume that

some and *all* share a scale, while *red* and *blue* do not, then [3d] would not be an alternative of [3a] to begin with. This move would certainly solve this particular challenge, but it's not without its own issues. First, scales become an additional, and core, part of the actual theory (sometimes called Horn scales), rather than just a useful metatheoretic device for describing the data. Second, it requires a specification of what counts as a scale and what doesn't. While many researchers have worked in this direction, the theory of Katzir, 2007, which is the basis of the recipe in [2] and which we ultimately build on, is intended to supersede the theory of Horn scales by replacing scales with structural operations alone. Lastly, such a move may be empirically incorrect: there may well be cases where *red* does implicate 'not black', such as *I have red cards in my hand*, which may suggest that I don't have black cards in my hand.[2] The solution we'll offer is that quantificational concepts like 'some' and 'all' are conceptual primitives, while color concepts are not (or some more graded notion of this idea), and that conceptually (more) primitive concepts are (more likely to be) recruited as alternatives.

Challenge 2: Some lexical properties are invisible. Consider the sentence in [4a]. By replacing *all* with *some*, we derive the alternative in [4b], which (to the extent that it's interpretable) has the reading 'there are some books x such that no one read x'. The negation of [4b] is thus equivalent to 'there are no books x such that no one read x', i.e. 'for all books x, someone read x'. As such, we expect [4a] to implicate that each book was read by someone or other; however, this implicature is intuitively unavailable. Rather, the implicature we typically draw from [4a] is that someone (at least one person) read some book (at least one book). We can derive this implicature by instead replacing *all* with *any*, yielding the alternative in [4c], which has the reading 'there is no person x such that x read some book'. The negation of [4c] is thus equivalent to 'there is some person x such that x read some book', precisely the implicature observed.

[4] a. No one read <u>all</u> of the books.
     b. No one read <u>some</u> of the books.   ✗
     c. No one read <u>any</u> of the books.   ✓

---

2 Thanks to an anonymous reviewer for raising this point about scales. The same reviewer also notes that, alternatively, we may assume that the implicature mechanism only excludes logically stronger alternatives, so that even if [3d] is an alternative of [3a], it won't be negated since the two sentences are logically independent. While this assumption would indeed solve this challenge, there are empirical reasons to think that the implicature mechanism actually does exclude non-stronger (and non-weaker, i.e. logically independent) alternatives. For instance, the example in the main text, *I have red cards in my hand*, may implicate the negation of the logically independent alternative *I have black cards in my hand*. Note that the logically stronger alternative *I have red and black cards in my hand* could play the same role, but this would require the potentially unnatural scale ⟨*red*, *red and black*⟩.

The problem, however, is that the recipe in [2], which simply makes reference to word-by-word replacement in strings, cannot distinguish [4c] (which we want to be an alternative of [4a]) from [4b] (which we don't). More generally, when *all* occurs in the syntactic scope of a negative quantifier like *no one*, it should be replaceable by *any*, but not by *some*. One possible solution is to revise the recipe in [2] to take syntactic structure into account, and to assume that *all* is replaceable by either *some* or *any*, depending on structure. Another possible solution, which we will explore, is to assume that *some* and *any* have roughly the same meanings — they're existential quantifiers (see Kadmon and Landman, 1993; Chierchia, 2013) — and that universal quantification is replaceable by existential quantification.

Even if the previous two challenges can be met by a suitably revised theory of alternatives as linguistic objects, the remaining challenges below suggest that the correct theory of alternatives must somehow go beyond plain language (which may, at the same time, solve the first two challenges, too).

Challenge 3: Alternatives may be ungrammatical/uninterpretable. This discussion recapitulates arguments from Buccola and Chemla, 2019, showing that some processes for creating alternatives (here about eliminating disjuncts) may output unusable alternatives. Sentence [5] intuitively licenses the inference that [6a] is false, and the inference that [6b] is false. How might we derive this pair of inferences on the basis of the recipe in [2]?

[5]   Every dad$_i$ called [his$_i$ daughter]$_j$ or her$_j$ dog.

[6]   a.   Every dad$_i$ called his$_i$ daughter.
      b.   Every dad$_i$ called his$_i$ daughter's dog.

One option is to assume that the recipe in [2] allows for replacement by the empty string (in other words, allows for deletion; cf. Katzir 2007). This would yield the alternatives in [7]. But while [7a] = [6a], it does not appear that [7b] = [6b]. Specifically, it's not immediately obvious how [7b] could mean [6b]: either j refers to some external individual (e.g. Alice), or else $i = j$; in neither case do we end up with the right meaning.

[7]   a.   Every dad$_i$ called his$_i$ daughter.
      b.   Every dad$_i$ called her$_j$ dog.

One option is to argue that, at the logical of logical form (LF), *her$_j$ dog* in [5] and [7b] receives a so-called paycheck interpretation, i.e. it means 'his$_i$ daughter's dog' after all (Geach, 1962; Karttunen, 1969).

A second option is to assume that the lexicon contains the special operators **L** and **R**, which are connectives that retain the structure of both disjuncts, but which

assert only the truth of the **L**eft and **R**ight disjunct, respectively (Sauerland, 2004). Then it's easy to show that [8a] = [6a], and that [8b] = [6b].

[8]   a.   Every dad$_i$ called [his$_i$ daughter]$_j$ **L** her$_j$ dog.
       b.   Every dad$_i$ called [his$_i$ daughter]$_j$ **R** her$_j$ dog.

Crucially, whatever option above is adopted, alternatives are best understood not as properly linguistic objects, but rather as objects at some more logico-conceptual level (LF). This is because, while [7b] or [8b] may yield the right inference, neither is a helpful English sentence in its own right.

Challenge 4: Alternatives may be inexpressible. This discussion rehearses arguments with a very similar flavor as the one above, here showing that lexical replacement, a simple process to obtain an alternative, is not available for lack of the relevant lexical item. The English sentence *John broke all of his arms* is odd: it implies that John has more than two arms. The oddity, and this inference, can be explained by competition with the alternative *John broke both of his arms*, which is a more appropriate description of a situation in which John broke his left arm and his right arm. That the English lexicon contains the terms *all* and *both* is crucial to this explanation, and the prediction is that in a language with only a lexical term meaning 'all', and no lexical term meaning 'both', the sentence corresponding to *John broke all of his arms* ought to not be odd. French is a case in point; however, the French sentence *Jean s'est cassé tous les bras* 'John broke all of his arms' is just as odd as its English counterpart (Chemla, 2007). This data point thus indicates that competition goes beyond just the lexicon of a given language: alternatives that are expressible in one language but not in another (due to a lexical gap) nevertheless get recruited in both languages at the level of competition, suggesting once again that alternatives are not totally linguistic.

In a very similar vein, Charlow (2019) examines the case of scalar implicatures arising from scalar expressions that have exceptional scope. For instance, the sentence *John repeated the rumor that Mary or Bill was expelled* invites the inference that it's not the case that Bill and Mary are *each* such that John repeated the rumor that he or she was expelled. Charlow argues that such utterances involve existential quantification over choice functions, and that their alternatives involve universal quantification over choice functions.[3] Similar to French above, these alternatives, as

---

3 Crucially, the alternative with *and* instead of *or* does not have the requisite wide-scope reading; that is to say, [i] below entails that there is *one* single rumor about both Mary and Bill being expelled; not two separate rumors, one about Mary and one about Bill.

[i]   John repeated the rumor that Mary and Bill were expelled.

Charlow notes, "are more abstract than we might have thought", and in particular "do not seem to correspond to any expressible lexical items" (in English).

A number of other researchers have recently provided similar arguments for the existence of alternatives that are inexpressible in the language under consideration (Smith, 2020; Szabolcsi, 2020). While space limitations prevent us from discussing all these phenomena, we note that, in principle, there may be cases where the relevant alternative is not expressible in *any* language, i.e. cases of a universal ineffability of an otherwise primitive meaning. Ideally, one should then find independent evidence for the primitiveness of those meanings. Our experiment (§5) is a first step towards demonstrating how such evidence might be sought.

Challenge 5: The quest for a deeper explanation. Even if all of the above challenges could be successfully met by a suitably revised lexical recipe, one may raise a more general, conceptual worry about linguistic approaches to alternatives. Why is the meaning 'all', but not 'some but not all', lexicalized in the first place, not just in English, but in language after language? A natural type of answer that has been sought is that 'all' is somehow more *natural*, or *primitive*, than 'some but not all'. Maybe this kind of answer can also provide an answer to the question of why *all*, but not *some but not all*, is an alternative to *some*. In other words, the observation that 'all' is lexicalized makes it a sensible candidate for being an alternative, but it leaves us with the question of why it's lexicalized in the first place. Once we understand why it's lexicalized, perhaps we won't need to explain why it's an alternative. It's an alternative for the same reason: because it's more natural/primitive.

3   Logical primitives and the language of thought

The challenges above arise only if competition is assumed to occur at the linguistic level, i.e. between properly linguistic objects (utterances). If, instead, competition is between 'primitives' at some more conceptual level, then the problems vanish. To be more concrete, we revise the recipe as in [9].

[9]   **Recipe (revised).** Take the conceptual representation of the sentence, and replace one primitive element with another primitive element.

This recipe may be seen as a 'neo-Katzirian' view of alternatives. We elaborate on this view in much greater detail in §4, but for now we want to discuss some initial consequences and how to render it concrete and testable. The main difference with Katzir, 2007 is that replacements are (preferably) drawn from a set of cognitive primitive elements, rather than from a linguistic lexicon. The full Katzirian set of alternatives, including those involving replacements made available by context, can

be recovered as well by assuming that some primitive elements can anaphorically refer to previously mentioned material (it seems easy to mentally point at or refer to thoughts or linguistic expressions).

Now reconsider Challenge 1, regarding the difference (in building alternatives) between quantifiers like *some* and *all*, on the one hand, and color terms like *red* and *black*, on the other. The idea behind the solution is that these quantifiers are conceptual primitives, while color terms are not (or some graded version of this idea; see §4). Thus, quantifier replacement is a (preferable) way to build alternatives, while color term replacement isn't.

As for Challenge 2 (*No one read all of the books*), if we assume that existential ($\exists$) and universal ($\forall$) quantification are primitives, and are the conceptual counterparts of *some* and *all*, respectively, then substitution of one for the other becomes possible regardless of syntactic environment. More precisely, $\forall$ can be replaced by $\exists$, yielding the attested implicature, despite the fact that replacing the word *all* by the word *some* (and keeping the structure the same) would yield an ungrammatical structure, due to the polarity sensitivity of *some*.

Challenges 3 and 4 are solved by assuming that the alternatives required are available at the conceptual level. In the case of *both* and its absence in the French lexicon, for example, the idea is that the concept 'both' is nevertheless relatively primitive and thus is used by French speakers to build the relevant alternative. As for the exceptional scope cases considered by Charlow (2019), we would again say that existential quantification ($\exists$) over choice functions is replaceable by universal quantfication ($\forall$) at the level of thought. Similar remarks hold for the other cases considered above.

Importantly, moving from a linguistic to a conceptual level provides new insight into the symmetry problem and into lexicalization (Challenge 5). Suppose that $\exists$ and $\forall$ are primitive, while the conceptual counterpart of *some but not all* is not; rather, it's composed of other primitives like $\exists$, $\wedge$ (conjunction), $\neg$ (negation), and $\forall$.

[10] **Hypothesis.** The conceptual counterparts of *some* and *all* are primitive, while that of *some but not all* is not.

Combined with the natural hypothesis that lexical items, across languages, are more likely to be primitive elements than non-primitive elements, this explains the cross-linguistic stability of lexicalization (Strickland, 2017). Moreover, if primitive elements are also privileged candidates for being alternatives, then this also addresses the symmetry problem. (In §4, we provide a more complete algorithm to derive alternatives in the language of thought.) Our experiment (§5) is intended to be a proof-of-concept of how to test hypothesis [10], with results that arguably support it.

3.1 What is a primitive element?

Here we try to give an intuitive sense of what we mean by a primitive element. The word *birds* describes an intuitively well-formed natural class (the set of all birds), and so does the word *red* (the set of all red things). Moreover, the combination *red (and) birds* is an intuitively well-formed natural class (the set that includes cardinals but not strawberries or blue jays). However, *red or birds* (the set that includes cardinals, strawberries, and blue jays, but not blue berries) is a rather unnatural class, and *red xor birds* (where *xor* means exclusive *or*) even less so (it includes strawberries and blue jays, but not cardinals).

These clear facts could be used to argue that (the concept) 'and' is, in some sense, more primitive than, or preferred to, 'or', and both of those more so than 'xor'. The specific question we aim to address is whether it can be shown that (the concept) 'all' is more primitive, or preferred, than 'some but not all'.

3.2 The logical primitives of thought

In the domain of content words, such as nouns, Gärdenfors, 2014 has most forcefully discussed constraints on lexicalized meanings, presumably originating from conceptual constraints on natural classes. These notions have been operationalized in psychology experiments since then (e.g. Xu and Tenenbaum, 2007), showing that specific properties of some word meanings were more readily accessible than others, in both children and adults. In the domain of logical words, which concerns us the most here, comparable investigations have been proposed from various traditions as well, e.g. Piantadosi et al., 2016; Katzir and Singh, 2013; Horn, 1972; Seuren and Jaspers, 2014 (see Chemla et al., 2019 for a unification of content and logical words).

4 A program for alternatives

We now flesh out in more detail what a move from linguistic alternatives to conceptual alternatives looks like, building on the seminal proposal in Katzir, 2007.

4.1 Katzirian view

Two lessons we learn from the Katzirian view of alternatives are that structure matters (and not just strings of words) and that different sorts of operations over structures exist. To be more precise, the alternatives of a given structure are best conceived as the result of applying operations to that structure. This view allows one to measure the relative complexity of a structure and its alternatives, by tracing the operations needed to obtain those alternatives. What operations, then, may transform a structure? Katzir proposes the following three:

(i) a part of the structure may be deleted;

(ii) a part of the structure may be replaced by a piece of structure made available in the current discourse (in virtue of its presence in the structure of the current sentence or of a previous sentence in the discourse); and

(iii) a part of the structure may be replaced by an individual lexical item.

4.2 Neo-Katzirian view

We propose an updated version of Katzir's algorithm, in which alternatives are located in the language of thought, and replacements occur only on the basis of single lexical items, albeit lexical items of the language of thought. Assume that a sentence and its structure can always be translated into the language of thought, a language which would thus be structurally similar to that of classical Logical Forms, but with potentially more lexical items. We propose that there is single operation available (recursively) to transform a structure: replacement of part of the structure by a lexical item of the language of thought. The lexical items are diverse, however, and may be:

(i′) a special empty element;

(ii′) a special pronoun capable of pointing at structures; and

(iii′) other elements that may or may not be lexicalized in the actual language.

There is a clear correspondence between (i′) and (i), and between (ii′) and (ii). We note, by the way, that (ii) is not very Katzirian in the first place, in that it may lead to more complex alternatives. However, its updated formulation (ii′) may be seen as Katzirian again (in the sense of not adding complexity), given that the resulting structure would definitely be simpler.[4]

The main innovation, then, is in (iii′), where replacements can occur even on the basis of elements (of the language of thought) that are not lexicalized in the actual language. This immediately gives us a handle on the challenge from French:

---

4 This may look like a technical trick, but the actual existence of a special pronoun of this form, capable of anaphorically referring to a piece of structure made available in the discourse, may actually be a sensible idea. For what it's worth, it seems appealing to say that we can refer to thoughts, at least for ourselves (as in, "*This* and *that* are the reasons why..."). This may even happen in a rather visual manner, and the analogy with pointing pronouns in sign language may be of relevance (Schlenker, 2017). Such considerations are not of central importance, but it's useful to note that even if quite abstract, they could lead to actual predictions: under the view that a pronoun is responsible for picking up the piece of structure from the context, then no further replacement should be possible, even if we assume that replacements can occur recursively.

the oddity of *tous les bras* 'all the arms' is due to the existence of an alternative 'both arms' at the level of thought, even though 'both' is not lexicalized in French.

However, our innovation in (iii') also raises the important question of what else the lexicon of the language of thought contains (besides the two special elements mentioned in (i') and (ii')), and we risk losing the Katzirian solution to the symmetry problem: if, for example, both 'all' and 'some but not all' were lexical items of the language of thought, then we would expect symmetric alternatives and hence fail to explain the observed implicatures. One straightforward way to interpret our experimental results, then, is as follows: they establish that 'all' is (more likely to be) a lexical item of the language of thought, while 'some but not all' is not (likely to be one), hence that 'all', but not 'some but not all', may be recruited during replacements by the structural rule in (iii'), thus solving the symmetry problem.

The resulting theory is a rather direct translation of the Katzirian view from natural language to the language of thought. However (foreshadowing a bit our graded experimental results), we may wish to incorporate gradedness into the theory. One way to do that would be to have either a graded conception of membership in the lexicon of the language of thought, or genuine epistemic uncertainty about what belongs in the lexicon of the language of thought. Both of these appear rather questionable, which suggests that something is wrong with the above formulation of (iii'). Instead, we may prefer a more graded notion of *alternativehood*, where rule (iii') involves replacement by potentially *any* element of the language of thought, but replacements are associated with costs that vary, among other things, according to the 'primitiveness' of the element being used.

4.3 A more nuanced view: Costs

In various current approaches to implicatures, alternatives come with relative costs, so that more costly alternatives are less likely to influence the inferential processes (Bergen et al., 2016). As a result, we may escape the symmetry between two alternatives A_1 and A_2 not by justifying that one is actually not an alternative, but simply from the fact that they come at different costs. In a Katzirian view, it's natural to try to obtain these costs from the way the alternative was derived, potentially by assigning costs to the transformations in (i), (ii), and (iii) above. For instance, alternatives obtained by type (i) may by design be simpler, and therefore produce a low cost.

From our neo-Katzirian perspective, we may simply say that the cost of an alternative varies inversely with the size of the structure being replaced, and directly with the complexity of the replacement material. The complexity of replacement material may then actually be the same thing as what we call primitiveness. Surely (i'), for instance, may incur a very low cost, but the cost of a replacement of the

other types may depend on how primitive these elements are.

Of course, we have no way to give precise numbers, but we may still sketch the general idea. If 'both' is relatively primitive, then, first of all, it's likely to be lexicalized in a variety of languages; and, second, even if it's not in a given language, replacing 'all' with 'both' (at the level of thought) is relatively cheap, hence could explain the challenge raised by French *tous les bras* 'all the arms'. Similarly, although our experimental results don't (and could never) definitively establish that 'all' is a lexical item of the language of thought, while 'some but not all' is not, they do support the hypothesis that the former is more primitive than the latter. As such, deriving alternatives via replacement with 'some but not all' is more costly than with 'all', which in turn gives us a considerable handle on solving the symmetry problem.

4.4 Connection with non-graded view of alternatives

The cost-based approach is compatible with a potentially more common, non-graded view of alternatives, according to which the inferential mechanism works from a set of alternatives which do not come with graded costs; a structure simply is or is not an alternative. Formally, one may still use the costs derived from a recipe as sketched above, and require that for a sentence to be in the set of (effective) alternatives, its cost has to pass some threshold. The threshold may also be variable from one occasion to the next, and the relative costs associated with different types of alternatives (i')–(iii') may then lead to generalizations of the form: when one disjunct is an alternative to a disjunction, then so is the other disjunction (because they come at the same cost). It would be interesting to see whether the well-described generalizations from Fox and Katzir, 2011 can be derived with appropriate parameters in such a model.

4.5 Influence of language on (the language of) thought

We have argued that at the level of thought, there is an inherent preference for the concept 'all' over the concept 'some but not all', and that such conceptual preferences help explain certain cross-linguistic lexicalization facts, such as the fact that language after language lexicalizes 'all', but not 'some but not all'. Thus, language proper (specifically, its lexicon) appears ready to host natural classes of concepts that fit whatever the constraints of the language of thought are.

However, if, as we have suggested, competition occurs at the level of thought, then nothing precludes the possibility that language proper (e.g. the lexicon) may also influence conceptual alternatives, either directly or indirectly. Directly, one could say that lexicalized alternatives get a cost reduction, simply because they are easily expressible both in the language of thought and presumably in language.

Indirectly, we may expect a bidirectional influence between language and language of thought, in a weak version of the Sapir-Whorf hypothesis. Concretely, assume, as we have argued, that competition occurs primarily at the level of thought, and that the concepts 'older brother' and 'younger brother' may well be universal (they appear in some languages, so they should be available at the level of the language of thought in the first place). Yet the fact that Japanese, but not English, lexicalizes them matters: as we have seen, the use of the word corresponding to the concept 'brother' in Japanese (*kyoodai*) implicates speaker uncertainty about 'older brother' vs. 'younger brother', whereas in English (*brother*) this is not the case. A possible translation of this in the current framework is that the costs associated with these concepts vary: they are moderate for all by default in the language of thought, but for Japanese speakers, the alternatives involving these concepts may benefit from a lower cost, simply because of the lexicalization. It's hard to describe, let alone test, the details of such an approach, but it's also hard to see reasons to discard this possibility.

4.6   Independent evidence for conceptual alternatives

Finally, we note that moving from alternatives as properly linguistic objects to alternatives as conceptual objects is potentially dangerous: we may lose the predictive power that the original Katzirian theory holds. In particular, without any independent constraints on the space of conceptual alternatives (or any independent measure of 'primitiveness'), we have little way of predicting which inferences should arise. In practice, it may be theoretically convenient to assume the existence of a conceptual alternative to explain some observed inference or other, but a lack of any independent evidence for such an alternative would undermine the explanatory power of the theory. Ideally, we should be able to turn the situation into an empirical question that can be tested empirically. Our practical recommendation, then, is to think about alternatives with an open mind, allowing non-lexicalized alternatives to play a role, and to seek independent evidence that these alternatives are indeed 'primitive'. Such independent evidence could come from standard linguistic investigations (e.g. finding their lexicalizations in different languages). Below we also propose the form of an experiment to gather such evidence.

5   Experiment

The goal of this experiment is to establish preferences between expressions, or rather between their counterparts at the level of thought (i.e. in a hypothesized language of thought). Using again the example from §2, we would like to know whether one can find an intrinsic preference for 'all' over 'some but not all', which would help solve the symmetry problem exposed above (as well as the

lexicalization facts, potentially). To do so, we used an implicit rule discovery task, very much inspired by Piantadosi et al., 2016 (as well as predecessors in the 1960s, e.g. Haygood and Bourne, 1965; King, 1966; Bourne, 1970). We adjusted the task, however, so as to be in a position to draw conclusions about pairwise differences between potential alternatives (while Piantadosi et al., 2016 had the wider-reaching ambition to evaluate as a whole the functional lexicon of the language of thought). The first two pairwise comparisons involve the 'some'/'all' scale of alternatives. Specifically, we will investigate whether 'all' is preferred over 'some but not all', thus explaining the paradigm in [11] (slightly modified from [3]). We will also explore whether 'no' is preferred over 'some but not all', thus explaining the parallel paradigm in [12] (a variant of [4]).[5]

[11] a. Some of the shapes are red.
      b. Not all of the shapes are red.   (observed inference)
      c. All of the shapes are red.   (actual alternative)
      d. Some but not all of the shapes are red.   (missing alternative)

[12] a. Not all of the shapes are red.
      b. Some of the shapes are red.   (observed inference)
      c. None of the shapes are red.   (actual alternative)

---

5 (Many thanks to Jacopo Romoli for raising the following issue clearly and forcefully.) In this experiment, we only address part of the problem, having to do with 'some but not all' not being an alternative neither in the direct case, (11), nor in the indirect case, (12). To be more complete, we would need to explain more, namely that:

  (i) Direct case: 'all' is an active alternative, while neither 'some but not all' nor 'not all' are.

  (ii) Indirect case: 'none' is an active alternative, while neither 'some but not all' nor 'some' are.

    The direct case doesn't seem to add much difficulty: the absence of 'not all' as an active alternative could simply be because 'not all' is intrinsically complex. This would need to be proven independently, but the ingredients seem to be there (on this view, 'not all' would not be an elementary primitive; it would have to be constructed from a negation and 'all'). Note that this would not be so easy to test because of the logical relations between 'all' and 'not all' and the constraints on the task (there needs to be an overlap between the items to be compared). The indirect case may be more challenging, however, because 'some' is likely to be a good primitive (given, e.g., that it's typically lexicalized), and so we need to explain why it does not play an active role in blocking the effect of the 'none' alternative. The experiment will not speak to this, but we mention here three possible solutions. First, even if both are primitives, it could be that 'some' is *less* primitive than 'none' (see §4.3). Second, if 'some' is a good primitive, then 'not some' is a very strong candidate in the context of 'not all'; that is, if single word replacements are favored in the first place, then this in fact makes 'not some' = 'none' a more salient candidate for being an alternative than 'some' on its own. Finally, there may well be additional constraints on alternative building not considered here, which could solve this issue. See, for instance, the 'Atomicity' constraint proposed by Trinh and Haida (2015) and their discussion of its application to indirect implicatures. It would be interesting, in fact, to investigate whether such a constraint is rooted in cognitive constraints or preferences of the kind explored here and elsewhere.

     d. Some but not all of the shapes are red.   (missing alternative)
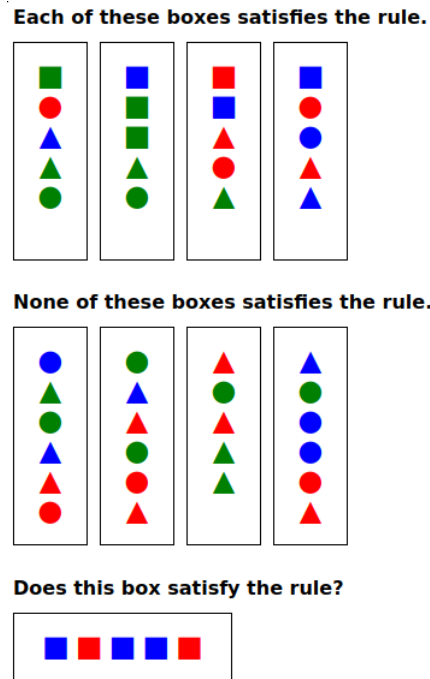
5.1 Participants

A group of 45 participants were recruited on Mechanical Turk and paid $3 each for their participation. IPs were restricted to the United States, and all participants reported to be native speakers of English. We precommitted to excluding participants with a median response time below 500ms on the assumption that they would not have been doing the task, but we found that this applied to none.

5.2 Design and task

Figure 1 presents an example of a stimulus that was shown to participants. Their task was to infer a rule that certain boxes of objects satisfy. They were presented with a group of boxes that they were told satisfy the rule (henceforth, *positive* boxes), and another group of boxes that they were told do not satisfy the rule (henceforth, *negative* boxes). At the bottom of the screen was another box, and the participants' task was to indicate whether they thought that this final box satisfied the rule, by clicking either "Satisfies the rule" or "Doesn't satisfy the rule".

    The experimental design involved two fully crossed factors: QUANTIFIERS and TYPE. The levels of the QUANTIFIERS factor are based on pairs of quantifiers, one of which is logically weaker than (entailed by) the other. In association with a particular property (random factor), a given quantifier determines a rule. For instance, in association with the property 'red' (see below for specific details about the stimuli), the quantifier 'all' determines a rule that can be stated as 'All the objects in the box are red'. Two quantifiers (in association with a given property) determine two rules which can be used to characterize the positive and negative boxes of items: positive boxes satisfied both rules, while negative boxes satisfied neither rule. The 8 levels of the QUANTIFIERS factor were: (some, all), (some, SBNA), (not all, no), (not all, SBNA), (at least 3, exactly 3), (at most 3, exactly 3), (at least 4, exactly 4), (at most 4, exactly 4). For the sake of the analysis, the last four were collapsed into two of the form (at least n, exactly n) and (at most n, exactly n). These numerical cases will allow us to check that the method can detect rather uncontroversial cases: the 'at least' rules should be easier than the 'at most' rules. This would align well with the known preference for reasoning with upward-monotonic over downward-monotonic quantifiers (Geurts and Slik, 2005) and with the fact that, in the absence of explicit modifier phrases, bare numerals more naturally acquire an 'at least' reading than an 'at most' reading (see Spector, 2013 for discussion).

    The TYPE factor, together with the two rules determined by the QUANTIFIERS factor, determined the final test box. It varied according to whether the final test

Each of these boxes satisfies the rule.

None of these boxes satisfies the rule.

Does this box satisfy the rule?

Figure 1: Example stimulus of the condition QUANTIFIERS=(some, SBNA), TYPE=TARGET, where the relevant property is 'square'. The positive boxes all satisfy the two rules 'some (or all) of the objects in the box are square' and 'some but not all of the objects in the box are square', while the negative boxes satisfy neither of those rules. The participants' task was to infer a rule and then to indicate whether the additional box at the bottom satisfies that rule. In this TARGET condition, the final box satisfies one but not the other of the two relevant rules.

box satisfied both of the rules determined by the QUANTIFIERS level (in the YES type), neither of them (in the NO type), or exactly one of them (in the TARGET type). As we will explain in further detail in the analysis, the YES and NO types serve as baselines to decide whether participants inferred a rule coherent with our goal, while the TARGET type helped us evaluate which of these two rules they favor: a positive response would correspond to the weaker rule (e.g. 'some'), a negative response to the stronger rule (e.g. 'SBNA').

The quantifiers in a pair had to allow for us to create boxes that satisfy both rules associated with them, none, or only one of them. Concretely, they always stood in an entailment relation. But we are also able to compare two quantifiers which do not stand in a logical entailment relation by comparing each of them to a third quantifier that they both entail and obtaining a measure of how strongly each of them is preferred (or dispreferred) with respect to that common reference point. For example, if we learn that 'all' is strongly preferred to 'some', while 'SBNA' is

only weakly preferred to 'some', then this indirectly indicates that 'all' is preferred to 'SBNA' (precisely the preference that would solve the symmetry problem).
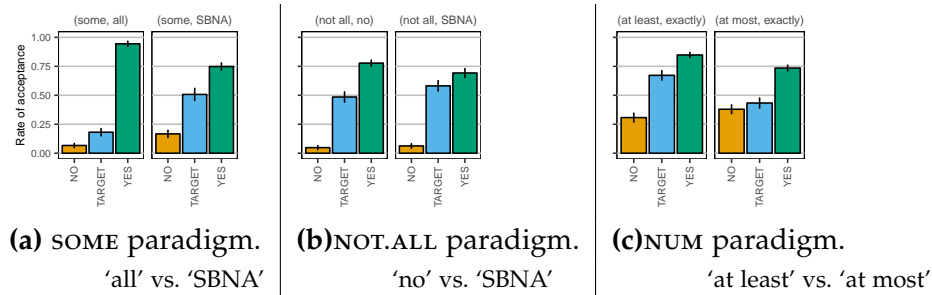
5.3 Stimuli

Figure 1 provides a full example of an item, with its characteristics described in the caption. Each item of a given condition was constructed by first randomly selecting the following superficial characteristics: the property required to determine the rules ('red', 'blue', 'green', 'triangle', 'square', or 'circle'), the number of positive boxes (3 or 4), the number of negative boxes (3 or 4), and for each box, the number of objects in that box (5 or 6). The shape and color of the objects in each box were then chosen so that each box would satisfy both, neither, or exactly one of the rules, depending on whether that box was a positive, negative, or test box. To simplify the task of the participants, and in particular the extraction of the relevant property, objects that satisfied the selected property were put together at the beginning (top or left) of their box.

We used a Python script (`https://osf.io/6qarq/`) and a Bash script (`https://osf.io/7dskm/`) to randomly construct the set of items that all participants saw, including 6 repetitions of each of the 24 combinations of the levels of TYPE (3 levels) and QUANTIFIERS (8 levels). The full set of 144 items was randomized for each participant. This randomized set was preceded by 4 practice items (the same ones, in the same order, for each participant). Their nature as practice items was not apparent to the participants, and they all involved a non-numerical QUANTIFIERS factor and a non-target TYPE factor. Our motivation for these choices was to ensure that participants did not hit a target item at the very beginning, and to ensure that participants did not see any discontinuity between the practice trials and subsequent trials. Once the list of items was generated offline, we used the online experimental platform created by Alex Drummond (`http://spellout.net/ibexfarm/`) to implement the task, presenting the (non-practice) items in a fresh random order to each participant.

5.4 Results

All data and an analysis script are available at `https://osf.io/uq3qp/`. Mean results are presented in Figure 2. First, note that the YES type (green bars) and NO type (orange bars) generate appropriate high and low acceptance rates, respectively (grand means of 79% and 17%), while the TARGET type (blue bars) generates intermediate response rates (average 50%). This shows that participants were paying attention and were able to satisfactorily perform what might seem like a difficult task (of a similar nature to the pattern recognition tasks in IQ tests, in that participants have to identify both the dimensions/properties to which the rule

(a) SOME paradigm.

'all' vs. 'SBNA'

(b) NOT.ALL paradigm.

'no' vs. 'SBNA'

(c) NUM paradigm.

'at least' vs. 'at most'

Figure 2: Proportion of 'yes' responses for each pair of QUANTIFIERS and each TYPE of item. NO cases made both rules attached to the quantifiers false; YES cases made them both true; and TARGET cases made one of them true and the other false. Error bars represent the standard error of the within-subject mean. Results are grouped by 'paradigm', i.e. pairs of QUANTIFIERS which meaningfully share one member ('some', 'not all', or 'exactly *n*') and are used for the comparison between their other members.

pertains and the rule itself).

Furthermore, observe the secondary but interesting difference between the QUANTIFIERS (some, all) and (not all, no). These two pairs are duals of each other (the first member of one is the negation of the second member of the other). Concretely, this means that the constraints on the positive and negative boxes were simply reversed from one to the other: what served as a potential YES test box for one was a NO test box for the other. Interestingly, we observe an asymmetry in response rates: the YES bar for (not all, no) is not as high as the NO bar for (some, all) is low. This suggests that participants did not perform the task via judgments of similarity (which of the two groups does the test box resemble the most?) — in which case the asymmetry would be unexpected — but rather engaged in a rule discovery task in which the positive exemplars played a different role from the negative exemplars. Overall, these first considerations suggest that participants solved the task as we were hoping.

Let us now move on to the question of how to interpret these data to shed light on the main point of interest. Imagine an ideal world in which participants always infer one of the two rules based on the quantifiers in the pair, and never any other rule. Then the YES items (where both rules are satisfied) and NO items (where neither rule is satisfied) should be at ceiling and floor, respectively. In this ideal world, the responses to TARGET items give us a direct measure of preference between the two rules. In actual fact, various biases, errors, and additional rules can play into the task, and hence YES and NO items are not at ceiling/floor. Still, the positioning of TARGET items within the range between the two extremes provided by YES and NO items provides a measure of preference: the closer TARGET is to YES, the more the weaker quantifier is preferred to the stronger. This rescaling of the

	ALL VS SBNA	NO VS SBNA	AT LEAST VS AT MOST
Δ_{elpd}	−32.0	−3.8	−11.7
$se(\Delta_{elpd})$	7.9	3.1	7.2

Table 1: Differences in estimated log pointwise predictive likelihood and their standard errors for models with one or two γ parameters for sets of quantifier pairs sharing a weaker member. Negative values of Δ_{elpd} indicate evidence in favor of the two-γ model.

range for TARGET items to the span between YES and NO eliminates noise created by general yes/no biases and rules that are weaker or stronger than both quantifiers in the pair (for more details, see the discussion in the appendix, §A).

For each pair p of quantifiers, we thus fitted models with three parameters: α_p determined the rate of acceptance of NO items, β_p determined the rate of acceptance of YES items, and γ_p determined where in the range between YES and NO we find the TARGET items. For details on the analysis to follow, see §A.

Recall, however, that what we are really interested in is comparing two quantifiers that are not ordered in a given pair. This we can now do by comparing the degree to which they are preferred or dispreferred relative to a third quantifier. We can compare two quantifiers y_1 and y_2 by fitting models for the data for the pairs (x, y_1) and (x, y_2), for some x, and seeing whether we obtain the same γ-parameter for both pairs. If there is a difference between the γ-parameters, this tells us that one of y_1 and y_2 is preferred to the other, in that it is more strongly preferred to x than the other is. The results of our Bayesian model comparison, done by leave-one-out cross-validation, are shown in Table 1.

We find evidence for a difference in all cases. Given the direction of the effect, preferences for 'at least' over 'at most' (our baseline case) and for 'all' over 'SBNA' are strongly supported, while there is more modest evidence for a preference of 'no' over 'SBNA'. (See details in §A, and Burnham and Anderson, 2002, §2.6 for heuristics about differences on the deviance scale; our Δ_{elpd} values are converted to that scale by multiplying them by −2.)

5.5 Discussion

The evidence we found for the preference of 'all' over 'some but not all' (and of 'no' over 'some but not all') provides a new solution to the symmetry problem, casts new light on the scalar inferences from 'some' to 'not all' (and the dual inference from 'not all' to 'some'), and gives new insight into the underlying source of cross-linguistic lexicalization facts. More generally, our work establishes a point of contact between linguistics and psychology, where the two domains seem to coincide: robust facts about language show the footprints of more general

properties of cognition.

We took inspiration from work in psychology to design a measure of 'primitivity' of logical elements, one that would apply both to non-linguistic rules and that could impact linguistic competition and linguistic lexicons. One may wonder whether the paradigm we used really taps into non-linguistic abilities, or whether its results merely reproduce the complexity facts at the linguistic (lexico-syntactic) level. We note that it's a complex issue. In particular, such an issue does not quite arise under a view of language as the ability to manipulate abstract representations, whether linguistic representations, non-linguistic representations, abstract rules, pronounced utterances, or unpronounced alternatives, etc. Concretely, it would be interesting to replicate our results under standard dual-tasks, which are supposed to block linguistic abilities (Norman and Bobrow, 1975), but failure to do so would not be very informative since this may simply show that the dual-tasks and the (linguistic) ability they block may already be at play for manipulating abstract rules in the first place. Investigating similar results with pre-linguistic infants or non-human animals could prove interesting too, allowing us to compare the logical repertoires of different animals, at various linguistic stages — that is, independently of the development of a pronounced language, at least as we know (hear/see) it.

Lastly, one may wonder whether the preferences we found truly constitute evidence for *a priori* conceptual preferences. An alternative interpretation would be that they are simply the result of Bayesian reasoning on the part of the participants about underlying rules, given the stimuli observed. Assume that participants have a flat prior over possible rules, and that they assume stimuli to be sampled from a uniform distribution over the space of possible stimuli that are compatible with or contradict a rule, respectively. Then the difference in likelihood of various stimuli, given the underlying rule, will result in different posterior probabilities for the rules in the same direction as found in our experiment. For example, consider the rules 'Some of the objects in the box are red' ('some') and 'All of the objects in the box are red' ('all'). The positive boxes that validate both rules contain only red objects. Imagine now that a participant sees an experimental item with positive boxes of all-red objects. If the rule is 'all', then this would necessarily be the case: the likelihood of all-red positive boxes, given the rule 'all', is 1. However, if the rule is 'some' and the positive boxes were randomly sampled from those possible boxes that validate the rule, then it would have to be quite the coincidence for all of them to be all-red. Hence, by Bayes' theorem, the underlying rule that generated the experimental item is more likely to be 'all' than it is to be 'some'. However, if we compare the two possible rules 'Some of the objects in the box are red' ('some') and 'Some but not all of the objects are red' ('SBNA'), then the situation is different: the set of possible boxes that validate both rules largely overlap, and while the 'some' rule would allow all-red boxes (which the 'SBNA'

rule does not), it would not require quite so great a coincidence for there to be no all-red boxes when positive examples are randomly sampled from the possible ones. Consequently, the observed preference for 'all' over 'some' is greater than any observed preference for 'SBNA' over 'some' (from which we inferred that 'all' is preferred over 'SBNA').[6] Thus, without further controls, our results don't tell us conclusively whether the preferences we observe reflect likelihood alone, *a priori* conceptual preferences alone (priors), or some combination of the two. Ideally, such an experiment would control for likelihood. Currently, we do not propose a way of doing so.[7] Nevertheless, we hope that future researchers will improve upon our proof-of-concept experiment so that these factors can be disentangled.

6 Conclusion

Alternatives are pervasive in natural language, and they lie at the heart of scalar inference, such as the inference from *some* to 'not all'. However, a complete theory of alternatives needs to explain how alternatives are constrained, e.g. why *all*, but not *some but not all*, is an alternative to *some*. That is, it needs to overcome the symmetry problem (Katzir, 2007; Fox and Katzir, 2011). We propose that alternatives are constrained both by the lexicon of a language, as standard wisdom has it, but also by more general restrictions on what counts as a primitive concept. Although the two sources remain necessary, we argued that the latter type of constraint deserves a better status, both for conceptual and practical reasons. Because this view relies on claims about the lexicon of an invisible language, it cannot be easily tested by standard means. But we propose that the relevant elements of this language can be inferred from regularities observed across languages (in language after language, we may find a word meaning 'all', but not necessarily one meaning 'some but not all'), or from general preferences in *a priori* non-linguistic tasks. Overall, a complete evaluation of the proposed approach will only be possible through a better integration of facts and methods that may typically come from psychology

---

6 We thank two anonymous reviewers who both pointed out this alternative interpretation of the results to us.

7 For example, one may try and reduce the number of objects in a box to just two, and the number of possible colors to just two (red and blue, say). In theory, then, a positive box that validates 'all' and one that validates 'SBNA' are equally constrained: both objects must be red in the former case, or one object must be red and the other blue in the latter. Thus, the likelihood of 'all' over 'some' is now equal to that of 'SBNA' over 'some'. There are at least two problems with this approach. First, it's unclear how one can ensure that the participants know that objects only come in pairs and can be only one of two possible colors. Even if they're told this information at the outset, it's difficult to determine whether they actually retain and use that information during the task. Second, even if that can be established or assumed, it's no longer clear whether the rules/concepts being tested are still 'some', 'all', and 'SBNA', rather than, say, 'some', 'both', 'exactly one', 'half', and so on, which thus complicates matters in other ways.

and formal linguistics independently of one another.

A Appendix: Statistical analysis

A.1 General description and notations

Semi-formally, our measure is based on three parameters for each quantifier pair p. First, $\overline{\alpha_p}$ represents our best estimate of the true rate of 'yes' responses in the NO cases, and $\overline{\beta_p}$ represents our best estimate of the true rate of 'yes' responses in the YES cases. In an ideal world, we should find $\overline{\alpha_p} = 0$ and $\overline{\beta_p} = 1$.[8] Critically, $\overline{\gamma_p}$ represents the preference between the weak and the strong quantifier, estimated as the position of the responses to the TARGET types not as an absolute rate, but as a proportion with $\overline{\alpha_p}$ and $\overline{\beta_p}$ as extreme points. In pseudo-formula, this means that the actual response in the TARGET type is predicted by $\overline{\alpha_p} + \overline{\gamma_p} \cdot (\overline{\beta_p} - \overline{\alpha_p})$. Concretely, a maximal preference for the weaker quantifier in p would correspond to $\overline{\gamma_p} = 100\%$, i.e. an expected response rate in the TARGET case as high as in the YES cases: $\overline{\alpha_p} + 100\% \cdot (\overline{\beta_p} - \overline{\alpha_p}) = \overline{\beta_p}$. Conversely, an extreme preference for the stronger quantifier would correspond to $\overline{\gamma_p} = 0\%$, i.e. an expected response rate in the TARGET case as low as in the NO cases: $\overline{\alpha_p} + 0\% \cdot (\overline{\beta_p} - \overline{\alpha_p}) = \overline{\alpha_p}$. Finally, a neutral preference corresponds to $\overline{\gamma_p} = 50\%$ and an expected response rate exactly in between $\overline{\alpha_p}$ and $\overline{\beta_p}$: $\overline{\alpha_p} + 50\% \cdot (\overline{\beta_p} - \overline{\alpha_p}) = (\overline{\alpha_p} + \overline{\beta_p})/2$.

To explain, imagine an ideal world in which participants always infer one of the two rules based on the quantifiers in the pair, and never any other rule. Then the YES items (where both rules are satisfied) and NO items (where neither rule is satisfied) should be at ceiling and floor, respectively. In this ideal world, the responses to TARGET items give us a direct measure of preference between the two rules. In actual fact, various biases, errors, and additional rules can play into the task, and hence YES and NO items are not at ceiling/bottom. Still, the positioning of TARGET items within the range between the two extremes provided by YES and NO items provides a measure of preference: the closer TARGET is to YES, the more the weaker quantifier is preferred to the stronger. This rescaling of the range for TARGET items to the span between YES and NO eliminates noise created by general yes/no biases and rules that are weaker or stronger than both quantifiers in the pair.

One source of simplification here is that we ignore the possibility that participants might infer a rule that is 'intermediate' between the two members of the pair;

---

8 In the upcoming implementation of this idea using logit models, the parameters will not directly represent rates of 'yes' responses, and they will not range between 0 and 1 but rather between $-\infty$ and $+\infty$. To convey the spirit of the analysis more simply, however, we ignore this technicality in this paragraph. To keep track of it, however, we mark the parameters in this informal description with a line above them, as in $\overline{\alpha_p}$.

for example, 'at least 3' and 'more than half' are intermediate between 'some' and 'all'. Depending on the specific rule participants infer and the particular test box showing, this could alter the rate of 'yes' responses, but not because participants opt for quantifiers *from our pair*. We thus work on the assumption that the two rules we compare are the most salient in their range. However, we also note that our main interpretations are based not only on pairs, but on the comparison between two pairs of quantifiers. We submit that the intermediate rules relevant for these two pairs are similar across the pairs we use, so that their role should be cancelled out in our analysis. For instance, when we compare 'all' and 'SBNA' through the pairs (some, all) and (some, SBNA), we assume that the relevant intermediate rules, which are compatible with the positive and negative examples being what they are, and which could make the test box true, are of the form 'at least x' in both cases.

A.2 Models and analyses

The logit models we fit thus had a slightly different structure from the generalized linear models commonly used in the analysis of binary response data. Their basic form was as follows:

$$Y_{pi} \sim \text{bernoulli}(\text{logit}^{-1}(\pi_{pi}))$$
$$\text{with: } \pi_{pi} = \alpha_p \cdot \text{NO}_i + \beta_p \cdot \text{YES}_i + (\alpha_p + \gamma_p \cdot (\beta_p - \alpha_p)) \cdot \text{TARGET}_i.$$

Given a pair of quantifiers $p = (x, y)$, the parameter $\gamma_{(x,y)}$ thus represents the preference for the weaker member over the stronger member. Given two pairs of quantifiers (x, y_1) and (x, y_2) that share a common member x, we can then compare y_1 and y_2 by fitting a model to the data from both pairs and seeing whether there is evidence for a difference between $\gamma_{(x,y_1)}$ and $\gamma_{(x,y_2)}$. If there is one, then one of the y's is preferred to the other, in that it is more distinct from x than the other is.

We partitioned our data into three sets based on the quantifier pairs that shared a member: (i) a set comprising data from (some, all) and (some, SBNA); (ii) a set comprising data from (not all, no) and (not all, SBNA); and (iii) a set comprising data from (at least n, exactly n) and (at most n, exactly n). On each of these three data sets, we fitted models that included either one or two γ parameters (allowing, or not allowing, for a relative preference between the two non-shared quantifiers in the two pairs, as discussed in the main text),[9] and different varying-effects structures: there was (or was not) a subject intercept u_{0s}, and either no modification of the parameters, parameter modifiers differing by subject, or parameter modifiers differing by both subject and pair. All models were fitted with MCMC methods using STAN through the rstan package in R,

---

9 α and β were always allowed to be distinct for the two pairs.

	ALL VS SBNA	NO VS SBNA	AT LEAST VS AT MOST
Δ_{elpd}	−32.0	−3.8	−11.7
$se(\Delta_{elpd})$	7.9	3.1	7.2

Table 2: Differences in estimated log pointwise predictive likelihood and their standard errors for models with one or two γ parameters for sets of quantifier pairs sharing a weaker member. Negative values of Δ_{elpd} indicate evidence in favor of the two-γ model.

with STAN's default uniform (improper) prior over α, β, and γ (in the latter case restricted to $[0, 1]$) as well as the standard deviation hyperparameters for the varying subject effects (which themselves were assumed to be normally distributed around 0). Models were evaluated by leave-one-out cross-validation, approximated by Pareto-smoothed importance sampling with the loo package (Vehtari et al., 2016) on the basis of 5,000 samples of the likelihood for each data point, drawn after 5,000 burn-in iterations.

A.3   Results

For all three data sets, the best model turned out to be the two-γ model with a subject intercept and parameter modifiers varying by subject, but not by pair:[10]

$$Y_{spi} \sim \text{bernoulli}(\text{logit}^{-1}(\pi_{spi}))$$
$$\text{with: } \pi_{spi} = (\alpha_p + u_{\alpha s}) \cdot \text{NO}_i + (\beta_q + u_{\beta s}) \cdot \text{YES}_i$$
$$+ (\alpha_p + u_{\alpha s} + (\gamma_p + u_{\gamma s})((\beta_p + u_{\beta s}) - (\alpha_p + u_{\alpha s}))) \cdot \text{TARGET}_i + u_{0s}.$$

In order to see whether the preference for two γ-parameters was meaningful, we compared the optimal models to the variants with identical varying effects structure, but only one (pair-independent) γ-parameter. Differences in estimated log pointwise predictive likelihood and their standard errors for models with one and two γ-parameters are given in Table 2.

In addition, posterior means and 95% credible intervals for the difference between γ-parameters are shown in Table 3:

With a Δ_{elpd} of −32.0, there is strong evidence for a preference for 'all' over

---

10 Models where parameter modifiers differed by subject and pair had this form:

$$Y_{spi} \sim \text{bernoulli}(\text{logit}^{-1}(\pi_{spi}))$$
$$\text{with: } \pi_{spi} = (\alpha_p + u_{\alpha sp}) \cdot \text{NO}_i + (\beta_q + u_{\beta sp}) \cdot \text{YES}_i$$
$$+ (\alpha_p + u_{\alpha sp} + (\gamma_p + u_{\gamma sp})((\beta_p + u_{\beta sp}) - (\alpha_p + u_{\alpha sp}))) \cdot \text{TARGET}_i + u_{0s}.$$

	$\gamma_{(\text{some,all})}$ $-\gamma_{(\text{some,SBNA})}$	$\gamma_{(\text{not all,no})}$ $-\gamma_{(\text{not all,SBNA})}$	$\gamma_{(\text{at least,exactly})}$ $-\gamma_{(\text{at most,exactly})}$
mean	-0.58	-0.19	0.52
95% CI	$(-0.75, -0.41)$	$(-0.32, -0.07)$	$(0.34, 0.70)$

Table 3: Posterior means and 95% credible intervals for differences between γ parameters, reflecting relative preference between quantifiers not shared between the pairs.

'SBNA'.[11] The direction of this preference is given by the negative difference in γ parameters. If x is stronger than y in a pair (x, y), then a low $\gamma_{(x,y)}$ indicates a stronger preference for x, and since $\gamma_{(\text{some,all})}$ is smaller than $\gamma_{(\text{some,SBNA})}$, we know that 'all' is preferred to 'some' more strongly than 'SBNA' is preferred to 'some', hence that 'all' is preferred to 'SBNA'.

The same reasoning applies to the comparison of 'no' with 'SBNA' via comparison of each with 'not all'. In this instance, however, the effect is smaller and the strength of the evidence, with a Δ_{elpd} of -3.8, much more modest.

For 'at least' and 'at most', a greater γ value indicates a stronger preference over 'exactly', because now the quantifiers of interest are the weaker members of the pair. Since $\gamma_{(\text{at least,exactly})}$ is greater than $\gamma_{(\text{at most,exactly})}$, we can infer a preference for 'at least' over 'at most'. The effect size here is in the same ballpark as with 'all' vs. 'SBNA', and the evidence quite strong.

References

Bergen, Leon, Roger Levy, and Noah D. Goodman (2016). Pragmatic Reasoning through Semantic Inference. In: *Semantics and Pragmatics* 9. DOI: `10.3765/sp.9.20`.

Bourne, Lyle E. (1970). Knowing and Using Concepts. In: *Psychological Review* 77.6, pp. 546–556. DOI: `10.1037/h0030000`.

Buccola, Brian and Emmanuel Chemla (2019). Alternatives of Disjunctions: When a Disjunct Contains the Antecedent of a Pronoun. In: *Snippets* 37, pp. 16–18. DOI: `10.7358/snip-2019-037-buch`.

Burnham, Kenneth P. and David R. Anderson (2002). *Model Selection and Multimodel Inference: A Practical Information-Theoretic Approach*. second. New York, NY: Springer-Verlag.

Charlow, Simon (2019). Scalar Implicature and Exceptional Scope. Rutgers University. URL: `https://ling.auf.net/lingbuzz/003181`.

---

11 See Burnham and Anderson, 2002, §2.6 for heuristics about differences on the deviance scale. Our Δ_{elpd} values can be converted to that scale by multiplying them by -2.

Chemla, Emmanuel (2007). French *both*: A Gap in the Theory of Antipresupposition. In: *Snippets* 15, pp. 4–5. URL: http://www.ledonline.it/snippets/allegati/snippets15001.pdf.

Chemla, Emmanuel, Brian Buccola, and Isabelle Dautriche (2019). Connecting Content and Logical Words. In: *Journal of Semantics* 36.3, pp. 531–547. DOI: 10.1093/jos/ffz001.

Chierchia, Gennaro (2013). *Logic in Grammar: Polarity, Free Choice, and Intervention*. Oxford, England: Oxford University Press.

Chierchia, Gennaro, Danny Fox, and Benjamin Spector (2012). Scalar Implicature as a Grammatical Phenomenon. In: *Semantics: An International Handbook of Natural Language Meaning*. Ed. by Claudia Maienborn, Klaus von Heusinger, and Paul Portner. Vol. 3. Berlin, Germany: Mouton de Gruyter, pp. 2297–2331.

Fox, Danny (2007). Free Choice and the Theory of Scalar Implicatures. In: *Presupposition and Implicature in Compositional Semantics*. Ed. by Uli Sauerland and Penka Stateva. Palgrave Studies in Pragmatics, Language and Cognition Series. New York, NY: Palgrave Macmillan. Chap. 4, pp. 71–120. DOI: 10.1057/9780230210752_4.

Fox, Danny and Roni Katzir (2011). On the Characterization of Alternatives. In: *Natural Language Semantics* 19.1, pp. 87–107. DOI: 10.1007/s11050-010-9065-3.

Gamut, L. T. F. (1991). *Logic, Language, and Meaning*. Vol. 1 (Introduction to Logic). Chicago, IL: University of Chicago Press.

Gärdenfors, Peter (2014). *The Geometry of Meaning: Semantics Based on Conceptual Spaces*. Cambridge, MA: The MIT Press.

Gazdar, Gerald (1979). *Pragmatics: Implicature, Presupposition, and Logical Form*. New York, NY: Academic Press.

Geach, Peter T. (1962). *Reference and Generality: An Examination of Some Medieval and Modern Theories*. Ithaca, NY: Cornell University Press.

Geurts, Bart and Frans van der Slik (2005). Monotonicity and Processing Load. In: *Journal of Semantics* 22.1, pp. 97–117. DOI: 10.1093/jos/ffh018.

Grice, H. P. (1975). Logic and Conversation. In: *Syntax and Semantics, Volume 3: Speech Acts*. Ed. by Peter Cole and Jerry L. Morgan. New York, NY: Academic Press, pp. 41–58.

Haygood, Robert C. and Lyle E. Bourne (1965). Attribute- and Rule-Learning Aspects of Conceptual Behavior. In: *Psychological Review* 72.3, pp. 175–195. DOI: 10.1037/h0021802.

Horn, Laurence R. (1972). On the Semantics of Logical Operators in English. PhD thesis. New Haven, CT: Yale University.

Horn, Laurence R. (1992). The Said and the Unsaid. In: *Semantics and Linguistic Theory (SALT)*. Ed. by Chris Barker and David Dowty. Vol. 40. Working Papers

in Linguistics. Columbus, OH: The Ohio State University, pp. 163–192. DOI: 10.3765/salt.v2i0.3039.

Kadmon, Nirit and Fred Landman (1993). Any. In: *Linguistics and Philosophy* 16.4, pp. 353–422. DOI: 10.1007/BF00985272.

Karttunen, Lauri (1969). Pronouns and Variables. In: *Chicago Linguistic Society (CLS)*. Vol. 5, pp. 108–116.

Karttunen, Lauri and Stanley Peters (1979). Conventional Implicature. In: *Syntax and Semantics*. Vol. 11. Leiden, Netherlands: Academic Press, pp. 1–56. DOI: 10.1163/9789004368880_002.

Katzir, Roni (2007). Structurally-Defined Alternatives. In: *Linguistics and Philosophy* 30.6, pp. 669–690. DOI: 10.1007/s10988-008-9029-y.

Katzir, Roni and Raj Singh (2013). Constraints on the Lexicalization of Logical Operators. In: *Linguistics and Philosophy* 36.1, pp. 1–29. DOI: 10.1007/s10988-013-9130-8.

King, William L. (1966). Learning and Utilization of Conjunctive and Disjunctive Classification Rules: A Developmental Study. In: *Journal of Experimental Child Psychology* 4.3, pp. 217–231. DOI: 10.1016/0022-0965(66)90022-1.

Magri, Giorgio (2009). A Theory of Individual-Level Predicates Based on Blind Mandatory Scalar Implicatures. In: *Natural Language Semantics* 17.3, pp. 245–297. DOI: 10.1007/s11050-009-9042-x.

Matsumoto, Yo (1995). The Conversational Condition on Horn Scales. In: *Linguistics and Philosophy* 18.1, pp. 21–60. DOI: 10.1007/BF00984960.

Norman, Donald A. and Daniel G. Bobrow (1975). On Data-Limited and Resource-Limited Processes. In: *Cognitive Psychology* 7.1, pp. 44–64. DOI: 10.1016/0010-0285(75)90004-3.

Piantadosi, Steven T., Joshua B. Tenenbaum, and Noah D. Goodman (2016). The Logical Primitives of Thought: Empirical Foundations for Compositional Cognitive Models. In: *Psychological Review* 123.4, pp. 392–424. DOI: 10.1037/a0039980.

Rooth, Mats (1985). Association with Focus. PhD thesis. Amherst, MA: University of Massachusetts Amherst.

Sauerland, Uli (2004). Scalar Implicatures in Complex Sentences. In: *Linguistics and Philosophy* 27.3, pp. 367–391. DOI: 10.1023/B:LING.0000023378.71748.db.

Schlenker, Philippe (2017). Sign Language and the Foundations of Anaphora. In: *Annual Review of Linguistics* 3, pp. 149–177. DOI: 10.1146/annurev-linguistics-011415-040715.

Seuren, Pieter A. M. and Dany Jaspers (2014). Logico-Cognitive Structure in the Lexicon. In: *Language* 90.3, pp. 607–643. DOI: 10.1353/lan.2014.0058.

Smith, Ryan Walter (2020). Similative Plurals and the Nature of Alternatives. PhD thesis. University of Arizona. URL: https://semanticsarchive.net/Archive/zZiMDljZ/.

Spector, Benjamin (2013). Bare Numerals and Scalar Implicatures. In: *Language and Linguistics Compass* 7.5, pp. 273–294. DOI: 10.1111/lnc3.12018.

Strickland, Brent (2017). Language Reflects "Core" Cognition: A New Theory About the Origin of Cross-Linguistic Regularities. In: *Cognitive Science* 41.1, pp. 70–101. DOI: 10.1111/cogs.12332.

Szabolcsi, Anna (June 25, 2020). Obviation in Hungarian: What is Its Shape, and is It Due to Competition? Colloquium talk given at Tel Aviv University. Tel Aviv, Israel.

Trinh, Tue and Andreas Haida (2015). Constraining the Derivation of Alternatives. In: *Natural Language Semantics* 23.4, pp. 249–270. DOI: 10.1007/s11050-015-9115-y.

Vehtari, Aki, Andrew Gelman, and Jonah Gabry (2016). Practical Bayesian Model Evaluation Using Leave-One-Out Cross-Validation and WAIC. In: *Statistics and Computing*, pp. 1–20. DOI: 10.1007/s11222-016-9696-4.

Xu, Fei and Joshua B. Tenenbaum (2007). Word Learning as Bayesian Inference. In: *Psychological Review* 114.2, pp. 245–272. DOI: 10.1037/0033-295X.114.2.245.