

Chapter 2

On the Locus of Expressivity

Deriving Parallel Meaning Dimensions from Architectural Considerations

Andrés Saab

1. Introduction

In this chapter I would like to defend the claim that certain non-truth conditional meanings are exclusively triggered by properties of vocabulary items (in the sense of Halle and Marantz 1993, and subsequent work). In other words, those meanings arise “late” and are not part of the syntactic-semantic derivation. Thus, I capture the notion of *parallel meaning dimension* from architectural considerations without the need for any metalogic operator (e.g., the symbol \bullet in Potts 2005 and McCready 2010, among others) especially designated to separate meaning dimensions. The basic expressive paradigm I am concerned with involves pairs of mixed words in Argentinian Spanish whose contribution to truth-conditional meaning is equivalent. They are only differentiated by register (e.g., ‘*comer*’ “to eat” vs. ‘*morfar*’ “to eat_{inf}”) or by register plus a derogative dimension (e.g., ‘*sudamericano*’ “South-American” vs. ‘*sudaca*’ “South-American_{pejorative}”). Here are two crucial properties of both informal and slur terms:¹

- (A) In the general case, expressives form doublets:² ‘*sudamericano/sudaca*’ “South-American”/“South-American(pejorative)”, ‘*comer*’/‘*morfar*’, “to eat”.
- (B) One of the members of these pairs contributes to a certain meaning dimension: style, tone or expressivity.

My initial conjecture is that at least some forms of expressivity are the direct result of free (lexical) variation; i.e., competition in the paradigmatic space among truth-conditionally identical terms gives rise to expressive meanings, an idea that has its roots in the functionalist tradition initiated by Trubetzkoy (1939) and others. Free variation implies competition in the paradigmatic (and phonetic) space and, at least in some cases, a non-truth conditional contribution to meaning. The final picture, I claim, results in a theory that integrates Fregean tone into semantics, against Frege's original considerations in this respect.

It makes no difference to the thought whether I use the word 'horse' or 'steed' or 'cart-horse' or 'mare'. The assertive force does not extend over that in which these words differ. What is called mood, fragrance, illumination in a poem, what is portrayed by cadence and rhythm, does not belong to the thought. (Frege 1918/1956, 23)

The idea that tone is not involved in the notion of truth in any relevant sense and, consequently, "does not belong to the thought" was revived, although not literally, by Capelen and Lepore (2013a, 2013b) and by Lepore and Stone (2018).

[...] tone, unlike meaning, does not seem to be a feature of language that speakers negotiate among one another and coordinate on. Indeed, tone, unlike meaning, does not seem to be something that speakers generally command in virtue of knowing their

language or universally respect in their choices of linguistic behavior [...] In short, Frege was right: tone “is not part of the thought expressed.” (Lepore and Stone 2018, 144)

In a few words, Lepore and Stone shares Frege’s view that tone is not part of the semantic agenda. Once that one acknowledges the basic fact that linguistic expressivity is subject to linguistic regularities (i.e., to a set of conventions that governs some sort of semantic competence), that view turns out to be somewhat mistaken. To be competent with the use of an informal term or a slur is to know the rules that make such uses expressively correct (see Kaplan 1999 and, more recently, Predelli 2013 and 2020). That the conditions behind expressive correctness don’t involve the notion of truth doesn’t make expressivity in natural language less “semantic” than those conditions that guarantee truth or falsity. Then, I agree with Kaplan and Predelli in that there is no obvious reason why the limits of semantics should be restricted in the narrow sense Frege suggested. Put differently, expressivity has the right to pertain to the research agenda of the theory of linguistic meaning or, perhaps more properly, I should say to the theory of linguistic *meanings*. Following also original insights in Trubetzkoy (1939), I call this research agenda “stylistic semantics”. Stylistic semantics is about a parallel dimension of conventional, non-truth conditional meaning. In this respect, the theory is (at least) two-dimensional, as in Kaplan (1999), Predelli (2013), Potts (2005), McCready (2010) and Gutzmann (2015), among others. However, as I have already advanced, I contend that expressivity is not represented in the Logical Form (LF) of a sentence, but it is deduced at Phonetic Form (PF), through a principle of lexical competition to be discussed in the next section. Assuming the Y-model of grammar, we can get a flavor of the division of labor in semantics in the following figure:




Figure 1.1: *Meaning dimensions and grammatical architecture*

Here, the symbol \bullet , used in Potts (2005) and others to separate meaning dimensions, is derived from the division in the Y-model. In this way, we dispense with this operator and propose that the division is given by the architecture of the grammar. The consequence of this is that I also dispense with mixed lexical entries like the following one for words like ‘*sudaca*’, which is the pejorative term for South-Americans in Argentinian Spanish:

$$(2) \llbracket \text{sudaca} \rrbracket = \lambda x. \text{South-American}(x) \blacklozenge \text{Bad}(\cap \text{South-American}): \langle e, t \rangle^a \times t^s$$

This is the lexical representation proposed in McCready (2010) for slur terms. It contains both truth-conditional and non-truth-conditional information. The latter is modeled as a conventional implicature (following Grice 1975 and, in particular, Potts 2005). The symbol \blacklozenge in this representation is changed for the symbol \bullet through a rule that allows interpreting conventional

implicatures in the semantic derivation. Putting technical details aside, compare the lexical entry in (2) to the following simplified lexical entry I propose is active in the syntax:

(3) Syntax-semantics: $\llbracket \text{sudaca} \rrbracket = \lambda x. \text{South-American}(x)$

According to (3), the denotation of ‘*sudaca*’ is equivalent to that of “South-American”; both denote the set of South-Americans. It is only when this word is sent to the PF component of the grammar (i.e., after syntax) that the expressive meaning is added to the phonological matrix of the relevant lexical item. This requires some particular implementation of lexical representation. In particular, it requires taking for granted that the phonetic makeup of lexical items is determined after syntax, concretely, at the PF interface.

The empirical argument in favor of this implementation of stylistic semantics comes from the behavior of expressives in the realm of ellipsis. The observation is that expressivity doesn’t survive ellipsis.³ This was already shown by Potts *et al* (2009) for ellipsis sites which take as antecedents predicates containing attributive adjectives like “fucking” in English (< > = ellipsis site):⁴

(4) A: I saw your fucking dog in the park.

B: No, you didn’t <see my dog> —you couldn’t have <seen my dog>. The poor thing passed away last week.

(Potts *et al* 2009, 364, ex. (32))

Here, the two ellipsis sites are modeled without expressives, an analysis justified by the clear intuition that B doesn't endorse A's hostility to the dog at hand. Potts *et al*'s conclusion is that ellipsis only requires recovery of descriptive content, an idea in consonance with the approach to ellipsis in Merchant (2001), according to which ellipsis identity boils down to mutual entailment between the antecedent and the ellipsis site. In principle, this seems to be correct when it comes to evaluating the distribution of mixed terms in ellipsis. Consider the following short dialogue:

- (5) A: Qué morfaste?
 what ate.2SG.INFORMAL
- B: Una pizza <?>, pero no tolero cuando
 a pizza but not tolerate.1SG when
 hablás tan informalmente. Yo nunca lo hago.
 speak.2SG so informally I never it do
 “A pizza. But I don't tolerate when you speak informally. I never do it.”

Here B's answer contains an ellipsis site that, unlike what happens in the antecedent, cannot be interpreted as containing the verb '*morf*', which is the informal term for the verb '*comer*' "to eat". In principle, we can generalize Potts *et al*'s argument and conclude that given that '*morf*' and '*comer*' can be substituted *salva veritate*, mutual entailment is satisfied and ellipsis applies. If this approach is correct, then ellipsis cannot help us distinguishing among different theories of expressivity in competition. However, I argue, following previous proposals, that ellipsis requires lexical identity as a necessary condition (Chomsky 1965, Chung 2006 and Saab 2008, among many others). If this is correct, then ellipsis can be indeed used as a relevant test, one that

seems to favor the approach to expressivity that I sketch in this chapter. I elaborate the argument from ellipsis in section 3 for informal terms and, in section 4, I extend the analysis to slurs. But before entering into the details concerning the behavior of mixed terms under ellipsis, I sketch my analysis for mixed terms. As I have already advanced, my implementation requires two core ingredients: (a) a model of late insertion for phonological information, and (b) a principle of lexical competition that, among other things, accounts for the fact that certain expressive terms require a non-expressive counterpart. These are the themes of the following section.

2. Implementation: Expressivity as a PF Property

2.1. Background Assumptions

I assume the Distributed Morphology framework (Halle and Marantz 1993, 1994 and, in particular, Embick 2000, 2007, Embick and Noyer 2001, Arregi and Nevins 2012, Marantz 2013, and Harley 2014, among many others). The general architecture is the following:

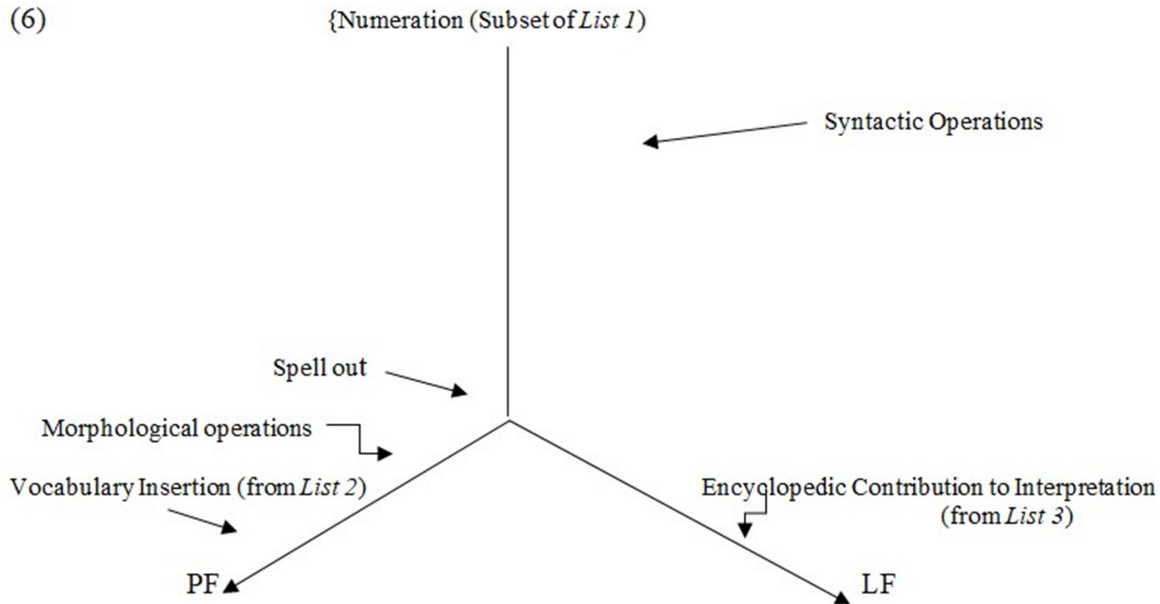


Figure 2.2: *Architecture of the grammar* (from Harley 2014, 228)

A crucial property of this conception of the grammar is form-meaning separation, i.e., the fact that meaning-form connections are determined by the syntax in an all-the-way-fashion (Halle and Marantz 1994). Syntax manipulates abstract objects (from List 1) that are supplied with a given phonological form after syntax through a set of lexical insertion rules that adds items from List 2. The primitives that syntax manipulates are Roots and abstract morphemes. Abstract morphemes are features drawn from a Universal Inventory and encode things like [past], [plural] and so on. Roots are represented by an index that is replaced at PF by a phonetic matrix (Chomsky 1995, Embick 2000, Saab 2008, Acquaviva 2008, and Harley 2014, among others). Finally, as for the LF side, syntax provides an abstract object built out from Roots and abstract morphemes, which is interpreted on the basis of the information available in List 3. Let me summarize the information contained in each of the lists in (6) (simplified from Harley 2014, 228):

List 1: Feature bundles: Syntactic primitives, both interpretable and uninterpretable, functional and contentful.

List 2: Vocabulary Items: Instructions for pronouncing terminal nodes.

List 3: Encyclopedia: Instructions for interpreting terminal nodes.

Thus, the full representation of an abstract morpheme like, say, the imperfect in Spanish (‘-ba/-ía’) requires information taken from the three lists:

List 1: [imperfect past]

List 2: [imperfect past] \leftrightarrow /-ía/ / Thematic Vowel_{2,3} ____

[imperfect past] \leftrightarrow /-ba/

List 3: [imperfect past] \leftrightarrow [[truth-conditional meaning]]

For a Root like ‘gat(o)’ “cat”, the relevant information, distributed in each list, can be summarized as follows:

List 1: $\sqrt{13}$

List 2: $\sqrt{13} \leftrightarrow$ /gát-/

List 3: $\sqrt{13} \leftrightarrow$ [[λx . Cat(x)]] / ____ n

In sum, on a DM conception, there is not a generative lexicon that syntax uses to construct syntactic objects; lexical objects are instead syntactic constructs derived by the same principles that allow us to derive phrasal objects. I contend now that expressivity, at least for the

cases I am concerned with, is an exclusive property of vocabulary items, that is, of objects taken from List 2.

2.2. Expressivity as a List 2 Property

With these background assumptions in mind, we can now state the basic hypothesis in the following way:

- (7) (At least) some forms of expressivity are exclusively encoded in List 2.

Let's illustrate the theory with the pair '*trabajar*' "to work" vs. '*laburar*' "to work_{informal}".

The crucial step is adding the informal flavor of ‘*laburar*’ as part of the vocabulary item:

- (8) a. *trabajar*

- b. *List 1:* $\sqrt{122}$

List 2: √122 ↔ /trabaj(ár)/ / ____ v

List 3: $\sqrt{122} \leftrightarrow \llbracket \lambda x. \text{Work}(x) \rrbracket / \quad \nu$

- (9) a. *laburar*

- b. *List 1:* $\sqrt{221}$

List 2: √221 ↔ /labur(ár)/_{BIAS} / ____ v

BIAS: $c \in CU(\text{laburar})$ only if, in c , c_a is a participant in register *informal*

List 3: $\sqrt{221} \leftrightarrow [\![\lambda x. \text{Work}(x)]\!] / __ v$

Let's assume that the expressive meaning of '*laburar*' is a bias on contexts of use (CU), which, in this particular case, characterizes the agent of the context as a participant in an informal register (Predelli 2013). We can illustrate the representation in (9) with the scheme in (10), assuming that the index for the Root of the informal term '*laburar*' is $\sqrt{35}$:




Figure 2.3: *Bias for 'laburar'*

As I advanced in the introduction, on this theory, the metalogic symbol ● is trivial; it is deduced from architectural considerations and a principle of lexical competition. A crucial question is, of course, how lexical competition in the paradigmatic space is determined. I conjecture that the following general principle of lexical competition is active in natural languages:

(11) *Principle of stylistic meaning (PSM)*: Given a pair of abstract nodes X and Y, taken from List 1, if X and Y are not semantically distinguishable at LF (List 3), they are “semantically” distinguished at PF (List 2).

The notion of *meaning* in this formulation requires some qualifications. If Y is semantically distinguished from X at LF, we say that the meaning contribution of both lexical items is truth-conditionally relevant. If Y is semantically distinguished from X at PF we say that the meaning contribution of both lexical items is stylistic or expressive, i.e., non-truth-conditionally relevant. Alternatively, we can characterize as *stylistic* those meanings that arise only at PF by virtue of a principle of lexical competition. In any case, the distinction is semantic in the favored sense of stylistic semantics. PSM is observed for the pairs of Roots ‘*laburar*’ and ‘*trabajar*’. And in this case in particular, we are led to conclude that the difference is purely stylistic. At any rate, note that PSM is not a principle of synonymy blocking: ‘*trabajar*’ and ‘*laburar*’ are still synonymous (truth-conditionally equivalent). I think this is an important result, but I will not explore its consequences here.

The considerations made here must be taken as a general sketch of how a theory of stylistic semantics should look like according to my assumptions. My main aim in this essay is just to argue that there is preliminary evidence that a radical dissociation between truth-conditional and expressive meaning is needed. If the general sketch turns out to be correct, then an explicit theory of stylistic semantics can be developed on a formal basis. There are, of course, many questions opened by this general approach. A pressing one is how the conceptual-intentional cognitive system accesses the information provided by the computational system of the language faculty. According to Chomsky (1995, 2000 and 2001), there are two cognitive systems that use the information provided by the PF and LF interfaces, namely, the sensorimotor, articulatory-perceptual system (S-M) and the conceptual-intentional system (C-I) related to systems of thought. In the classical Y model, the C-I system only uses information from the LF

interface. According to Reinhart (2006), it seems that such a cognitive system relates lexical and syntactic information to a set of disparate interpretation processes, involving, at least, inference, context and concepts. Reinhart (2006) proposes, then, associating these processes to three distinct cognitive systems, namely, context, concept and inference. For the purposes of this chapter the context system is the most relevant one and also the hardest to define:

The hardest to define given our present state of knowledge are the context systems that narrow the information transmitted through the derivation (coded in the relevant representation), and select the information that is useful for the context of use. (Reinhart 2006, 4)

To be competent with a slur, for example, requires knowing the conditions that make it correct in a certain context of use. This supposes manipulation of information other than purely inferential or conceptual. Both are required, yet. A pair of terms opposed in register is constituted as such because the concepts expressed by each member are identical. Yet, they are used in different contexts for different purposes, and this is something that requires linking lexical and inferential content to particular contexts of use. It is in this sense that I assume here that the context system must have access both to the PF and LF interfaces. I leave these issues open and turn my attention to the empirical argument in favor of the proposed meaning division.

3. The Argument from Ellipsis

3.1. Expressive Mismatches under Ellipsis

For many kinds of ellipses in natural languages, the following generalization seems to be relevant:

*Vehicle Change Generalization (Barros and Saab 2016):*⁵

- (12) Recoverability conditions in ellipsis make reference to content not character.

This observation generalizes over different types of well-known ellipsis mismatches, Vehicle Change and indexical mismatches being two prominent ones:

Classic Vehicle Change (Fiengo and May 1994 and Merchant 2001):

- (13) a. They arrested the man₃, but he₃ doesn't know why.
b. They arrested the man₁, but he₁ doesn't know why <they arrested *the man₁/him₁>.

Indexical mismatches (Thoms 2013, 2015):

- (14) A: Can you help me? [requesting help]
B: Yes, I can <help you>.

This type of mismatches follows if identity in ellipsis is stated in purely semantic terms. Merchant (2001) famously proposed a semantic condition based on the crucial notion of *e*(ellipsis)-GIVENness. Here is Merchant's condition:

Focus condition:

- (15) A constituent α can be deleted only if α is *e*-GIVEN.
- (16) An expression E counts as *e*-GIVEN iff E has a salient antecedent A and, *modulo* \exists -type shifting,
- i. A entails the Focus closure of E (written F-clo(E)), and
 - ii. E entails F-clo(A)
- (17) F-clo(α) is the result of replacing F-marked parts of α with \exists -bound variables.

Take as an example a simple case of VP-ellipsis in English like (18a) and its associated (and simplified) semantic derivation in (18b):

- (18) a. [_F Ann] loves Peter and [_F Mary] does <love Peter> too.
- b. F-(clo)([A]) = $\llbracket \exists x. x \text{ loves Peter} \rrbracket$
- F-(clo)([E]) = $\llbracket \exists y. y \text{ loves Peter} \rrbracket$
- Therefore, [A] entails F-clo([E]) and [E] entails F-clo([A]).

When it comes to the mismatches in (13) and (14), the crucial point is that descriptive properties of indexicals or R-expressions do not alter the mutual entailment relation under some variable assignment. If this is on the right track, then other mismatches should be allowed beyond indexicals and R-expressions. Consider, for instance, pairs of words opposed to each other only by the bias they express. As I have already noted, in Argentinian Spanish, for instance, the “neutral” verb ‘*comer*’ “to eat” is semantically undistinguishable from the biased verb ‘*morfar*’. This can be demonstrated by well-known substitution tests: any occurrence of the verb ‘*comer*’ can be replaced (*modulo* metalinguistic and sociolinguistic tones) by an occurrence

of the verb ‘*morfar*’ and vice versa. The predictions for mutual entailment as formulated in (15) are more or less clear. In principle, register mismatches between A and E should be allowed, in a way such that modeling the following E(lipsis)-sites as indicated should be possible in fragments like the following ones:

- (19) A: Qué comiste?
 what ate.2SG.NEUTRAL
 “What did you eat?”
 B: Una pizza <morfé>.
 a pizza ate.1SG.INFORMAL
 “a pizza.”
- (20) A: Qué morfaste?
 what ate.2SG.INFORMAL
 B: Una pizza <comí>.
 a pizza ate.1SG.NEUTRAL

Of course, without any discursive clue it would be just impossible to know whether such E-sites are possible or not. Consider in this respect the following discourse:⁶

- (21) A: Qué morfaste?
 what ate.2SG.INFORMAL
 B: Una pizza <?>, pero no tolero cuando
 a pizza but not tolerate.1SG when

hablás tan informalmente. Yo nunca lo hago.
 speak.2SG so informally I never it do

“A pizza. But I don’t tolerate when you speak informally. I never do it.”

At first sight, the metalinguistic comment introduced by B allows us to reject an E-site modeled as containing the informal counterpart of the verb “to eat”. Note that a non-elliptical version of (22) is infelicitous here:⁷

- (22) A: Qué morfaste?
 what ate.2SG.INFORMAL
- B: #Una pizza morfé, pero no tolero cuando
 a pizza ate.2SG.INFOR but not tolerate.1SG when
 hablás tan informalmente. Yo nunca lo hago.
 speak.2SG so informally I never it do

“#I ate(informal) a pizza. But I don’t tolerate when you speak informally. I never do it.”

The mutual entailment approach apparently provides the right answer to the problem, as the neutral form ‘*comer*’ could take the informal form ‘*morfar*’ as antecedent and outputs a legitimate E-site.

- (23) A: Qué [_{TP} *pro* morfaste *t*]?
 what ate.2SG.INFORMAL
- B: Una pizza <[_{TP} *pro* comí *t*]>, pero no tolero

a	pizza	ate.1SG.NEUTRAL	but	not	tolerate.1SG
cuando	hablás	tan	informalmente.		
when	speak.2SG	so	informally		

I call this “Bias Vehicle Change”, cases where the change between A and E is produced in the particular bias of some lexical expression. For (21), and assuming that short answers are derived as cases of TP ellipsis (Merchant 2004), mutual entailment between A and E should be permitted under Focus Closure:⁸

- (24) a. $F\text{-clo}(\llbracket A \rrbracket) = \llbracket A \rrbracket = \exists x[g(1) \text{ morfar } x]$ entails $\llbracket TP_E \rrbracket = \exists y[g(1) \text{ comer } y]$
b. $F\text{-clo}(\llbracket E \rrbracket) = \llbracket E \rrbracket = \exists y[g(1) \text{ comer } y]$ entails $\llbracket TP_A \rrbracket = \exists x[g(1) \text{ morfar } x]$
c. $\llbracket A \rrbracket$ entails $F\text{-clo}(\llbracket E \rrbracket)$ and $\llbracket E \rrbracket$ entails $F\text{-clo}(\llbracket A \rrbracket)$.

On the basis of the grammaticality of (21), we can then state the following generalization:

Generalization 1 (G1):

- (25) Bias Vehicle Change is licensed under TP-ellipsis.

As far as I can tell, examples similar to (21) can be constructed for myriads of pairs of words contrasting only in register. Here is a non-exhaustive list of pairs of verbs from Argentinian Spanish:

Neutral

Informal

(26)	tomar	chupar	“to drink”
	sudar/transpirar	chivar	“to sweat”
	eyacular	acabar	“to ejaculate/to come”
	pagar	garpar	“to pay”
	trabajar	laburar	“to work”
	escapar	rajar	“to escape”
	defecar	cagar	“to defecate/to shit”
	delatar	buchonear	“to betray”
	molestar	joder	“to bother”
	pasarse	sarparse	“to cross the limits”

However, G1 does not hold for every type of ellipsis. Consider in this respect Spanish NP-ellipsis (see Saab 2019 for detailed discussion on NP-ellipsis). In this language, words like ‘*culo*’ “ass” and ‘*cola*’ “tail” when applied to humans refer to the same body part, the difference being, once again, in the bias dimension of each word. Thus, ‘*culo*’ is coarse language and ‘*cola*’ is the polite form at least in some dialects (Argentinian Spanish, for instance). Interestingly, both nouns differ in gender: ‘*culo*’ is masculine, but ‘*cola*’ is feminine. This allows us to test their behavior in NP-ellipsis contexts. As shown below, bias mismatches are fully ungrammatical in any direction:

(27)	a.	El	<i>culo</i>	de	Juan	es	más	grande
		the.MASC.SG	ass.MASC.SG	of	J.	is	more	big
		que	el	< <i>culo</i> >	de	María.		

- that the.MASC.SG ass.MASC.SG of M.
- b. La *cola* de Juan es más grande
the.FEM.SG tail.FEM.SG of J. is more big
que la <*cola*> de María.
that the.FEM.SG tail.FEM.SG of M.
- c. *El *culo* de Juan es más grande
the.MASC.SG ass.MASC.SG of J. is more big
que la <*cola*> de María.
that the.FEM.SG <tail.FEM.SG > of M.
- d. *La *cola* de Juan es más grande
the.FEM.SG tail.FEM.SG of J. is more big
que el <*culo*> de María.
that the.MASC.SG <ass.MASC.SG > of M.

Note that, as far as mutual entailment under focus closure is concerned, the fact that ‘*culo*’ and ‘*cola*’ differ in gender is irrelevant, since gender is semantically arbitrary here. We arrive then at our second generalization:

*Generalization 2 (G2):*⁹

- (28) Bias Vehicle Change is not licensed under NP-ellipsis.

Here is the problem: either we have a dissociated identity condition or one of the two generalizations is spurious. The impossibility of Bias Vehicle Change in Spanish NP-ellipsis

casts doubts on the alternative of extending mutual entailment (or relatives) to this particular elliptical construction. We can of course dissociate the identity condition in one semantic condition for TP-ellipsis and one lexical-syntactic identity condition for NP-ellipsis, after all we already know that different types of ellipsis are subjected to different conditions beyond identity (e.g., parallelism conditions regulating the correlate/remnant distribution, discourse conditions regulating the legitimacy of some types of ellipsis but not others, and so on). This would be compatible with generalizations 1 and 2, but such a solution would be unappealing under uniformity considerations. Alternatively, we can make an attempt to solve semantic identity for ellipsis in general by adding some non-truth-conditional conditions to the theory. In what follows, I argue that such an alternative cannot handle the entire set of facts and, then, I show that a uniform identity condition, with lexical identity as a necessary condition, plus my theory of expressivity, accounts for why expressivity doesn't survive ellipsis sites in the TP-ellipsis cases. The NP-ellipsis cases are, in turn, ruled out as a simple failure of lexical identity.

3.2. A Pragmatic Alternative: No Telepathy

Let's assume that the choice of remaining silent (i.e., making use of ellipsis) nullifies the bias of whatever expression the silence is replacing. This would have obvious consequences for any theory of information recoverability in ellipsis involving the communicative force of non-elliptical expressions regarding bias, metalinguistic uses and register. By remaining silent, such information is unrecoverable, no matter what your favorite theory of identity is (semantic or syntactic). Put differently, we are not telepaths. Let's assume that the phenomenon under discussion falls under some No Telepathy Condition:

No Telepathy (NT):

- (29) Information regarding lexical choices is unrecoverable under ellipsis without additional discourse clues.

By “lexical choices” I understand some competition among lexical items in the same semantic space (de Saussure’s paradigmatic relations), where the selection of such and such lexical item would depend, among other things, on the speaker’s attitude towards the content of the speech act. Such choices, arguably, also convey information regarding socio-economic, the speaker’s gender and age. No Telepathy then prevents us from introducing new lexical material in the E-site that would require some sort of telepathic activity. Crucially, NT correctly rules out standard cases of Vehicle Change like (13), repeated below:

- (30) They arrested [Alex]₃, but he₃ doesn’t know why <they arrested *Alex₃/him₃>.

One reasonable way to approach (30) under NT is just claiming that the use of the pronoun is unavoidable, i.e., there is no lexical choice to make, given that one of the possible choices provided by the paradigmatic space introduces a flagrant violation of Principle-C of Binding Theory (although see Johnson 2012). As for (23), the situation is different: we can guess a lexical choice made in the E-site on the basis of the metalinguistic comment made by B. So, NT seems to be sufficiently flexible as to allow some bias change between A and E and sufficiently constrained as to rule out mismatches in lexical choices.

Thus far, it seems that the NT is a good alternative theory that would account for Bias Vehicle Change. The NT Condition, however, would have to explain the impossibility of coarse mismatches in NP-ellipsis as illustrated in (27). The question here is why gender information does not suffice to recover the missing information that the speaker has chosen the coarse ‘*culo*’ or the polite ‘*cola*’ in the relevant cases. The answer would be that gender information cannot resolve the issue of how to know whether in (27c) we are talking, for instance, about María’s nose (‘*nariz*’ “nose” is feminine in Spanish). Yet, this problem should be avoided by contextual and discursive factors. So, suppose we are walking behind Juan and María, and pointing out to Juan’s relevant body part, I say:

- (31) *El culo de Juan es grande pero...
the.MASC.SG ass.MASC.SG of J. is big but
[now pointing out to María’s relevant body part]
la < cola > de María es más grande.
the.FEM.SG tail.FEM.SG of M. is more big

Even when we are no telepaths, we should be able to recover the relevant information in this case, but we cannot. Compare with cases of pragmatic recoverability of empty nouns in general:

[pointing to some toy; the word ‘toy’, ‘*juguete*’, is masculine in Spanish]

- (32) a. Yo quiero ese con rueditas.
I want this.MASC with wheels
“I want that one with wheels.”

[pointing to a bike. The word “bike”, ‘*bicicleta*’, is feminine in Spanish]

- b. Cuando era chico, tenía una como
when was.1P.SG boy had. 1P.SG one.FEM like
esa.
that.FEM

“When I was a child, I had one like that.”

The contrast between (31) and (32) constitutes a strange state of affairs under the NT Condition. It seems that some degree of telepathy should be permitted for (32) but not for (31). In other words, the addressee in (32) is able to guess that the speaker is talking about toys or bikes, but the same guessing capacity is impossible in (31). This is connected to the fact that recoverability conditions are different for surface and deep anaphora in Hankamer and Sag’s (1976) terms. As Merchant (2010) has shown, whenever ellipsis (i.e., surface anaphora) and deep anaphora compete, ellipsis is always preferred. This is exactly what seems to be happening in (31), where there is a linguistic antecedent for the E-site that blocks pragmatic recoverability. If this is on the right track, there is some basis to conclude that NT is a suspicious condition at least as a recoverability condition for surface anaphora.

As it stands, the entire paradigm discussed in this section remains as a puzzle for this attempt to formulate a uniform semantic identity condition supplied with a pragmatic constraint. Yet, in the following section I show that Generalization 1 is a spurious observation. I argue that lexical/syntactic identity is the right condition for several types of surface anaphora, in particular, for the type of ellipsis I am concerned with.

3.3. A Syntactic Solution

Generalization 1 is spurious. Ellipsis does not allow for the type of mismatches that mutual entailment predicts. A uniform syntactic identity condition applying in narrow syntax plus the theory of expressivity I am defending are enough to make the right predictions. Let's assume that identity in ellipsis applies in narrow syntax and makes reference to the lexical content of both Roots and abstract morphemes, although other structural or semantic recoverability conditions must also apply:




Figure 2.4.: *Lexical-syntactic identity*

If lexical identity (strictly speaking, List 1 identity) is a necessary condition for ellipsis to apply, then the impossibility of Bias Vehicle Change in NP-ellipsis follows without further ado as a general violation of lexical identity. In order to see this, let's reconsider the ungrammatical example in (27c) and assume that $\sqrt{38}$ is the index for '*cola*' and $\sqrt{83}$ for '*culo*':

- (34) *El $[\text{NP } \sqrt{83} + \text{masculine}]$ de Juan es más grande
the.MASC.SG of J. is more big

que	la	[_{NP} √38 + feminine]	de	María.
that	the.FEM.SG		of	M.

As shown in (34), the strong ungrammaticality of (27c) follows now as an extreme violation of the identity condition, where both the abstract morpheme for gender and the Root in the E-site do not match the features of the corresponding phrases in the antecedent. We derive thus Generalization 2.

This analysis opens the question of why TP-ellipsis in fragments seems to behave different to NP-ellipsis. Recall the analysis sketched in (23B) and its associated (putative true) empirical observation:

- (35) A: Qué [_{TP} *pro* morfaste *t*]?
 what ate.2SG.INFORMAL
- B: Una pizza <[_{TP} *pro* comí *t*]>, pero no tolero
 a pizza ate.1SG.NEUTRAL but not tolerate.1SG
 cuando hablás tan informalmente.
 when speak.2SG so informally

Generalization 1:

- (36) Bias Vehicle Change is licensed under TP-ellipsis.

It is clear now that taking for granted that the same identity condition is operative in TP and NP-ellipsis, such an analysis should be regarded as impossible, given that ‘*comer*’ and ‘*morfar*’ are

the PF realizations of distinct Roots. So, (36) should be regarded as false. Under uniformity considerations, we expect (37) to hold:

Bias Vehicle Change Generalization (Final):

(37) Bias Vehicle change is neither licensed under TP-ellipsis nor NP-ellipsis.

Fortunately, the theory of expressivity I am defending here gives rise to an alternative analysis for (35) according to which the same Root for ‘*morfar*’, say $\sqrt{102}$, is generated both in the antecedent and the elided TP. Given that all other syntactic features are identical *in the syntax* (agreement information being supplied at PF), we conclude that both TPs are strictly identical:

(38) A: Qué [_{TP} *pro* $\sqrt{102}$ *t*]?

B: Una pizza <[_{TP} *pro* $\sqrt{102}$ *t*]>...

Thus, bias mismatches in TP-ellipsis are just illusory, i.e., just the surface reflex of the basic fact that bias is encoded in List 2 at PF. Such information, however, cannot be accessed under ellipsis, whose PF effect is precisely blocking vocabulary insertion, the operation that introduces information from List 2.

Additional evidence comes from related facts. Other cases from the list in (26) point to the conclusion that, at least in some instances, the bias/conventional implicature of a given word is not present at LF. The relevant cases are the so-called ‘*vesre*’ “reverse” words illustrated in the list in (26) with words like ‘*garpar*’ “to pay_{informal}” and ‘*sarparse*’ “to cross the limits_{informal}”. *Vesre* speech is an informal way to speak that consists of reversing the syllable structure of a

given word. Thus, ‘*pagar*’ is converted into a *vesre* word by inverting the two syllables it contains (e.g., *pa – gar* → *gar – par*). *Vesre* speech uses to be productive in informal registers in a way such that even non-lexicalized *vesre* words are subject to the rule (e.g., *mesa – same* ‘table’). From a DM perspective, *vesre* speech can only be obtained after syntax (see Bohrn 2019). A rough derivation for the verb ‘*garpar*’ could be as follows:




Figure 2.5: Derivation of a ‘*vesre*’ word at PF

If this analysis is on track a parallel example to (35) containing the verb ‘*garpar*’ in the antecedent should be analyzed as in (40):

(40) A: Qué [_{TP} *pro* garpaste *t*]? (*garpar* = √75 in the syntax)

what payed.2P.SG.INFORMAL

“What did you pay?”

B: Una pizza <[_{TP} *pro* √75 *t*]>, pero no tolero cuando

a pizza but not tolerate. 1P.SG when

hablás tan informalmente. Yo nunca lo hago.

speak.2P.SG so informally I never it do

Like in the case of ‘*morfar*’, we also have here the same syntactic Root both in the antecedent and in the E-site. It is only after a process of syllabic inversion applying at PF that we obtain informal register for the inverted *phonological* word. Such a process is obviously blocked by ellipsis.

4. Slurs

Paradigmatic slurs are expressions *prima facie* associated with the expression of a contemptuous attitude concerning a group of people identified in terms of their origin or descent (“spic”), race (“nigger”), sexual orientation (“faggot”), ethnicity or religion (“kike”), gender (“whore”), etc. They meet the two criteria I initially proposed to determine the type of expressivity I am concerned with:

- (A) In the general case, expressives form doublets: ‘*sudamericano*’/‘*sudaca*’ “South-American”/”South-American (pejorative)”, ‘*comer*’/‘*morfar*’, “to eat”.
- (B) These doublets contribute to a certain meaning dimension: style, tone or expressivity.

That the first criterion is satisfied is shared by most extant accounts of slurs. This comes in the form of the so-called *Identity Thesis*, the idea that the representational dimension of a slur is equivalent to the representational dimension of its neutral counterpart.¹⁰ Thus, according to this conception, the two sentences in (41) are extensionally equivalent:

- (41) a. Juan es sudaca.
 J. is South-American_{PEJORATIVE}
 “Juan is South-American (pejorative).”
- b. Juan es sudamericano.
 J. is South-American
 “Juan is South-American.”

In several accounts, the additional meaning slurs contribute is conceived of as independent of the at-issue content. This content is modeled as a bias on contexts of use (Predelli 2013), a conventional implicature (McCready 2010), or a stereotype (Orlando and Saab 2020a,b), among other options. Regardless of particular implementations, the expressive meaning of slurs projects out of at-issues operators (see McCready 2010):

- (42) a. Juan no es sudaca.
 Juan not is South-American_{PEJORATIVE}
 “Juan is not South-American (pejorative).”
- b. Juan cree que Ana es sudaca.
 Juan believes that Ana is South-American_{PEJORATIVE}
 “Juan believes that Ana is South-American (pejorative).”
- c. Juan puede ser sudaca.
 Juan might be South-American_{PEJORATIVE}
 “Juan might be South-American (pejorative).”
- d. ¿Es sudaca Juan?

is South-American_{PEJORATIVE} Juan

“Is Juan South-American (pejorative)?”

Part of the debate is whether these projection effects are presuppositional or not.¹¹ The problem is subtle as it is not easy to construct the relevant tests. Yet, I think that the behavior of slurs under ellipsis favors a non-presuppositional account of slurs as much as it favors the particular approach to stylistic semantics I am pursuing here. In order to see how the argument works, let’s consider an example of NP-ellipsis in Spanish for which we have enough evidence that the ellipsis site contains a full-fledged NP structure which is deleted at PF:

(43) A: ¿A cuántos sudacas viste en la fiesta?
to how-many South-Americans_{PEJ} saw.2SG in the party
“How many South-American_{PEJ} did you see at the party?”

B: Vi a tres <sudacas>, pero podrías evitar
saw.1SG to three South-American_{PEJ} but could.2SG avoid
ese modo de hablar de los sudamericanos.
that way of speaking of the South-American_{PEJ}
Yo nunca hablo así de ellos.
I never speak so of them

“Three, but you could avoid talking that way about South-Americans. I never talk that way.”

By eliding the noun phrase, the speaker B provides the relevant information asked in A's question and, at the same time, objects A's xenophobic way of speaking. Not eliding the noun, obviously, leads to infelicity, unless air quotations or a similar strategy is employed:

- (44) A: ¿A cuántos sudacas viste en la fiesta?
to how-many South-Americans_{PEJ} saw.2SG in the party
“How many South-Americans_{PEJ} did you see at the party?”
- B: #Vi a tres sudacas, pero podrías evitar
saw.1SG to three South-Americans_{PEJ} but could.2SG avoid
ese modo de hablar de los sudacas.
that way of speaking of the South-Americans_{PEJ}
Yo nunca hablo así de ellos.
I never speak so of them

‘I saw three South-Americans_{PEJORATIVE}, but you could avoid that way of speaking about South-Americans. I never talk that way.’

This is an interesting state of affairs. On the one hand, it shows that ellipsis is an apt strategy to nullify the bias encoded in some lexical items. This follows from the present account that requires lexical insertion in order to make the expressivity salient in the discourse. Alternative accounts that encode the bias in the syntax don't offer a good explanation for why B's answer in (43B) is a non-biased answer. Recall that McCready (2010) proposes lexical entries like the

following ones, where a conventional implicature is directly encoded in full lexical items that are manipulated by the syntax and LF:

$$(45) \llbracket \text{sudaca} \rrbracket = \lambda x. \text{South-American}(x) \blacklozenge \text{Bad}(\cap \text{South-American}): \langle e, t \rangle^a \times t^s$$

Without further ado, ellipsis should not block the conventional implicature that B has a negative attitude towards South-Americans. A way of avoiding the problem is assuming with Potts *et al* (2009) that identity in ellipsis only refers to descriptive properties of antecedents (see section 1). But note that this claim makes sense if ellipsis doesn't require lexical identity and it is calculated, say, over the focus closures of antecedents and elliptical sites. But if lexical identity is indeed required such an assumption remains as a controversial stipulation. In this respect, the proposal of deriving expressivity at PF gives better results.

On the other hand, as I have already advanced, the facts in (43) also seem to give additional support for the hypothesis that the expressive dimension in slurs is non-presuppositional. In effect, standard presuppositions are not cancelled by ellipsis, as witnessed in cases like the following:¹²

- (46) A. Juan dejó de fumar.
 J. stopped of smoking
- B. Pedro también <dejó de fumar>,
 P. too stopped of smoking
- #aunque nunca fumó.
 although never smoked

In sum, slurs behave under ellipsis as predicted by the hypothesis of “late” expressivity. The moral to be extracted from this particular behavior of informal terms and slurs is that expressivity must be expressed (uttered). Putting the type of non-truth conditional meanings that these terms have in the syntactic-semantic derivation leaves this basic observation unexplained.

A final question is how the stylistic meaning of a given slur can be modeled in the present framework. An important similarity with informal terms is that, in the general case, slurs are also informal. This means that we want to model part of their stylistic meaning as we did for informal terms in section 2.2 (see Díaz Legaspe *et al* 2020 on the relation between register and slurs):¹³

(47) *Bias for sudaca*: $c \in CU(\text{sudaca})$ only if, in c , c_a is a participant in register *informal*

Of course, this is not enough. A concrete use of a slur involves much more than just informal register. Adapting ideas from Orlando and Saab (2020a,b), I would like to suggest that slurs characterize contexts by making salient a certain stereotype of the target group. Concretely, I conjecture that stereotypes can be conceived of as semantic objects ranging over sets of propositions which, taken together, constitute a certain misinformed theory of a given human group. In model-theoretic terms, stereotypes are similar to Kratzerian Modal Bases, *i.e.*, functions of the $\langle\langle s, t \rangle, t \rangle$ type that take propositions as inputs and return sets of propositions of a certain form:

(48) $\llbracket \text{Stereotype} \rrbracket^{g,w} = \lambda p. \exists P_{\langle e, t \rangle} [P \in C \ \& \ p = [\lambda w. P(\text{Group})(w)]]$

Thus, the stereotype for ‘*sudaca*’ would contain, among other propositions, the propositions that South-Americans are prone to laziness, that South-Americans are undocumented migrants and so on. The content of each predicate P is contextually determined by its membership to C , i.e., by the large socio-cultural context in which a certain stereotype is in force. With this in mind, we can proceed to define the bias dimension of ‘*sudaca*’ as follows:

- (49) *Bias for ‘sudaca’*: $c \in CU(sudaca)$ only if, in c , c_a is a participant in register *informal* and a stereotype about South-Americans is in force in c

On this view, slurs serve to the end of conventionally characterizing a certain type of context of use, in which a certain hate ideology is in force. Alternatively, we could describe the dual meaning of a given slur in more neutral terms, using Potts’s bullet only to separate the LF (i.e., truth-conditional) and PF (i.e., stylistic) meaning of a certain dual expression:

- (50) LF: $\lambda x. \text{South-American}(x)$ • PF: $\lambda p. \exists P[P \in C \ \& \ p = [\lambda w. P(\cap \text{South-American})(w)]]$

At any rate, in both cases, and this is what I would like to stress, the particular stereotype that a given slur makes salient is a property of the vocabulary item in the favored sense of this chapter, i.e., a property of a form.

To sum up, I have discussed two important empirical properties of slurs: (i) they are in lexical competition with their corresponding neutral counterparts, and (ii) their expressive content can be eliminated by ellipsis. These two properties make slurs analogous to other expressive parts of the lexicon (lexical pairs opposed only in register), for which I have already

argued that the best account is that of constructing the expressive content at PF and not LF. If this analogy turns out to be substantial, we can conclude that the same PF explanation is the best account of slurs, particularly, of how their expressive content is built.

5. Conclusion

I have sketched a general project for what I have called “stylistic semantics”. The main goal of this research agenda is to offer a PF account of the expressive content of a subset of biased words, in particular, register terms and slurs. I have proposed that this can be done to the extent the PF interface is capable of introducing stylistic meanings in the favored sense. I have conjectured that PSM, repeated below, is responsible for generating PF meanings on the basis of semantic vacuity at LF.

(51) *Principle of stylistic meaning (PSM)*: Given a pair of abstract nodes X and Y, taken from List 1, if X and Y are not semantically distinguishable at LF (List 3), they are “semantically” distinguished at PF (List 2).

The proposal captures the two basic properties of a certain type of expressive terms, repeated below, and receives robust evidence from ellipsis and *vesre* speech:

- (A) In the general case, expressives form doublets: ‘*sudamericano/sudaca*’ “South-American”/“South-American(pejorative)”, ‘*comer*’/‘*morfar*’, “to eat”.
- (B) One of the members of these pairs contributes to a certain meaning dimension: style, tone or expressivity.

The hope is that a proper and more explicit formulation of the conjectures I have discussed here will be extended to other empirical domains in the realm of stylistic semantics.

Acknowledgements

I have presented different versions of this material at the Facultad de Letras y Ciencias Humanas, Pontificia Universidad Católica del Perú (Lima, 2019), *Going Romance XXXII*, Utrecht University (Utrecht, 2018), the *IV Congreso Internacional de la Sociedad de Filosofía del Uruguay* (Montevideo, 2018), the *IV Issues in Contemporary Semantics and Ontology: Modals and Evaluatives: Linguistic and Philosophical Perspectives*, SADAF (Buenos Aires, 2018), and the *Comparative Syntax Meeting*, Leiden University (Leiden, 2017). I am thankful to these audiences for many interesting and fruitful discussions, in particular, to Ramiro Caso, Nicolás Lo Guercio, Anikó Lipták, Eleonora Orlando, Ana Clara Polakoff, Stefano Predelli, Marco Ruffino and Ludovic Soutif. I am also in debt with Matthew Barros, with whom I have discussed the hypotheses presented here extensively. I am also grateful to the linguistics and philosophy of language team at the Sociedad Argentina de Análisis Filosófico (SADAF): Eduarda Calado, Ramiro Caso, Sofía Checchi, Eduardo García Ramírez, Federico Jaimes, Justina Díaz-Legaspe, Nicolás Lo Guercio, Alfonso Losada, Laura Skerk, Matías Verdecchia and, in particular, to Nicolás Lo Guercio and Eduardo García-Ramírez for many comments and suggestions and to Eleonora Orlando for many years of mutual collaboration and friendship.

Bibliographical References

- Acquaviva, Paolo. 2008. "Roots and Lexicality in Distributed Morphology". Paper given at the *5th York-Essex Morphology Meeting*. Available at ling.auf.net/lingBuzz/000654.
- Anderson, Luvell, and Ernie Lepore. 2013a. "Slurring Words." *Noûs* 47(1): 25-48.
- Anderson, Luvell, and Ernie Lepore. 2013b. "What did You Call Me? Slurs as Prohibited Words." *Analytic Philosophy* 54(3): 350-363.
- Arregi, Karlos, and Andrew Nevins. 2012. *Morphotactics: Basque Auxiliaries and the Structure of Spellout. Studies in Natural Language and Linguistic Theory*. Berlin: Springer.
- Ashwell, Lauren. 2016. "Gendered Slurs." *Social Theory and Practice* 42(2): 228–239.
- Barros, Matthew, and Andrés Saab. 2016. "Implicatures from Silence. A Remark on Chung's Puzzle." Ms., Yale University, CONICET-UBA.
- Barros, Matthew, and Hadas Kotek. 2019. "Ellipsis Licensing and Redundancy Reduction: A Focus-Based Approach." *Glossa: A Journal of General Linguistics* 4(1): 1–36.
<https://doi.org/10.5334/gjgl.811>.
- Bohrn, Andrea. 2019. "Inversión silábica y realización de género y número: el caso del *vesre* rioplatense." *Rasal* 2015: 29-48.
- Bolinger, Renee Jörgensen. 2017. "The pragmatics of slurs." *Noûs* 51(3): 439-462.
- Caso, Ramiro, and Nicolás Lo Guercio. 2016. "What Bigots Do Say: A Reply to DiFranco." *Thought: A Journal of Philosophy* 5(4): 265-274.
- Chomsky, Noam. 1965. *Aspects of the Theory of Syntax*. Cambridge, Mass.: MIT Press.
- Chomsky, Noam. 1995. *The Minimalist Program*. Cambridge, Mass.: MIT Press.

- Chomsky, Noam. 2000. "Minimalist Inquiries: The Framework." In *Step by Step. Essays on Minimalist Syntax in Honor of Howard Lasnik*, edited by Roger Martin, David Michaels, and J. Uriagereka, 89–156. Cambridge, Mass.: MIT Press.
- Chomsky, Noam. 2001. "Derivation by Phase." In *Ken Hale: A Life in Language*, edited by Michael Kenstowicz, 1–52. Cambridge, Mass.: MIT Press.
- Chung, Sandra. 2006. "Sluicing and the Lexicon: The Point of no Return." In *Proceedings of the Annual Meeting of the Berkeley Linguistics Society (31)*, edited by Rebecca T. Cover, and Yuni Kim, 73–91. Berkeley: University of California, Berkeley Linguistics Society.
- Díaz Legaspe, Justina. 2018. "Normalizing Slurs and Out-Group Slurs: The Case of Referential Restriction." *Analytic Philosophy* 59(2): 1–22.
- Díaz Legaspe, Justina, Liu, Chang and Stainton, Robert. 2020. "Slurs and Register: A Case Study in Meaning Pluralism". *Mind and Language* 35(2): 156-182.
<https://doi.org/10.1111/mila.12236>.
- Embick, David. 2000. "Features, Syntax, and Categories in the Latin Perfect." *Linguistic Inquiry* 31(2): 185-230.
- Embick, David, and Rolf Noyer. 2001. "Movement Operations after Syntax." *Linguistic Inquiry* 32: 555-595.
- Embick, David. 2007. "Linearization and Local Dislocation: Derivational Mechanics and Interactions." *Linguistic Analysis* 33(3-4): 2-35.
- Embick, David. 2015. *The Morpheme. A Theoretical Introduction*. Berlin: de Gruyter Mouton.
- Fiengo, Robert, and Robert May. 1994. *Indices and Identity*. Cambridge, Mass.: MIT press.
- Frege, Gottlob. 1918/1956. "The Thought: A Logical Inquiry." *Mind* 65: 289-311.

- Grice, H. Paul. 1975. "Logic and Conversation." In *Syntax and Semantics: Speech Acts*, edited by Peter Cole and Jerry Morgan, 43–58. New York: Academic Press.
- Gutzmann, Daniel. 2015. *Use-Conditional Meaning. Studies in Multidimensional Semantics*. Oxford: Oxford University Press.
- Halle, Morris, and Alec Marantz. 1993. "Distributed Morphology and the Pieces of Inflection." In *The view from Building 20*, edited by Hale, Ken and Samuel Keyser, 111-176. Cambridge, Mass.: MIT Press.
- Halle, Morris, and Alec Marantz. 1994. "Some Key Features of Distributed Morphology." In *Papers on Phonology and Morphology (MITWPL 21)*, edited by Andrew Carnie, Heidi Harley, and Tony Bures, 275-288. Cambridge, Mass.: MIT Press.
- Hankamer, Jorge, and Ivan A. Sag. 1976. "Deep and Surface Anaphora." *Linguistic Inquiry* 7:391–428.
- Harley, Heidi. 2014. "On the Identity of Roots." *Theoretical Linguistics* 40: 225–276.
- Hom, Christopher. 2008. "The Semantics of Racial Epithets". *The Journal of Philosophy* 105: 416-440.
- Hom, Christopher. 2010. "Pejoratives." *Philosophy Compass* 5(2): 164-185.
- Jeshion, Robin. 2013a. "Slurs and Stereotypes." *Analytic Philosophy* 54: 314–325.
- Jeshion, Robin. 2013b. "Expressivism and the Offensiveness of Slurs." *Philosophical Perspectives* 27: 231-259.
- Johnson, Kyle. 2012. "Pronouns vs. Definite Descriptions." In *Generative Linguistics and Acquisition: Studies in Honor of Nina M. Hyams*, edited by Misha Becker, John Grinstead, and Jason Rothman, 157–184. Amsterdam: John Benjamins Publishing Company.

- Kaplan, David. 1989. "Demonstratives." In *Themes from Kaplan*, edited by Joseph Almog, John Perry, and Howard Wettstein, 481–563. Oxford: Oxford University Press.
- Kaplan, David. 1999. "What Is Meaning? Explorations in the Theory of Meaning as Use." Brief Version, Draft 1. Ms.
- Lepore, Ernie, and Matthew Stone. 2018. "Pejorative Tone." In *Bad Words: Philosophical Perspectives on Slurs*, edited by David Sosa, 132-154. Oxford: Oxford University Press.
- Lipták, Anikó. 2020. "Fragments with Correction." *Linguistic Inquiry* 51(1): 154-167.
- Lo Guercio, Nicolás. 2020. "Slurs and Antipresuppositions." Ms., SADAF.
- Losada, Alfonso. 2020. "*Sudaca*. Slurs and Typifying." This volume.
- Marantz, Alec. 2013. "Locality Domains for Contextual Allomorphy Across the Interfaces." In *Distributed Morphology Today. Morphemes for Morris Halle*, edited by Ora Matushansky, and Alec Marantz, 95-116. Cambridge, Mass.: MIT Press.
- McCready, Elin. 2010. "Varieties of Conventional Implicatures." *Semantics and Pragmatics* 3: 1-57.
- Merchant, Jason. 2001. *The Syntax of Silence: Sluicing, Islands, and the Theory of Ellipsis*. Oxford: Oxford University Press.
- Merchant, Jason. 2004. "Fragments and Ellipsis." *Linguistics and Philosophy* 27. 661-738.
- Merchant, Jason. 2010. "Three kinds of ellipsis." In *Context-Dependence, Perspective, and Relativity*, edited by François Recanati, Isidora Stojanovic, and Neftali Villanueva, 141-192. Berlin: Walter de Gruyter.
- Nunberg, Geoff. 2018. "The Social Life of Slurs." In *New Work on Speech Acts*, edited by Daniel Fogal, Daniel Harris, and Matt Moss, 237-295. Oxford: Oxford University Press.

- Orlando, Eleonora and Andrés Saab. 2020a. "Slurs, Stereotypes and Insults." *Acta Analytica* (Forthcoming). (DOI: 10.1007/s12136-020-00424-2)
- Orlando, Eleonora and Andrés Saab. 2020b. "A Stereotype Semantics for Syntactically Ambiguous Slurs." *Analytic Philosophy* (forthcoming). (DOI: 10.1111/phib.12184)
- Potts, Christopher. 2005. *The Logic of Conventional Implicature*. Oxford: Oxford University Press.
- Potts, Christopher, Ash Asudeh, Seth Cable, Yurie Hara. 2009. "Expressives and Identity Conditions." *Linguistic Inquiry* 49(2): 356-366.
- Predelli, Stefano. 2013. *Meaning Without Truth*. Oxford: Oxford University Press.
- Predelli, Stefano. 2020. "Taboo: The Case of Slurs." This volume.
- Reinhart, Tanya. 2006. *Interface Strategies. Optimal and Costly Computations*. Cambridge, Mass.: MIT Press.
- Saab, Andrés. 2008. "Hacia una teoría de la identidad parcial en la elipsis." PhD diss., University of Buenos Aires.
- Saab, Andrés. 2019. "Nominal Ellipsis." In *The Oxford Handbook of Ellipsis*, edited by Jeroen van Craenenbroeck, and Tanja Temmerman, 526-561. Oxford: Oxford University Press.
- Schlenker, Philippe. 2007. "Expressive Presuppositions." *Theoretical Linguistics* 33(2). 237–245.
- Thoms, Gary. 2013. "Lexical Mismatches in Ellipsis and the Identity Condition." In *Proceedings of North Eastern Linguistic Society (vol. 42)*, edited by Keine, Stefan, and Shayne Sloggett, 559-572. Amherst, Mass.: CSLI Publications.
- Thoms, Gary. 2015. Syntactic Identity, Parallelism and Accommodated Antecedents. *Lingua* 166:172-198.

Trubetzkoy, Nikolai. 1939. *Grundzüge der Phonologie*. Prague: Travaux du Cercle Linguistique de Prague.

Whiting, Daniel. 2013. “It’s not What You Said, It’s the Way You Said It: Slurs and Conventional Implicatures.” *Analytic Philosophy* 54(3): 364-377.

¹ Even if I am only mentioning, not using, slurs and coarse language, I apologize in advance for any unintentional offense I might cause to the reader.

² As Eduardo García-Ramírez (pers. comm.) observes, the set of items in competition may be more than two. For instance, in Argentina the word ‘*homosexual*’ has more than one slur as expressive counterpart. This is consistent with all what I will say in what follows. On this point, see also Losada (2020).

³ As Eduardo García-Ramírez pointed out me, this seems to be a too strong claim in view of the fact that there are many possible scenarios in which B could endorse the negative evaluation A conveys with respect to the relevant dog in her utterance. This is correct, but I would say that in those scenarios such an endorsement (i) is not a lexical entailment, and, consequently, (ii) only arises in virtue of some pragmatic mechanism.

⁴ Note that, in principle, (4) could be analyzed as a standard case of Vehicle Change, where the E-sites are modeled as “see it”. However, Thoms (2013) provides the following example that seems to rule out such a possibility:

(i) A: You should *fucking* fire that asshole John!

B: I know you think I should, but I won’t as I like him.

Here, “fucking” is an adverbial modifier and, as such, it would not allow for Vehicle Change.

However, I am not convinced that (iB) cannot be modeled as a special case of Vehicle Change, where the E-site is modeled as “do it”. Whatever the right analysis is, this is orthogonal to the main point I am making here.

⁵ Following Kaplan (1989), I take the character of any expression E as a function from context to content and the content itself as function from circumstances of evaluation to truth values (i.e., to $\langle s, t \rangle$ objects on some accounts). Standardly, a Kaplanian context is a tuple consisting at least of the following parameters: $\langle w, t, a, h, l \rangle$, where w is a possible world, t is a time, a is the agent of the utterance, h the hearer, and l the location.

⁶ This present test is modeled after Lipták (2020), who discusses a different type of putative mismatch in ellipsis. Here is one of her examples:

(i) A: What are you devouring?

B: A pizza, but I am not devouring it.

B: # I am devouring a pizza, but I am not devouring it!

At first glance, this example does not parallel (21), since in (i) one is tempted to conclude that whatever verb one postulates within the E-site (e.g., “ate”) it should contribute a different truth-conditional profile from the one contributed by the verb in the antecedent.

⁷ Of course, (21B) improves if speaker B adds air quotations or other metalinguistic gestures/devices.

⁸ The same results obtain in other semantic proposals, like Barros and Kotek’s (2019) identity condition.

Redundancy reduction:

-
- (i) XP_E may be reduced (elided or deaccented) provided that it has a salient antecedent, XP_A , and $U[XP_E]^f = U[XP_A]^f$. (Barros and Kotek 2019, 4)

And the same with more flexible approaches (Q-equivalence approaches in Barros and Kotek's terms). I use Merchant's condition just because of the influence it had in the last twenty years.

⁹ Obviously, with nouns in which there is no gender mismatch one can have the illusion that NP-ellipsis allows bias mismatches (Muñoz Pérez, pers. comm.):

- (i) A. ¿Fuiste al laburo de Pedro hoy?
 went.2SG to.the work of P. today
 “Did you go to Juan's work (informal) today?”
- B. No, al <trabajo> de Juan, pero no tolero cuando
 No, to.the work of J. but not tolerate.1SG when
 hablás tan informalmente. Yo nunca lo hago.
 Speak.2SG so informally I never it do.1SG
 “No, to John's, but I don't tolerate when you speak informally. I never do it.”

¹⁰ This is the core thesis of Neutral Counterpart Theories, supported by Anderson and Lepore (2013a,b), Jeshion (2013a,b), Predelli (2013), Whiting (2013), and McCready (2010), among others. For objections to that thesis, see Hom (2008, 2010), Ashwell (2016), and Losada (2020), among others. See also Díaz Legaspe (2018) for insightful considerations in favor of restricting the thesis in the case of some kind of slurs, the so-called “normalizing” ones. For a general defense of the Identity Thesis, see Caso and Lo Guercio (2016).

¹¹ See Schlenker (2007) for a presuppositional account, and McCready (2010) for a non-presuppositional approach.

¹² But see Schlenker (2007) for arguments in favor that expressive presuppositions are not standard presuppositions, and Lo Guercio (2020) for a recent criticism to presuppositional accounts. One interesting point raised by Lo Guercio (pers. comm.) is that pragmatic accounts of expressivity like those defended in Bolinger (2017) or Nunberg (2018) are immune to the this type of criticism. Indeed, they seem to have a good explanation of why ellipsis can block expressive effects.

¹³ On the relation between register and slurs, see Díaz Legaspe *et al* (2020).