Assessing gain modulation as a cortical principle of natural language computation

Elliot Murphy[1,2]

1. Vivian L. Smith Department of Neurosurgery, McGovern Medical School, University of Texas Health Science Center

2. Texas Institute for Restorative Neurotechnologies, University of Texas Health Science Center

Recent work in language and computation in neural systems has begun to develop a model of linguistic computation relying on gain modulation, or the process whereby neurons combine multiple sources of information. In this brief article, I will review these proposals and provide two conclusions: First, gain modulation models of language face a number of internal conceptual and empirical difficulties. Second, our current understanding of gain modulation does not support its use as a core component of linguistic computation.

Language and gain modulation: The proposal

With ambitious interdisciplinary scope, Martin (2020) defends a multidimensional coordinate system for language based on neurophysiological models of sensory processing. Martin maintains that neural trajectories (of these coordinates) encode *sensory*, *motor* and *abstract linguistic states*. Crucial to Martin's model is the assumption that gain modulation tunes the path of these trajectories in accordance with behavior, and that this is "how latent structure is inferred" (2020). Similar to how Chomsky famously points to the paradox behind language use being appropriate to circumstance but not caused by circumstance (*Descartes' problem*), Martin points to a more recent, apparent paradox between the "marvelous expressive

23 capacity" of language (i.e. its generative nature) and the "general remit" of the human brain

24 to be "driven by statistical regularities in its environment". Martin adds that, lacking a

25 neurophysiological account of core linguistic properties such as hierarchical structure,

26 function application and scope, and compositionality, theories of the human brain "seem

27 startingly incomplete".

28 Attempting to move beyond this situation, Martin (2020) proposes that linguistic structure

29 building is "a form of perceptual inference". She adds that increasingly abstract linguistic

30 structures (from syllables to morphemes to words) are inferred via gain modulation which

31 she defines as "the way neurons combine information from two or more sources", with, for

32 instance, inhibition and signal amplification being notable ways that the brain seems to

33 combine and separate information.

34 Martin (2020) reports how gain modulation underlies "coordinate transform between

35 sensory modalities and between sensory and motor systems". She cites classic work on

36 cellular dynamics to support this (Salinas & Abbott 2001), but does not cite other important

37 work by the same co-author (most importantly, Dayan & Abbott 2001) where the limitations

38 of gain modulation are stressed, e.g. gaze-dependent gain modulation of retinotopic visual

39 receptive fields is well-reported, but the relevance of gain modulation to higher cognition is

40 not established; further, "[t]he mechanism by which sensory and modulatory inputs combine

41 in a multiplicative way in gain-modulated neurons is not known" (Dayan & Abbott 2001: 17;

42 although see discussion of Ferguson and Kardin 2020 below). Dayan and Abbott (2001: 48)

43 note the potential that linear and non-linear feedforward and recurrent networks have to act

44 as content-addressable memories *alongside* the exhibition of gain modulation, but it is not

45 clear that gain modulation itself carries any explanatory scope certainly to the level of

46 sentential representations. Martin (2020) does provide an interesting speculation that

47 "[c]oordinate transform may be a computational requirement of any system with multiple

48 data types or formats from multiple perceptors", yet the risk of centering gain modulation in

49 these processes is that the genericity of this system-wide process leads to a situation in which

50 one cannot easily test or falsify Martin's claims. Martin (2020) notes that "how [linguistic]

51 units combine together via gain modulation over time is hypothetical and must be tested",

52 but no empirical suggestions for testing are discussed. Martin (2020) later claims that "the

53 role of gain modulation in the brain is likely to be much more broad than as a neuronal

54 instantiation of attention" (Martin 2020). Unfortunately, Martin provides no empirical

55 support for this claim.

56 Martin (2020) makes an additional claim that "gain modulation also offers a built-in system

57 for predictive coding", citing Friston (2005); more recent, updated accounts (e.g. Friston 2018)

58 rely on an interplay of a range of other processes, not gain modulation. Moreover, Friston's

59 ventures into language (most prominently in Friston et al. 2017, exploring epistemic foraging

60 and Bayesian belief updating) do not seem compatible with Martin's (2020) extensive use of

61 gain modulation (see Shipp 2016 for discussion of some possible, motivated neural

62 implementations of predictive coding, and also Heilbron & Chait 2018 for mechanisms more

63 in line with the spirit of Martin's proposal).

64 One interesting advantage of Martin's model is the multiple-realisability of gain modulation,

65 which is a deceptively more abstract process than its textbook definition would imply.

66 Martin's (2020) supplied pseudocode for her model (Table 2) is instructively clear, yet it also

67 begs the question as to the neuronal substrates of the proferred representations; "lexico-

68 syntactic" relations are assumed *ab initio*, and one cannot simply stipulate that they align with

69   "gain-field trajectories" without cashing out these higher-level abstractions within some form

70   of compatible lower-level mechanism. In addition, Martin's (2020) account of minimal pair

71   interpretation relies on inhibitory implementations of gain modulation, whereby activation

72   of one representation inhibits the other, an interesting feature which might also extend to

73   broader linguistic processes like polysemy. Still, as mentioned, the broader architectural

74   setting of these operations is unclear. Relatedly, Martin only addresses in a footnote a core

75   question at hand concerning the underlying non-linear vs linear (and Euclidean vs non-

76   Euclidean) structure of the brain dynamics, but this is an essential issue to sustain a model

77   relying on gain modulation.[1]

78

79   **Language and gain modulation: A reassessment**

80   How do these assumptions about the role of gain modulation fare when we examine the

81   apparent mechanisms underlying cortical gain modulation? Ferguson and Kardin (2020)

82   provide a comprehensive review of gain modulation, pointing to a common set of

83   mechanisms: GABAergic inhibition, synaptically driven fluctuations in membrane potential,

84   and changes in cellular conductance. Ferguson and Kardin (2020) note that diverse cortical

---

[1] Some other issues serve to make the exposition less clear. In Table 1 of Martin's (2020) paper, 'like an arrow' is presented as both a PP and an AdvP (incorrectly). The 'Role-filler binding calculus' in Table 1 is not informative: Representing 'Time flies like an arrow' as {'flies(time)'; 'like(an arrow)'} does not approach the underlying richness in thematic relations. Turning to Figure 1, the cartoon representation of words and morphemes is incorrectly matched, such that morphemes are represented as being composed of words, rather than the other way round. It is also unclear how Figure 2 provides conceptual or psycholinguistic clarity beyond the level of stating that 'Time flies like an arrow' contains five words with separate phrases.

85   functions such as information integration across cognitive, sensory and motor systems seem

86   to be performed through gain modulation. Further, regulation of neural gain can provide "an

87   integration mechanism whereby information from multiple sources can be non-linearly

88   combined via multiplicative modulation of the cell's response to inputs" (2020: 81).

89   Morphological features of dendrites also regulate the extent to which gain control is neurons

90   is possible; e.g. moderate branching in pyramidal neurons potentially promotes the greatest

91   possible range of gain modulation (2020: 84).

92   Martin (2020) also briefly defends an additive, rather than a tensor-based multiplicative,

93   formulation of natural language compositionality. On this issue of additive versus

94   multiplicative gain modulation, Ferguson and Kardin (2020) review how visual properties that

95   must be *decoded separately* are combined multiplicatively, whereas parts that must be

96   *integrated* are combined additively. If gain modulation is to play a central role in linguistic

97   computation, then one would assume that Martin's additive-only model will be able to

98   capture integrative functions of natural language compositionality, but that processes

99   involved in parsing discrete linguistic (and embedded) representations that require structural

100   separation may require multiplicative gain modulation.

101   What is more in line with Martin's specific proposal is Ferguson and Kardin's (2020: 88)

102   summary that gain modulation serves the "rapid adaptation to varying ranges of input" in

103   addition to "enhancing the salience of relevant information", with the notion of salience being

104   relevant to a range of linguistic processes. In addition, Ferguson and Kardin (2020: 88) note –

105   using terminology familiar to syntacticians – that gain modulation may be able to *increase*

106   *information transmission* and provide *computational efficiency* within a given neural network.

107 Nevertheless, it remains currently unknown how gain modulation at the single-cell level

108 contributes to population coding that can be read out by downstream targets.

109 The authors conclude that "gain is regulated by a wide range of influences, including

110 attention, learning, locomotion, arousal and neuromodulatory activity, and that these may

111 act through a common set of cellular and circuit mechanisms" (2020: 89). As such, while the

112 role of gain modulation in sensory-to-motor conversions is well-established, gain modulation

113 has not currently been implicated in any sub-component of language-relevant cognition, with

114 *attention* very much constituting the periphery of the B in FLB ("Faculty of Language Broad";

115 Hauser et al. 2002). Further, "the precise relationship between gain modulation of single

116 neurons and the encoding and transmission of information at the population level is not well

117 understood", while "the contribution of neural gain control to perceptual and cognitive

118 performance remains to be fully explored" (Ferguson & Kardin 2020: 89). As such, it seems

119 premature to use gain modulation as a core component of any model of higher-level linguistic

120 computation.

121 Turning to related concerns, Ferguson and Kardin (2020: 89) note that "the precise

122 relationship between gain modulation of single neurons and the encoding and transmission

123 of information at the population level is not well understood", and that "the reliability and

124 repeatability of gain modulation of single neurons and cortical networks is unknown". Given

125 this, it seems not well-motivated to stipulate (as Martin 2020 does) a close connection

126 between gain modulation and the full range of oscillatory frequency interactions which seem

127 to index the processing of linguistic representations. Since gain modulation does in fact

128 appear relevant to enabling networks of neurons to produce distributed representations of

129 stimulus features (as opposed to generating inferences to higher abstract structures from

130 such features) , it may be that gain modulation goes the way of entrainment: Initially viewed

131 as a crucial mechanism for all manner of linguistic processes (again, from syllables to

132 sentences) before being relegated to a role in auditory cortex for *some* aspects of lower-level

133 speech perception (Murphy 2020).

134 Keeping to these systems-level concerns, while Martin has previously sided against

135 multiplexing as a viable mechanism for grounding linguistic computation, in Martin (2020) she

136 now proposes that cortical oscillations structure speech input into linguistic representations

137 via gain-modulated multiplexing. Multiplexing models of language have been proposed

138 elsewhere (e.g. Benítez-Burraco & Murphy 2019, Murphy 2016a, 2016b, 2018, 2020 utilize

139 multiplexing as operationalised via a range of distinct cross-frequency couplings; see also

140 Gross et al. 2013) and, given the above-mentioned problems with relating our current

141 understanding of gain modulation to higher-order language processing, it is unclear how

142 Martin's (2020) *gain-modulated multiplexing* is a step forward. In addition, Martin puts

143 forward one oscillation-related hypothesis, that "[t]he theta oscillation may be a likely carrier

144 signal for linguistic sensory input, but more carrier signals and coherence between them must

145 exist for the perceptual inference of linguistic structure, which itself is not recoverable from

146 the sensory codes alone". Which sensations in particular are carried by theta, and how such

147 a multiplexing mechanism plays a role in the broader architecture Martin alludes to but does

148 not develop, are questions not explored but presumably need to be addressed in any given

149 model of *language* as a system, rather than its sub-systems. In addition, while Martin's

150 intuition of gain modulation being an elegant algorithmic-level description of

151 compositionality is surely novel and appealing (putting aside the objections I have made

152 here), and while Martin's ideas are also not incompatible with a range of oscillatory-models

153 of language in the literature, Martin does not discuss any limitations of her model.

154   Finally, one potential successful application of gain modulation models is to learning, which

155   appears to be more amenable to implementations via gain modulation (Ferguson & Kardin

156   2020). Nevertheless, testing Martin's (2020) gain modulation model will require a much

157   deeper understanding of the relationship between gain modulation and processes that we

158   can more easily record in the human brain.

159

160

161

162

163

164

165

166

167

168

169

170

171

172

173

174 **References**

175 Benítez-Burraco, A., & Murphy, E. (2019). Why brain oscillations are improving our

176 understanding of language. *Frontiers in Behavioral Neuroscience* 13: 190.

177 Dayan, P., & Abbott, L.F. (2001). *Theoretical Neuroscience: Computational and Mathematical*

178 *Modeling of Neural Systems.* Cambridge, MA: MIT Press.

179 Ferguson, K.A., & Cardin, J.A. (2020). Mechanisms underlying gain modulation in the cortex.

180 *Nature Neuroscience* 21: 80–92.

181 Friston, K. (2005). A theory of cortical responses. Philosophical Transactions of the Royal

182 Society of London, Series B: Biological Sciences 360: 815–836.

183 Friston, K.J., Rosch, R., Parr, T., Price, C., & Bowman, H. (2017). Deep temporal models and

184 active inference. *Neuroscience and Biobehavioral Reviews* 77: 388–402.

185 Gross, J., Hoogenboom, N., Thut, G., Schyns, P., Panzeri, S., Belin, P., & Garrod, S. (2013).

186 Speech rhythms and multiplexed oscillatory sensory coding in the human brain. *PLoS Biology*

187 11: e1001752.

188 Hauser, M., Chomsky, N., & Fitch, W.T. (2002). The faculty of language: what is it, who has it,

189 and how did it evolve? *Science* 298(5598): 1569–1579.

190 Heilbron, M., & Chait, M. (2018). Great expectations: is there evidence for predictive coding

191 in auditory cortex? *Neuroscience* 389: 54–73.

192 Martin, A. (2020). A compositional neural architecture for language. *Journal of Cognitive*

193 *Neuroscience* 1–20. doi.org/10.1162/jocn_a_01552.

194   Murphy, E. (2016a). A theta-gamma neural code for feature set composition with phase-
195   entrained delta nestings. *UCL Working Papers in Linguistics* 28: 1-23.

196   Murphy, E. (2016b). The human oscillome and its explanatory potential. *Biolinguistics* 10: 6-
197   20.

198   Murphy, E. (2018). Interfaces (travelling oscillations) + recursion (delta-theta code) =
199   language. Luef, E., & Manuela, M. (Eds). *The Talking Species: Perspectives on the Evolutionary,*
200   *Neuronal and Cultural Foundations of Language.* Graz: Unipress Graz Verlag. 251-269.

201   Murphy, E. (2020). *The Oscillatory Nature of Language*. Cambridge: Cambridge University
202   Press.

203   Salinas, E., & Abbott, L.F. (2001). Coordinate transformations in the visual system: how to
204   generate gain fields and what to compute with them. *Progress in Brain Research* 130: 175–
205   190.

206   Shipp, S. (2016). Neural elements for predictive coding. *Frontiers in Psychology* 7: 1792.