

Lexicon against Naturalness: Unnatural Gradient Phonotactic Restrictions in Tarma Quechua

Gašper Beguš Aleksei Nazarov Anna Björklund Blas Puente Baldoceda

Abstract

It has been shown in separate studies that phonological grammar can operate probabilistically and in phonetically unnatural directions. This paper examines whether phonological grammar can be both probabilistic and unnatural at the same time. We create a new corpus of Tarma Quechua vocabulary (based both on published and unpublished data), and argue that the unnatural probabilistic phonotactic trends in the Tarma Quechua lexicon are statistically significant and show signs of productivity, with evidence from loanword phonology and from morphophonological alternations. We also perform an acoustic analysis of existing recordings to confirm the phonetic status of the the unnatural gradient restriction. To our knowledge, this is the first report of a fully unnatural gradient phonotactic restriction on segmental structure. The existence of unnatural probabilistic phonology has broad theoretical consequences because it requires probabilistic approaches to phonology to derive unnatural patterns. While weighted constraint approaches are more powerful than non-probabilistic approaches, we argue that models without unnatural constraints are not sufficient for deriving Tarma Quechua data. We propose a framework for evaluating phonological analyses using the goodness of fit measurement. This quantitative evaluation approach supports analyses with some, but not all of the relevant conceivable unnatural constraints. We argue that the proposed approach can serve for evaluating other models and approaches.

Keywords: naturalness in phonology, analysis evaluation, goodness of fit, gradient phenomena, phonotactics, maximum entropy grammar

1 Introduction

The question of how phonetic naturalness shapes language is relevant both for theories of synchronic grammar (Chomsky & Halle, 1968; Stampe, 1973; Hellberg, 1978; Donegan & Stampe, 1979; Anderson, 1981; Westbury & Keating, 1986; Archangeli & Pulleyblank, 1994; Hayes, 1999; Wilson, 2006; Carpenter, 2010; Becker et al., 2011; White, 2013; Hayes & White, 2013; Beguš, 2018) and for a discussion on how phonetic, cognitive, and historical factors influence phonological patterns in the world’s languages (Moreton, 2008; Blevins, 2004; Beguš, 2019, 2020, 2022). We present here a case of phonotactic restrictions on segmental features (specifically, the feature $[\pm \text{voice}]$) that are both lexically gradient (Coetzee & Pater, 2008) and phonetically unnatural (Buckley, 2000; Yu, 2004; Coetzee & Pretorius, 2010). Such cases are interesting because they suggest that these phonetically arbitrary or unnatural restrictions extend not only to categorical bans of certain sequences (Coetzee & Pater, 2008), but even to manipulating the relative frequencies of sequences in the lexicon.

1.1 Background

Two questions have received increased attention in the phonological literature: (i) how to represent gradient phonotactic restrictions in the grammar (Frisch et al., 2004; Anttila, 2008; Coetzee &

Pater, 2008; Wilson & Obdeyn, 2009), and (ii) whether and how to represent unnatural processes in the grammar (Hayes, 1999; Hyman, 2001; Blevins, 2004, 2008; Wilson, 2006; Hale & Reiss, 2008; Samuels, 2009; Carpenter, 2010; Coetzee & Pretorius, 2010; Becker et al., 2011; White, 2013; Hayes & White, 2013; Beguš, 2022). To our knowledge, however, there exists no systematic treatment of the intersection of these two topics: unnatural gradient phonotactics, i.e. phonotactic restrictions that, given a particular environment, target a single (segmental) feature and gradiently favor the value of this feature that is unnatural in that environment.

Phonetically unmotivated (see Section 1.2) phonotactics are primarily discussed in light of the “surfeit of the stimulus” problem: there exist unmotivated statistically significant trends in the lexicon, but speakers often fail to generalize these unmotivated trends to nonce words (see, for instance, Becker et al. 2011, 2012b; Hayes & White 2013). Hayes & White (2013), for example, identify ten potential “unnatural” phonotactic restrictions based on systematic gaps in the English lexicon and test in a behavioral experiment whether those restrictions are internalized by native speakers. Using Hayes & Wilson (2008)’s Phonotactic Learner to identify systematic gaps, they find that, for instance, the English lexicon features a restriction against /3/ before a stressed vowel followed by an obstruent (*[+continuant, +voice, –anterior] [+stress] [–son] in constraint formalism; Hayes & White 2013), yet learners did not generalize this to nonce words as strongly as they generalize phonetically motivated constraints. Most studies on this problem agree that speakers do not generalize unmotivated phonotactic restrictions to nonce words, with the suggestion that they might not be represented in grammar (Becker et al., 2011, 2012b; Hayes & White, 2013; Becker et al., 2017; Wilson & Gallagher, 2018). Other studies (for instance, Hayes et al. 2009; Hayes & White 2013) suggest instead that unmotivated processes can actually be generalized to nonce words, but to a lesser extent than natural processes.

Beside its unnaturalness, the other aspect of our data that bears implications for phonological theory is the fact that restrictions are *gradient* rather than categorical. The necessity of encoding gradient phonotactics in the grammar has recently been motivated on the basis of gradient lexicalized co-occurrence restrictions (Berkley, 2000; Frisch et al., 2004; Anttila, 2008; Coetzee & Pater, 2008; Wilson & Obdeyn, 2009; Zuraw, 2010) as well as other processes (Martin, 2007). Coetzee & Pater (2008), for example, discuss two cases of homorganic consonant co-occurrence restrictions: Arabic and Muna. In both cases, an Obligatory Contour Principle constraint (Leben 1973, McCarthy 1986) creates a pressure that is strong, but fails to reach categorical status because of lexical variation: homorganic consonants are dispreferred within the same root, but not completely ruled out. This dispreference, however, operates in the *natural* direction: a restriction against the co-occurrence of homorganic consonants is strongly motivated by both articulatory (Garrett & Johnson, 2013) and perceptual factors (Gallagher, 2010).

As we will argue in Section 1.2, most cases of what literature thus far calls “unnatural” phonotactic restrictions in the literature are in fact either phonetically motivated or phonetically unmotivated, including the restrictions that were tested in the *surfeit of stimulus* experiments. To distinguish our cases from these, we will adopt a new subdivision of naturalness (Beguš, 2019) and focus on gradient processes that are not only phonetically unmotivated, but that operate against universal, phonetically motivated tendencies.

1.2 On naturalness

Traditionally, the naturalness of phonological processes has been conceptualized as a binary split between natural and unnatural processes (see, for instance, Kiparsky 2006, 2008). Natural processes are phonetically well-motivated and typologically frequent. Their high typological frequency is driven by the existence of phonetic precursors, which Yu (2013, 398) defines as “systematic

contextually-induced variations that parallel some cross-linguistically recurring sound patterns”. Unnatural processes, on the other hand, are typologically rare or non-existent and lack phonetic precursors.

However, such a division is insufficient, as it fails to identify an important distinguishing feature within the “unnatural” group itself. All processes in the group traditionally called “unnatural” indeed lack phonetic motivation or phonetic precursors. However, a subset of these processes not only lacks phonetic motivation, but also operates against universal phonetic tendencies — that is, against phonetic precursors responsible for a large number of (natural) phonological processes across many languages. Beguš (2019) proposes a new division of phonological processes with respect to naturalness which we follow in this paper.¹

Consistent with the literature, natural processes are defined as having clear articulatory or perceptual motivations (phonetic precursors; Yu 2013) and being typologically prevalent. Unmotivated and unnatural processes are similar in that they both lack articulatory/perceptual motivation, i.e., they have no phonetic precursors. However, in contrast to unnatural processes, unmotivated processes have no universal phonetic tendencies that militate against them — see the examples in (2), taken from Blevins (2008). On the other hand, (true) unnatural processes not only lack articulatory/perceptual motivation, but also operate against universal phonetic tendencies. In other words, if $A \rightarrow B / X$ is a universal phonetic tendency, $B \rightarrow A / X$ is an unnatural process.

To be precise, a *universal phonetic tendency* is a phonetic pressure that has a clear articulatory or perceptual motivation, and operates passively (without active articulatory control) across different languages (Beguš, 2019).

(1) *Universal Phonetic Tendency* (UPT) (Beguš, 2019):

“UPTs are phonetic pressures motivated by articulatory or perceptual mechanisms (for an overview of phonetic mechanisms, see Garrett & Johnson 2013) that passively operate in speech production cross-linguistically and result in typologically common phonological processes.”

Unlike simply ‘unmotivated’ phonological processes, the tendencies described in this paper are ‘unnatural’ in the sense that they not only lack phonetic motivation, but operate against universal tendencies. It is in this sense that the phenomena in this paper—such as lack of obstruent voicing post-nasally—are not only simply phonetically unmotivated but phonologically unexpected. The use of ‘unnatural’ to describe such a process is chosen especially with reference to this latter property, following usage by Beguš (2019) and others.

For example, post-nasal voicing, intervocalic voicing, and word-final devoicing are universal phonetic tendencies: they have clear articulatory motivations, are amply attested cross-linguistically as typologically common phonological alternations, and operate passively even in languages in which such processes are not among the synchronic alternations (for discussion, see Section 2).

The literature on (un)naturalness has primarily focused on unmotivated processes. Most of the 26 listed examples in the survey of “unnatural” alternations in Blevins (2008) and the ten “unnatural” gradient phonotactic restrictions in Hayes & White (2013) are unmotivated rather than unnatural, according to the definition in Beguš (2019). Consider, for example, the following processes from Blevins (2008) and Hayes & White (2013).²

¹Similar, but slightly different distinctions have been proposed before: see Morley (2014). For Morley (2014), for example, “anti-natural” processes need to violate implicational universals and they are defined as “unattested patterns that do not conform with posited language universals”.

²Unmotivated processes are often also complex (i.e. they manipulate multiple features). Avoidance of complex alternations has long been established as complexity bias Moreton (2008). For a discussion on the complex interplay of naturalness and complexity, see Beguš (2022).

- (2) *Some processes labeled as “unnatural” in Blevins (2008) and Hayes and White (2013)*
- a. /p/ → [s] / __i
 - b. /i/ → [u] / d__
 - c. * $\begin{bmatrix} +COR \\ +cont \\ -strid \end{bmatrix} \begin{bmatrix} -stress \\ +round \end{bmatrix}$ (“No [θ, ð] before stressless rounded vowels”)
 - d. * $\begin{bmatrix} +cont \\ +voice \\ -ant \end{bmatrix} [+stress] [-son]$ (“No [ʒ] before stressed vowel + obstruent”)

All these processes lack phonetic motivation, but they do not operate against universal phonetic tendencies. In other words, it is not the case that [θ, ð] in the context before unstressed rounded vowels are universally preferred: there exists no cross-linguistic passive phonetic tendency that would prefer fricatives before unstressed rounded vowels.

Some gradient restrictions that operate in the unnatural direction have recently been reported in light of the surfeit of the stimulus experiments mentioned in Section 1.2, including (i) a higher rate of alternation in monosyllabic compared to polysyllabic words (Becker et al., 2012b), and (ii) a stronger preference for antepenultimate stress in LLL words compared to HLL words (Garcia, 2017). Strictly speaking, the first restriction does not contradict any universal phonetic tendency (perhaps just a universal phonological tendency). The second restriction targets non-segmental features and while the restriction is significant, it is also very subtle in magnitude (see Garcia 2017).

Few truly unnatural active alternations are reported (for an overview, see Beguš 2018; Dabkowski & Beguš 2024; Beguš & Dabkowski 2025); there is additional debate as to whether unnatural processes are, in fact, possible as productive synchronic alternations. The most compelling case of an unnatural synchronic alternation is post-nasal devoicing (PND) in Tswana and Shekgalagari (Hyman, 2001). PND qualifies as unnatural according to all criteria laid out in this section: it operates against the well-motivated, typologically common, and passive phonetic tendency of post-nasal voicing. Coetzee & Pretorius (2010) have shown that PND is phonetically real and fully productive: speakers generalize it to nonce words (although it does not appear in all varieties and it might have alternative analyses; Gouskova et al. 2011; Downing & Hamann 2021). On the other hand, one of the most convincing cases against unnatural alternations in synchronic grammars is that of final (de)voicing. While final devoicing is a highly common and phonetically motivated process (Steriade, 1997; Iverson & Salmons, 2011) its unnatural counterpart — final voicing — is not attested in any language as a synchronic process (Kiparsky 2006, 2008; for a possible exception, see Yu 2004; de Lacy 2002; Blevins et al. 2020). This has led some scholars to conclude that final voicing is an impossible synchronic alternation (Kiparsky, 2006, 2008). The question of whether fully productive synchronic alternations can be unnatural remains open.

1.3 Aims

To our knowledge, no systematic treatment of unnatural gradient phonotactic restrictions exists in the literature. This paper presents a novel analysis of the lexicon of the Tarma dialect of Quechua (data from Adelaar 1977; Puente Baldoceda 1977 and novel unpublished data) which has a highly unnatural trend in the lexicon that targets a single segmental laryngeal feature [$\pm$ voice] at a considerably greater magnitude than any cases of unnatural gradient restriction reported thus far. In our reported case, the distributions are approaching categorical levels (for example in the post-nasal and post-obstruent position). We show that these trends run counter to specific universal

phonetic tendencies, and argue that the trends are statistically significant, phonetically real, and morphophonologically productive. Based on these findings, we interpret unnatural trends as a gradient phonotactic restriction (in line with Coetzee & Pater 2008 and others).

The existence of unnatural gradient phonotactic restrictions has several implications. One of the most consequential debates in phonology is to what degree it is influenced by phonetics and whether it can accommodate unnatural processes (Hyman, 2001; Hayes & Steriade, 2004; Blevins, 2004; Kiparsky, 2006, 2008; Reiss, 2018; Beguš, 2022). Our data speaks to this debate on two fronts: first, we argue that phonotactic restrictions can extend in phonetically unnatural directions. Second, phonological processes are not only categorical, but often operate probabilistically. Our presented case suggests that not only categorical, but also gradient phonological processes, can operate in the unnatural direction.

The analysis in this paper raises a theoretical question for the theory of Markedness (see, e.g., de Lacy 2002; de Jong 2004; de Lacy 2006; de Lacy & Kingston 2013) in weighted-constraint theories like Harmonic Grammar or MaxEnt (Goldwater & Johnson, 2003; Coetzee & Pater, 2008, 2011; Pater, 2008, 2009; Albright, 2009; Potts et al., 2010). While weighted constraint models are able to derive gradient phonotactics (see, e.g., Coetzee & Pater 2008), we show that, without unnatural Markedness constraints, they are not always able to derive phonotactic systems in which the unnatural element in that context is more frequent than the natural element in a given context. We call this effect the “Natural Gradience Bias”. Since the cases presented here are of this exact nature, this opens up the possibility that Markedness constraints might, after all, be able to counter universal phonetic tendencies (cf. Hayes & White 2013; Zhang 2010 and other work in that direction).

We also apply, for the first time to our knowledge, a model comparison paradigm to investigate the necessity of unnatural constraints to model our data in a grammar framework. We model Tarma Quechua data with natural and selected unnatural constraints, then compare the goodness of fit of various models. Such simulations, as well as the proposed evaluation technique, bring several advantages. Firstly, we can test whether phonetically unnatural gradient restrictions can be derived with natural-only constraints. Secondly, we can quantitatively evaluate different analyses and argue in favor of unnatural constraints based on a quantitative metric. Finally, the technique also facilitates interpretability of MaxEnt models: we provide explanations for various predictions made by the MaxEnt simulations.

2 Universal tendencies for voicing

The case of unnatural phonotactics presented in this paper targets a single phonological feature: $[\pm\text{voice}]$. Alternations and phonotactic restrictions targeting this feature are among the most well-studied phenomena in phonology and we have a substantial body of research on the typology as well as the phonetics of voicing cross-linguistically. The attention that voicing has received in the literature helps us determine in which environments the voice feature is universally (dis)preferred, i.e. which processes targeting $[\pm\text{voice}]$ are “unnatural” according to the definition above (see (1)). Table 1 summarizes these tendencies.

A body of research has established that voicing is universally articulatorily dispreferred in word-initial and word-final or coda positions (Westbury & Keating, 1986; Iverson & Salmons, 2011); it is universally preferred intervocally (Westbury & Keating, 1986; Davidson, 2016) and after a nasal (Rothenberg, 1968; Kent & Moll, 1969; Ohala & Ohala, 1993; Ohala, 1983; Hayes & Stivers, 2000). Voicing is also dispreferred in clusters and geminates: the longer the closure, the more antagonistic articulatory forces are to voicing (Westbury & Keating, 1986). Finally, the voicing feature is also

	[+voice] preferred	[+voice] dispreferred
a. Intervocalic voicing	between two vowels	
b. Postnasal voicing	after [+nasal] consonant	
c. Voicing agreement	adjacent to [+voice] consonant	adjacent to [-voice] consonant

Table 1: Some universal phonetic tendencies for [+voice].

universally dispreferred before another obstruent with a different value in voicing: clusters that disagree in $[\pm\text{voice}]$ are universally dispreferred (Myers, 2010).

Tarma Quechua phonotactic restrictions violate naturalness in two main contexts: post-nasally and in consonant clusters. Post-nasal voicing fulfills all three requirements for a universal phonetic tendency from section 1.2. To begin, it is a common productive phonological alternation and sound change (Kümmel, 2007): as reported in Hayes & Stivers (2000), Locke (1983)’s study identifies 15 out of 197 (or 8%) languages surveyed with post-nasal voicing as a synchronic alternation. The articulatory explanation for post-nasal voicing is described in detail in Hayes & Stivers (2000) and Coetzee & Pretorius (2010), building on previous work by Rothenberg (1968), Kent & Moll (1969), Ohala (1983), Ohala (1993), and others. Two basic factors are identified: in the transition from the nasal stop into the oral stop of nasal-plosive clusters, the velum has to rise from its lowest position to a complete closure. During this transition, air can still escape through the nasal cavity (“nasal leakage”), which makes it more difficult to achieve the increase in supraglottal pressure necessary to cease voicing (cf. Coetzee & Pretorius 2010). Secondly, as the velum rises, the volume of the oral cavity increases. This also promotes voicing, as the increased oral cavity volume translates to more time where voicing is possible (Hayes & Stivers, 2000; Coetzee & Pretorius, 2010). Post-nasal voicing is also well-motivated from a perceptual perspective: cues for voicelessness, such as the release burst of a stop, or cues for voicing, such as low frequency energy, are reduced in post-nasal position (Coetzee & Pretorius 2010, p. 405). Finally, voicing is a passive phonetic tendency in post-nasal position, i.e. stops universally feature more voicing into closure post-nasally than in other positions (Hayes & Stivers, 2000; Davidson, 2016).

The second universal phonetic tendency that Tarma Quechua’s phonotactics operate against is voicing agreement in obstruent clusters, which, again, fulfills all three conditions mentioned above. Clusters that agree in voicing are typologically very common and voicing assimilation is one of the most common processes cross-linguistically (Myers, 2010). Myers (2010) lists at least 28 such languages; this list is neither exhaustive, nor does it result from a survey. The phonetic motivation is straightforward: laryngeal features “overlap and blend” in obstruent clusters, which results in the passive voicing of preceding voiceless stops due to laryngeal coarticulation Myers (2010, 164ff). Perceptual factors have also been proposed to promote the agreement of laryngeal features (Myers, 2010). Finally, agreement in laryngeal features fulfills the third condition of a universal phonetic tendency: voicing before another voiced stop is a passive phonetic tendency (Barry & Teifour, 1999) and is attested in languages in which voicing assimilation is not a complete synchronic process. As Myers (2010) claims, “[a]coustic studies have shown that there is a longer voiced interval in an obstruent before a voiced consonant than before a voiceless consonant in English (Haggard, 1978; Docherty, 1992; Smith, 1997; Jansen, 2004), French (Snoeren et al., 1992), and Syrian Arabic (Barry & Teifour 1999; Myers 2010, p. 164).

3 Unnatural trends in the lexicon

This section shows that the identified trends in the lexicons that operate against the two universal phonetic tendencies are statistically significant, phonetically real, and point to evidence in favor of the productivity of these processes.

3.1 Tarma Quechua

Tarma Quechua is a dialect of Quechua (Quechua I) spoken in the Tarma district of the Junín province of Peru. The number of speakers is difficult to establish, as the dialect is rapidly being replaced by Spanish (Adelaar, 1977). Adelaar (1977) and Puente Baldoceda (1977) report approximately 30,000-40,000 inhabitants of the Tarma district and an additional 3,500 inhabitants of La Unión Leticia. However, the number of speakers of the dialect with the particular unnatural phonotactics that are of interest here is much smaller and difficult to estimate. The unnatural phonotactic restriction presented here is also reported in the dialect of Paccho (Adelaar & Muysken, 2004). Because there are no descriptions or recordings of Paccho Quechua available, we leave this dialect out of our discussion.

3.1.1 Stop voicing

The Quechua dialect continuum almost uniformly has only voiceless stops in native vocabulary (Adelaar & Muysken, 2004). Adelaar (1977) and Puente Baldoceda (1977) report that some of these voiceless stops have become voiced in Tarma Quechua (henceforth: TQ), or more precisely, the Quechua dialects spoken in Tarma, Huaricolca, Palcamayo, and La Unión Leticia. Voiceless velar and labial stops (to the exclusion of alveolars) are reported to undergo voicing in intervocalic and post-consonantal positions, but not after a nasal consonant. Adelaar (1977) and Puente Baldoceda (1977) note that voicing does not apply categorically (and that it does not apply post-nasally), but no further analyses on the lexicon are performed. Below we present results of a statistical analysis of the TQ lexicon that reveals a highly unnatural lexical trend. We show that, in addition to limitations on voicing after nasals, the relative rates of voicing in different environments, including intervocalic and post-consonantal positions, contradict the universal phonetic tendencies from Section 2. This paper, to our knowledge, is the first report of this unnatural distribution in TQ.

The obstruent voicing observed in this paper does not reflect a distinction of phones and allophones. It is indeed true that some voiced/voiceless obstruent minimal pairs exist in Tarma Quechua (e.g. $\widehat{tjimba}$ - ‘to put a bag on someone’s back [colocar un bulto en la espalda de una persona]’ and $\widehat{tjimpa}$ ‘opposite side (of a watercourse) [orilla contraria]’, despite a historical lack of voiced obstruents. However, the phenomenon discussed here is that of gradient lexical phonotactics: the preference within the lexicon for phonetically dispreferred outcomes. Thus, the question is not whether voiced and voiceless obstruents form distinct phonemes in Tarma Quechua, but rather why the phonemes within Tarma Quechua as a whole are distributed in a phonetically unnatural way.

3.2 Data

For the purpose of the analysis, we collected all tokens of prevocalic labial and velar stops from the vocabulary list in Adelaar (1977) and an unpublished lexicon by Puente Baldoceda. Altogether 2,097 entries were thus collected. As a comparison, Ernestus & Baayen (2003) analyze 1,697 entries from the CELEX database; Hayes & Londe (2006) use 10,974 unique stem for their analyses. The entries were annotated for place of articulation (labial or velar), voicing, part of speech, loanword status, derived status, left and right context, source dictionary, and syllable position. The root

	#__	N__	V__V	R__	S__	T__
voiced	14	11	158	82	71	29
voiceless	417	83	202	28	20	3
% voiced	3.2	11.7	43.9	74.5	78.0	90.6

Table 2: Voiced vs. voiceless labial and velar stops in Tarma Quechua native vocabulary across contexts.

structure of Quechua is predominantly disyllabic. Some entries in our data include words analyzed as non-derived that involve more than two syllables (Muysken, 2011; Gouskova, 2023). The question of whether a lexical item is analyzed as derived or non-derived by the speakers is non-trivial. In order to avoid subjective analyses of the derived status of lexical entries, we follow the analysis of roots in the two dictionaries. While it is possible that some entries marked as non-derived might be analyzed as derived by the speakers, we do not expect these mismatches to alter the results substantially. For details on data collection and annotation, see Section A.1 in the Supplementary Materials.

3.3 Statistical analysis

3.3.1 Individual segments

The raw data analysis of non-loan non-derived vocabulary in Table 2 reveals a surprising trend: voicing surfaces very rarely post-nasally (11.7%) (e.g. *anku-bina* vs more naturally expected *angu-bina*), in almost half of the lexicon intervocalically (43.9%) (e.g. *a:pa*), and almost always post-consonantally, including in positions after a voiceless fricative (78.0%) (e.g. *alaymoska*) or voiceless stop (90.6%) (e.g. *pikpi*).³ Raw counts suggest that voiced stops are least frequent word-initially and postnasally, and most frequent after voiceless obstruents. To test the statistical significance of this trend, we fit a logistic regression model to the data with the R statistical software (R Core Team) using the *glm()* function. The dependent variable was binary: presence or absence of voicing (voicing coded as success); the independent variables were PLACE of articulation, and LEFT CONDITION (i.e., the phonological class of the immediately preceding segment; additionally, when Left Condition = V, only intervocalic segments are selected). The model also includes the PART OF SPEECH predictor (with four levels; nouns, verbs, dual noun/verbs, and other) to control for potential effects of parts of speech on the voicing pattern and the DICTIONARY source predictor (to test whether the described voicing patterns differ across the two dictionaries).

The final model includes PLACE (sum-coded with two levels: labial and velar, with labial as the reference level)⁴ and LEFT CONDITION (treatment-coded with five levels: initial, post-nasal, intervocalic, post-sonorant, and post-fricative and post-stop, with intervocalic as the reference level). While their interaction is not warranted by AIC, we leave it in the model per Barr et al. (2013)’s work on the benefits of keeping models maximal. The significance of the estimates of interest does not differ between the best-fitting model (which includes the DICTIONARY predictor) and the final model.

The estimates of the final model in Figure 1 and Table 3 illustrate the unnatural trends in the Tarma Quechua lexicon. In the intervocalic position (V_V), voiced and voiceless stops are

³Unless noted otherwise, we will henceforth abbreviate phonological classes as follows: T – voiceless stop, D – voiced stop, S – voiceless fricative, N – nasal, R – non-nasal sonorant, including glides, V – vowel.

⁴For a discussion on contrast coding, see Brehm & Alday (2022). For a discussion on how the final model was chose, see Section A.2 in the Appendix.

equally frequent ($\beta = -0.18, z = -1.62, p = 0.11$). Voiced stops are significantly less frequent word-initially ($\beta = -3.40, z = -9.79, p < 0.0000$) and postnasally than intervocally ($\beta = -1.96, z = -4.61, p < 0.0000$). On the other hands, voiced stops are significantly more frequent after non-nasal sonorants and glides ($\beta = 1.52, z = 5.20, p < 0.0000$), after voiceless fricatives ($\beta = 1.67, z = 5.16, p < 0.0000$), and after voiceless stops ($\beta = 2.49, z = 3.83, p = 0.0001$) than they are intervocally. No overdispersion was detected in the models (ratio = 1.01; tested with *overdisp_fun()* in Bolker 2019).




Figure 1: Percentage of voiced stops according to position based on a logistic regression model in Table 3. All bars in figures represent 95% confidence intervals based on model estimates.

These trends are significant even if we include loanwords and derived words in the analysis. Data with the entire collected vocabulary including loanwords and derived words (total N=1,793) were fit to the same model with LEFT CONDITION, PLACE and their interaction as predictors. Estimates in Table 12 suggest all differences remain significant in this model.⁵

Furthermore, if we add LOANWORD status as a predictor to our model (including loan and derived vocabulary), the fit improves significantly, but it is not clear whether loanword status as a predictor is justifiable in a cognitive model. If speakers are aware of the loanword status of the vocabulary, the distribution of *[pmvoice]* for loanwords ceases to be of interest (in this case speakers may use two different grammars that govern native and loanword phonologies; see, e.g., Itô & Mester (2002); Itô & Mester (2003)). The model with three independent variables was fit to data: LEFT CONDITION, PLACE of articulation, and LOANWORD status (sum-coded with native words as the reference level). Two two-way interactions are significant (LEFT CONDITION:PLACE and LEFT CONDITION:LOAN) and are included in the model. The significance of all main effects remains the same as before (at means of other predictors), except that the V__V vs. S__ and V__V vs. T__ cease to be significant as a main effect, likely due to data scarcity. There is only one loanword with labials or velars in post-stop position, and only twelve in post-fricative position. Other frequency differences remain significant: there is still more voicing intervocally than post-nasally.

If we isolate loanwords from the native vocabulary, we do not observe any unnatural patterns,

⁵The interaction is significant in this model.

	Estimate	Std. Error	z value	Pr(> z)
(Intcpt)=V	-0.1791	0.1105	-1.62	0.1051
#	-3.4002	0.3474	-9.79	0.0000
N	-1.9615	0.4254	-4.61	0.0000
R	1.5191	0.2921	5.20	0.0000
S	1.6683	0.3232	5.16	0.0000
T	2.4916	0.6502	3.83	0.0001
Place1	0.5358	0.1105	4.85	0.0000
#:Place1	-0.0433	0.3474	-0.12	0.9007
N:Place1	-0.7466	0.4254	-1.76	0.0792
R:Place1	0.2760	0.2921	0.95	0.3447
S:Place1	0.1150	0.3232	0.36	0.7219
T:Place1	-0.7082	0.6502	-1.09	0.2760

Table 3: Logistic regression model with PLACE (sum-coded) and LEFT CONDITION and their interaction as predictors (non-loan non-derived vocabulary)

which is not surprising as the donor language, Spanish, does not feature any of the unnatural patterns in TQ. In a model that excludes derived words, but includes loanwords with three predictors (LEFT CONDITION, PLACE, and LOAN with all interactions), the loanword vocabulary does not reflect the unnatural pattern from the non-loan vocabulary (see Figure 8). Adding loanword status as a predictor also introduces a problem of data scarcity, since there are only a few loanwords with labials or velars in post-obstruent position. As a whole, loanwords prove inconclusive to our analysis. However, the unnatural voicing pattern in TQ does apply to a subset of loanwords.

To test the effect of syllable structure, we fit a model with LEFT CONDITION and SYLLABLE STRUCTURE as predictors with their interaction (on non-loan non-derived vocabulary). Adding SYLLABLE STRUCTURE introduces highly sparse data in some conditions, which is why it is not added to the main model and which is why we do not include PLACE as a predictor here. We divide data into four groups according to the syllable structure: onset of the ultimate syllable in disyllabic words (the most common position given TQ root structure), onset of the ultimate syllable in trisyllabic words, onset of the penultimate syllable in trisyllabic words, and all other syllabic positions (elsewhere). The estimates of the model in Figure 9 suggest that in the ultimate syllable of disyllabic words, all trends are as observed in the full model. Occurrences of labial and velar stops in trisyllabic words are rather rare, and the effects are not significant. Raw data, however, suggest that the same tendencies do also operate in trisyllabic words. For example, in the postnasal position in trisyllabic words, 13 (76.5%) stops are voiceless and 4 are voiced (23.5%). In the position after a voiceless stop (T), 2 stops are voiced (66.7%) and 1 is voiceless (33.3%). A notable deviation from the general trend is that in trisyllabic words, stops in a position after a voiceless fricative are voiceless in 5 cases and voiced in 1 case. In sum, the observed unnatural differences in voicing are confirmed when they all appear in the most common position given TQ’s syllable structure: the ultimate syllable of disyllabic words. In trisyllabic words, the data are sparse, but at least the distinction in voicing between post-nasal and post-stop positions operate in the expected unnatural direction.

As shown in this analysis, [+voice] in labial and velar stops is significantly less frequent word-initially and post-nasally compared to in intervocalic position. [+voice] is significantly more frequent in post-sonorant and post-obstruent position compared to intervocalic position in the TQ native vocabulary. These trends hold across parts of speech, syllable structures, and source dictio-

1 st member	2 nd member	
	Labial	Velar
t	lutbi	mutgi
tʃ	/	aʃga
tʃ̥	aʃba	maʃga
k	takba	/
s	tʃasbu	tʃasgi
ʃ	kaʃbi	ifgi
x	saxbi	manexax-gunas
l	tʃilbi	tʃilgi
r	karba	argu
j	ajba	ajga
w	kawbu	awgis

Table 5: Obstruent clusters by the first member with an example for each combination in TQ (from Adelaar 1977).

naries.

As argued in Section 2, post-nasal and intervocalic positions universally prefer voicing, while voiced stops after voiceless obstruents are universally dispreferred. The fact that TQ exhibits less voicing post-nasally and intervocalically than after voiceless obstruents (labial and velar stops never occur after a voiced obstruent in our data)⁶ is thus highly unnatural. TQ voicing thus operates in a direction opposite to two universal phonetic tendencies: it operates more frequently where it is universally dispreferred (post-consonantly) and less frequently where it is universally preferred (post-nasally and intervocalically). These findings are summarized in the table in 4.

Universal tendencies for [+voice]	Observed significant trends in TQ
T/S__ < V__V	V__V < T/S__
T/S__ < N__	N__ < V__V < T/S__

Table 4: Unnatural distribution of [+voice]

3.3.2 Disagreement of voicing in clusters

Another aspect of gradient unnaturalness emerges in TQ if we look into the within-context distribution of voicing in consonant clusters involving obstruents: obstruent clusters that agree in voicing are gradiently dispreferred in TQ — clusters that disagree in voicing are significantly more frequent.

In TQ, labial and velar stops surface as voiced in non-nasal post-consonantal position (Table 2). The following consonants are attested as triggering voicing: [t, tʃ, tʃ̥, k, s, ʃ, x, l, l̥, r, j, w]. The list includes voiceless fricatives, affricates, and even voiceless stops. The following clusters of two stops are attested: [kb, tb, tg]. Table 5 presents examples of clusters that disagree in voicing after each consonant (data from Adelaar 1977).

A statistical analysis of this trend reveals that obstruent clusters that disagree in voicing are much more frequent than clusters that agree in voicing if the second consonant is either a labial

⁶The analysis shows post-nasal < intervocalic and intervocalic < post-obstruent, from which post-nasal < post-obstruent can be derived by transitivity.

	TT	TD	DT	DD
Count	23	100	2	0
Percent	18.4%	80.0%	1.6%	0%

Table 6: Voice feature in obstruent clusters in which the first element is an obstruent and the second element is a labial or a velar stop in non-derived non-loan TQ vocabulary.

or a velar stop. Table 6 shows the number of occurrences of obstruent clusters (both stops and fricatives; for the purposes of this particular analysis, we mark all voiceless obstruents with T, and all voiced obstruents with D) in which the first element is an obstruent and the second element is a labial or a velar stop in the non-derived non-loan TQ vocabulary. To test the statistical significance of this distribution, a logistic regression model was fit to data with voicing as the dependent variable, PLACE of articulation and CLUSTER POSITION (first vs. second) without the interaction (which is not significant). In the second position in obstruent clusters (after the first obstruent which in all but two cases are voiceless), voicing is significantly more frequent than non-voicing ($\beta = 2.25, z = 5.21, p < 0.00001$). Voicing is significantly less frequent when the obstruent is in the first position compared to the second position ($\beta = -5.48, z = -6.91, p < 0.00001$). Tested differently, the proportion of clusters that disagree in voicing is significantly higher than 50% (according to the Exact binomial test; $p < 0.00001$).

TQ thus features a statistically significant trend that restricts clusters agreeing in voicing in favor of disagreeing clusters. This trend is both gradient and unnatural.

The trend against agreeing obstruent clusters in TQ is unnatural in one two additional respects. Table 6 shows a preference for TD clusters, compared to DT clusters — which goes against yet another phonetic tendency. Voicing is articulatorily easier to maintain in the initial parts of closure than it is to initiate voicing after a period of voiceless closure (Ohala & Riordan, 1979; Ohala, 1997). The reason for this articulatory dispreference is straightforward and has been identified as the Aerodynamic Voicing Constraint: airflow and a subglottal-supraglottal pressure difference, necessary for voicing, are sufficient during vowel articulation, but decrease into closure. The reason why voicing is articulatorily difficult to initiate after a period of voiceless closure is that it is difficult to reinstantiate the necessary airflow and pressure difference — once the closure has caused them to decrease — without releasing the stop closure completely. In addition, there is a typological tendency towards respecting the Syllable Contact Law (Vennemann, 1988), which states that the coda of the preceding syllable should be more sonorous than the onset of the following syllable and which also prefers DT over TD clusters. Decreasing phonation into closure is observed as a passive tendency in several languages (see, for instance, Möbius 2004; Davidson 2016). In other words, voicing has a universal tendency to decrease rather than increase during the closure. The restriction in TQ against DT (decreasing in voicing) clusters in favor of TD clusters (increasing in voicing) is thus unnatural: it operates against the universal phonetic tendency that decreases voicing into closure. Finally, the fact that labials and velars have voiced counterparts in TQ, but alveolars do not also goes against the phonetic tendency whereby labials are most likely to be voiced and velars least likely (Everett, 2018).

3.4 Phonetics

The segmental inventory of Tarma Quechua includes 22 consonant and 10 vowel segments. This includes five voiceless stops (/p/, /t/, /tʰ/, /c/, /k/), three voiced stops (/b/, /d/, /g/), four voiceless fricatives (/f/, /s/, /ç/, /x/), four nasals (/m/, /n/, /ɲ/, /ŋ/), four liquids (/l/, /ʎ/, ɭ, ʟ),

and two semivowel/glides (/w/, /j/). The vowel inventory consists of long/short pairs of /a/, /i/, /e/, /o/, and /u/. The canonical syllable shape is CVC(C), where VV sequences are not permitted and V: may only occur word-finally. Consonant clusters are permitted in coda position but not coda onset (Adelaar (1977)).

No previous detailed phonetic analyses of the system of voicing in TQ exist: Adelaar (1977) and Puente Baldoceda (1977) are based on qualitative descriptions of recordings and are not supported by acoustic analyses. Our analysis confirms the phonetic reality of the TQ voice system as described above, making the case for true unnaturalness in the TQ data. In this analysis, a segment is considered voiced if it has a clear voicing bar in the spectrogram.

The analyzed recordings were obtained online⁷ in .wav format, sampled at 90 kHz⁸ with 16-bit quantization and analyzed with Willem Adelaar’s permission in the Praat software (Boersma and Weenink 2016). The recordings were made by Willem Adelaar in 1970 in Tarma, in the Junín province of Peru. The informant was a 35 year old male speaker of TQ. The recordings are noisy with considerable reverberation (for a detailed acoustic analysis of reverberation, see Section A.3 in the Supplementary Materials), but the analysis nevertheless reveals important aspects of the unnatural gradient phonotactics and of the phonetic system of TQ in general. The following discussion includes waveforms and spectra that were chosen as best representing the described voicing pattern. Given the quality of the recordings and a substantial presence of echo noise, a more comprehensive quantitative phonetic analysis is difficult to undertake.

Figure 2 shows four waveforms and spectrograms of two TD clusters: [tb] and [kb]. All four spectrograms clearly show that the initial stop of the cluster is voiceless with almost no phonation into closure and that phonation does not start until the onset of the second stop’s closure.




Figure 2: Waveforms and spectrograms of four TD clusters: [atbi], [atbi], [akba], and [ukba].

⁷ Accessible online at: <https://corpus1.mpi.nl/ds/asv/?0&0%5C&openpath=node:1483874>

⁸ Although not specified by the archive from which the audio was retrieved, it can be assumed from their age that the recordings were originally analogue, and therefore have no original sampling frequency.

The same situation holds when the first element is an obstruent other than a stop, such as in the [sb] or [sg] sequences in Figure 3. The lack of any low frequency energy in the fricative portion of the cluster confirms that the first element, [s], is voiceless. Phonation starts at the onset of the stop.




Figure 3: Waveforms and spectrograms of two TD clusters: [sb] and [sg].

The exact realization of voiced stops in clusters may vary. A precise distribution is difficult to establish with limited data, but a short transitional vocalic element is occasionally found between the voiceless and voiced obstruent, indicating a smaller degree of gestural overlap (Figure 5).⁹ Occasionally, the voiced element surfaces as a continuant (though whether it is an approximant or fricative is harder to distinguish). These occur in apparent free variation. Figure 4 presents waveforms and spectrograms of phonemic voiced stops that surface as voiced continuants, as indicated by the presence of formant structure during the closure. These formants are present throughout the consonantal part, even in cases where formants do not result from reverberation see Section A.3 in the Supplementary Materials) — after a voiceless stop.




Figure 4: Waveforms and spectrograms of voiced stops that surface as voiced fricatives: intervocalically (left) and after a voiceless stop (right).

⁹Occasionally, the second element is found to surface as voiceless or deleted.




Figure 5: Waveforms and spectrograms of voiced labial stops in post-consonantal position with a short vocalic element between the voiceless and voiced element: [xb] (left) and [fb] (right).

After nasals, on the other hand, voiceless stops are the preferred variant, as detailed in Table 3. Figure 6 shows spectrograms with voiceless [p] and [k] after nasals. Also note that voiceless stops in TQ are unaspirated, which means that the phonotactic restriction in fact targets the feature $[\pm\text{voice}]$ rather than the feature $[\pm\text{spread glottis}]$.




Figure 6: Waveforms and spectrograms of voiceless stops in post-nasal position: [mp] (left) and [ŋk] (right).

3.5 Productivity

Corroborating its status as a phonotactic restriction, there exists evidence that the gradient phonotactic restriction is synchronically active in some morphophonological alternations. Creider (1968) and Adelaar (1977) identify four suffixes with an initial voiced labial stop that feature morphophonemic alternation:

- (3) *Alternating suffixes*
 - a. -ba/-pa ‘genitive’
 - b. -bax/-pax ‘purposive’
 - c. -bita/-pita ‘procedentive’
 - d. -bis/-pis ‘even, too’

The allomorph with voiced initial stops is selected after vowels and non-nasal consonants, including voiceless obstruents; the allomorph with voiceless initial stop is selected after nasals (Creider, 1968). The distribution is illustrated in (4).

- (4) a. *Intervocalic*
 wawxi-gi-**ba** wayi-n
 ‘the house of your brother’
 b. *Post-nasal*
 wayi-n-**pa** pasa-fun
 ‘we’re going to walk by way of his house’
 c. *Post-obstruent*
 tamya-ya-n nuqa-ntʃik-**baq**
 ‘it is raining now for us’ (Creider, 1968, p. 12-13)

This process is productive for a subset of suffixes. Other suffixes do not enter the alternation. For example, the highly frequent plural suffix /-guna/ and other suffixes /-bura/, /-gama/, and /-gasga/ have no voiceless allomorphs in post-nasal position (Adelaar, 1977, p. 59). The productivity of this morphophonemic alternation also differs across the dialects. (Adelaar, 1977, p. 59) reports that voiceless allomorphs are required in Vicora Congas, whereas in Huanuquillo the rate of application varies, i.e. is gradient.

The behavior of loanwords provides further evidence for the productivity of unnatural gradient phonotactics. Most Spanish loanwords retain their original voicing. Sporadically, however, voicing of an originally unvoiced segment or devoicing of an originally voiced segment does occur (data from Adelaar 1977). It is possible that voicing is facilitated because Spanish voiceless stops are at least partially voiced in these positions (Hualde et al., 2011).

- (5) a. Sp. *cuculi* > kuguli: ‘white-winged dove’
 b. Sp. *cotpe* > kutbi ‘an animal from the mountains’
 c. Sp. *sauco* > sawgu ‘magic tree’
 d. Sp. *vaca* > wa:ga ‘cow’

In two loanwords, a Spanish voiced intervocalic stop devoices to a TQ voiceless stop (data from Adelaar 1977).

- (6) a. Sp. *taruga* > taruka ‘deer’
 b. Sp. *dios se lo pague* > jusulpa:ki ‘thank you’

The two loanwords with devoicing of intervocalic stops are especially relevant for the discussion on the productivity of TQ unnatural gradient phonotactic restrictions. Devoicing in TQ [jusulpa:ki] from Sp. *dios se lo pague* is hardly a result of an earlier borrowing, when TQ voicing supposedly did not operate yet. The historical development of TQ involves only voicing of voiceless stops, not devoicing of voiced stops. Regardless of when Spanish *dios se lo pague* was borrowed, sound change could not have produced TQ [jusulpa:ki]. Similar claims can be made about [taruka], but the situation there is more complex as the word is not of Spanish origin and has the devoiced variant even in Spanish. The loanwords suggest that the gradient phonotactic restriction was potentially productive.¹⁰

While morphophonological and loanword data suggest the process in TQ is at least to some degree productive, there still exists a possibility that morphophonological alternations in (4) are lexicalized or governed by non-phonological parts of grammar and that speakers would not productively extend the distributions to wug-tests. The possibility that a phonological process is not productive is always present when studies do not perform wug-tests (which represents the majority

¹⁰This assumes, of course, that these loanwords were not borrowed into TQ via some other Quechuan dialect without this particular voicing pattern after the historical development of this pattern was completed in TQ.

of phonological literature). We present several pieces of evidence that the process might be at least partially if not fully productive, but the possibility of the process not being extended to wug-tests is always present.

4 Grammar learning simulations

The results from the statistical and acoustic analysis of the unnatural trends in TQ in Section 3 suggest that the grammar needs to derive not only phonetically motivated, but also phonetically unnatural gradient processes in phonology. Whether phonetically unnatural constraints are needed in the grammar has been a long-standing discussion in phonology (Hyman, 2001). In this section, we tackle this question by simulating grammars with natural and unnatural weighted constraints within the Maximum Entropy (MaxEnt) framework (Goldwater & Johnson, 2003; Hayes & Wilson, 2008). We use model selection techniques to investigate which natural/unnatural constraints are necessary for deriving gradient phonotactic restrictions. In doing this, we demonstrate the existence of a “natural gradient bias” in MaxEnt grammars: the presence of only natural constraints cannot motivate all unnatural gradient processes, *pace* Magri & Anttila (2018).¹¹ Such a bias is to be expected: natural constraints prefer natural configurations, so that zero weight on any natural constraint gives unnatural and natural candidates equal probabilities, and non-zero weight would give natural candidates a greater probability than unnatural ones, so that unnatural candidates should never prevail over natural ones. Our approach also facilitates interpretation of what various MaxEnt models are learning and where they fail to learn phonotactic generalizations.

In other words, we compare phonotactic MaxEnt learning models with different constraint sets to find how the presence of particular unnatural constraints affect the model’s success in capturing the data (see also literature like Hayes et al. 2012). We also consider a MaxEnt model that includes faithfulness, learned based on the weights of the model without faithfulness, which is able to produce an even better fit for the data by accounting for the fact that each word has a uniform pronunciation, rather than a probabilistic distribution over pronunciations.

4.1 Setup of MaxEnt Models

All our MaxEnt models are trained on the same training data, which consist of all 1321 native TQ words found in the corpus (in Section 3) together with their attested pronunciations and competing surface candidates that explore all alternative possibilities of assigning voicing to the labial and velar stops in the word. For example, $\widehat{t\text{sim}pajgu}$ - ‘to (make) kneel [arodillar(se)]’ will have attested candidate $[t\text{sim}pajgu-]$ and competing candidates $[t\text{sim}pajku-]$, $[t\text{sim}bajku-]$, $[t\text{sim}bajgu-]$. An illustrative sample from the training data is given in section A.4 in the Supplemental Materials.

Our MaxEnt models assign probabilities to surface candidates that sum to 1 for each underlying representation (UR), as in Goldwater & Johnson (2003) and unlike Hayes & Wilson (2008). We work here with a setup where the underlying voicing of all labial and velar stops is unspecified, e.g., $/\widehat{t\text{sim}PajKu}/$ (with P and K being labial and velar stops, respectively, unspecified for [voice]). In this dialect of Quechua, voicing is contrastive, and we do believe that it is ultimately underlying. However, we make this pragmatic modeling decision to simplify learning (see the difficulties with UR learning mentioned by Tesar 2004, 2006, 2014; Jarosz 2006; Nyman & Tesar 2019, amongst others). In section 4.4, we show that adding voicing contrast to the model while retaining desirable predictions is possible.

¹¹Anttila (2008) alludes to this generalization, but does not provide a more elaborate account. For an earlier treatment of “natural gradient bias”, see Beguš & Nazarov (2018).

Family	Natural constraints	Unnatural constraints
Context-free	*[+voice,-son], *b, *g	*[-voice,-son], *p, *k
Word-initial	*#[+voice,-son], *#b, *#g	*#[-voice,-son], *#p, *#k
Voicing agreement	AGREE(voice)/[-son], AGREE(voice)/[+labial,-son], AGREE(voice)/[+dorsal,-son]	DISAGREE(voice)/[-son] DISAGREE(voice)/[+labial,-son], DISAGREE(voice)/[+dorsal,-son]
Coda condition	CODACOND([+voice])/[-son], CODACOND([+voice])/[+labial,-son], CODACOND([+voice])/[+dorsal,-son] ¹³	CODACOND([-voice])/[-son], CODACOND([-voice])/[+labial,-son], CODACOND([-voice])/[+dorsal,-son]
Intervocalic	*VTV, *VpV, *VkV	*VDV, *VbV, *VGv
Postnasal	*NT, *Np, *Nk	*ND, *Nb, *Ng
Post-R	*RT, *Rp, *Rk	*RD, *Rb, *Rg

Table 7: Constraints considered in our MaxEnt models.

In this setup, there are no feature mismatches between URs and surface candidates, which means that IDENT constraints like IDENT(voice) would have no violations for any candidate. Since we do not include any unfaithful mappings except voicing changes in labial/velar stops, meaning that MAX, DEP, and other faithfulness constraints would also yield zero violations, we only include markedness constraints in our simulations, but this is only provisional: if a broader candidate set were considered (one with deletions, insertions, and other unfaithful mappings), various faithfulness constraints would be necessary, so faithfulness must be included in a fuller analysis of the data set. Each resulting grammar is still compatible with the Rich Base (Prince & Smolensky, 1993/2004), since underlying /p/, /b/, /k/, /g/ etc. would still surface as a voiced or voiceless stop, depending on phonotactic considerations.¹²

We consider 42 constraints: 21 typologically and phonetically natural ones and 21 unnatural mirror images. These can be split up into 7 families based on the context in which they operate (Table 7). The constraints whose name starts with * assign one violation mark for every occurrence of the sequence of sounds in the constraint name (as in Section 3, T stands for a voiceless obstruent, D stands for a voiced obstruent, N stands for a nasal, and R stands for a nonnasal sonorant consonant); definitions of (DIS)AGREE and CODACOND will be given below.

The natural constraint families can be justified as follows. There is a general bias against voiced obstruents both across and within languages (Ohala, 1983) — hence the constraint *[+voice,-son] or, in short, *D (Lombardi 2001). Voicing is also generally avoided word-initially (Westbury & Keating, 1986), given processes like word-initial strengthening: hence the constraint *#[+voice]. Voicing agreement within clusters is a widely attested phenomenon (Myers, 2010); we use the AGREE constraint family (Baković, 2000).¹³ AGREE([F]) assigns one violation mark for every pair of adjacent segments that disagree on [F]. Alternative proposals could be used, instead, e.g., Agreement by Correspondence (Rose & Walker, 2004) or Share (McCarthy, 2011).

To account for the fact that marked feature values are preferred in onset position or in a coda assimilated to an onset, Itô (1986) proposed the Coda Condition, later adapted as the CODA-

¹²Surface candidates with unspecified voicing are not included: it is assumed that a high-weighted constraint HAVEVOICE leads to very low probabilities on such candidates.

¹³We have defined the (DIS)AGREE constraints over obstruents only. One of the reviewers remarked that defining this family over sonorants and obstruents together would allow DISAGREE to trigger postnasal devoicing as well as post-obstruent voicing. Since postnasal devoicing is already motivated by *[+voice,-son], as defined above, and the unnatural *ND constraint, to be defined below, we have not pursued this option.

COND constraint family (e.g., Itô & Mester, 1994; Zoll, 1998; McCarthy, 2007, 2008b). CODA-COND([F])/[G] assigns one violation mark for every instance of [F] in a coda segment that bears feature [G] such that [F] is not shared with the following onset; e.g., CODA-COND([+voice])/[-son] assigns one violation mark for every instance of [-sonorant] in a coda (i.e., a coda obstruent) that is voiced itself but is not followed by a voiced onset.

Finally, constraints against intervocalic (Westbury & Keating, 1986; Kaplan, 2010; Davidson, 2016) and postnasal (Pater, 1999) voiceless obstruents have been proposed to account for the avoidance of such configurations across languages. The *RT constraint family is an extension of these constraints, based on the general idea that a sonorant prefers to be followed by a voiced rather than voiceless obstruent.

The place-specific versions of these constraints (specific to [+labial] or [+dorsal] consonants) are violated whenever the target obstruent in the constraint definition is labial or velar, respectively. The place-specific CODA-COND constraints target codas that are labial or velar. The place-specific AGREE constraints are violated whenever at least one obstruent in a cluster is labial (or velar, respectively) and the obstruents disagree in voicing.

The unnatural counterparts of these constraints are obtained by reversing the value of [voice] in the constraint, if said constraint only refers to one of the values (+ or -). As the unnatural counterpart of AGREE, we devised the constraint DISAGREE, which assigns a violation mark every time that two adjacent obstruents have equal values of [voice]. Its place-specific variants are defined the same as for the AGREE family: for instance, DISAGREE(voice)/[+labial,-son] is violated once for every obstruent cluster that contains a labial obstruent and in which the obstruents agree in voicing.¹⁴

With these constraints, we built 11 models in total. These comprise 2 basic models: Full (all constraints), Natural (only natural constraints); 8 expanded models, which have a tailored mix of natural and unnatural constraints; and 1 model with faithfulness constraints, which is considered in Section 4.4.

4.1.1 Description of models

The Full model has all 42 constraints in Table 7: the learner weights both natural and unnatural constraints to account for the pattern. This model gives the learner the greatest chance of succeeding and is tested to see what roles the learner assigns to natural vs. unnatural constraints.

The Natural model has all 21 natural constraints and provides a baseline of how well a model with only universal, natural constraints fares on the TQ data.

For each of the 7 unnatural constraint series in Table 7, we consider an expanded model that adds the unnatural constraints from one family to the Natural model. For instance, the *ND model contains the unnatural constraints from the Postnasal family in addition to all natural constraints. Each of these models tests whether including the unnatural constraints from one of the families improves the fit of the resulting optimized model to the data compared to the Natural model, and whether the addition of this constraint is warranted is checked with the likelihood ratio test (which is applicable here since the Natural model is nested inside all the expanded models).

Given the results in Section 3.3, we predict that the addition of unnatural constraints from the

¹³The definition of these place-specific CODA-COND-constraints is: for every labial coda (or dorsal coda, respectively), if it has the feature [+voice], this [+voice] feature is also linked to a segment in onset position.

¹⁴For an alternative solution, DISAGREE(voice)/[-son] could be replaced by the combination of OCP(voice) (see Itô & Mester 1986, 2003) and a constraint against two consecutive obstruents sharing the same autosegmental voicing specification. However, we chose the one-constraint option so that grammars with natural and unnatural constraints could have matching numbers of constraints.

Postnasal and Voicing agreement families (*ND, DISAGREE) will lead to an improved fit to the data. If true, this would mean that the gradient phonotactic pattern cannot be modeled without these unnatural constraints. This is what we mean by a “natural gradient bias”: if only natural constraints are present, the unnatural gradient in our data cannot be accounted for.

Based on the latter prediction, we also trained one last model: the DISAGREE & *ND model, which contains unnatural constraints from the Voicing agreement and Postnasal families in addition to the constraints of the Natural model. This model tests whether adding constraints in favor of voicing disagreement and postnasal voicing constitutes an improvement over a model that only has voicing agreement and postnasal voicing constraints.

4.2 Learning setup

All models are trained with Staubs’ (2011) MaxEnt training package, which optimizes constraint weights given an array of tableaux with pre-defined OT-style constraints. Since the data are categorical, the objective of the learner is to maximize data log-likelihood (or, in practice, minimize negative data log-likelihood). The models are optimized with the built-in *optim* function in R (R Core Team, 2018) with the L-BGFS-B method (Byrd et al., 1995), L2 optimization with $\mu=0$ and $\sigma=1,000,000$, and a maximum number of iterations of 100,000. Since the models have no hidden structure (cf. Tesar & Smolensky 2000), there is no reason to assume that different runs and different initializations of this algorithm will yield different optima. Nevertheless, each model is run 5 times with all constraint weights initialized at 0. No differences were observed among the 5 runs for any of the models, so the results presented in the next subsection will not refer to distinct runs of the same model.

4.3 Results

The MaxEnt grammars optimized by this method are evaluated quantitatively in terms of the (conditional) log-likelihood on the training data (using the likelihood ratio test and related techniques, which also take the number of constraints into account), and qualitatively in terms of their performance on a test dataset and in terms of their constraint weights. The test dataset consists of labial and velar stops in all relevant environments, as shown in Table 8. Like in the training dataset, surface candidates differ in the voicing of labial and/or velar stops, which allows us to gauge whether the model prefers a voiced or voiceless stop in the specified context.

The purpose of the test dataset is to see how accurately the statistical tendencies in the training dataset are predicted by the grammar (since the training data do not show variation within the same word, we cannot directly predict individual data points). We have no available data on generalization to new items, but, lacking any evidence to the contrary, we assume that matching the frequencies in the lexicon is desirable (e.g., Zuraw 2000; Ernestus & Baayen 2003; Hayes et al. 2009; Moore-Cantwell 2016; Zymet 2018). Obtaining actual generalization data to test this assumption would be an important next step.

4.3.1 Results for basic and expanded models

Table 9 shows the results for all basic and expanded models, including the number of parameter(s) (constraints to weight), the log-likelihood ($\log\lambda$) of the training data for the models (closer to zero means a better fit to the data), and the Akaike Information Criterion (AIC, Akaike 1974) and Bayesian Information Criterion (BIC, Schwarz 1978), two measures that score models based on log-likelihood and number of parameters; the lower the AIC/BIC value, the closer the model is to a perfect fit to the training data while being able to generalize. *** indicates a significant difference

Context	Data points		
Word-initial	/Pata/	→	[pata, bata]
	/Kata/	→	[kata, gata]
Post(/pre)-obstruent	/atPa/	→	[atpa, atba]
	/atKa/	→	[atka, atga]
	/aKPa/	→	[akpa, agpa, akba, agba]
Intervocalic	/aPa/	→	[apa, aba]
	/aKa/	→	[aka, aga]
Postnasal	/amPa/	→	[ampa, amba]
	/aŋKa/	→	[aŋka, aŋga]
Post-R	/alPa/	→	[alpa, alba]
	/alKa/	→	[alka, alga]

Table 8: Test dataset.

($p < .001$) between a model and the Natural model according to the likelihood ratio test (LRT). This test, performed according to NIST (2012), chapter 8.2.3.3, is applied to a pair of models in a superset-subset relationship and asks whether the superset model’s log-likelihood significantly improves upon that of the subset model. We take the difference in number of constraints between the models as the number of degrees of freedom. As can be seen, the Full model and the models that add either DISAGREE or *ND constraints to the Natural model improve on the Natural model on both the LRT and on the AIC/BIC criterion. The weights of the Natural model and these four improving models are given in Table 14 in the Supplementary Materials, section A.5.

Model	Number of parameters	Log-likelihood	AIC	BIC
DISAGREE	24	609***	1266	1391
DISAGREE/*ND	27	609***	1272	1412
Full	42	609***	1302	1520
*ND	24	714***	1476	1601
Natural	21	736	1514	1623
*[-voice,-son]	24	736	1520	1644
*#[-voice,-son]	24	736	1520	1644
CODACOND[-voice]	24	736	1520	1644
*VDV	24	736	1520	1644
*RD	24	736	1520	1644

Table 9: Results for basic and expanded models.

To gauge the necessity of adding the DISAGREE and/or *ND constraint families to a Max-Ent model of the TQ data, we compare the DISAGREE and *ND models to the combined DISAGREE/*ND model. The DISAGREE/*ND model constitutes an improvement over the *ND model ($\chi^2(3)=210$, $p < .001$), but not over the DISAGREE model ($\chi^2(3)=0$, $p > .5$).¹⁵ Thus, adding the DISAGREE family to a model that already has the *ND family leads to a significant improvement, whereas adding the *ND family to a model that already has the DISAGREE family does not. This

¹⁵The DISAGREE/*ND model actually performs marginally worse than the DISAGREE model, which would yield a negative χ^2 statistic, but since the latter is impossible, we report 0. The worse performance of the DISAGREE/*ND model most likely stems from the regularization term: there are more constraints in this model, and the regularization term minimizes constraint weights stronger when there are more constraints, meaning that the actual constraint weights are slightly farther removed from the true optimal constraint weights.

means that the *ND is not necessary: it helps to account for the data compared to the Natural model, but it does not add anything to a model that already has the DISAGREE family. On the other hand, the DISAGREE family is necessary, as its presence significantly improves a model that already has the *ND family.

Table 9 shows that the DISAGREE model has the same log-likelihood as the full model when rounded to whole numbers, and, indeed, the addition of all the other unnatural constraint families in the full model does not yield a significant improvement ($\chi^2(15)=0.09$, $p>.5$). This means that the DISAGREE family is, in fact, sufficient to capture the improvement that unnatural constraints can bring to fitting the TQ data.¹⁶

These findings are corroborated by inspecting test data performance. Table 10 shows how well the predictions of the best four models on the test dataset match the training dataset using four of the tendencies in the TQ data; the Natural model is included for comparison. Values that differ from the training data column by ≥ 0.05 are shaded). The attested proportions in the training data are calculated as the proportion of words that start with [p/k] out of all words that start with a labial or velar stop, the proportion of occurrences of a post-obstruent [b/g] out of all obstruent+[p,k,b,g] clusters, the proportion of [mp/ŋk] clusters out of all nasal+[p,k,b,g] clusters, and the proportion of all [Rb/Rg] clusters out of all R+[p,k,b,g] clusters, respectively. The proportions in the test data are calculated as follows: given a context (e.g., word-initial), for the two words for that context (e.g., /Pata/, /Kata/) the model-assigned probabilities of the candidates that obey the tendency active in that context (e.g., [pata], [kata] obey word-initial devoicing) are averaged.

The DISAGREE and DISAGREE/*ND models produce results that are practically identical to the full model. As can be seen in the shaded cells, the Natural and *ND models fail to capture the prevalence of voicing-disagreeing clusters, which are predicted to occur less than 50% of the time, while the Natural model predicts a much lower rate of postnasal devoicing than observed in the lexicon.

Context	Trend-obeying forms	Attested	Natural	DisAgree	*ND	DisAgree/*ND	Full model
Word-initial	[pata, kata]	0.98	0.98	0.98	0.98	0.98	0.98
Post-obstruent	[atba, atga, a{k,g}ba]	0.80	0.39	0.83	0.47	0.83	0.83
Postnasal	[ampa, aŋka]	0.90	0.67	0.91	0.91	0.91	0.91
Post-R	[alba, alga]	0.73	0.74	0.74	0.74	0.74	0.74

Table 10: Strength of phonotactic tendencies in the best four models, as gauged on the test set.

The behavior of the Natural and *ND models can be interpreted as follows. Since the Natural and *ND models do not have access to the DISAGREE family of constraints, which prefers voicing mismatch in obstruent clusters, the preference for voicing in the first members of clusters cannot be explained by any constraint. This means that there is unexplained voicing variability in the data (voicing after obstruent, devoicing otherwise), which leads the the *[+voice,-son] family receiving lower weights (see Table 14 in the Supplementary Materials, section A.5) and not being able to capture postnasal devoicing. This means that, in the *ND model, the postnasal devoicing pattern

¹⁶At the suggestion of one of the reviewers, we applied stepwise backwards model selection (e.g., Hayes et al. 2012) with significance level $p < .1$ to find, in an unbiased way, a smallest set of constraints that can account for the pattern. The results, see Supplementary Materials section A.6, also include members of the DISAGREE family, but exclude the *ND family. Interestingly, the unnatural constraint CODACOND([-voice])/[+dorsal,-son] is also included. However, this does not mean that CODACOND([-voice])/[+dorsal,-son] is essential to accounting for the pattern, as the DISAGREE model already sufficiently accounts for the data. In fact, it appears that this constraint is only there to counteract the effects of DISAGREE in /kt/ or /k/ + affricate or fricative clusters, where the absence of a candidate in which the second member of the cluster is voiced would mean that DISAGREE would favor [g] in the coda.

is captured by high-weighted *Nb and *g, instead, while the Natural model can explain postnasal voicing to a lesser extent by using the lower-weighted *[+voice,-son] family.

4.4 Accounting for voicing stability within lexical items

4.4.1 Morphophonological learning setup

One important limitation of the models discussed above is that they do not predict categorical voicing in the attested lexical items of TQ, but rather, predict that most labial and velar stops will vacillate between [p/b] and [k/g], respectively, while this is not attested in the dataset: for instance, *t̪imba*- ‘to put a bag on someone’s back [colocar un bulto en la espalda de una persona]’ is only attested with [b] and *t̪impa* ‘opposite side (of a watercourse) [orilla contraria]’ is only attested with [p]).

This can be remedied by learning URs with a fixed voicing value and including IDENT, so that words with categorical [b] have underlying /b/, words with categorical [p] have underlying /p/, etc. However, the process of learning URs is a complex one (see, e.g., Tesar 2004, 2006, 2014; Jarosz 2006; Nyman & Tesar 2019). Here, we use a two-phase model (cf. Jarosz 2006), in which learning a purely phonotactic grammar, as described in the preceding subsections, is followed by a learning stage that involves specific URs.

For the second stage, we assume that URs are equal to the attested forms in all cases, which creates an incentive for the learner to promote faithfulness constraints over markedness constraints, leading to a lack of generalization (Prince & Tesar, 2004). By starting the optimization on the previously learned weights of the best faithfulness-free model (the DISAGREE model), the markedness constraints get a head start. By adding a bias that favors higher markedness constraint weights in relation to the weight of IDENT, the learner is incentivized to make structured predictions about new forms. This markedness-over-faithfulness bias (Prince & Tesar, 2004) is implemented by penalizing, during the optimization process, the weight of IDENT minus the combined weights of the markedness constraints, a standard option in the hgR software (Staubs, 2011); the degree to which this is done is determined by a coefficient we will call M/F.

4.4.2 Results with faithfulness

The model described above, which may be called DISAGREE/Faithfulness, is trained and tested in the same way as the other models. The *optimal* (best generalizing) version of this model is obtained by setting the initial weights of the markedness constraints to the weights learned by the DISAGREE model, while the initial weight of IDENT is set to 0; and setting the M/F coefficient (controlling the strength of the markedness-over-faithfulness bias) to 0.0001.¹⁷ This model yields a log-likelihood of about -0.0001 for the entire dataset, which is a nearly perfect performance.

A *baseline* version of the DisAgree/Faithfulness model is also run with initial weights at 0 and without a markedness-over-faithfulness bias, to explore what happens without the additional assumptions of two-stage learning and the markedness-of-faithfulness bias. On the training data, this model also has a near-perfect performance (with a log-likelihood of -0.0003), but it differs in its performance on the test data, as indicated below.

The test dataset is kept identical to the one for the previous simulations, including URs in which the voicing of labial and velar plosives is unspecified, so that IDENT has no effect. This is done to maximize the chance that the model with faithfulness would yield a match with the pattern in the

¹⁷Higher M/F coefficient values tested did not lead to convergence, while lower values 0.00001 and 0.000001 yielded a decreased performance on the test data. Not using the weights of the initial DISAGREE model or not using a markedness-over-faithfulness bias also yielded degraded test data performance.

lexicon (and to ensure that the model with faithfulness was tested in the exact same way as the models without faithfulness).

As seen in (11), the DISAGREE/Faithfulness model run with optimal parameters yields a close match (within a tolerance of 0.05) to the lexical tendencies for the phonetically natural word-initial devoicing and post-nonnasal sonorant voicing tendencies, whereas the unnatural post-obstruent voicing and postnasal devoicing tendencies are still active, but weaker than in the data or the predictions of the DISAGREE model without faithfulness. The DISAGREE/Faithfulness model run with baseline parameters matches the lexical tendencies much less closely than the version with optimal parameters (with the exception of the postnasal devoicing tendency), but even in this “worst case scenario”, where the markedness constraints are not boosted, each of the tendencies is still observed with a chance above 50%.

The observed discrepancy between the faithfulness-free and faithfulness-included models is to be expected, as the DISAGREE/Faithfulness model does not require markedness constraints to be active to account for the training data, so despite the markedness-over-faithfulness bias in the optimal parameter version, we expect the markedness constraints to be less active in encoding the lexical tendencies (Prince & Tesar, 2004). However, this setup does seem to predict that, despite the availability of mechanisms to account for the two unnatural tendencies in the grammar, the unnatural tendencies might not match the frequencies in the lexicon as strongly as the natural tendencies.

Context	Trend-obeying forms	Attested	DisAgree (no faith)	DisAgree/Faith optimal	DisAgree/Faith baseline
Word-initial	[pata, kata]	0.98	0.98	0.93	0.80
Post-obstruent	[atba, atga, a{k,g}ba]	0.80	0.83	0.72	0.54
Postnasal	[ampa, aŋka]	0.90	0.91	0.70	0.69
Post-R	[alba, alga]	0.73	0.74	0.68	0.54

Table 11: Strength of phonotactic tendencies in the Disagree/Faithfulness model (with the DISAGREE model without faithfulness as comparison), as gauged on the test set.

These results demonstrate a few interesting points. Firstly, they show the feasibility and learnability of a model that combines the representation of gradient unnatural voicing tendencies (across words) along with lexically fixed voicing values (for individual words); see Moore-Cantwell & Pater (2016), amongst others, for similar work on natural tendencies. Secondly, they show that, even when such models do not match the trends in the lexicon, more strongly misrepresenting the phonetically unnatural trends, these trends are still represented in the grammar at more than 50%, indicating the DISAGREE model’s capability of capturing the data in a generalizable way.

4.5 General discussion of grammar simulations

The Maximum Entropy simulations shown in this section are aimed to test to what extent, if any, the Tarma Quechua data require unnatural constraints in the grammar. These simulations have pre-defined natural and unnatural markedness constraints. Based on the literature reviewed in Section 4.1, it is assumed that there are natural constraints against obstruent voicing across the board and word-initially, and in favor of obstruent voicing after nasals, non-nasal sonorants, and intervocalically. Furthermore, we assume natural constraints for voicing agreement between adjacent segments (AGREE(voice)/[-son]) and against voicing in codas unless followed by a voiced obstruent (CODACOND([+voice])/[-son]). These constraints are inverted in ways specified in Section

4.1 to yield unnatural counterparts. We compare a model with only natural constraints to various models that also contain unnatural constraints.

The model with unnatural voicing disagreement constraints and only natural constraints otherwise (the DISAGREE model) comes out as the best one: it yields the maximal increase in fit to the training data compared to the Natural model (equivalent to the one achieved by adding all unnatural constraints) with the addition of just one unnatural constraint family. The model that includes both voicing disagreement and postnasal devoicing constraints performs equivalently, while a model that includes only postnasal devoicing constraints performs worse. This implies that an unnatural constraint promoting voicing disagreement is necessary, while an unnatural constraint promoting postnasal devoicing is not: postnasal devoicing can be accounted for by giving high weight to the context-free natural constraint $*[+voice, -son]$, and the best possible performance on the training data can be achieved without adding a postnasal devoicing constraint.

Finally, including faithfulness can ensure that there is no within-word variation while still maintaining probabilistic phonotactic tendencies in the grammar. Our DISAGREE/Faithfulness model accounts for the data near-perfectly, assuming that the URs are equal to the observed surface forms. However, in terms of generalization, this model does not represent the unnatural tendencies as strongly as the best faithfulness-free model. Nevertheless, even with non-optimal learning settings, both natural and unnatural trends in the lexicon are generalized to some extent by this DISAGREE/Faithfulness model.

Thus, while TQ exhibits two statistically significant phonetically unnatural tendencies –post-obstruent voicing and postnasal devoicing (Section 3.3)– our simulations suggest that only post-obstruent devoicing requires unnatural constraints in the grammar. Gradient postnasal devoicing in our dataset can be modeled as an interaction between a low-weighted constraint in favor of postnasal voicing ($*NT$) and a high-weighted constraint against obstruent voicing ($*[+voice, -son]$), since voiced obstruents are still very much a minority in the native vocabulary. This means that, against our expectations, the Tarma Quechua phonotactic data do not provide evidence for a postnasal devoicing constraint in the phonological grammar. If the proportion of voiced obstruents had been higher across the board in this language while the proportion of voiced obstruents after nasals had remained the same, there might have been phonotactic evidence for a phonologically unnatural constraint $*ND$. While TQ phonotactics do not support the $*ND$ constraint, the *morphophonological* alternations shown in Section 3.5 that are limited to a subset of morphemes might provide independent evidence for $*ND$. Testing the productivity of these alternations and their theoretical implications is left for future work. Additionally, there are other languages that provide strong evidence in favor of the $*ND$ constraint: Hyman 2001; Coetzee & Pretorius 2010; Beguš 2019.

The post-obstruent voicing pattern, however, does require an “unnatural” constraint in the grammar, DISAGREE(voice)/[-son], which penalizes adjacent obstruents that agree in voicing. According to the definition in Section 1.2, an unnatural pattern is one that runs counter to a universal phonetic tendency. By extension, a constraint that consistently favors such unnatural patterns, in this case, voicing-disagreeing obstruent clusters, is seen as an unnatural constraint.

One uniquely unnatural aspect of the effects of DISAGREE(voice)/[-son] in TQ is that it targets two adjacent obstruents. Phonetic pressures towards voicing uniformity in these conditions operate as strong phonetic tendencies across languages (see discussion in Section 1.2). While voicing dissimilation is attested (cf. Dahl’s Law in Gikuyu, Lombardi 1995 and references cited therein, and Rendaku in Japanese, e.g., Itô & Mester 2003), it typically operates non-locally, where the articulatory phonetic tendencies towards voicing agreement are less strong. Another uniquely unnatural aspect arises from the interaction of DISAGREE(voice)/[-son] with the (natural) CODA-COND([+voice])/[-son] constraint, namely, a preference for output clusters that contain voicing in the second member (TD) rather than in the first member (DT; see Table 6 and Section 3.3.2 for

why this distribution is phonetically unnatural). As far as we can see, DISAGREE(voice)/[-son] does not drive any phonetically natural patterns in TQ in addition to the phonetically unnatural voicing disagreement pattern, meaning that this is truly an unnatural constraint according to our definition.

Despite its clear phonetic unnaturalness, a constraint like DISAGREE may have a functional motivation, as it promotes dissimilation. For instance, Bennett (2015) provides an account of rhotic dissimilation in Sundanese motivated by the combination of pressures to have surface rhotics correspond to one another (Rose & Walker 2004) and to limit such correspondences when the two rhotics are not onsets of adjacent syllables. However, it can be argued that such pressures are rather of a structural or cognitive character, not contradicting the phonetic unnaturalness of the process. Moreover, TQ exhibits segmentally local dissimilation, which Bennett (2015):9.2.2 argues shows different typological patterns from long-distance dissimilation and may have a different motivation.

We do not take a specific stance on the origin of unnatural constraints here. In the literature, there is doubt about the productivity of (some) unnatural patterns (e.g., Hayes et al. 2009; Becker et al. 2012a). If unnatural constraints are indeed learned, not universal, then the question of constraint induction comes into play (e.g., Hayes 1999; Hayes & Wilson 2008). Unfortunately, a full investigation of this topic is outside the scope of this pattern, but the current simulations do show that adding DISAGREE constraints to a model with natural constraints significantly increases the fit to the training data, which shows that the existing data provide evidence for a learner to induce (a) constraint(s) in favor of post-obstruent voicing.

Finally, despite the fact that it has been shown that Maximum Entropy models are not restricted to generating phonetically natural patterns when the grammar only contains natural constraints (see, e.g., Magri & Anttila 2018), the results of the current simulations show that there is still a “natural gradient bias” in these types of models: not all (gradient) phonetically unnatural patterns can be modeled by interactions of natural constraints alone, and the presence of certain unnatural constraints leads to a significantly better fit to the training data when modeling datasets such as the TQ one. See also Hayes (2016) and O’Hara (2022) for similar results showing that an all-natural constraint set imposes restrictions on possible (gradient) patterns in Maximum Entropy models.

5 Conclusion

This paper argues that phonological grammar can be probabilistic and operate in phonetically unnatural directions at the same time. We present a case of unnatural gradient phonotactic restrictions in the lexicon: Tarma Quechua stop voicing. We show that the phonotactic restrictions operate in an unnatural direction (as defined in Section 1.2): in the environments presented (post-nasal, post-voiceless-obstruent, and intervocalic), the articulatorily or perceptually preferred value of $[\pm\text{voice}]$ is less frequent than the dispreferred value. Based on a new corpus (which includes both published and unpublished data) we argue that the unnatural phonotactic restriction is statistically significant and phonetically real. This appears to be the first report of unnatural gradient phonotactics that target segmental features in a segmental context. In Section 3.5, we provide evidence for the productivity of these phonotactic restrictions, which suggests that they are actually present in the grammar (see Coetzee & Pater 2008 and others for arguments that gradient phonotactic patterns must be present in the synchronic grammar).

We furthermore use grammar learning simulations to test whether some phonetically unnatural gradient phonotactic restrictions can be derived as a by-product of natural constraint interactions without any unnatural constraints. We argue that at least some phonetically unnatural constraints are necessary in the grammar to model the Tarma Quechua data to their full extent. Interestingly,

our results (with and without faithfulness constraints) provide evidence for a post-obstruent devoicing constraint, but not for a postnasal devoicing constraint. This is because postnasal devoicing can be parasitically accounted for by a general devoicing constraint, given that the data show a general dispreference for voiced obstruents, while post-obstruent voicing goes against the general devoicing trend in the data.

This result is obtained using a new approach to evaluating the necessity of unnatural constraints by measuring goodness of fit of individual models. This approach allows us to evaluate analyses with natural and unnatural constraints or models with and without indexed constraints. The results of this approach can serve as a diagnostic for which constraints are needed for a model that provides the best fit to the data. The proposed approach also helps us interpret how the MaxEnt learners derive the data. For example, it has helped us understand how post-nasal devoicing can be derived without unnatural constraints. The proposed approach can be used both in further discussions on natural and unnatural constraints as well as for evaluating other phonological analyses.

In future work, we hope that our evidence for the gradient unnatural phonotactic patterns presented here will be corroborated with experimental data. Because of the relative inaccessibility of the Tarma speech community, obtaining such experimental data was outside the scope of this research. However, this type of evidence would greatly enhance the strength of our claim that phonological grammars should at the very least allow some phonetically unnatural constraints.¹⁸ Another related direction for future work would be to test the productivity of the limited morphophonological alternations mentioned in Section 3.5 and computationally model these patterns, to see whether they might provide evidence for the activity of a postnasal devoicing constraint (*ND) in the grammar.

Furthermore, we hope this paper will inspire exploration of other cases of gradient unnatural phonotactics, particularly ones that involve a feature other than $[\pm\text{voice}]$. Finding phonotactic restrictions that gradiently favor the phonetically unnatural variant is a challenging task which has thus far not been a focal area within phonological theory. This paper draws attention to the importance of such processes, points to their theoretical implications, and provides a quantitative framework for evaluating and interpreting weighted-constraint approaches to unnatural phonotactics. We hope this work will spark further interest in phonological processes that are simultaneously probabilistic and unnatural, and we expect that more such cases will be discovered in the future.

References

- ADELAAR, WILLEM F. H. 1977. *Tarma Quechua: Grammar, texts, dictionary*. Lisse: The Peter de Ridder Press.
- ADELAAR, WILLEM F. H., and PIETER C. MUYSKEN. 2004. *The languages of the Andes*. Cambridge: Cambridge University Press.
- AKAIKE, HIROTSUGU. 1974. A new look at the statistical model identification. *IEEE Transactions on Automatic Control* AC-19.716723.
- ALBRIGHT, ADAM. 2009. Feature-based generalization as a source of gradient acceptability. *Phonology* 26.9–41.
- ANDERSON, STEPHEN R. 1981. Why phonology isn't "natural". *Linguistic Inquiry* 12.493–539.
- ANTTILA, ARTO. 2008. Gradient phonotactics and the complexity hypothesis. *Natural Language & Linguistic Theory* 26.695–729. Online: <https://doi.org/10.1007/s11049-008-9052-2>.

¹⁸Whether unnatural constraints are universal/innate or learned from the data is a separate question, one that we do not take a stance on in this paper.

- ARCHANGELI, DIANA, and DOUGLAS PULLEYBLANK. 1994. *Grounded phonology*. Cambridge, MA: MIT Press.
- BAKOVIĆ, ERIC. 2000. *Harmony, dominance, and control*. Rutgers University dissertation.
- BARR, DALE J.; ROGER LEVY; CHRISTOPH SCHEEPERS SCHEEPERS; and HARRY J. TILY. 2013. Random effects structure for confirmatory hypothesis testing: Keep it maximal. *Journal of Memory and Language* 68.255–278.
- BARRY, MARTIN, and RYAD TEIFOUR. 1999. Temporal patterns in Syrian Arabic voicing assimilation. *14th international congress of phonetic sciences*, 24292432.
- BECKER, MICHAEL; LAUREN EBY CLEMENS; and ANDREW NEVINS. 2017. Generalization of French and Portuguese plural alternations and initial syllable protection. *Natural Language & Linguistic Theory* 35.299–345. Online: <https://doi.org/10.1007/s11049-016-9343-y>.
- BECKER, MICHAEL; NIHAN KETREZ; and ANDREW NEVINS. 2011. The surfeit of the stimulus: Analytic biases filter lexical statistics in Turkish laryngeal alternations. *Language* 87.84–125.
- BECKER, MICHAEL; NIHAN KETREZ; and ANDREW NEVINS. 2012a. Asymmetries in generalizing alternations to and from initial syllables. *Language* 88.84–125.
- BECKER, MICHAEL; ANDREW NEVINS; and JONATHAN LEVINE. 2012b. Asymmetries in generalizing alternations to and from initial syllables. *Language* 88.231–268. Online: <http://www.jstor.org/stable/23251831>.
- BEGUŠ, GAŠPER. 2018. *Unnatural phonology: A synchrony-diachrony interface approach*. Harvard University dissertation. Online: <http://nrs.harvard.edu/urn-3:HUL.InstRepos:40050094>.
- BEGUŠ, GAŠPER. 2019. Post-nasal devoicing and the blurring process. *Journal of Linguistics* 55.689–753.
- BEGUŠ, GAŠPER. 2020. Estimating historical probabilities of natural and unnatural processes. *Phonology* 37.515–549.
- BEGUŠ, GAŠPER. 2022. Distinguishing cognitive from historical influences in phonology. *Language* 98.1–34. Online: <https://muse.jhu.edu/article/849525>.
- BEGUŠ, GAŠPER, and MAKSYMILIAN DABKOWSKI. 2025. The blurring history of intervocalic devoicing. *Journal of Linguistics* 61.211230.
- BEGUŠ, GAŠPER, and ALEKSEI NAZAROV. 2018. Gradient trends against phonetic naturalness: The case of Tarma Quechua. *Proceedings of the 48th Meeting of the North East Linguistic Society*, ed. by Sherry Hucklebridge and Max Nelson. Amherst, MA: GLSA.
- BENNETT, WILLIAM. 2015. *The phonology of consonants : harmony, dissimilation and correspondence*. Cambridge: Cambridge University Press.
- BERKLEY, DEBORAH. 2000. *Gradient Obligatory Contour Principle effects*. Northwestern University dissertation.
- BLEVINS, JULIETTE. 2004. *Evolutionary phonology*. Cambridge: Cambridge University Press.
- BLEVINS, JULIETTE. 2008. Natural and unnatural sound patterns: A pocket field guide. *Naturalness and iconicity in language*, ed. by Klaas Willems and Ludovic De Cuypere, 121–148. Amsterdam: Benjamins.
- BLEVINS, JULIETTE; ANDER EGURTZEGI; and JAN ULLRICH. 2020. Final obstruent voicing in Lakota: Phonetic evidence and phonological implications. *Language* 96.294–337.
- BOLKER, BEN. 2019. GLMM FAQ. Retrieved on June 22, 2022. Online: bbolker.github.io/mixedmodels-misc/glmmFAQ.html.

- BREHM, LAUREL, and PHILLIP M. ALDAY. 2022. Contrast coding choices in a decade of mixed models. *Journal of Memory and Language* 125.104334. Online: <https://www.sciencedirect.com/science/article/pii/S0749596X22000213>.
- BUCKLEY, EUGENE. 2000. On the naturalness of unnatural rules. *Proceedings from the Second Workshop on American Indigenous Languages. UCSB Working Papers in Linguistics* 9.1–14.
- BYRD, RICHARD H.; PEIHUANG LU; JORGE NOCEDAL; and CIYOU ZHU. 1995. A limited memory algorithm for bound constrained optimization. *SIAM Journal on Scientific and Statistical Computing* 16.1190–1208.
- CARPENTER, ANGELA C. 2010. A naturalness bias in learning stress. *Phonology* 27.345–392.
- CHOMSKY, NOAM, and MORRIS HALLE. 1968. *The sound pattern of English*. New York: Harper & Row.
- COETZEE, ANDRIES W., and JOE PATER. 2008. Weighted constraints and gradient restrictions on place co-occurrence in Muna and Arabic. *Natural Language & Linguistic Theory* 26.289–337. Online: <http://www.jstor.org/stable/20533006>.
- COETZEE, ANDRIES W., and JOE PATER. 2011. The place of variation in phonological theory. *The handbook of phonological theory*, ed. by Jason Riggle John Goldsmith and Alan C. L. Yu., 401–434. Blackwell.
- COETZEE, ANDRIES W., and RIGARDT PRETORIUS. 2010. Phonetically grounded phonology and sound change: The case of Tswana labial plosives. *Journal of Phonetics* 38.404–421.
- CREIDER, CHET A. 1968. *A morphological sketch of Tarma Quechua*. Master's thesis, University of Minnesota.
- DABKOWSKI, MAKSYMILIAN, and GAŠPER BEGUŠ. 2024. Complex diachronies of final nasalization in austronesian and dakota. *Glossa: a journal of general linguistics* 9.
- DAVIDSON, LISA. 2016. Variability in the implementation of voicing in American English obstruents. *Journal of Phonetics* 54.35–50.
- DOCHERTY, GERARD J. 1992. *The timing of voicing in British English obstruents*. Berlin: Foris Publications.
- DONEGAN, PATRICIA J., and DAVID STAMPE. 1979. The study of natural phonology. *Current approaches to phonological theory*, ed. by Daniel Dinnsen, 126–173. Bloomington: Indiana University Press.
- DOWNING, LAURA J., and SILKE HAMANN. 2021. Why phonetically-motivated constraints do not lead to phonetic determinism: The relevance of aspiration in cueing nc sequences in tumbuka. *Phonological Data and Analysis* 3.139. Online: <https://phondata.org/index.php/pda/article/view/49>.
- ERNESTUS, MIRJAM, and R. HARALD BAAYEN. 2003. Predicting the unpredictable: Interpreting neutralized segments in Dutch. *Language* 79.5–38.
- EVERETT, CALEB. 2018. The global dispreference for posterior voiced obstruents: A quantitative assessment of word-list data. *Language* 94.e311–e323.
- FRISCH, STEFAN A.; JANET B. PIERREHUMBERT; and MICHAEL B. BROE. 2004. Similarity avoidance and the OCP. *Natural Language & Linguistic Theory* 22.179–228. Online: <https://doi.org/10.1023/B:NALA.0000005557.78535.3c>.
- GALLAGHER, GILLIAN. 2010. Perceptual distinctness and long-distance laryngeal restrictions. *Phonology* 27.435–480.
- GARCIA, GUILHERME D. 2017. Weight gradience and stress in Portuguese. *Phonology* 34.4179.
- GARRETT, ANDREW, and KEITH JOHNSON. 2013. Phonetic bias in sound change. *Origins of sound change: Approaches to phonologization*, ed. by Alan C. L. Yu, 51–97. Oxford: Oxford University Press.

- GOLDWATER, SHARON, and MARK JOHNSON. 2003. Learning OT constraint rankings using a Maximum Entropy model. *Proceedings of the workshop on variation within Optimality Theory*, ed. by Jennifer Spenader, Anders Eriksson, and Östen Dahl, 111–20. Stockholm: Stockholm University.
- GOUSKOVA, MARIA. 2023. *Phonological asymmetries between roots and affixes*, 1–44. Wiley. Online: <https://onlinelibrary.wiley.com/doi/abs/10.1002/9781119693604.morphcom007>.
- GOUSKOVA, MARIA; ELIZABETH ZSIGA; and ONE TLALE BOYER. 2011. Grounded constraints and the consonants of setswana. *Lingua* 121.2120–2152. Online: <https://www.sciencedirect.com/science/article/pii/S0024384111001732>.
- HAGGARD, MARK. 1978. The devoicing of voiced fricatives. *Journal of Phonetics* 6.95102.
- HALE, MARK, and CHARLES REISS. 2008. *The phonological enterprise*. Oxford: Oxford University Press.
- HAYES, BRUCE. 1999. Phonetically-driven phonology: The role of Optimality Theory and inductive grounding. *Functionalism and formalism in linguistics, volume I: General papers*, ed. by Michael Darnell and Edith Moravcsik, 243–285. Amsterdam: John Benjamins.
- HAYES, BRUCE. 2016. Varieties of Noisy Harmonic Grammar. Talk given at the 2016 Annual Meeting on Phonology, University of Southern California, Los Angeles, CA, October 21–23, 2016.
- HAYES, BRUCE, and ZSUZSA CZIRKY LONDE. 2006. Stochastic phonological knowledge: the case of hungarian vowel harmony. *Phonology* 23.59104.
- HAYES, BRUCE, and DONCA STERIADE. 2004. Introduction: The phonetic bases of phonological markedness. *Phonetically based phonology*, ed. by Bruce Hayes, Robert Kirchner, and DoncaEditors Steriade, 133. Cambridge University Press.
- HAYES, BRUCE, and TANYA STIVERS. 2000. Postnasal voicing. Ms., University of California, Los Angeles. <http://linguistics.ucla.edu/people/hayes/Phonet/NCPhonet.pdf> (accessed 9 May 2018).
- HAYES, BRUCE, and JAMES WHITE. 2013. Phonological naturalness and phonotactic learning. *Linguistic Inquiry* 44.45–75.
- HAYES, BRUCE, and COLIN WILSON. 2008. A Maximum Entropy model of phonotactics and phonotactic learning. *Linguistic Inquiry* 39.379–440. Online: <https://doi.org/10.1162/ling.2008.39.3.379>.
- HAYES, BRUCE; COLIN WILSON; and ANNE SHISHKO. 2012. Maxent grammars for the metrics of Shakespeare and Milton. *Language* 88.691–731.
- HAYES, BRUCE; KIE ZURAW; PETER SIPTAR; and ZSUZSA LONDE. 2009. Natural and unnatural constraints in hungarian vowel harmony. *Language* 85.822–863.
- HELLBERG, STAFFAN. 1978. Unnatural phonology. *Journal of Linguistics* 14.157–177.
- HUALDE, JOSÉ IGNACIO; MIQUEL SIMONET; and MARIANNA NADEU. 2011. Consonant lenition and phonological recategorization. *Laboratory Phonology* 2.30e1–329. Online: <https://doi.org/10.1515/labphon.2011.011>.
- HYMAN, LARRY M. 2001. The limits of phonetic determinism in phonology: *NC revisited. *The role of speech perception in phonology*, ed. by Elizabeth Hume and Keith Johnson, 141–186. San Diego, CA: Academic Press.
- ITÔ, JUNKO. 1986. *Syllable theory in prosodic phonology*. University of Massachusetts Amherst dissertation.
- ITÔ, JUNKO, and ARMIN MESTER. 1986. The phonology of voicing in Japanese: Theoretical consequences for morphological accessibility. *Linguistic Inquiry* 17.4973.

- ITÔ, JUNKO, and ARMIN MESTER. 1994. Reflections on CodaCond and alignment. *Phonology at santa cruz 3*, ed. by Jason Merchant, Jaye Padgett, and Rachel Walker, 27–46. Santa Cruz: University of California, Santa Cruz.
- ITÔ, JUNKO, and ARMIN MESTER. 2002. *One phonology or many? Issues in stratal faithfulness theory*. Tokyo: Kaitakusha.
- ITÔ, JUNKO, and ARMIN MESTER. 2003. *Japanese morphophonemics: Markedness and word structure*. Cambridge, MA: MIT Press.
- IVERSON, GREGORY K., and JOSEPH C. SALMONS. 2011. Final devoicing and final laryngeal neutralization. *The Blackwell companion to phonology, volume II: Suprasegmental and prosodic phonology*, ed. by Marc van Oostendorp, Colin J. Ewen, Elizabeth Hume, and Keren Rice, 1622–1643. Malden, MA: Wiley-Blackwell.
- JANSEN, WOUTER. 2004. *Laryngeal contrast and phonetic voicing: A laboratory phonology approach to English, Hungarian, and Dutch*. University of Groningen dissertation.
- JAROSZ, GAJA. 2006. *Rich lexicons and restrictive grammars - Maximum Likelihood learning in Optimality Theory*. Johns Hopkins University dissertation.
- DE JONG, KENNETH. 2004. Stress, lexical focus, and segmental focus in English: Patterns of variation in vowel duration. *Journal of Phonetics* 32.493–516.
- KAPLAN, ABBY. 2010. *Phonology shaped by phonetics: The case of intervocalic lenition*. University of California, Santa Cruz dissertation.
- KENT, RAYMOND D., and KENNETH L. MOLL. 1969. Vocal tract characteristics of the stop consonants. *Journal of the Acoustical Society of America* 46.1549–1555.
- KIPARSKY, PAUL. 2006. Amphichronic program vs. Evolutionary Phonology. *Theoretical Linguistics* 32.217–236.
- KIPARSKY, PAUL. 2008. Universals constrain change, change results in typological generalizations. *Linguistic universals and language change*, ed. by Jeff Good, 23–53. Oxford: Oxford University Press.
- KÜMMEL, MARTIN. 2007. *Konsonantenwandel*. Wiesbaden: Reichert.
- DE LACY, PAUL. 2002. *The formal expression of markedness*. University of Massachusetts, Amherst dissertation.
- DE LACY, PAUL. 2006. *Markedness: Reduction and preservation in phonology*. Cambridge Studies in Linguistics. Cambridge University Press.
- DE LACY, PAUL, and JOHN KINGSTON. 2013. Synchronic explanation. *Natural Language and Linguistic Theory* 31.287–355.
- LOCKE, JOHN. 1983. *Phonological acquisition and change*. New York: Academic Press.
- LOMBARDI, LINDA. 1995. Dahl’s law and privative [voice]. *Linguistic Inquiry* 26.365–372.
- LOMBARDI, LINDA. 2001. Why place and voice are different: Constraint-specific alternations in Optimality Theory. *Segmental phonology in Optimality Theory: Constraints and representations*, ed. by Linda Editor Lombardi, 1345. Cambridge University Press.
- MAGRI, GIORGIO, and ARTO ANTILA. 2018. Does MaxEnt overgenerate? Implicational universals in Maximum Entropy Grammar. *Proceedings of the 2017 Annual Meeting on Phonology*, ed. by Gillian Gallagher, Maria Gouskova, and Sora Heng Yin. Washington, DC: Linguistic Society of America.

- MARTIN, ANDREW T. 2007. *The evolving lexicon*. University of California, Los Angeles dissertation.
- MCCARTHY, JOHN J. 2007. Slouching toward optimality: Coda reduction in OT-CC. *Phonological studies 10*, ed. by Phonological Society of Japan, 89–104. Tokyo: Kaitakusha.
- MCCARTHY, JOHN J. 2008b. The gradual path to cluster simplification. *Phonology* 25.271–319.
- MCCARTHY, JOHN J. 2011. Autosegmental spreading in Optimality Theory. *Tones and features: Phonetic and phonological perspectives*, ed. by John Goldsmith, Elizabeth Hume, and Leo Wetzels, 195–222. Berlin: Mouton de Gruyter.
- MÖBIUS, BERND. 2004. Corpus-based investigations on the phonetics of consonant voicing. *Folia Linguistica* 38.5–26.
- MOORE-CANTWELL, CLAIRE. 2016. *The representation of probabilistic phonological patterns: Neurological, behavioral, and computational evidence from the English stress system*. University of Massachusetts Amherst dissertation.
- MOORE-CANTWELL, CLAIRE, and JOE PATER. 2016. Gradient exceptionality in Maximum Entropy Grammar with lexically specific constraints. *Catalan Journal of Linguistics* 15.53–66.
- MORETON, ELLIOTT. 2008. Analytic bias and phonological typology. *Phonology* 25.83–127.
- MORLEY, REBECCA L. 2014. Implications of an exemplar-theoretic model of phoneme genesis: A velar palatalization case study. *Language and Speech* 57.3–41.
- MUYSKEN, PIETER. 2011. Root/affix asymmetries in contact and transfer: case studies from the andes. *International Journal of Bilingualism* 16.22–36, First published online May 10, 2011. Online: <https://doi.org/10.1177/1367006911403211>.
- MYERS, SCOTT. 2010. Regressive voicing assimilation: Production and perception studies. *Journal of the International Phonetic Association* 40.163–179. Online: <http://www.jstor.org/stable/44527013>.
- NIST. 2012. *NIST/SEMATECH e-handbook of statistical methods*. Gaithersburg, MD. Online: <http://www.itl.nist.gov/div898/handbook/>.
- NYMAN, ALEXANDRA, and BRUCE TESAR. 2019. Determining underlying presence in the learning of grammars that allow insertion and deletion. *Glossa: A Journal of General Linguistics* 4.1–41.
- OHALA, JOHN J. 1983. The origin of sound patterns in vocal tract constraints. *The production of speech*, ed. by Peter F. MacNeilage, 189–216. New York: Springer.
- OHALA, JOHN J. 1993. The phonetics of sound change. *Historical linguistics: Problems and perspectives*, ed. by Charles Jones, 237–278. London: Longman.
- OHALA, JOHN J. 1997. Aerodynamics of phonology. *Proceedings of the 4th Seoul International Conference on Linguistics (SICOL)*, 84–91. Seoul: Linguistic Society of Korea.
- OHALA, JOHN J., and MANJARI OHALA. 1993. The phonetics of nasal phonology: Theorems and data. *Nasals, nasalization, and the velum*, ed. by Marie K. Huffman and Rena A. Krakow, 225–249. San Diego, CA: Academic Press.
- OHALA, JOHN J., and CAROL J. RIORDAN. 1979. Passive vocal tract enlargement during voiced stops. *Speech communication papers*, ed. by J. J. Wolf and D. H. Klatt, 89–92. New York: Acoustical Society of America.
- O’HARA, CHARLIE. 2022. MaxEnt learners are biased against giving probability to harmonically bounded candidates. *Proceedings of the Society for Computation in Linguistics* 5.24.

- PATER, JOE. 1999. Austronesian nasal substitution and other NC effects. *The prosody-morphology interface*, ed. by René Kager, Harry van der Hulst, and Wim Zonneveld, 310–343. Cambridge: Cambridge University Press.
- PATER, JOE. 2008. Gradient phonotactics in Harmonic Grammar and Optimality Theory. Ms., University of Massachusetts, Amherst.
- PATER, JOE. 2009. Weighted constraints in generative linguistics. *Cognitive Science* 33.999–1035.
- POTTS, CHRISTOPHER; JOE PATER; KAREN JESNEY; RAJESH BHATT; and MICHAEL BECKER. 2010. Harmonic Grammar with linear programming: From linear systems to linguistic typology. *Phonology* 27.77–117.
- PRINCE, ALAN, and PAUL SMOLENSKY. 1993/2004. *Optimality Theory: Constraint interaction in generative grammar*. Malden, MA: Blackwell. First published in Tech. Rep. 2, Rutgers University Center for Cognitive Science.
- PRINCE, ALAN, and BRUCE TESAR. 2004. Learning phonotactic distributions. *Fixing priorities: Constraints in phonological acquisition*, ed. by René Kager, Joe Pater, and Wim Zonneveld, 245–291. Cambridge: Cambridge University Press.
- PUNTE BALDOCEDA, BLAS. 1977. *Fonología del quechua tarmaño*. Lima: Universidad Nacional Mayor de San Marcos, Centro de Investigación de Lingüística Aplicada.
- R CORE TEAM. 2018. *R: A language and environment for statistical computing*. R Foundation for Statistical Computing, Vienna, Austria. Online: <https://www.R-project.org/>.
- REISS, CHARLES. 2018. Substance free phonology. *The Routledge handbook of phonological theory*, 1 edn., 425–452. Routledge.
- ROSE, SHARON, and RACHEL WALKER. 2004. A typology of consonant agreement by correspondence. *Language* 80.475–531.
- ROTHENBERG, MARTIN. 1968. *The breath-stream dynamics of simple-released-plosive production*. Basel: S. Karger.
- SAMUELS, BRIDGET. 2009. *The structure of phonological theory*. Harvard University dissertation.
- SCHWARZ, GIDEON. 1978. Estimating the dimension of a model. *Annals of Statistics* 6.461–464.
- SMITH, CAROLINE L. 1997. The devoicing of /z/ in American English: Effects of local and prosodic context. *Journal of Phonetics* 25.471 – 500. Online: <http://www.sciencedirect.com/science/article/pii/S009544709790053X>.
- SNOEREN, NATALIE; PIERRE HALLÉ; and JUAN SEGUI. 1992. A voice for the voiceless: Production and perception of assimilated stops in French. *Journal of Phonetics* 34.241268.
- STAMPE, DAVID. 1973. *A dissertation on natural phonology*. University of Chicago dissertation.
- STAUBS, ROBERT. 2011. *Harmonic Grammar in R (hgR)*. Software package. Online: <http://blogs.umass.edu/hgr/>.
- STERIADE, DONCA. 1997. Phonetics in phonology: The case of laryngeal neutralization. Ms., University of California, Los Angeles.
- TESAR, BRUCE. 2004. Using inconsistency detection to overcome structural ambiguity. *Linguistic Inquiry* 35.219–253.
- TESAR, BRUCE. 2006. Faithful contrastive features in learning. *Cognitive Science* 30.863–903.

- TESAR, BRUCE. 2014. *Output-driven phonology*. Cambridge: Cambridge University Press.
- TESAR, BRUCE, and PAUL SMOLENSKY. 2000. *Learnability in Optimality Theory*. Cambridge, MA: Massachusetts Institute of Technology Press.
- VENNEMANN, THEO. 1988. *Preference laws for syllable structure and the explanation of sound change: With special reference to German, Germanic, Italian, and Latin*. Berlin: Mouton de Gruyter.
- WESTBURY, JOHN R., and PATRICIA A. KEATING. 1986. On the naturalness of stop consonant voicing. *Journal of Linguistics* 22.145–166.
- WHITE, JAMES C. 2013. *Bias in phonological learning: Evidence from saltation*. UCLA dissertation.
- WILSON, COLIN. 2006. Learning phonology with substantive bias: An experimental and computational study of velar palatalization. *Cognitive Science* 30.945–982.
- WILSON, COLIN, and GILLIAN GALLAGHER. 2018. Accidental gaps and surface-based phonotactic learning: A case study of South Bolivian Quechua. *Linguistic Inquiry* 49.610–623. Online: https://doi.org/10.1162/ling_a.00285.
- WILSON, COLIN, and MARIEKE OBDEYN. 2009. Simplifying subsidiary theory: Statistical evidence from Arabic, Muna, Shona, and Wargamay. Ms., Johns Hopkins University. <https://linguistics.ucla.edu/people/hayes/219/papers/WilsonObdeyn2009.pdf> (accessed 22 June 2022).
- YU, ALAN C. L. 2004. Explaining final obstruent voicing in Lezgian: Phonetics and history. *Language* 80.73–97.
- YU, ALAN C. L. 2013. Typologizing phonetic precursors to sound change. *Language typology and historical contingency: In honor of Johanna Nichols*, ed. by Balthasar Bickel, Lenore A. Grenoble, David A. Peterson, and Alan Timberlake, 395–414. John Benjamins: Amsterdam/Philadelphia.
- ZHANG, JIE. 2010. Issues in the analysis of chinese tone. *Language and Linguistics Compass* 4.1137–1153. Online: <https://compass.onlinelibrary.wiley.com/doi/abs/10.1111/j.1749-818X.2010.00259.x>.
- ZOLL, CHERYL. 1998. Positional asymmetries and licensing. Ms., Massachusetts Institute of Technology.
- ZURAW, KIE. 2000. *Patterned exceptions in phonology*. University of California Los Angeles dissertation.
- ZURAW, KIE. 2010. A model of lexical variation and the grammar with application to Tagalog nasal substitution. *Natural Language & Linguistic Theory* 28.417–472. Online: <https://doi.org/10.1007/s11049-010-9095-z>.
- ZYMET, JESSE. 2018. *Lexical propensities in phonology: corpus and experimental evidence, grammar, and learning*. University of California Los Angeles dissertation.

A Supplementary Materials

A.1 Data collection and annotation

Each target stop’s syllable position was annotated right-to-left, with the first counted syllable therefore being the rightmost (thus backwards to English reading order) . Syllable position was furthermore annotated in two ways: firstly via syllabifying each word according to normal Tarma Quechua syllable structure, and secondly via syllabifying each word in order to maximize onsets, regardless of whether these are actually licensed in the language. The latter technique is used out of an abundance of caution, in order to simplify the process of analyzing target CC clusters. Tarma Quechua does not typically tolerate within-syllable CC clusters, although a stop followed

by a glide or liquid seems to be permitted, at least in loans. This can be seen in examples such as *kwarda* [Sp. *cuarto*] ‘room’ and *grat:is* [Sp. *gratis*] ‘thanks’. Because of examples such as these, consonant + liquid or consonant + glide clusters were kept when syllabifying according to the actual phonotactics of the language. Consonant + nasal clusters were not syllabified together, as evidence is marginal that these are permitted. The data also contained several glide + vowel or vowel + glide sequences; these sequences were treated as one vowel V.

Part of speech was also annotated in two ways. The first annotation simply recorded the tagging of each source text; however, each author’s tagging system did not exactly align. In order to provide consistency across part-of-speech tagging, we collapsed the tagging systems of both authors into a simple four-way system of nouns, verbs, dual noun/verbs, and “other” parts of speech (e.g. adjectives, adverbs, interjections, etc). Closed-class nouns (e.g. pronouns, numbers) were not tagged as nouns for the purpose of this analysis; instead, they were tagged as “other.”

Duplicate entries between the Adelaar and Baldoceda dictionaries were treated in several ways. If multiple identical-looking entries appeared together in the same dictionary entry, they were treated as unique and not consolidated together, as in the case of homophones. If a word was entered multiple times under the same entry in a single dictionary, all were considered a single entry (e.g. a single word with both verbal and adjectival senses, listed separately, would be marked as one entry). If two entries occurred in separate dictionaries, but one appeared to be a hypernym of the other, the two were consolidated into one entry (e.g. *parba* ‘food offering’, and *parba* ‘sweet corn’ [‘dulce de maiz’]). Lastly, entries for which one definition seemed like a plausible semantic extension of the other were still marked as separate entries, in the absence of explicit confirmation of their relationship.

Because alveolars never undergo voicing, they were omitted from the analysis; only labials and velars were kept. In addition, word-final stops and the first members of consonant clusters almost always surface as voiceless in native vocabulary, which is why they were also excluded from the analysis. For the purpose of statistical analysis, we only test consonants in prevocalic position. Morpheme boundary and word-boundary are collapsed into one category (lexical boundary, marked with #). To test significance of the observed trends in the lexicon, we include only native vocabulary in the analysis and exclude words marked as derived in order to avoid overrepresenting patterns in those lexical items that are frequently derived ($N = 1,118$). We also perform a significance test on vocabulary that includes loanwords and derived items to test whether observed patterns persist when the entire vocabulary is included.

A.2 Models

The initial full model was fit with all predictors and their interactions. The best fitting model was chosen with the step-wise backwards model selection technique: higher order interactions were removed step-wise from a full model. If the Akaike Information Criterion (AIC) determined an interaction or predictor does not improve fit significantly, they were removed until all predictors in the model significantly improved the fit. The best-fitting model includes PLACE, LEFT CONDITION, and DICTIONARY as predictors and two interactions: LEFTCOND:DICTIONARY and PLACE:DICTIONARY. The PART OF SPEECH predictor does not substantially improve model fit according to the AIC test.

The DICTIONARY predictor has two levels: Adelaar includes all data entries that are in Adelaar or in both dictionaries; Baldoceda includes those entries that appear in the Baldoceda dictionary only and are absent from Adelaar. This predictor thus indicates whether either data source deviates significantly from the other. Figure 7 in Section A.2 in the Supplementary Materials shows that the major trends in the Adelaar’s dictionary hold in those lexical items in Baldoceda that are absent

from Adelaar. This suggests that the voicing pattern is independently confirmed in two sources. However, because the predictor **DICTIONARY** is not linguistically justifiable (lexical items that are absent from Adelaar 1977, but present in Baldoceda have no linguistic function; their only function is to test whether the trends persist across dictionaries), we exclude this predictor from the final model.




Figure 7: Estimates of a logistic regression model with **PLACE**, **LEFT CONDITION**, and **DICTIONARY** as predictors and two interactions: **LEFTCOND:DICTIONARY** and **PLACE:DICTIONARY** (non-loan non-derived vocabulary only).

	Estimate	Std. Error	z value	Pr(> z)
(Intercept)=V	-0.1623	0.0852	-1.91	0.0567
#	-2.0487	0.1671	-12.26	0.0000
N	-1.7818	0.3061	-5.82	0.0000
R	1.2640	0.2147	5.89	0.0000
S	1.1647	0.2368	4.92	0.0000
T	2.3277	0.6120	3.80	0.0001
Place1	0.5264	0.0852	6.18	0.0000
#:Place1	0.2004	0.1671	1.20	0.2302
N:Place1	-0.8849	0.3061	-2.89	0.0038
R:Place1	0.2589	0.2147	1.21	0.2278
S:Place1	0.0582	0.2368	0.25	0.8059
T:Place1	-0.8460	0.6120	-1.38	0.1669

Table 12: Estimates of a logistic regression model with **PLACE** (sum-coded) and **LEFT CONDITION** and their interaction as predictors with loans and derived words included in the analysis.




Figure 8: Estimates of a logistic regression model with three predictors, LEFT CONDITION, PLACE, and LOAN with all interactions (derived words excluded).




Figure 9: Estimates of a logistic regression model with LEFT CONDITION and SYLLABLE STRUCTURE as predictors with their interaction.

A.3 Analysis of reverberation

Due to generally sub-optimal recording conditions, reverberation is present throughout the audio analyzed here. However, we can still reliably distinguish between energy caused by the reverberation of the preceding vowel and voicing. This is done by comparing the suspected reverberation in two environments: word-medially and utterance-finally.

Figure 2 shows four utterances beginning with a VT sequence. The initial vowel is either [a] or [u]. We can see from these examples that following the ending of the initial vowel, energy is still visible in the waveform and spectrogram. This energy seems to last for approximately 75-100 ms, before the signal reverts to reflecting the basic background noise of the room.

If this energy were reverberation, we would expect that it would still be present utterance-finally; that is, when there is no possible following voiced sequence. If this energy were voicing, we would not expect to see it utterance-finally for the same reason. Figures 10 show the vowels from Figure 2 utterance finally. It is evident from these utterance-final vowels that some vibration is still present after the end of the vowel, again for around 75-100 ms. This supports the interpretation that this energy in both cases is reverberation, and not voicing.




Figure 10: Waveform and spectrogram of utterance-final [a] and [u], demonstrating reverberation.

That the energy present in the initial consonant segment of TD clusters in Figure 2 is not voicing can also be shown through comparing each putative T segment with its putative D segment. The voiced segments each have a noticeably darker bar at the bottom, low-frequency end of their respective spectrograms, indicative of voicing. Clear periodicity can also be seen on their respective waveforms throughout the segment, again indicative of voicing. While the T segments do show some periodicity in the beginning of their waveforms, this is not unexpected, and likely represents the periodic reverberation of the previous vowel during the articulatory transition to a voiceless stop.

As a comparison, we also add a voiceless in Figure 11.

A.4 Fragment of training data for grammar learning simulations

Table 13 shows a fragment of the training data with a few of the 21 natural constraints. This is to illustrate what the training data look like. Two inputs are included, and the attested candidates are marked with 1.

UR	Output	Attested	*D	AGREE(vce)/[-son]	CODACOND(+vce)/[-son]	*NT	*VTV
/t̥simPajKu-/	[t̥simpajgu-]	1	-1	0	0	-1	0
/t̥simPajKu-/	[t̥simpajku-]	0	0	0	0	-1	-1
/t̥simPajKu-/	[t̥simbajgu-]	0	-2	0	0	0	0
/t̥simPajKu-/	[t̥simbajku-]	0	-1	0	0	0	-1
/qiPt̥u/	[qipt̥u]	1	0	0	0	0	0
/qiPt̥u/	[qibt̥u]	0	-1	-1	-1	0	0




Figure 11: A waveform and spectrogram of a voiceless [kt] cluster in a relatively silent recording.

Table 13: Fragment of the training data. Output candidates explore all possible ways of making underlying /P,K/ voiced or voiceless. 5 of the 21 natural constraints are shown here.

A.5 Constraint weights in selected models

The table below shows the weights obtained for the Full model, the Natural model, the faithfulness-free DISAGREE/*ND (D.A./*ND), *ND, and DISAGREE (D.A.) models, and the DISAGREE/Faithfulness (D.A./Faith) model.

Constraint	Full	Natural	D.A./*ND	*ND	D.A.	D.A./Faith
Ident						15.996
<i>Natural markedness constraints</i>						
*[+voice,-son]	3.177	0.164	2.446	0	2.451	1.601
*b	1.006	0.129	1.034	0	1.029	0.724
*g	2.171	0.974	1.847	0.816	1.842	1.399
*#[+voice,-son]	1.753	2.076	0	2.280	0	0
*#b	1.128	1.095	0	1.182	0	0.004
*#g	0.625	0.981	0	1.098	0	0.005
AGREE(voice)/[-son]	0.157	0	0.142	0	0.233	0.856
AGREE(voice)/[+labial,-son]	0.442	0	0.555	0	0.688	0.999
AGREE(voice)/[+dorsal,-son]	0.111	0	0.553	0	0.600	1.083
CODACOND([+voice])/[-son]	4.701	2.326	4.331	2.520	4.423	3.847
CODACOND([+voice])/[+labial,-son]	0.001	0	0.093	0.075	0	0.033
CODACOND([+voice])/[+dorsal,-son]	5.700	9.618	6.614	11.044	8.758	8.771
*VTV	2.792	0.300	2.394	0.032	2.394	1.685
*VpV	1.520	0.252	1.344	0.227	1.344	0.837
*VkV	1.417	0.101	1.161	0.046	1.162	1.066
*NT	1.478	0	1.278	0.282	0.711	0.813
*Np	0.237	0	0.325	0.475	0.228	0.431
*Nk	1.376	0	1.275	0.278	1.516	1.510
*RT	3.727	1.243	3.357	1.038	3.357	2.274
*Rp	2.075	0.857	1.930	0.770	1.931	1.296
*Rk	1.651	0.386	1.427	0.268	1.427	1.085
<i>Unnatural markedness constraints</i>						
*[-voice,-son]	0.484					
*p	0.175					
*k	0.309					
*#[-voice,-son]	1.882					
*#p	1.053					
*#k	0.857					
DISAGREE(voice)/[-son]	4.317		4.716		4.803	4.339
DISAGREE(voice)/[+labial,-son]	1.831		1.215		1.351	1.232
DISAGREE(voice)/[+dorsal,-son]	1.702		1.198		1.249	0.960
CODACOND([-voice])/[-son]	0.270					
CODACOND([-voice])/[+labial,-son]	0					
CODACOND([-voice])/[+dorsal,-son]	0.953					
*VDV	0.286					
*VbV	0.243					
*VgV	0.106					
*ND	0.186		0.209	1.685		
*Nb	0.548		0.454	1.559		
*Ng	0.181		0.117	0.126		
*RD	0.111					
*Rb	0.358					
*Rg	0.221					

Table 14: The best performing models and the weights assigned to them by the learner.

A.6 Results of stepwise backwards model selection

Applying a stepwise backwards model selection procedure (e.g., Hayes et al. 2012) to the training data with significance level $p < .1$ yields the following model with a training data log-likelihood of -610:

Constraint	Weight
<i>Natural markedness constraints</i>	
*[+voice,-son]	3.268
*g	1.206
CODACOND[+voice]/[+labial,-son]	4.393
CODACOND[+voice]/[+dorsal,-son]	15.273
*VpV	3.526
*VkV	3.736
*Nk	2.406
*RT	5.000
<i>Unnatural markedness constraints</i>	
DISAGREE(voice)/[+labial,-son]	5.021
DISAGREE(voice)/[+dorsal,-son]	5.400
CODACOND[-voice]/[+dorsal,-son]	4.757

Table 15: Selected constraints and their weights resulting from applying stepwise backwards model selection.