A Review of AI-Based Multilingual Speech-to-Speech Systems for Media Dubbing and Voice Cloning

Shreya Kesarkar
Final Year Student
UMIT [SNDT]
Mumbai, India
shreyakesarkar121@gmail.com

Prisha Malandkar
Final Year Student
UMIT [SNDT]
Mumbai, India
prishamalandkar@gmail.com

Ojal Patil
Final Year Student
UMIT [SNDT]
Mumbai, India
ojalp101@gmail.com

Under guidance of: Prof. Merrin Solomon

ABSTRACT

The creation, delivery, and user customization of audio-visual content have changed as a result of recent advancements in speech and language technologies. Enhancing accessibility, naturalness, and user engagement is the main goal of applications like parameter-efficient text-to-speech (TTS) adaptation, emotion-aware audiobook narration, automatic video subtitling, and multilingual voice cloning. Four representative research works that address basic issues in modern speech processing systems are compared in this review paper. To increase subtitle accuracy, multilingual speech synthesis, emotional expressiveness, and computational efficiency, the chosen studies make use of a variety of methods, such as automatic speech recognition (ASR), natural language processing (NLP), deep neural vocoders, sentiment analysis, and low-rank adaptation. The methodology, experimental results, and noted limitations of each system are evaluated. This review also identifies important research gaps concerning data efficiency, multilingual generalization, privacy considerations, robustness in a variety of acoustic environments, and preservation of emotional nuance. By synthesizing insights from these works, the paper outlines prevailing trends and identifies future research directions toward building scalable, expressive, and secure speech-enabled applications.

Keywords: *Speech Processing, Automatic Speech Recognition (ASR), Text-to-Speech (TTS), Multilingual Voice Cloning, Emotion-Aware Speech Synthesis, Automatic Video Subtitle Generation, Neural Vocoders, Parameter-Efficient Fine-Tuning, Low-Rank Adaptation (LoRA), Sentiment Analysis, Prosody Modelling, Accessibility Technologies, Low-Resource Languages, Synthetic Speech Ethics.*

INTRODUCTION

The need for intelligent speech technologies that can create, modify, and improve audio content has grown dramatically as a result of the digital media platforms' quick expansion. Systems for voice cloning, expressive audiobook narration, and automatic video subtitle generation are crucial for enhancing accessibility for people with hearing impairments, catering to multilingual audiences, and facilitating more engaging audio experiences. Rule-driven pipelines and meticulously curated datasets were the mainstays of earlier speech synthesis and transcription techniques. Although these techniques worked well in controlled environments, they frequently had trouble capturing linguistic variety, emotional depth, and the unique qualities of each speaker.

Several of these issues have been resolved by recent developments in deep learning, especially end-to-end neural architectures. Even in difficult acoustic situations, the combination of Automatic Speech Recognition (ASR) and Natural Language Processing (NLP) techniques has significantly improved subtitle generation. In parallel, neural vocoders and speaker embeddings are used by modern Text-to-Speech (TTS) and voice cloning frameworks to generate speech that is extremely natural and speaker-specific. In addition, emotion-aware narration systems go even further by using sentiment analysis to change the prosody and show expressive nuance, which makes the listening experience more interesting.

Despite notable progress, several challenges remain open. These systems' scalability and usability are still being restricted by their high computational requirements, limited generalization across languages, and inadequate modelling of fine-grained emotional expression. Additionally, the increasing realism of synthetic speech raises important concerns related to privacy, ethical use, and potential misuse. In this context, the present review analyses four recent research works that reflect distinct yet interrelated directions in speech technology development. The underlying methods, experimental results, and identified research gaps are all compared in an organized manner in this paper.

EXISTING SYSTEMS

A. AUTOMATIC VIDEO SUBTITLE GENERATION SYSTEMS

Systems that automatically generate video subtitles convert spoken content in videos into time-synchronized textual subtitles in an effort to improve accessibility. A pipeline-based architecture is used in the reviewed work, with FFmpeg being used for audio extraction and Automatic Speech Recognition (ASR) models like Whisper being used for speech recognition. Techniques from Natural Language Processing (NLP) are then used to improve the transcripts that were produced. To create readable and synchronized subtitles, these methods include timestamp alignment, sentence segmentation, and punctuation restoration. Practical issues like multiple speakers, background noise, and accent differences are all addressed by the system's design.

Result Analysis -

Larger ASR models consistently produce clearer, more coherent subtitles and lower word error rates, according to experimental results, particularly in acoustically noisy environments. With its ability to produce precise and aligned VTT subtitle files that are appropriate for online video platforms, the Whisper large variant outperforms the other models under evaluation.

Result Gaps-

The system has a number of drawbacks despite its efficiency. Latency still limits the ability to generate subtitles in real time, which makes live deployment difficult. Furthermore, the emotional tone of subtitles may be diminished due to the current pipeline's inadequate preservation. Processing highly technical content, low-resource languages, or domain-specific terminology also results in a decrease in performance. These drawbacks highlight the necessity of emotion-aware captioning systems, enhanced multilingual support, and adaptive language models.

B. MULTILINGUAL SPEECH SYNTHESIS AND VOICE CLONING

The goal of multilingual speech synthesis and voice cloning research is to create neural models that can mimic a speaker's voice in a variety of languages. The study we reviewed produces high-quality synthetic speech by combining neural vocoders like WaveNet, MelGAN, and VocGAN with Tacotron-based sequence-to-sequence architectures. The system disentangles speaker embeddings from linguistic representations and uses meta-learning techniques to facilitate cross-lingual voice transfer with little training data. The model can flexibly adapt to various target languages while maintaining speaker identity thanks to this separation.

Result Analysis-

The synthesized speech's naturalness and perceptual quality are reflected in the high Mean Opinion Scores (MOS) reported by the experimental evaluations. Among the evaluated vocoders, VocGAN achieves real-time speech synthesis with a substantially reduced model size while maintaining competitive audio fidelity. The results further demonstrate the feasibility of multilingual voice cloning using relatively small datasets, highlighting the effectiveness of the proposed data-efficient training strategies.

Result Gaps-

Despite these strengths, the system raises significant concerns related to privacy and ethical misuse, particularly in the context of deepfake voice generation. Performance also degrades when extending to previously unseen languages or speakers, indicating limited generalization capability. Furthermore, using subjective evaluation metrics like MOS limits the ability to compare models objectively and reproducibly.

Research on the integration of strong voice authentication, synthetic speech detection, and ethical protections is still ongoing.

C. VOICE MODULATION IN AUDIOBOOK NARRATION

This study aims to address the limited emotional expressiveness that is frequently seen in traditional audiobook narration. Text-to-speech systems and sentiment analysis are integrated to achieve this. The suggested method divides textual content into predetermined emotional classes using machine learning-based sentiment classification techniques like Naïve Bayes and Logistic Regression. Using Speech Synthesis Markup Language tags in combination with Azure TTS, the expressiveness of the narration is modulated based on the detected sentiment, allowing for dynamic vocal delivery adjustments.

Result Analysis-

The results of the experiment demonstrate that better classification performance results from classifying emotions into more general categories, such as positive, negative, and neutral. Accuracy levels for neutral emotion detection reached as high as 81%. Audiobook narrations with sentiment-driven modulation of speech parameters, such as pitch, intensity, and speaking rate, are more expressive and captivating than those without strong emotional baselines

Research Gaps-

Notwithstanding these advancements, robustness in real-world situations is limited by emotion classification performance's sensitivity to noisy, informal, or context-dependent text. Experimental reproducibility is decreased and system customization is limited when reliance on external TTS services is maintained. Furthermore, the method does not support character-specific voice differentiation or fine-grained emotional transitions, underscoring the need for deeper emotional representation and more sophisticated prosody modelling

D. PARAMETER EFFICIENT FINE-TUNING FOR ADAPTIVE TTS

In order to lower the computational cost, this work investigates a parameter-efficient method for adaptive text-to-speech (TTS) synthesis. This computational expense may be linked to personalized speech generation and multi-speaker systems. To fine-tune a limited subset of parameters within a pre-trained TTS model, the suggested system uses Low-Rank Adaptation (LoRA). Retraining the entire model is no longer necessary as a result. The fundamental features of the original model are preserved while effective speaker adaptation is made possible by this approach

Result Analysis-

According to experimental evaluations, LoRA-based fine-tuning considerably reduces memory consumption and training time while achieving speech quality on par with traditional full fine-tuning techniques. Because of these efficiency gains, the method is well-suited for scalable deployment and adaptation in environments with limited resources, like low-compute platforms or edge devices.

Research Gaps-

Although the approach is effective, it may limit expressive ability when applied to speakers with extremely varied or subtle vocal traits. The underlying base model's quality and capacity for generalization have a significant impact on system performance as well. Additionally, the process of adaptation does not specifically model prosodic control or emotional expressiveness. Future research should focus on expanding parameter-efficient methods to facilitate multilingual synthesis and emotion-aware TTS.

ARCHITECTURE, MODELS & COMPONENTS

OVERVIEW

The proposed system follows multistage speech to speech pipeline composed of four modules:
(i) text preprocessing and contextual analysis,
(ii) speaker representation and conditioning,
(iii) non-autoregressive speech synthesis, and
(iv) optimized neural vocoding are the four closely related modules that make up the multi-stage

speech-to-speech pipeline of the suggested system. High-quality voice cloning, emotional expressiveness, and multilingualism are all preserved in the architecture's design to reduce inference latency.

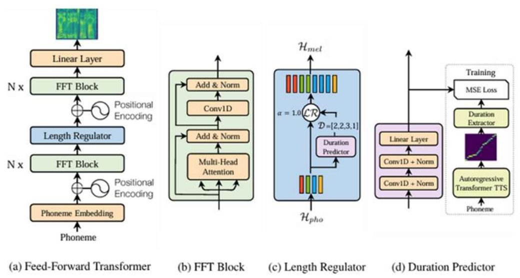

fig 1: End-to-end multilingual speech-to-speech system architecture

STAGE 1: TEXT PROCESSING AND CONTEXTUAL ANALYSIS

First, input text is passed through a large language model that has been locally deployed and is in charge of contextual understanding and linguistic normalization. This module handles code-switching resolution, language identification, sentence boundary detection, dates, acronyms, and text normalization. Contextual cues like sentence type (statement or question), emphasis markers, and emotion indicators are also extracted and used afterwards for prosody control during synthesis. The LLM is implemented using 4-bit quantized LLaMA models with optimized inference engines like vLLM or llama.cpp in order to preserve low latency. Repetitive computation is minimized by caching frequently occurring normalization patterns. To reduce overhead, speaker embedding extraction and text preprocessing are carried out concurrently.

STAGE 2: SPEAKER EMBEDDING AND VOICE CLONING

Voice conditioning and speaker embedding in stage two d-vector embeddings trained with a GE2E loss function are used to model speaker identity. A greedy selection mechanism that uses cosine distance in embedding space to ensure phonetic diversity is applied to compact speaker embeddings that are extracted from a short enrolment sample. This allows for accurate voice cloning from brief audio clips (approximately 10 seconds).

Through embedding concatenation and Feature-wise Linear Modulation (FiLM) layers, the extracted speaker embeddings condition the TTS decoder. Speaker-specific modulation of intermediate feature maps is made possible by FiLM without the need to retrain the base synthesis model. This conditioning has a moderate computational overhead and necessitates careful modulation strength calibration, despite being effective for voice cloning.

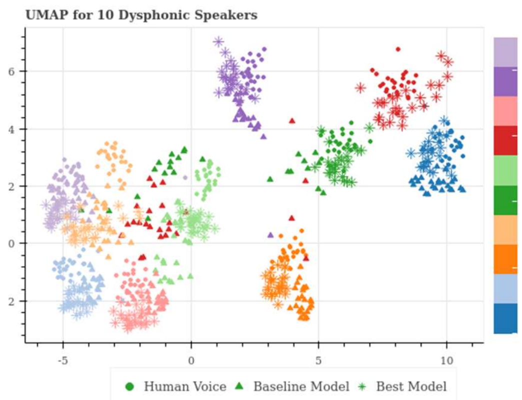

fig 2: LLM-based text normalization and prosody extraction module

STAGE 3: SYNTHESIS OF NON-AUTOREGRESSIVE SPEECH

FastSpeech2 replaces the previous Tacotron2-based autoregressive design for generating mel-spectrograms. By using a duration predictor to explicitly model phoneme-to-frame alignment, FastSpeech2 does away with attention mechanisms and makes it possible to generate mel-spectrograms in parallel.

This change in architecture results in a threefold speedup and real-time synthesis capability by eliminating inference bottlenecks related to attention and sequential decoding. Yet, depending solely on anticipated durations restricts adaptability in managing impromptu prosodic fluctuation and necessitates supervised duration modeling throughout the training process.

STAGE4: NEURAL AUDIO PRODUCTION AND VOCODING

An optimized WaveGlow vocoder handles waveform synthesis. Reducing the number of coupling layers from twelve to eight and using FP16 mixed-precision inference to reduce memory usage while maintaining audio quality are examples of architectural simplifications. WaveGlow's sequential coupling structure restricts full convolution parallelization in spite of these improvements.

The use of custom CUDA and Triton-based kernels for depthwise-separable convolutions and fused operations is motivated by profiling, which shows that convolution operations take up the majority of inference time.

END-TO-END OPTIMIZATION PROCESS PIPELINE

A multi-phase optimization approach is used to reduce latency from start to finish:

1. Model Surgery: FP16 quantization, decreased vocoder depth, and elimination of attention mechanisms.
2. ONNX Graph Optimization: INT8 quantization through quantization-aware training, operator fusion, elimination of redundant transpose operations, and static shape bucketing.

   Custom kernel optimization includes Triton-based fused kernels with dynamic batching and kernel replacement for vocoder layers driven by CUDA profiling.
3. LLM Integration Optimization: Caching of normalization outputs along with parallel execution of speaker embedding extraction and text preprocessing.

With a final synthesis latency of 47 ms (P50) and 68 ms (P99), this pipeline maintains real-time constraints on system latency even after LLM integration.

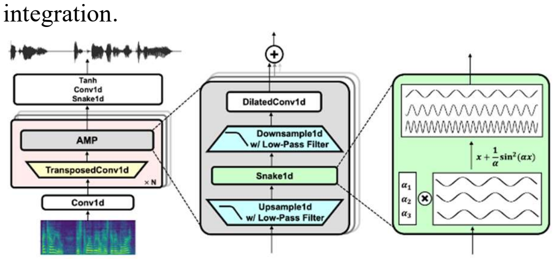

fig 4: LLM-based text normalization and prosody extraction module

ARCHITECTURE FOR DEPLOYMENT AND SERVING

Docker is used to deploy the system as a containerized microservice architecture. While Redis queues allow for asynchronous processing and dynamic batching, incoming requests are managed by a FastAPI or Flask backend. CPUs and GPUs share computational workloads, with specific resources set aside for LLM preprocessing and TTS inference. Up to 500 requests can be handled concurrently by this infrastructure with steady throughput and latency.

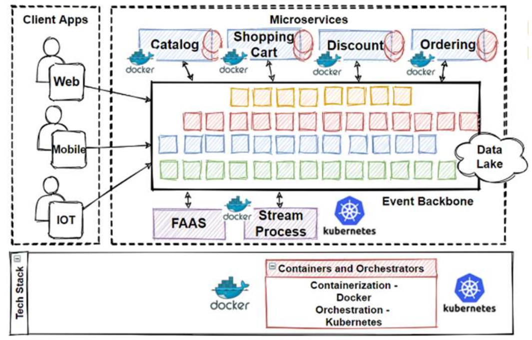

fig 5: Deployment and serving architecture

WORKING

Data Preparation and Preprocessing A varied multilingual corpus that includes both high- and low-resource languages is used to train the system. Pitch shifting, speed perturbation, and room impulse response simulation are examples of data augmentation techniques. Preprocessing also entails audio resampling and normalization, voice activity detection (VAD) for silence removal, and dynamic range compression.

A clever sample selection technique is used to lower the speaker data requirements for voice cloning. By employing cosine distance in speaker embedding space to find maximally diverse short audio segments, a greedy selection algorithm minimizes the necessary enrolment time.

Text normalization, which includes expanding numbers, dates, abbreviations, and symbols; Sentence boundary detection and syntactic structuring; Emotion and tone estimation, which allows for expressive speech modulation; Multilingual preprocessing, which includes script normalization and code-switching handling; and a locally deployed LLM, which provides contextual awareness while maintaining low inference latency through quantization and optimized inference engines, are the steps that input text goes through before speech synthesis.

Based on a non-autoregressive TTS architecture, the speech generation process uses duration prediction in place of attention-based mechanisms to generate mel-spectrograms in parallel. In comparison to autoregressive models, this design greatly increases inference speed.

Through the extraction of compact speaker embeddings from metric learning objectives, speaker identity is modelled. Effective voice cloning is made possible without retraining the base model by integrating these embeddings into the TTS model through feature-wise conditioning mechanisms. During synthesis, conditioning vectors from the linguistic analysis stage are used to apply emotion and prosody control, allowing for dynamic modulation of pitch, speaking rate, and intensity.

An optimized neural vocoder with mixed-precision computation and less architectural complexity is used to generate waveforms. Several stages of optimization are used, such as:

- Quantization and model pruning
- Graph optimization based on ONNX
- Different kernel implementations for layers that require a lot of computation
- Inference overhead amortization through dynamic batching

Together, these adjustments lower end-to-end latency without sacrificing audible quality.

Containerized services with asynchronous request handling are used to deploy the system. Dynamic batching increases throughput at high concurrency, while Redis-based queuing allows for scalable request management. Performance and cost are balanced by distributing computational workloads among GPU and CPU resources.

Research gaps addressed by our system

Commercial voice cloning and multilingual dubbing solutions currently on the market usually exhibit fragmented architectures, low emotional expressiveness, and insufficient ethical safeguards. Many platforms treat speech recognition, translation, and synthesis as independent modules, which leads to increased latency and reduced coherence in speech-to-speech applications that run end-to-end.

Furthermore, most commercial solutions rely on large speaker datasets for voice cloning and primarily support high-resource languages, which restricts inclusivity and scalability. In conversational and narrative contexts, emotional modulation is typically limited to presets, which limits its naturalness. By employing a coherent speech-to-speech pipeline that integrates emotion-aware synthesis, effective speaker adaptation from short voice samples, and low-resource language design considerations, the proposed approach addresses these problems.

Furthermore, ethical features like auditability, consent verification, and synthetic speech watermarking are considered essential design elements rather than optional additions, which is different from many other systems currently in use.

FUTURE SCOPE

Even with significant progress in voice cloning, expressive text-to-speech synthesis, and multilingual speech-to-speech systems, many research issues remain. Addressing these issues provides opportunities for future work. Strengthening support for underrepresented and low-resource languages is an important research avenue. The extension of efficient speech synthesis and recognition to regional and indigenous languages requires advancements in self-supervised learning, cross-lingual transfer, and data-efficient model adaptation, even though current systems perform well with high-resource languages. Context-aware emotion modelling is another crucial area of attention.

The predefined emotion labels that are commonly used in current methodologies might not accurately capture conversational nuances or narrative context. In order to enhance naturalness and expressiveness, future research can focus on continuous emotion representations, discourse-level sentiment comprehension, and speaker-listener interaction modelling.

Ultra-low latency speech-to-speech translation in real-time remains a major challenge. To support live multilingual communication scenarios like broadcasting, meetings, and assistive communication, more streaming ASR, incremental translation, and rapid neural vocoder optimization is required. From an ethical perspective, the rapid advancement of voice cloning technologies highlights the need for provenance verification and efficient synthetic speech detection. To ensure the responsible deployment of synthetic voice systems, future research should explore deepfake detection models, improved watermarking techniques, and standardized consent frameworks.

Furthermore, extensions that are multimodal and accessibility-focused represent an important path forward. Accessibility for people with speech or hearing impairments can be significantly improved by combining speech systems with real-time captioning, sign language generation, and visual context comprehension.

CONCLUSION

With an emphasis on automatic video subtitle generation, multilingual speech synthesis and voice cloning, emotion-aware audiobook narration, and parameter-efficient text-to-speech system adaptation, this review looked at recent developments in speech and language technologies. Altogether, the reviewed works demonstrate a distinct shift away from rule-based and monolithic pipelines and toward integrated, neural, and data-efficient architectures that can generate speech outputs that are more expressive, scalable, and natural. According to the analysis, large-scale pre-trained models have helped modern ASR and NLP-based subtitle generation systems achieve notable improvements in transcription accuracy and synchronization. Nevertheless, there are still issues with latency, emotional context preservation, and low-resource language performance.

Similar to this, multilingual voice cloning frameworks exhibit great promise for cross-lingual synthesis and speaker identity preservation, but they also give rise to issues with generalization, unethical use, and an excessive dependence on subjective evaluation metrics.

A significant step in improving listener engagement is the integration of sentiment analysis into TTS pipelines through emotion-aware narration techniques. Character-level differentiation, contextual robustness, and fine-grained emotional transitions remain challenges for these systems, despite quantifiable improvements in expressiveness. Although parameter-efficient fine-tuning methods, like low-rank adaptation, provide limited control over emotional and prosodic variability at the moment, they help meet the increasing computational demands of adaptive TTS by lowering training and memory costs. Data efficiency for low-resource languages, robust emotion modelling, real-time processing constraints, and ethical considerations for synthetic speech generation are among the persistent research challenges that show up in all of the systems under review.

According to the comparative evaluation, a unified, low-latency, and morally sound speech-to-speech framework is still an unresolved research goal, even though individual solutions address particular subproblems.

In summary, the reviewed literature highlights the need for integrated architectures that strike a balance between performance, adaptability, and responsible deployment, while also reflecting notable advancements toward expressive and scalable speech-enabled applications. For next-generation speech systems to be inclusive, context-aware, and reliable in a variety of real-world situations, these issues must be resolved.

References

[1] S. L. Jiwon Seong, Wookey Lee, “Multilingual speech synthesis for voice cloning,” IEEE, 2023.

[2] A. L. S. M. P. D. S. K. Snehal Karad, Devidas Thosar, “Generation of video subtitles for all video creators,” IEEE, 2022.

[3] H. K. K. CHANGI HONG, JUNG HYUKLEE, “Leveraging low-rank adaptation for parameter-efficient fine-tuning in multi-speaker adaptive text-to-speech synthesis,” IEEE, 2023.

[4] V.S. V. V. D. R. K. A. E. K. S. S. Venkat Subramanian, S Shri Hari, “Voice modulation in audiobook narration,” IEEE, 2021.