

Persona Framing in Human–LLM Interaction

Edervaldo José de Souza Melo 

Independent Research, Campo Grande/MS, Brazil
edersouzamelo@gmail.com, edervaldo.melo@sou.ufmt.br

March 2026

Abstract

Large Language Models (LLMs) exhibit a high degree of stylistic adaptability, yet they frequently produce unstable responses in terms of discursive register, enunciative position, and pragmatic coherence across extended interactions. A common practice among users is *persona prompting*, in which the model is instructed to assume a specific discursive role. Although widely used in practice, this technique has been treated primarily as a prompt engineering strategy, with limited analysis from the perspective of linguistics and discourse pragmatics.

This paper proposes interpreting *persona prompting* as a mechanism of discursive framing (*persona framing*) in human–LLM interaction. It argues that persona instructions function as pragmatic devices that establish discursive registers, delimit enunciative positions, and constrain the stylistic space of generated responses. At the same time, the concept of *persona framing* is proposed not only as a linguistic interpretation of an existing technique, but also as a useful conceptual lens for the design of structured interactions with language models.

As an exploratory study, we analyze structured interactions with a system of discursive personas called *Nemosine*, in which different linguistic roles have defined discursive vocations and are maintained throughout dialogue. Based on this analysis, we propose a preliminary typology of framing levels in human–LLM interaction: (1) *ephemeral persona framing*, in which the persona appears only in a single isolated turn; (2) *persistent persona framing*, in which the persona remains active across multiple turns; (3) *environment framing*, in which interaction takes place within a symbolic space of interpretation that allows the invocation of different personas; and (4) *system-level persona framing*, in which the architecture of the interaction system is organized around multiple personas with specific discursive functions.

The qualitative analysis suggests that persona-based framing contributes to greater consistency of register, greater stylistic stability, and greater pragmatic predictability in the model’s responses. The results indicate that *persona framing* may be understood as an emerging linguistic interface for interacting with LLMs, connecting prompt engineering practices with classical theories of discursive framing, polyphony, and pragmatic organization of discourse.

Keywords: Persona Framing; Persona Prompting; Large Language Models; Human–AI Interaction; Discourse Framing; Polyphony; Prompt Engineering; Discursive Roles

Contents

1	Introduction	2
2	Persona Prompting in Current LLM Practice	5
3	Discursive Framing and Enunciative Position	6
4	Persona Framing	6
5	Levels of Framing in Human–LLM Interaction	8
5.1	Ephemeral Persona Framing	8
5.2	Persistent Persona Framing	8
5.3	Environment Framing	9
5.4	System-Level Persona Framing	9
6	Case Study: The Nemoline System	9
6.1	Research Materials and Program	10
7	Interaction Analysis	10
8	Implications for Human–AI Interaction	11
8.1	Future Research Directions	12
9	Conclusion	12

1 Introduction

Large Language Models (LLMs) have, in recent years, become one of the dominant paradigms for interaction between humans and computational systems. Systems based on LLMs are

capable of producing coherent text, adapting discursive registers, and responding to a wide range of linguistic tasks. This stylistic flexibility, however, also presents a recurring side effect: discursive instability across extended interactions.

In many usage scenarios, the same model may rapidly alternate between different registers, enunciative positions, or argumentative styles, even when the user attempts to maintain a consistent discursive line. This variability arises partly from the probabilistic nature of these systems, whose textual generation depends heavily on the immediate prompt context and on statistical patterns learned during training.

To address this instability, users and researchers have developed various textual instruction strategies commonly referred to as *prompt engineering*. Among these strategies, one of the most widely used consists of explicitly instructing the model to assume a particular discursive role, such as “act as a teacher”, “respond as a lawyer”, or “explain as a scientist”. This practice is commonly referred to in the technical literature as *persona prompting* or *role prompting*. Recent work in prompt engineering and interaction with LLMs has explored different forms of textual instruction to improve performance and control model behavior (Reynolds and McDonell, 2021; Wei et al., 2022).

Despite being widely used in practice, most discussions of persona prompting have been conducted within the domains of prompt engineering and human–computer interaction, focusing primarily on task performance or response quality. Relatively little attention has been devoted to the phenomenon from a broader linguistic perspective, particularly with regard to the pragmatic and discursive effects produced by this type of instruction.

This article starts from the hypothesis that the practice of persona prompting can be interpreted as a specific case of *discursive framing* in human–LLM interaction. We propose the term *persona framing* to describe the pragmatic effect by which a persona instruction establishes an interpretive frame that organizes discursive register, enunciative position, and communicative expectations within the interaction.

By introducing this concept, we aim to contribute in two complementary directions. First, we offer a linguistic interpretation of a practice already widespread in interactions with LLMs. Second, we suggest that the concept of persona framing may serve as a conceptual tool for the design of more structured discursive interfaces between humans and artificial language systems.

To explore these ideas, we also present a case study based on the *Nemosine system*, an interaction environment structured around multiple discursive personas. The system was conceived as an experiment in the polyphonic organization of interaction with a language model, in which different discursive voices possess distinct functions and analytical vocations within the dialogue.

Based on this exploratory study, we propose a preliminary typology of framing levels in human–LLM interaction, distinguishing different ways in which persona instructions and interaction contexts may structure the discursive space of a dialogue with language models.

Contributions of this Work

This article makes three main contributions:

1. Introduction of the concept of *persona framing*, proposed as a linguistic interpretation of persona prompting in LLM interaction.
2. Proposal of a typology of discursive framing levels in human–LLM interaction, including *ephemeral persona framing*, *persistent persona framing*, *environment framing*, and *system-level persona framing*.
3. Presentation of the *Nemosine system* as an exploratory case study illustrating how multiple layers of framing may be combined in an architecture of interaction based on discursive personas.

Conceptual Figure of the Proposed Model




Figure 1: Conceptual relationship between prompting, persona framing, framing levels, and structured interaction architectures.

Prompting (operational instruction) → Persona Framing (discursive effect) → Levels of Framing (ephemeral, persistent, environment, system-level) → Structured interaction architectures (e.g., Nemosine).

2 Persona Prompting in Current LLM Practice

The use of discursive role instructions has become a common practice in everyday interaction with language models. Users frequently rely on commands such as “act as”, “explain as”, or “respond in the role of”, seeking to guide the model toward a particular style of response.

This technique, known as *persona prompting* or *role prompting*, essentially consists in adding to the prompt an instruction that defines a discursive identity for the model. By assuming that identity, the system is expected to adopt not only a specific vocabulary, but also a broader set of pragmatic expectations associated with the requested role.

For example, when a model is instructed to respond “as a teacher”, it tends to produce more explanatory and didactic answers. When asked to act “as a lawyer”, it may adopt a more argumentative or normative tone. This behavior suggests that persona instructions affect not only the content of the response, but also the discursive frame within which the response is produced.

In the technical literature on LLMs, persona prompting is usually discussed primarily as a strategy for improving performance in specific tasks. Studies on *prompt engineering* frequently examine how different formulations of instructions can influence response quality, factual accuracy, or contextual appropriateness (Reynolds and McDonell, 2021; Wei et al., 2022).

However, this approach tends to treat the phenomenon predominantly as a problem of prompt optimization. The linguistic dimension of the process — that is, the way persona instructions reorganize the discursive space of the interaction — remains relatively underexplored.

In this sense, although the practice of persona prompting is widely recognized, its interpretation as a mechanism of discursive framing still lacks a clearer conceptual formulation. This article proposes that the notion of *persona framing* can partially fill this gap by offering an analytical lens that connects the phenomenon to traditions in pragmatics and discourse analysis.

3 Discursive Framing and Enunciative Position

The idea that human communication takes place within specific interpretive frameworks has a long tradition in linguistics and the social sciences. One of the most influential concepts in this field is that of the *frame*, developed by Erving Goffman to describe the interpretive schemas that guide the understanding of social situations (Goffman, 1974).

According to this perspective, participants in an interaction do not interpret utterances in isolation, but rather within contextual structures that define “what is going on” in a given communicative situation. These frames influence expectations regarding the appropriate type of discourse, the roles assumed by participants, and the possible modes of interpretation.

Another concept relevant to understanding the organization of voices in discourse is *polyphony*, developed in the work of Mikhail Bakhtin. For Bakhtin, discourse is not monological but often composed of multiple voices and discursive positions that coexist and interact within the same textual or communicative space (Bakhtin, 1981).

In addition, theories of discursive register emphasize that different social and communicative contexts demand distinct linguistic styles. The register used in a classroom lecture, for instance, differs from that employed in an informal conversation or in a legal opinion. These differences involve not only vocabulary, but also argumentative structures, explanatory modes, and forms of speaker positioning.

When these concepts are applied to the context of interaction with LLMs, it becomes possible to reinterpret the practice of persona prompting as a mechanism that explicitly introduces a discursive frame into the dialogue. When a user requests that the model respond “as a teacher” or “as an expert”, they are not merely asking for a different answer, but also establishing a set of expectations regarding the type of discourse that will be produced.

In this sense, what we call *persona framing* can be understood as the process by which a persona instruction establishes a pragmatic frame that guides the development of the interaction. This framing operates by delimiting enunciative positions, discursive registers, and forms of discourse organization.

From this perspective, the phenomenon ceases to be viewed merely as a technique of textual instruction and instead becomes part of a broader dynamic of role construction and discursive framing within human-machine interaction.

4 Persona Framing

Based on the considerations presented in the previous sections, we propose the concept of *persona framing* to describe the discursive effect produced when a persona instruction

establishes an interpretive frame within the interaction with a language model.

While *persona prompting* refers to the operational action performed by the user — that is, the textual instruction requesting the model to assume a particular role — the term *persona framing* describes the pragmatic and discursive state that emerges from that instruction. In other words, prompting constitutes the technical intervention, whereas framing corresponds to the reorganization of the discursive space that subsequently guides the interaction.

From this perspective, persona instructions do not function merely as superficial modifiers of style. Instead, they introduce pragmatic expectations regarding:

- the appropriate discursive register
- the enunciative position assumed by the model
- the argumentative structure of responses
- the type of discursive relationship established with the user

For example, when a model is instructed to respond “as a teacher”, the interaction becomes interpreted within a didactic frame. In such a context, the discourse is expected to include progressive explanations, pedagogical examples, and a logical organization oriented toward the transmission of knowledge. Similarly, when the model is instructed to respond “as a lawyer”, the discursive frame tends to privilege normative argumentation, rule interpretation, and formal justification.

In this sense, the concept of *persona framing* allows these transformations to be understood not merely as technical effects of prompting, but as processes of *pragmatic configuration of dialogue*. The persona functions as a framing device that delimits the interpretive space within which subsequent utterances are produced and interpreted.

Beyond offering a linguistic interpretation of the practice of persona prompting, this perspective also suggests that persona-based framing may be deliberately explored as a design principle for interaction systems with LLMs. By structuring dialogue around explicit discursive roles, it becomes possible to organize the textual behavior of the model in a more predictable way.

To avoid conceptual ambiguity, it is useful to clarify the central distinction proposed in this work. *Persona prompting* refers to the operational instruction provided by the user to the model — typically a textual formulation assigning the system a discursive role or identity. *Persona framing*, by contrast, describes the discursive configuration that emerges from that instruction. In other words, prompting acts as a mechanism of intervention, while framing describes the linguistic and pragmatic effect produced by that intervention within

the interaction. This distinction allows the phenomenon to be understood not merely as a technique of prompt engineering, but as a reorganization of the discursive space in which the dialogue unfolds.

5 Levels of Framing in Human–LLM Interaction

Based on the observations presented in the previous sections, we propose a preliminary typology of *discursive framing levels* in interaction with language models. This typology is not intended to be exhaustive, but rather to offer an initial framework for understanding different ways in which the discursive space of dialogue may be structured in interactions with LLMs.

5.1 Ephemeral Persona Framing

The simplest level corresponds to *ephemeral persona framing*, in which the persona instruction is used only for a single isolated turn of interaction.

In this case, the discursive frame introduced by the persona has a local and temporary character. The requested discursive identity influences the model’s immediate response but does not necessarily persist across subsequent turns of the dialogue.

This type of use is common in punctual interactions, in which the user requests specific explanations or reformulations. For example, a user may ask the model to “explain a concept as a teacher” in a single response without intending to maintain this frame throughout the entire conversation.

5.2 Persistent Persona Framing

A second level corresponds to *persistent persona framing*, in which the persona remains active across multiple turns of the interaction.

In this case, the discursive frame ceases to be merely local and instead organizes the continuity of the dialogue. The persona assumes a relatively stable enunciative position, influencing not only individual responses but also the overall coherence of the discourse.

The persistence of the persona tends to produce greater consistency of register, greater pragmatic predictability, and greater argumentative continuity throughout the interaction.

5.3 Environment Framing

A third level involves what we call *environment framing*. Unlike the previous cases, this form of framing does not directly define a specific persona but instead establishes a *symbolic space of interpretation* within which the interaction takes place.

In this context, the environment functions as a “discursive setting” that organizes how different personas may be invoked and interpreted. The frame is therefore not associated with a single role, but with the symbolic contextualization of the interaction itself.

This type of framing may be understood as a parallel layer of discursive organization: while personas define specific enunciative positions, the environment defines the interpretive context within which these voices operate.

Furthermore, the environment does not need to remain static. It may evolve over time, incorporating new symbolic references, implicit interaction rules, or new forms of dialogical organization. In this sense, environment framing also allows the plural participation of personas, including the possibility of simultaneous invocation within the same interpretive space.

5.4 System-Level Persona Framing

The most comprehensive level corresponds to *system-level persona framing*, in which the interaction architecture itself is organized around multiple personas with distinct discursive functions.

In this configuration, personas are not merely occasional roles introduced by the user but structural elements of the dialogue system. Each persona may possess a specific discursive vocation, contributing different perspectives or analytical styles within the interaction.

This type of organization creates a polyphonic environment in which multiple discursive voices coexist and collaborate in the construction of dialogue.

6 Case Study: The Nemosine System

To explore empirically the ideas presented in the previous sections, we analyze the *Nemosine system* as a case study of discursive organization based on multiple personas.

The Nemosine system was introduced as an experimental architecture for structured interaction with language models (Melo, 2025) in which different discursive personas are explicitly defined and maintained throughout the dialogue. Each persona represents a relatively stable discursive position associated with a specific vocation of analysis, interpretation, or textual production.

Within this system, interaction does not occur solely between the user and the model, but within a structured discursive architecture. The user may invoke different personas to address the same problem, allowing distinct perspectives to coexist within the same interpretive space.

Beyond the presence of multiple personas, the system also incorporates elements of *environment framing*, insofar as certain interactions occur within specific symbolic contexts that guide the type of analysis being performed. These environments function as shared interpretive spaces within which different discursive voices can be mobilized.

This environment may also evolve over time, incorporating new symbolic elements, implicit rules of interaction, and new forms of dialogical organization. Within a single environment framing, it is even possible to invoke multiple personas simultaneously, allowing different voices to participate in parallel in the analysis of the same problem.

From a conceptual standpoint, Nemotine may be understood as an example of *system-level persona framing*, since personas are not merely occasional prompt instructions but integral components of the interaction architecture. In this sense, the system resembles architectures in which cognitive activity is distributed across multiple semi-autonomous agents or perspectives (Minsky, 1986).

6.1 Research Materials and Program

The ideas presented in this paper form part of an ongoing research program on discursive architectures for human–LLM interaction. The Nemotine system functions as an experimental environment within this program, allowing the exploration of structured interaction models based on discursive personas and layered framing mechanisms.

Documentation of the conceptual framework, architectural descriptions, and related research materials are progressively being made available through open scientific repositories. These materials include technical documentation, conceptual models, and experimental interaction protocols.

7 Interaction Analysis

The analysis of interactions conducted within the Nemotine environment suggests several discursive effects associated with the systematic use of persona-based framing.

First, a greater stability of discursive register can be observed when a persona is maintained across multiple turns of interaction. In contrast to dialogues without explicit framing—where textual style may vary considerably between responses—the presence of a persistent

persona tends to produce greater stylistic continuity.

Second, the use of multiple personas within the same interpretive environment may facilitate the diversification of analytical perspectives. By invoking different discursive voices to examine the same topic, the user can explore alternative or complementary interpretations within a single dialogical process.

Another observed effect concerns the pragmatic predictability of responses. When the discursive position of the persona is clearly established, it becomes easier to anticipate the type of argument, explanation, or analytical structure that the model is likely to produce.

These observations are exploratory in nature and do not constitute, at this stage, a systematic quantitative evaluation. Nevertheless, they suggest that persona-based framing may function as a relevant mechanism of discursive organization in human–LLM interaction.

8 Implications for Human–AI Interaction

Interpreting persona prompting as a form of persona framing opens new possibilities for the development of more structured interaction interfaces with language models.

Rather than treating persona instructions merely as local prompt engineering strategies, future systems could explicitly explore multiple layers of discursive framing — including persistent personas, interpretive environments, and polyphonic dialogue architectures.

In this context, language models may function not only as generators of isolated responses, but as participants within organized dialogical structures. Such structures could support activities such as complex problem analysis, exploration of alternative perspectives, and reflective organization of thought.

Architectures based on multiple personas also resonate with earlier ideas in cognitive science that describe intelligence as emerging from the interaction of multiple semi-autonomous processes (Minsky, 1986). Similarly, the use of structured external systems to support reasoning and analysis can be interpreted through the lens of the extended mind theory, in which external artifacts function as forms of cognitive scaffolding for human cognitive processes (Clark and Chalmers, 1998).

Systems inspired by this type of architecture may have potential applications across several domains, including:

- decision-support environments
- dialogue-based educational systems
- conceptual analysis tools

- interfaces for assisted cognitive reflection

From this perspective, environments such as Nemoline may be understood as structures of reflective dialogue in which different personas represent distinct discursive perspectives mobilized to examine a problem or organize a process of thought.

8.1 Future Research Directions

As the present work has an exploratory character, several questions remain open for future investigation.

These include:

- systematic empirical studies on discursive stability under different framing configurations
- quantitative analysis of interactions involving persistent personas
- comparative experiments between different dialogue architectures

To facilitate replicability and further development of this research line, future experimental datasets, architectural descriptions, and research materials associated with the Nemoline system may be made available through open science repositories such as Open Science Framework (OSF) or Zenodo, as mentioned above.

9 Conclusion

This article proposed the concept of *persona framing* as a linguistic interpretation of the practice of persona prompting in interactions with large language models.

We argued that persona instructions do not merely modify the superficial style of responses, but instead introduce pragmatic frames that organize discursive register, enunciative position, and interpretive expectations within the dialogue.

From this perspective, we presented a preliminary typology of framing levels in human–LLM interaction, including *ephemeral persona framing*, *persistent persona framing*, *environment framing*, and *system-level persona framing*.

The case study of the Nemoline system illustrates how these layers of framing can be combined within an interaction architecture structured around multiple discursive personas.

Although the analysis presented in this work is exploratory in nature, it suggests that discursive framing may play an important role in organizing interactions with language models that are more stable, predictable, and analytically productive.

Future research may investigate more systematically the effects of different forms of framing in human–AI interaction environments, as well as explore new dialogue architectures based on multiple discursive voices.

In this sense, persona framing may be understood as an emerging linguistic interface layer in the interaction between humans and large language models.

References

- Bakhtin, M. (1981). *The Dialogic Imagination: Four Essays*. University of Texas Press.
- Clark, A. and Chalmers, D. (1998). The extended mind. *Analysis*, 58(1):7–19.
- Goffman, E. (1974). *Frame Analysis: An Essay on the Organization of Experience*. Harvard University Press.
- Melo, E. (2025). A modular cognitive architecture for assisted reasoning: The nemosine framework.
- Minsky, M. (1986). *The Society of Mind*. Simon and Schuster.
- Reynolds, L. and McDonell, K. (2021). Prompt programming for large language models. *arXiv preprint arXiv:2102.07350*.
- Wei, J. et al. (2022). Chain-of-thought prompting elicits reasoning in large language models. *arXiv preprint arXiv:2201.11903*.