

You Can't Be Neutral on a Moving Train

Jeffrey Heinz

Department of Linguistics and Institute for
Advanced Computational Science
Stony Brook University
jeffrey.heinz@stonybrook.edu
jeffreyheinz.net

Jonathan Rawski

Department of Linguistics and Language
Development
San Jose State University
jon.rawski@sjsu.edu
jrawski.info

Commentary on Futrell, R. & Mahowald, K. (in press). How Linguistics Learned to Stop Worrying and Love the Language Models. Behavioral and Brain Sciences.
<http://doi.org/10.1017/S0140525X2510112X>

Abstract

The commentary argues the authors employ misdirection and strawmanning to cast others as polarized extremes and themselves as the reasonable centrists. We argue that these patterns of misrepresentation ultimately damage any consensus and middle ground they claim to hope to reach.

Futrell & Mahowald (F&M) want readers to see them as reasonable centrists, but do so by casting other positions as polarized extremes via misdirection and strawmen. Their centrism is an illusion, and comes at the cost of saying anything substantive or advancing the conversation. By the end, the sheep becomes a wolf: Claudette is inviting Norm for dinner, and Norm is on the menu. With an invitation like that, no wonder linguists would rather continue to focus on the human capacity for language in all its glorious variations.

F&M frame radical positions as a middle ground by strawmanning the debate. For example, regarding learnability constraints, they state *"There is no logical necessity for a restrictive formalism to explain language learning."* and *"The key mistake was ... that inductive bias has to come from restrictions on the learner's hypothesis space."* They highlight Wilson's (2025) position paper describing several deep learning phenomena using the PAC-Bayes framework via 'soft' biases. In reality these are a set of hidden 'hard' biases: a fixed computable hypothesis class, a fixed description-length prior tied to the architecture, I.I.D. sampling assumptions, bounded-loss assumptions, task-dependent representational primitives (by fallaciously appealing to Li & Vitanyi's (2008) 'additive-constant' invariance theorem), and other capacity constraints. While description-length priors are coherent, Wolpert (2021) shows any bias only justifies a particular hypothesis class to the extent it matches the (unknown) distribution over tasks (see James et al 2023 and Curth et al. 2023 on task-dependent complexity and double-descent).

Without explicitly formalizing that alignment, moving from ‘real-world language data is inherently compressible’ to endorsing biases like those above is meaningless. Whether hypotheses are restricted, admissible data presentations are restricted, or learning resources in time and space are unrestricted, there are fundamental tradeoffs in solving learning problems (Angluin 1988, Goyal and Bengio 2022, a.o.).

F&M’s own candidate constraints like sensitivity and information locality are also model-dependent (Hahn & Roßin 2024, Someya et al 2025), task-dependent (Hahn et al. 2021) and ultimately ‘hard’. Even Wilson and colleagues agree that *“it is largely the design of the architecture that makes generalizable solutions more easily accessible”* (Goldblum et al 2024). Like all ‘domain-general’ biases, F&M’s candidates turn out to be just ‘domain-specific’ biases of a different sort (Rawski & Heinz 2019). In other words, the domains and biases F&M want are real, the others are not. F&M can only achieve centrism from such an extreme position by misassigning people to trivially false claims, like that Rawski & Baumont (2023)’s 1-page paper says LLMs *“can learn to approximate anything”* when it does not. Strawmanning debates by mischaracterizing one perspective to position oneself as a middle ground, while ignoring the substance of the issues, is disingenuous and ultimately leads nowhere.

F&M cover these realities with doublespeak and fallacious arguments. They hype LLM success, concede the successes are disagreed upon and may always emerge from shallow heuristics (Liu et al 2022, McCoy et al 2019, Petty et al 2025, Vafa et al 2025, Vázquez Martínez et al. to appear, a.o.), backtrack to *“this is not the right place to adjudicate these issues,”* then reassert LLMs have *“learned ‘the real thing’”* and are *“not purely relying on heuristics or memorization or cheap tricks”*. Even their interpretability ‘thought experiment’ about an idealized ‘c-command neuron’ is a revamp of standard cognitive neuroscience techniques with the same reification fallacies and circularity (Buzsáki 2019). Again they admit *“LMs just don’t represent anything in such human-interpretable ways,”* yet insist they are *“linguistically interesting”*. This is not a serious discussion.

Beyond the banal observation that *“there are increasingly researchers using language models to ask questions that are informed by and which can inform linguistic theory,”* what is F&M’s actual point? Simultaneously insisting on the Eliminativist fallacy (Smolensky 1988) that restrictive symbolic formalisms are *“dogma”*, while claiming *“we are not eliminativists about grammatical structure,”* is incoherent. Taking a Limitivist view (Smolensky 1988) where symbolic formalisms are an ‘approximate explanation’ and *“graded, probabilistic, and fuzzy”* subsymbolic primitives are cognitively real, they bizarrely omit 30 years of neural networks in generative linguistics via Optimality Theory and Harmonic Grammar (Prince & Smolensky 1997, Smolensky and Legendre 2006). Advocating for ‘domain-general’ learning biases, which end up being domain-specific, paints them into a Dennettian corner to conclude the ‘troll’s

truism’ (Schackel 2005) that “*linguistic structure is real*”. Little else remains, certainly nothing new.

Here we see the tragedy of the centrist, ruled by “*the zeitgeist*” in F&M’s terms, at the cost of ideas or content. Like most self-described centrists, however, their strategies serve to cover another goal. F&M seem to want linguistics to effectively become a niche within psychology (Núñez et al. 2019), where replicability and theory crises (Eronen & Bringmann 2021) mean “*theories rise and decline, come and go, more as a function of baffled boredom than anything else; and the enterprise shows a disturbing absence of that cumulative character that is so impressive*” (Meehl 1978). Linguistic theory is so robust and successful at phenomena discovery that LLM performance benchmarks (Gauthier et al 2020, a.o.) and evaluations (Hale & Stanojević 2024, a.o.) crucially rely on them. When neural models are used to trivialize the substance of linguistic theory, as F&M do with competence and performance (Berwick 2018, Firestone 2020, West et al 2024), linguists have no reason to engage. Replacing theory-driven science with model-driven science, and phenomena with effects-chasing, is a profound retreat from the core insights which gained linguistics membership in cognitive science in the first place.

F&M wonder why many linguists remain uninterested in LLMs, which have not yet discovered virtually any new phenomena nor explained anything of theoretical significance. Yes, LLMs are a technological advance and engineering achievement. Nonetheless, language models, in whatever guise, will continue to be a scientific embarrassment to language scientists as long as they do not offer robust phenomena discovery, counterfactual support, explanatory depth, and other scientific insight into the fundamental properties of the human capacity for language.

Competing Interests: None

Funding: This work is supported by NSF Award #2416183 to JH

References

- Angluin, D. (1988). Identifying languages from stochastic examples. Technical Report 614. Yale University. Department of Computer Science.
- Berwick, R. (2018). Revolutionary New Ideas Appear Infrequently. In *Syntactic Structures after 60 Years*. Norbert Hornstein, Howard Lasnik, Pritty Patel-Grosz and Charles Yang, eds. de Gruyter Mouton.
- Buzsáki, G. (2019). The brain from inside out. Oxford University Press.
- Curth, A., Jeffares, A., & van der Schaar, M. (2023). A u-turn on double descent: Rethinking parameter counting in statistical learning. In *Advances in Neural Information Processing Systems* 36, 55932-55962.

- Eronen, M. I., & Bringmann, L. F. (2021). The Theory Crisis in Psychology: How to Move Forward. *Perspectives on Psychological Science*, 16(4), 779-788.
<https://doi.org/10.1177/1745691620970586>
- Firestone, C. (2020). Performance vs. competence in human-machine comparisons. *Proceedings of the National Academy of Sciences*, 117(43), 26562-26571
- Gauthier, J., Hu, J., Wilcox, E., Qian, P., Levy, R. 2020. SyntaxGym: An Online Platform for Targeted Evaluation of Language Models. In *Proceedings of the 58th Annual Meeting of the Association for Computational Linguistics: System Demonstrations*, pages 70–76, Online. Association for Computational Linguistics.
- Goldblum, M., Finzi, M. A., Rowan, K., & Wilson, A. G. (2024, July). Position: the no free lunch theorem, Kolmogorov complexity, and the role of inductive biases in machine learning. In *Forty-first International Conference on Machine Learning*.
- Goyal, A., & Bengio, Y. (2022). Inductive biases for deep learning of higher-level cognition. *Proceedings of the Royal Society A: Mathematical, Physical and Engineering Sciences*, 478(2266).
- Hahn, M., Jurafsky, D., & Futrell, R. (2021). Sensitivity as a complexity measure for sequence classification tasks. *Transactions of the Association for Computational Linguistics*, 9, 891-908.
- Hahn, M., & Rojin, M. (2024, August). Why are Sensitive Functions Hard for Transformers?. In *Proceedings of the 62nd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)* (pp. 14973-15008).
- Hale, J. and Stanojević, M. 2024. Do LLMs learn a true syntactic universal?. In *Proceedings of the 2024 Conference on Empirical Methods in Natural Language Processing*, pages 17106–17119, Miami, Florida, USA. Association for Computational Linguistics.
- James, G., Witten, D., Hastie, T., Tibshirani, R., & Taylor, J. (2023). *An introduction to statistical learning: With applications in Python*. Cham: Springer International Publishing.
- Kearns, M. J., & Vazirani, U. (1994). *An introduction to computational learning theory*. MIT press.
- Li, M., & Vitányi, P. (2008). *An introduction to Kolmogorov complexity and its applications*. New York: Springer.
- Liu, B., Ash, J. T., Goel, S., Krishnamurthy, A., and Zhang, C. Transformers learn shortcuts to automata. In *The Eleventh International Conference on Learning Representations*, 2022.
- McCoy, R. T., Pavlick, E., & Linzen, T. (2019, July). Right for the Wrong Reasons: Diagnosing Syntactic Heuristics in Natural Language Inference. In *Proceedings of the 57th Annual Meeting of the Association for Computational Linguistics* (pp. 3428-3448).
- Meehl, P. E. (1978). Theoretical risks and tabular asterisks: Sir Karl, Sir Ronald, and the slow progress of soft psychology. *Journal of Consulting and Clinical Psychology*, 46(4), 806–834. <https://doi.org/10.1037/0022-006X.46.4.806>
- Núñez, R., Allen, M., Gao, R., Miller Rigoli, C., Relaford-Doyle, J., & Semenuks, A. (2019). What happened to cognitive science?. *Nature human behaviour*, 3(8), 782-791.

- Petty, J., Hu, M., Wang, W., Ravfogel, S., Merrill, W., and Linzen, T. 2025. RELIC: Evaluating compositional instruction following via language recognition. In NeurIPS 2025 Workshop on Evaluating the Evolving LLM Lifecycle: Benchmarks, Emergent Abilities, and Scaling.
- Prince, A., & Smolensky, P. (1997). Optimality: From neural networks to universal grammar. *Science*, 275(5306), 1604-1610.
- Rawski, J., & Heinz, J. (2019). No free lunch in linguistics or machine learning: Response to Pater. *Language*, 95(1), e125-e135.
- Rawski, J., & Baumont, L. (2023). Modern language models refute nothing. Lingbuzz Preprint.
- Shackel, N. (2005). The vacuity of postmodernist methodology. *Metaphilosophy*, 36(3), 295-320.
- Someya, T., Svete, A., DuSell, B., O'Donnell, T., Giulianelli, M., and Cotterell, R. (2025). Information Locality as an Inductive Bias for Neural Language Models. In *Proceedings of the 63rd Annual Meeting of the Association for Computational Linguistics (Volume 1: Long Papers)*, pages 27995–28013, Vienna, Austria. Association for Computational Linguistics.
- Smolensky, P. (1988). On the proper treatment of connectionism. *Behavioral and brain sciences*, 11(1), 1-23.
- Smolensky, P. and Legendre, G.. 2006. *The Harmonic Mind: From Neural Computation to Optimality-Theoretic Grammar*. Volumes I and II. MIT Press.
- Vafa, K., Chang, P. G., Rambachan, A., & Mullainathan, S. (2025). What Has a Foundation Model Found? Using Inductive Bias to Probe for World Models. In *International Conference on Machine Learning* (pp. 60727-60747).
- Vázquez Martínez, H. J., Heuser, A., Yang, C., and Kodner, J. (to appear). Evaluating the Existence Proof: LLMs as Cognitive Models of Language Acquisition. In *Artificial Knowledge of Language*. Lingbuzz preprint: <https://lingbuzz.net/lingbuzz/008277>.
- West, P., Lu, X., Dziri, N., Brahman, F., Li, L., Hwang, J. D., Jiang, L., Fisher, J., Ravichander, A., Chando, K., Newman, B., Koh, P.W., Ettimner, A., Choi, Y. The Generative AI Paradox: “What It Can Create, It May Not Understand”. In *The Twelfth International Conference on Learning Representations*
- Wolpert, D. H. (2021). What is important about the no free lunch theorems?. In *Black box optimization, machine learning, and no-free lunch theorems* (pp. 373-388). Cham: Springer International Publishing.