

No Free Lunch: Copredication Beyond Predicate Extension

Elliot Murphy^{1,2}

Affiliations: 1. Vivian L. Smith Department of Neurosurgery, McGovern Medical School, UTHealth, Houston, TX, USA

2. Texas Institute for Restorative Neurotechnologies, UTHealth, Houston, TX, USA

Permanent Address: 1133 John Freeman Blvd, Houston, TX 77030, USA

Disclosure of Interest: The author reports there are no competing interests to declare.

Abstract: Under Liebesman and Magidor's Property Versatility account, predicates such as 'on the shelf' and 'by Tolstoy' are more versatile than theorists have assumed. I argue that approach turns a genuine problem about how linguistic representations coordinate incompatible conceptual formats into a permissive doctrine about predicate extension, and does not align with reported empirical patterns found in the psycholinguistics literature. I defend an alternative internalist account on which polysemous nominals make available organized combinations of senses, while predicates act as probes that select and reweight dimensions of these representations. Copredication succeeds when multiple senses can be incrementally coordinated and stabilized under tolerable costs of order, complexity, coherence, and individuation. This perspective also explains why copredication cannot be dismissed as merely an illusion or unrecognized zeugma.

Keywords: Predicate order; coherence relations; semantic complexity; polysemy

Liebesman and Magidor (2026) (L&M) begin from a familiar observation: sentences such as ‘Lunch was delicious but took hours’, ‘The school was vandalized after expelling Bob’, and ‘The red book on the shelf is enthralling and was written by Tolstoy’ can be used to express truths, despite joining predicates that select for different kinds of things. Their diagnosis is that much of the literature has over-blown the consequences of these kinds of ‘copredication’ structures, and they defend a Property Versatility account (PV). According to PV, copredication does not require special semantic machinery: a single entity can satisfy the relevant noun and predicates because ordinary predicates are more versatile than we typically give them credit for.

Below, I offer some critical commentary and defend an alternative semantic internalist perspective. While I concur that copredication may not need any special semantic machinery, I will argue that L&M’s PV account does not align with the current empirical terrain: predicate order, coherence relations, semantic complexity, and individuation effects all modulate copredication acceptability in ways that cannot be predicted by simply expanding predicate extensions. I defended instead a cognitively grounded internalist account supported by a range of work not cited by L&M (nor discussed in their other recent publications), according to which polysemous nominals provide structured combinations of senses, and predicates act as probes into these senses. Copredication is therefore neither evidence for bizarre external objects nor a harmless artifact of predicate flexibility, but a controlled form of conceptual coordination at the grammar-pragmatics interface (Murphy 2016).

1. The paper was boring and long: Evaluating Liebesman and Magidor’s thesis

L&M succeed methodologically mainly by intentionally narrowing the explanatory target: Copredication is taken to be a problem about the truth conditions of selected sentences, and they then ask whether familiar extensional resources can be preserved. From the perspective of the psycholinguistic (Murphy 2017, 2021a, 2021b, 2022, 2024a) and internalist (Murphy 2023) approach I have developed elsewhere, this loses sight of some other critical phenomena: not just how a sentence can be made true, but how comprehenders coordinate conflicting representational formats in real time, why some predicate orderings are better than others, and why the same nominal-predicate materials vary systematically with order, coherence, context, and conceptual complexity.

L&M state, in reference to *lunch*-type copredications, that “Nothing can be delicious and take hours”. There apparently are things that can be delicious and take hours: Lunches turn out to be quirkier and more conceptually complex than we initially would intuit, they just don’t exist as some ‘physical(ist)’, external entity or set of truth-conditions that human languages come to meet

on neutral ground. The metaphysical horror is produced by philosophers like the Dual-Nature theorists first imposing a clean external typology, and then acting surprised when a vast number of ordinary lexical items refuse to respect it. The lesson of copredication should not be that ordinary speakers mysteriously ascribe impossible properties to impossible objects. It should be that ordinary lexical items are instructions for building conceptual representations whose admissible inferences do not line up cleanly with the ontology favored by externalist truth-conditional semantics (Pietroski 2018). The conceptual landscape of the human mind is multi-faceted, with many sub-domains offering representational types that often conflict in their inferential resources. But that is the very nature of mental life – it is precisely in the tensions and conflicts that different conceptual modules/resources offer that compositional semantic objects can be constructed and used to update multiple aspects of our world model. Objects and ‘things’ are built in our minds precisely by a process of *negotiation and compromise* across distinct conceptual domains.

L&M write that it is indeed “puzzling” that structures involving copredication can be evaluated for truth, given that they call upon the comprehender to bind together categorially otherwise incompatible conceptual features. However, this is only a puzzle if we assume that truth-evaluability needs to be in relation to an externalist framing, whereby linguistic inferences are somehow evaluated in relation to some stable external world, rather than an internal world model that can readily admit for conceptual incompatibilities.

Some cognitive modules and core knowledge systems will be able to export ‘un-readable’, non-compatible forms of information to each other. This is a remarkably common, yet under-appreciated, aspect of mental life. We can try and export certain musical, linguistic, or visual forms of information to undergo affective inferences, or moral inferences, or olfactory inferences – yet, often this will be barred from closure through the sheer fact of the variability in conceptual representational formats.

So too for the syntax-semantics and syntax-pragmatics interfaces with copredication. Certain grammatical configurations can be exported for evaluation by various (and presently ill-understood) conceptual systems, but often this will result either in conceptual paradoxes, or certain pragmatic and performance systems will serve to disguise the underlying logico-semantic contradictions and prevent us from worrying too much about the meaning of ‘lunch’. But, as professional philosophers, worry we shall.

2. Property Versatility serves to relocate, not explain, the puzzle

The central concern about PV that I would like to raise here is that it threatens to redescribe the explanandum as the explanation. If ‘The bestselling book is on the shelf’ is acceptable, PV says that ‘on the shelf’ is versatile enough to apply to informational books. If ‘The bestselling book is slightly wrinkled’ is infelicitous, PV says that ‘wrinkled’ is not versatile enough to apply to informational books. In individual cases this may sound plausible. But as a general theory it risks a form of *ex post facto* lexical bookkeeping: acceptable copredications are those in which at least one predicate turns out to be quote-unquote “versatile” in the required way; unacceptable ones are those in which it does not.

The issue here is not that predicates cannot be versatile, but that PV needs a principled theory of which extensions are available, why they are available, how they are learned, and why certain nearby extensions fail. If informational books can be on the shelf because they are instantiated by shelf-bound copies, why can they not also be wrinkled when all their copies are wrinkled? Some of these are acceptable in special contexts; others are degraded; many are unstable. PV can describe these contrasts only *after* consulting the very judgments it is meant to predict. A theory of copredication should predict this distribution before consulting speaker judgments. PV instead risks treating versatility as a label assigned to whichever predicate extensions speakers happen to tolerate.

I have raised similar methodological concerns in relation to other approaches to copredication (Murphy 2022), but the same logic applies here. A useful theory should make clear predictions about impossible or degraded combinations (Murphy 2024a). It should say, before we ask speakers, which property combinations are likely to be licensed across languages, which require strong discourse support, which are order-sensitive, and which will remain zeugmatic. A maximally permissive theory of property versatility can always be rescued by expanding or contracting the extension of the predicate. But that flexibility is purchased at the cost of explanatory precision.

The same concern arises in L&M’s treatment of category mistakes. L&M are right to deny that predicates impose only crude type requirements such as *physical* or *abstract*. They propose more specific constraints: ‘on the shelf’ requires something like locatability; ‘red’ requires having a color. Yet, the more specific the constraint becomes, the more it appears to be reverse-engineered from the data. To say that informational books are located, but numbers are not, may capture ‘War and Peace is on the shelf’ versus ‘The number two is on the shelf’. But it does not explain how the conceptual system determines the relevant kind of location, why one form of mediated instantiation counts and another does not, or why discourse can temporarily rescue otherwise degraded cases. A theory of copredication should not merely list the presuppositions

predicates happen to tolerate; it should explain the cognitive and discourse mechanisms by which those tolerances are computed (Murphy 2021a).

L&M are right that category mistakes cannot be explained by crude type labels such as PHYSICAL and EVENT. But replacing crude types with more specific predicate presuppositions does not yet explain copredication. The relevant question is not whether 'on the shelf' requires physicality or locatability; the question is why informational books are locatable in the relevant mediated sense while numbers, jokes, symphonies, and legal principles are not – except under special representational or institutional settings. The missing notion is a licensed connection between inferred senses. *Lunch* conventionally links food and event structure; *sandwich* does not. *Book* conventionally links physical copy, informational work, authorship, and edition/token structure; *number* does not. Copredication succeeds when a predicate can access one feature space of a nominal's senses, while the discourse can stabilize its relation to the remaining feature spaces. Category mistakes arise when no such bridge is lexically or contextually supplied.

In prior work (Liebesman & Magidor 2017), L&M claim that nominals like 'book' are not ambiguous but encode a single sense (a form of univocity) simultaneously designating both physical and informational entities, and that it is via contextual restriction that multiple readings are permitted. For example, an argument they make against the ambiguity of 'book' is that cases like 'Three colours are on the canvas' – involving three distinct shades of red – cannot automatically be seen as supporting the ambiguity of the nominal 'colours'. Rather, it is other words in the sentence like 'three' which constrain the meaning of 'colour', or it is 'red' which is ambiguous; in addition, 'colour' is clearly contextually restricted (2017: 136). What to make of this? It might simply be that contextual restrictions are active for colour-type scenarios while genuine ambiguity is at play for book-type nominals – in particular given that 'book' putatively encodes senses of distinct types, while colour can only encode different gradations of the same physical type. Another core aim of L&M is to show that copredication-licensing nominals are not the only linguistic entities that can instantiate semantically distinct properties. To illustrate their point, they discuss properties that can be instantiated by both events and non-events, such as being 'tolerable': A certain condescending explanation can be deemed intolerable (event reading) by virtue of the mental state of the individual listening to it, while an individual can also be judged as intolerable (non-event reading) because of their penchant for delivering condescending explanations. As such, "properties can be instantiated by different objects in different ways" (Liebesman & Magidor 2017: 143), and they claim that copredications can be derived from this basic metaphysical fact. Yet, the crucial difference between event and non-event readings of 'tolerable' and event and non-event readings of 'lunch' is that the latter involves the simultaneous

instantiation of these readings; one cannot combine both readings of 'tolerable' or attribute them to a single entity.

Polysemous nominals do not denote external objects (Chomsky 2000), nor do they merely invite arbitrarily flexible predicates. They provide structured sense combinations: internally organized packages of semantically distinct but related senses. A noun such as *book* makes available physical-copy, informational-work, authorial, edition, and token/type dimensions; *lunch* makes available food, event, and social-occasion and dimensions; *school* makes available building, institution (staff, pupils), and educational-function dimensions. Predicates act as probes into these frames. They select a sense, impose integration demands, and require discourse to stabilize the resulting representation. Copredication is licensed when this stabilization succeeds under tolerable costs of order, coherence, complexity, and individuation (and likely many other psycholinguistic variables yet to be empirically tested).

3. Copredication is not homonymy with better predicates

A second general problem is L&M's assimilation of copredication to ordinary semantic variability. The authors repeatedly compare 'book' to 'bat': just as 'bat' has multiple readings, 'book' can range over physical books or informational books. But inherent polysemy is not homonymy plus a relation between extensions. The physical, informational, eventive, and institutional aspects of nouns such as 'book', 'lunch', 'school', 'newspaper', 'bank', and 'city' are not merely alternative dictionary entries. Copredication is not 'homonymy + tolerant predicates'; it involves organized representational packages whose senses can be jointly activated, ranked, suppressed, and stabilized.

This matters because the authors infer too quickly from a possible truth-conditional uniformity to a *cognitive* uniformity. Even if a given reading of 'The lunch was delicious but took hours' assigns lunch a single event-like denotation, that does not show that the food sense plays no role in interpretation. 'Deliciousness' is not processed as a free-floating versatile property of events; it recruits the food-associated dimensions of the lunch representation. Likewise, even if 'The school was vandalized after expelling Bob' is ultimately represented as involving one institution-like entity, the building sense is not irrelevant. Certainly, a school's necessary features involve its institutional/organizational meaning, and the building sense is not necessary (a school can exist and practice in the middle of a field with no building required; discussed in Murphy 2021a), but this does not mean that 'school' is not polysemous. It is precisely the availability of building-associated inferences that makes vandalized appropriate.

During copredication, a predicate searches the nominal's compatible senses for an admissible attribute binding, then imposes constraints on how that attribute must be integrated with the rest of the discourse. I would argue that the mistake that PV makes is that it treats the outcome as an extensional fact about the predicate: 'delicious' or 'on the shelf' just has a broader extension. More things turn out to be delicious than we had hitherto considered! But the psychologically serious explanation is procedural: the predicate recruits one sense of a structured nominal representation and forces the system to integrate that with other senses, naturally being constrained by language-external cognitive demands on what our core knowledge systems will permit.

This is one domain where the internalist framing is useful: linguistic expressions instruct language-external conceptual systems to construct representations, and those systems may differ in format, granularity, and inferential access. Some conceptual information is readily exported across systems; other information is unreadable or only partially compatible. Copredication is the grammatical and conceptual site at which these formats collide. The puzzlement arises not because 'lunch' must pick out one external thing that is both food and event, but because the language faculty can generate a configuration that asks the conceptual system to close over partially incompatible modes of representation.

Yet, humans routinely entertain conceptual compromises, useful fictions, idealizations, category-crossing representations, and counterfactual or mythic entities. Lunches, books, banks, universities, and cities are not less real for being conceptually complex, or licensing eccentric interpretations. They are real as any objects entertained by our representational practices. The real question is how the grammar-pragmatics interface delivers stable enough representations for judgment, inference, and action. PV treats this stability as an extensional fact about predicates; the approach I would favor treats it as an achieved state of interpretation.

When discussing 'sandwich' versus 'lunch', L&M argue that "if 'lunch' picks out some food then the predicate 'took place on Tuesday' should give rise to a similar anomaly (and similarly, if 'lunch' picks out an event, then the parallel worry arises with respect to 'is delicious')." However, 'lunch' and 'sandwich' are, quite simply, different words. One intrinsically licenses both abstract and concrete readings, while the other only permits certain forms of deferred reference (e.g., 'the ham sandwich at table 12 needs his check'). As L&M also note, we do not even need to invoke the usual internalist tactics here at all. 'Lunch' provides instructions to access eventive and physical senses, whereas 'sandwich' does not.

Many of L&M's intuitions about acceptable copredication licensing are also odd to me (hence why I am a strong supporter of empirically assessing these properties of copredication).

With three physical copies of *War and Peace* on the shelf, with no other books, to my mind it would *not* be wholly acceptable for someone to turn to me and say ‘Exactly one book is on the shelf’. Clearly, there are three physical books on the shelf, and the *partial* acceptability in fact supports the polysemy view whereby both senses are available for inference but one is saliently triggered but not to the extent that it automatically *negates and suppresses* the other. Gotham’s (2022) work has explored this question in far superior detail, but certainly it would also only be *partially* acceptable to say ‘Exactly three books are on the shelf’ if there was only one physical volume containing a trilogy of novels. In fact, to my mind both structures are equally as acceptable (namely, only partially so). The symmetrical nature of this partial acceptability between three-informational and three-physical strongly complies with the polysemy approach (Murphy 2021a).

Things become more complicated with *newspaper*-type copredications, where we have informational, physical *and* institutional senses to count (discussed at length in Murphy 2021a), and we can have situations where none of these three senses align in terms of how many there are (e.g., 2 institutions, 5 physical copies, 4 informational sources – for instance, multiple copies of Thursday’s *Daily Telegraph*) – which is perfectly in line with the idea that *newspaper* is polysemous. L&M correctly observe that ‘The newspaper fell off the table and fired its editor’ is not particularly acceptable. Yet ‘The newspaper that was sued can be found on the table’ is mostly acceptable, as is ‘The newspaper purchased more ink and has five colorful pages’, for reasons discussed in Murphy (2021a) to do with coherence relations, predicate ordering, and syntactic structure. The polysemy account does not imply that *all* possible combinations of senses must be acceptable – only that semantic compositions also have to pass assessments from pragmatic and conceptual systems, which have their own idiosyncratic criteria. The problem here is not “what objects exist?” but “which internal individuation criteria are active, and can they be jointly satisfied?” The individuation criteria discussed in the copredication literature do not need to be addressed by externalist mereological models, nor via versatile predicates, but simply by a model of the syntax-semantics interface that takes into consideration how nominals offer our faculties for quantification multiple options – much as how our sense of smell can offer our lexicon various options for selection, which often lead to a sense of paradox or internal contradiction (for example, when sampling fine wines).

L&M seem to think that manipulating types and tokens of nouns yielding unacceptable copredications (‘The book was translated to many languages and has a small coffee stain on its cover’) carries with it serious implications. Alternatively, it might merely show that reducing or compressing a token to its general type yields an odd reading. This is true not just for copredication, but for many kinds of structures:

- (1) The average American has 1.9 children and is waiting outside.
- (2) The young bachelor is unmarried and would like to know who ate his yogurt.

L&M's discussion about words at the end of their article is also not in keeping with assumptions in modern psycholinguistics concerning the form, syntax, and semantics of lexical items, and how these need not stand in a 1:1:1 relation (Murphy 2024b).

Some of L&M's other claims about the alleged implications of certain copredication structures being unacceptable also strike me as peculiar: for instance, why are we not also stroking our chins about sentences like 'The newspaper sued the French government, rolled down the hill, joked about the FIFA World Cup, and was then collectively burned in seven different households'? The answer seems rather obvious – polysemy approaches *do not* require all possible sense combinations to be valid, any more than any other form of complex compositional semantic processes need to meet some arbitrary acceptability criteria set by philosophers in order to be psycholinguistically plausible elements of linguistic cognition.

In any event, to return the above theme, 'Exactly one book is on the shelf' (when we have three physical copies of the same informational book) may force us to count differently, but while this is a serious problem for truth-conditional externalist semantics, it is not an immediate issue for semantic internalism. 'John read and then burned every book in the library', in a situation in which a library has multiple copies of certain books, also ensures that 'every book' picks out distinct quantities, depending on the sense we attend to. But what are the implications of this? The function of human language is not to go around counting things and making sure that all the logically possible quantities inferred from sentences that can yield *multiply true* meanings are somehow all aligned.

This also reframes the double-distinctness debate. If 'Granger memorized three yellow books' requires three distinct volumes expressing three distinct texts, that is not because the sentence quantifies over exotic dual-nature objects. It is because the predicate pair imposes two simultaneous individuation pressures: 'yellow' selects physical-token distinctness, while 'memorized' selects informational-content distinctness. Where those criteria diverge, acceptability becomes graded and context-sensitive. This is not an embarrassment for a polysemy-based internalism; it follows from a system in which lexical items provide multiple internal criteria of identity.

4. Missing empirical terrain: order, coherence, complexity

Perhaps the largest omission in L&M's approach is any remote engagement with empirical terrain. Predicate order (Murphy 2019, 2024a), predicate coherence (Murphy 2021a) and predicate semantic complexity (Murphy 2022) are all relevant empirical factors that need to be entertained (and explained). I will briefly entertain each of these here.

First, PV has no natural place for the kind of asymmetry found in predicate order preferences. If the explanation of a specific case of acceptability is that one predicate is versatile enough to apply to the relevant entity, then why should reversing the predicates affect acceptability? A truth-conditional account can perhaps say that both have a true reading, or that one reading is pragmatically harder. But without a theory of incremental conceptual construction (e.g., the 'Incremental Semantic Complexity' parsing bias triggering a general Simple-Complex ordering preference, able to be overridden by coherence relations; Murphy 2021a), it cannot predict the direction and size of the effect. The significant predicate order effects reported elsewhere (Murphy 2021a, 2024a) are not cosmetic – they reveal subtle preferences in the directionality of permissible sense construction (see the 'Copredication Licensing Effect' in Murphy 2024a).

Second, coherence relations matter hugely here. Copredications improve when the predicates stand in causal, explanatory, temporal, or common featural relations. 'The newspaper was sued because it was offensive' is more natural than a flat conjunction ('and it was offensive') lacking any inferable relation. 'The city outlawed smoking in bars and has 500,000 smokers' is more acceptable than the same sentence with '500,000 inhabitants', because the predicates share a salient conceptual feature. Coherence can weaken or override predicate-order penalties (Murphy 2024a). This means that copredication licensing is not solely a matter of which predicate can apply to which entity. It is a matter of how discourse binds the predicates into a meaningful relation. *Metaphysics, versatility, truth, ontology* – these are almost beside the point entirely.

Third, complexity is not reducible to frequency or prediction (Murphy 2021b). The best predictor of the persistence conditions of copredications involving institutional entities (i.e., which sense of an institutional nominal 'survives' certain narrative interventions) is not sense frequency, but semantic complexity, or approximation to the institutional sense. This naturally undermines explanations that treat copredication as merely a matter of dominant senses, expectations, or familiar truth-conditional readings (Murphy 2022). It also constrains whatever version of PV one might want to still consider. If speakers preserve institutional senses under degradation scenarios, then their judgments are sensitive to the internal organization of the nominal's conceptual package, not just to the extension of predicates. A city, bank, school, or newspaper is not a neutral

node to which predicates attach. It is a structured cluster of senses whose components differ in complexity, stability, and inferential centrality (Pietroski 2018; Pustejovsky & Batiukova 2019).

Importantly, L&M and various other accounts lean heavily on expectations, priors, discourse, and world knowledge. As discussed elsewhere (Murphy 2022), sense frequency does not explain predicate order effects: sense frequency has no statistically significant relation with sense-order acceptability, and the evidence instead points toward semantic complexity. Context does not determine copredication acceptability by itself; it supplies bridging material that can reduce the cost of sense integration. But, as empirically documented (Murphy 2021a), context cannot rescue everything equally: the parser still pays costs for complexity inversion, sense complexity distance, individuation conflict, and low coherence.

Copredication acceptability therefore appears to be a function of: sense availability + coherence + contextual bridging + anaphoric/syntactic support – complexity inversion – sense distance – individuation conflict. A copredication is licensed when its predicates can be incrementally integrated into a stable combination of polysemous senses under tolerable costs of complexity, coherence, and individuation.

5. Curb your externalism

In general, philosophers enjoy writing papers about copredication because it makes them feel like they are traversing serious linguistics terrain without necessitating much expertise in any given domain of language (syntax, formal semantics, pragmatics). Likewise, linguists enjoy writing papers about copredication because it makes them feel like they are contributing to the most pressing philosophical questions without having to wade too deeply into the philosophy of language/mind literature. How might we evaluate PV from the perspective of one of these linguists/philosophers who might be interested in preserving something of valuable to researchers of all backgrounds?

PV is surely right to resist the thought that every successful copredication requires a baroque dual-nature object. It is also right that many predicates are more flexible than philosophers sometimes assume. But these points should be embedded inside a richer model rather than offered as a unified proposal. Predicate versatility is one ingredient of copredication; it is not the phenomenon itself. A more adequate view would distinguish at least four levels. At the lexical level, complex polysemous nominals encode structured sense packages: physical, informational, eventive, institutional, and other components. At the compositional level, predicates impose demands on these packages and may activate, suppress, or re-weight particular senses. At the discourse level, coherence relations and other pragmatic processes help determine

whether the resulting configuration can be integrated. At the broader cognitive level, the resulting representation is evaluated against an internal model rich enough to host apparent paradoxes, mediated instantiation, metonymic transfer, and pragmatic repair.

Under this view, it is certainly true that some predicates really are truly versatile; some cases can be analyzed without novel semantic machinery; some dual-nature accounts overgenerate. But to constitute a model, let alone a theory, PV should be required to specify, in advance, which predicate orderings should degrade acceptability, which discourse relations and pragmatic forms of assistance should rescue them, and how the internal representational make-up of certain lexical items guides the formation of copredications. This is where philosophers of language need to learn from philosophers of science, in terms of what actually constitutes a serious model or theory.

L&M's statements about internalism also require some comment. Mental representations are *useful fictions* (Ramstead et al. 2020) for achieving particular inferential goals, which can very often conflict with our sentential assessments of truth. Chomsky further rebuts the externalist position (personal communication, April 2022): "What notion of extension permits an object X to be within or outside the extension of a term depending on how we think about it? I've never heard of any, and can't imagine what it would be".

Similarly, if the externalist object we choose to formulate is incoherent, that is not evidence that copredication is "illusory" (as per Collins 2024). It is evidence that externalist object-based semantics is the wrong level of analysis. Indeed, Collins (2024: 381-382) offers examples of copredication that violate ISC (Murphy 2021a, 2024a) without also noting that the acceptable copredication cases he offers (e.g., 'The school was ranked first in the county despite being dilapidated') are acceptable *precisely because of an established coherence relation*, as per the Copredication Licensing Effect (Murphy 2024a). If copredication is merely unrecognized zeugma (as Collins argues), then the central empirical burden is explaining why some "illusions" are better than others, why they vary with predicate order, and why causal/featural coherence improves acceptability. Collins can redescribe good copredications as easy illusions and bad copredications as failed illusions, but he lacks an independent account of what makes the illusion easy or difficult. Once he introduces order, coherence, and semantic complexity to explain this (as he partly does; Collins 2024: 379-382), he has already conceded the central terrain put forward in Murphy (2021a, 2024a).

Overall, an alternative to L&M's PV view and the illusion view (Collins 2024) is the thesis that copredication is not the discovery that ordinary predicates have unexpectedly broad extensions, nor is it the production of a merely associated content that disguises semantic

incoherence. It is the online stabilization of a single discourse index across multiple internally structured combinations of senses. Polysemous nominals can make available organized representational packages. Predicates act as probes into these packages, selecting senses and imposing integration demands. Predicate order matters because the first predicate anchors the discourse index, while later predicates must enrich or reconstruct that anchor (Murphy 2024a). Coherence matters because causal, functional, narrative, and featural relations bind otherwise divergent facets into a usable conceptual object (Murphy 2021a). Predicate semantic complexity matters because predicates differ in the amount of structure they import and in the granularity of the facets they select (Murphy 2022). Copredication is therefore neither metaphysically uniquely instructive nor semantically illusory. It is a controlled form of conceptual coordination at the grammar-pragmatics interface, with a range of interesting implications for psycholinguistic, neurolinguistic, and artificial models of semantics, some of which are already being explored (Baggio & Murphy 2026; Dentella et al. 2024; Leivada et al. 2026; Murphy 2025; Murphy et al. 2023, 2025, 2026).

This point is continuous with the empirical picture developed in Murphy (2021a, 2022, 2024a), not cited by L&M. As mentioned, predicate order effects are not reducible to dominance or sense frequency: Complex-Simple orderings are generally degraded relative to Simple-Complex orderings even where frequency-based expectations predict no such asymmetry. Likewise, coherence relations improve acceptability not by changing the extension of predicates, but by making the transition between senses easier to stabilize. The resulting picture is not a sense-enumeration theory on which all senses are equally available, nor a PV theory on which predicates simply stretch. It is a constrained activation model: senses belonging to a complex nominal are internally related, but they differ in complexity, accessibility, mutual priming strength, and discourse stability.

6. Conclusion

L&M's account of copredication is ultimately too conservative. It saves externalist compositionality by allowing predicates to stretch, while leaving unexplained why they stretch in some directions and not others, why certain stretches are easier in one order than another, and why coherence can repair otherwise degraded cases. Their account therefore captures the truth-evaluable residue of copredication but not its generative core.

Copredication shows that natural language meanings are instructions for building internal conceptual objects whose parts can be incompatible yet usable, unstable yet truth-evaluable, paradoxical yet communicatively ordinary. Lunches really can be delicious and take hours. Books

really can be heavy and profound. Cities really can vote, burn, relocate, and remember. The task is not to domesticate these facts by stretching predicates until the puzzle disappears, but to explain the internal architecture that makes such representations possible, graded, and constrained.

References

- Baggio, Giosuè, & Elliot Murphy. (2026). On the referential capacity of language models: an internalist rejoinder to Mandelkern & Linzen. *Computational Linguistics* <https://doi.org/10.1162/COLI.a.601>.
- Chomsky, Noam. (2000). *New Horizons in the Study of Language and Mind*. Cambridge: Cambridge University Press.
- Collins, John. (2024). Copredication as illusion. *Journal of Semantics* 40: 359–389.
- Dentella, Vittoria, Fritz Günther, Elliot Murphy, Gary Marcus, & Evelina Leivada. (2024). Testing AI on language comprehension tasks reveals insensitivity to underlying meaning. *Scientific Reports* 14: 28083.
- Gotham, Matthew. (2022). Property inheritance, deferred reference and copredication. *Journal of Semantics* 39(1): 87–116.
- Leivada, Evelina, Gary Marcus, Fritz Günther, & Elliot Murphy. (2026). A sentence is worth a thousand pictures: can large language models understand hum4n L4ngu4ge and the W0rld behind W0rds? *Philosophical Transactions of the Royal Society A* 384: 20250008.
- Liebesman, David, & Ofra Magidor. (2017). Copredication and property inheritance. *Philosophical Issues* 27(1): 131-166.
- Liebesman, David, & Ofra Magidor. (2026). What copredication can(not) teach us. *Theoretical Linguistics* <https://doi.org/10.1515/tl-2026-2002>.
- Murphy, Elliot. (2016). Phasal eliminativism, anti-lexicalism, and the status of the unarticulated. *Biolinguistics* 10: 21–50.
- Murphy, Elliot. (2017). Acquiring the impossible: developmental stages of copredication. *Frontiers in Psychology* 8: 1072.
- Murphy, Elliot. (2019). Acceptability properties of abstract senses in copredication. Bolognesi, M., & Steen, G. (Eds). *Perspectives on Abstract Concepts: From Cognitive Processing to Semantic Representation and Linguistic Expression*. Human Cognitive Processing Series. Amsterdam: John Benjamins. 145–165.
- Murphy, Elliot. (2021a). *Linguistic Representation and Processing of Copredication*. PhD thesis. University College London.

- Murphy, Elliot. (2021b). Persistence conditions of institutional entities: investigating copredication through a forced-choice experiment. *Frontiers in Psychology* 12: 528862.
- Murphy, Elliot. (2022). Copredication and complexity revisited: a reply to Loehr and Michel. *Cognitive Science* 46(10): e13207.
- Murphy, Elliot. (2023). The citadel itself: defending semantic internalism. *Global Philosophy* 33: 7.
- Murphy, Elliot. (2024a). Predicate order and coherence in copredication. *Inquiry* 67(6): 1744-1780.
- Murphy, Elliot. (2024b). What is a word? arXiv:2402.12605.
- Murphy, Elliot. (2025). ROSE: A universal neural grammar. *Cognitive Neuroscience* 16(1-4): 49–80.
- Murphy, Elliot, Kiefer J. Forseth, Cristian Donos, Kathryn S. Snyder, Patrick S. Rollo, & Nitin Tandon. (2023). The spatiotemporal dynamics of semantic integration in the human brain. *Nature Communications* 14: 6336.
- Murphy, Elliot, Evelina Leivada, Vittoria Dentella, Raquel Montero, Fritz Günther, & Gary Marcus. (2025). Fundamental principles of linguistic structure are not represented by ChatGPT. *Biolinguistics* 19: e19021.
- Murphy, Elliot, Oscar Woolnough, Cale W. Morse, Xavier Scherschligt, & Nitin Tandon. (2026). Frontotemporal network interactions causally support rapid concreteness judgments during reading. *PLoS Biology* 24(3): e3003723.
- Ramstead, Maxwell J.D., Karl J. Friston, & Inês Hipólito. (2020). Is the free-energy principle a formal theory of semantics? From variational density dynamics to neural and phenotypic representations. *Entropy* 22: 889.
- Pietroski, Paul. (2018). *Conjoining Meanings: Semantics Without Truth Values*. Oxford: Oxford University Press.
- Pustejovsky, James, & Olga Batiukova. (2019). *The Lexicon*. Cambridge: Cambridge University Press.