Profiles, Not Prosodies: A Weighted-Feature Diagnostic for Connotative Meaning in Lexical Semantics

Tariq Al Mozahid

Jahangirnagar University

Author Note

I have not received any fundings and have no conflicts of interest.

Large Language Models have been used to collect references. References have been independently verified later.

Ethical Approval: Not Applicable (no human participant is involved)

House#288, Road#06, Mohammadia Housing Society, Mohammadpur, Dhaka-1207, Bangladesh

Email: talmozahid@gmail.com

Abstract

The theory of semantic prosody holds that words acquire evaluative coloring through repeated collocation with charged contexts, such that connotation is absorbed from usage rather than encoded in the lexicon. This paper argues that the phenomena motivating that theory are real but have been correctly measured and incorrectly attributed. It treats lexical meaning as a weighted bundle of semantic features and distinguishes two independent readouts of that bundle: a *structural axis*, tracking coverage of a word's non-evaluative content, and a *valence axis*, tracking its evaluative content specifically. The prosody tradition has read collocational skew almost exclusively onto the valence axis; this paper argues that for a substantial class of frequently studied words the skew belongs on the structural axis instead, and that this misattribution, not any failure to detect real pattern, is the source of the tradition's most persistent unsolved problems. Through analysis of *cause*, *utterly*, *commit*, *set in*, *foster*, and *regime*, alongside the contrasting case of *murder*, the paper develops a falsification procedure — the *contradiction test* — for locating a feature's conventionalization: features admitting structured, unmarked counterexamples are diagnosed as weak; those admitting none, as strong. The six classical words encode strongly conventionalized structural profiles alongside weak valence content, while *murder*'s valence sits strong, fused with its descriptive content. The account answers two problems named in the prosody literature — the introspection problem and the causal-direction problem — and distinguishes its diagnostic claim from gradient absorption accounts, which tolerate exceptions but lack a mechanism for locating them.

Keywords: lexical semantics, connotation, semantic profiles, corpus linguistics, collocation, valence, cancellability, distributional semantics

Profiles, Not Prosodies: A Weighted-Feature Diagnostic for Connotative Meaning in Lexical Semantics

When a corpus linguist reports that *cause* collocates overwhelmingly with *problems*, *damage*, and *death*, or that *commit* collocates overwhelmingly with *crime*, *fraud*, and *suicide*, two questions are typically run together as though they were one. The first is empirical: is the collocational skew real, replicable, and statistically robust? The second is interpretive: what does the skew tell us about the word's meaning? The semantic-prosody tradition inaugurated by Sinclair (1991) and developed by Louw (1993), Stubbs (2001), and Partington (2004) has answered the second question by treating the skew itself as diagnostic of an evaluative component absorbed into the word through repeated exposure to charged contexts — *cause* has "gone negative," in Louw's phrase, in more or less the way a fabric absorbs dye.

This paper accepts the first finding without reservation and takes the second question more seriously than a simple rejection would allow. The skews are real, well-documented, and replicated across corpora, registers, and decades; the corpus linguists who found them found something. Denying that a pattern exists in the frequency data would be an odd position for a theory of meaning to hold, since the existence of stable, learnable pattern in usage is precisely what makes shared linguistic communication possible in the first place. The claim developed here is narrower and, for that reason, harder to dismiss: the pattern is real, but it has been read onto the wrong feature. A word's profile, on the account developed below, is a bundle of several features, only one of which is evaluative valence; frequency data is genuine evidence about that bundle, but a given frequency skew is not automatically evidence about the valence slot specifically. For *cause*, *commit*, *foster*, and the other classical cases examined here, the skew is best read as evidence for a *structural* feature — an *agentive* profile for *cause*, a *volitional-binding* profile for *commit*, a *system-control* profile for *regime* — interacting with an independent fact about discourse: some categories of event are more reportable, more newsworthy, and more likely to be described using these words than others. The skew is real. Its source has been correctly located in the data and incorrectly assigned within the word's own semantic architecture.

This is not a new observation in isolation. A minority tradition within lexical semantics, most notably associated with force-dynamic and aspectual analyses of causative and inchoative predicates (Talmy, 2000; Vendler, 1957), has long argued that verbs like *cause* encode structural rather than evaluative content. Nor is the present paper the first to press on the absorption model itself: Hunston (2007) revisits semantic prosody from within the tradition and grants that the phenomenon is better treated as a matter of degree than as a fixed, binary attachment, allowing that a word's prosodic association can be strong in some contexts and weak or absent in others. That concession matters, and this paper takes it seriously rather than arguing past it. But gradience on a single evaluative dimension is a different move from splitting the dimension itself. A gradient-absorption account can accommodate the existence of *cause a sensation* by saying the negative prosody is simply weak or partially overridden there; what it has no machinery to do is ask whether the exception pattern tracks a specifiable non-evaluative feature with its own independent motivation, because on a single-axis account there is nowhere else for the pattern to be tracked to. The two-axis architecture developed here is offered as what gradience on its own does not supply: a location for the exception to belong to, and a procedure for testing whether it belongs there. What is new in this paper is a principled, falsifiable procedure for adjudicating, word by word and feature by feature, which slot in a word's profile a given piece of frequency evidence actually supports — and a formal notation for representing the difference as a matter of degree along two independent dimensions rather than as a single scalar, however gradiently that scalar is construed.

The stakes of getting this right are not narrowly terminological, and they connect to problems the prosody literature has already identified in its own house without resolving. Stubbs (2001, p. 105) names what he calls "the introspection problem": the robust collocational patterns revealed by corpora frequently surprise native speakers, who intuitively regard words like *cause* as evaluatively neutral. Under an absorption account, in which the frequency pattern simply constitutes the word's evaluative meaning, this is an uncomfortable finding with no principled explanation on offer beyond declaring introspection generally unreliable — an expedient that concedes more than it resolves, since the same speakers' introspective judgments are otherwise treated as reliable evidence throughout the rest of semantic theory. A second, related problem is causal direction: correlation between a word and its evaluatively charged collocates is clear in the corpus record, but establishing that the word absorbed the charge from those collocates, rather than the charge being explained by some third factor, remains, on Stubbs's own admission, theoretically unresolved. Absorption offers no intermediate causal mechanism connecting exposure to meaning-change; it asserts the connection and stops. Both problems dissolve, rather than requiring separate ad hoc treatment, once frequency is treated as multiply ambiguous evidence — capable of supporting a structural-feature attribution as readily as a valence attribution — and once a procedure exists for determining, word by word, which attribution the evidence actually supports.

The remainder of the paper proceeds as follows. Section two develops the theoretical architecture: meaning as a weighted feature bundle, the distinction between the bundle's structural axis and its valence axis as two readouts of a single representational object, the synchronic/diachronic distinction that governs how such bundles change, and the grounding of the framework in meaning as a matter of dispositional usage. Section three introduces the contradiction test as the paper's central diagnostic and situates it relative to gradient formulations of absorption. Section four applies the diagnostic across seven case words, six drawn from the classical semantic-prosody literature and one, *murder*, introduced as a strong-valence contrast. Section five examines diachronic evidence, section six examines the behavior of large language models under this account, and section seven briefly situates the argument relative to the lexical-encyclopedic distinction. Section eight discusses implications and significance, and section nine addresses limitations and outlines a program of empirical validation, including a proposed test of the decorrelation prediction that the case studies motivate but do not, on their own, confirm.

Meaning as a Weighted Profile

The Problem with a Single-Feature Account

The absorption model of semantic prosody, even in its gradient formulations, treats connotative meaning as a single evaluative dimension attached to a word: a word is more or less negative, more or less positive, with the degree of attachment itself allowed to vary by context. This is a parsimonious model, and its parsimony is part of its appeal. It is also, on inspection, empirically strained, because it has difficulty accommodating two facts that are jointly true of every word examined in this paper: a robust majority skew toward one evaluative pole, and a structured, non-marginal minority of felicitous uses at the opposite pole. A single evaluative dimension, however gradiently construed, predicts differences of degree in how strongly a word leans one way — words that are "very negative" versus "mildly negative," or contexts in which the lean is stronger or weaker — but it has no independent mechanism for predicting *where* the exceptions to a skew will fall, why they are unmarked rather than ironic, or why they persist rather than eroding under continued exposure to the majority pattern. An account with only one moving part, however finely graded that part is allowed to be, cannot explain a distribution with two describable regions without treating one of those regions as noise.

The alternative developed here treats a word's meaning not as a single dimension but as a bundle: a finite set of semantic features, each holding some degree of conventionalization. Formally,

$$\text{Profile}(w) = \{(f_1, \alpha_1), (f_2, \alpha_2), \ldots , (f_n, \alpha_n)\}$$

where each f_i is a semantic feature associated with the lexical item w, and each α_i is that feature's *agreement score*: a value indexing the degree to which the feature is fixed, non-cancelable, and convergently represented across speakers of the relevant speech community at a given synchronic moment. A feature whose agreement score sits near the strong end of the scale behaves like a classical entailment — non-negotiable, resistant to contextual override, immediately available to explicit judgment. A feature in the moderate range behaves as a weak entailment: it generates systematic expectations that influence production and comprehension, is detectable in aggregate behavior such as corpus frequency and priming, and is genuinely cancelable by sufficient contextual support in a way that leaves a trace of contextual work having been done. A feature near the weak end of the scale is a different kind of case again: not a defeasible default that gets overridden in some uses, but a weak, largely idiosyncratic association that was never strong enough to require defeat in the first place, present in some speakers or registers without having achieved real conventionalization across the community.

The critical architectural claim of this paper is that evaluative valence is not a separate kind of thing standing outside this bundle, waiting to be absorbed into or subtracted from the "real" meaning. Valence is simply one more feature, f_v, occupying one more slot in Profile(w), with its own α_v. For some words that slot holds a weak value; for others, a moderate value; for still others, a value near the strong end. Nothing about the architecture privileges valence as either present or absent by default — the question of where a given word's valence sits is answered empirically, word by word, using the diagnostic developed in section three.

Two Axes, One Bundle

A word's profile can be interrogated along two distinct axes, and distinguishing them explicitly resolves a confusion that has attended every case study of this kind. The *valence axis* asks a narrow question: what is the agreement score of the single feature f_v, the word's evaluative-polarity feature, considered on its own? The *structural axis* asks a different question of the same bundle: what is the coverage of the word's remaining, non-evaluative features — how completely do they account for the full range of attested and elicitable uses, without requiring a structured class of counterexamples to be set aside? These are not two different kinds of representation competing for the same explanatory role, in the way a categorical lexical-versus-encyclopedic split would propose, and they are not the same move as allowing a single evaluative dimension to vary gradiently by context. They are two different readouts of one flat bundle, in the sense developed further in section seven: the same formal object, Profile(w), answers both questions, differing only in which slot or slots of the bundle the question is directed at.

Keeping the two axes distinct matters because a single frequency statistic — a seventy percent skew toward negative direct objects, for instance — is not, by itself, evidence exclusively about either axis. It is compatible with a strong valence-α (the word genuinely encodes robust evaluative content, and the skew simply reflects that content operating as expected) and equally compatible with a weak valence-α alongside a high-coverage structural feature that happens to intersect unevenly with which events are salient enough to be discussed at all. Distinguishing the axes does not, by itself, tell a reader which explanation is correct for a given word; it only clarifies that the two explanations are genuinely distinct hypotheses requiring separate tests, rather than the same claim stated two ways. The contradiction test developed in section three is precisely the procedure for adjudicating, feature by feature and axis by axis, which reading a given frequency pattern actually supports.

Stated this way, the diagnosis this paper offers of the semantic-prosody tradition can be put precisely. The tradition has not failed to find pattern in the data — finding the pattern is its principal empirical achievement, and nothing in this paper disputes it. What the tradition's standard methodology lacks, even in gradient formulations such as Hunston's, is any procedure for adjudicating which axis a given pattern belongs to. Coding a collocate as evaluatively positive or negative is a comparatively cheap operation, available by inspection; identifying the structural feature — agentivity, volitional binding, gradual onset — that a word's full distribution instantiates requires the kind of force-dynamic and aspectual analysis that corpus methodology, as standardly practiced, was not built to run. The taxonomy through which the field has typically organized collocational relationships makes this default explicit rather than incidental: semantic preference is defined relationally to any definable semantic field a word's collocates might belong to, while semantic prosody, on the standard characterization, "is always either for positive or for negative evaluation" (McEnery & Hardie, 2012, as cited in Ben Ghozlen, 2022, p. 71), leaving no third option within the taxonomy itself for a structural, non-evaluative regularity to be coded onto once semantic preference has been set aside. Allowing that evaluation, so coded, comes in degrees — as a gradient-absorption view does — does not add the missing third option; it only adds gradations within the same two bins. Given only these two coding options, a genuine structural-axis regularity that fails to fit tidily within a single semantic field will be coded, by elimination, onto the valence axis — not because the analyst mistook the data, but because the available categories offered nowhere else for a real pattern to go. This paper's diagnosis is therefore a claim about a methodological default rather than an error confined to any one analyst, one word, or one formulation of absorption, gradient or otherwise: wherever a structural feature's own distribution happens to correlate with a reportability-skewed discourse domain, the valence axis will absorb credit that the structural axis was owed, and it will do so systematically, across independently chosen words, for exactly the reason just given.

Synchrony and Diachrony

A recurring objection to profile-based accounts is that meaning obviously does change over time through use — words drift, narrow, broaden, and acquire new evaluative colorings across centuries, and any theory that denies this contradicts well-established findings in historical linguistics (Traugott & Dasher, 2002; Bybee, 2010). This objection conflates two questions that the weighted-bundle architecture is built to keep separate: what determines a feature's α_i value across historical time, and what determines whether a given token of use, produced now, is felicitous.

The answer to the first question is genuinely usage-based. Over centuries, the accumulated pattern of how a feature is deployed, reinforced, and reinterpreted across a speech community shapes where its α-value settles and how that value shifts. This is conventionalization in the ordinary diachronic sense, and nothing in this paper denies it. The answer to the second question is not usage-based in the same way. At a given synchronic cross-section, a word's profile — the specific set of features and their current α-values — is already fixed, the deposit of that diachronic process rather than an ongoing negotiation with each new token. A speaker producing or comprehending *cause* today is not renegotiating the word's agentive profile against the current sentence; that profile is already settled, and it is what filters which uses will sound felicitous, unmarked, or strange. The direction of explanation therefore runs in opposite directions at the two timescales: diachronically, accumulated usage shapes the profile; synchronically, the already-shaped profile constrains usage. Cooccurrence data collected today is evidence about the current state of the profile, not raw material from which the profile is assembled on the spot.

This distinction directly answers a version of the circularity objection that might otherwise be raised against profile characterizations: if profiles are, in the end, the residue of historical usage, does explaining today's usage by appeal to the profile not simply explain usage by usage? The answer is no, for the same reason predicting tomorrow's weather from today's barometric pressure is not circular even though both quantities belong to "the same" atmospheric system. What is required for non-circular explanation is not that the explanans be drawn from some usage-independent realm — no such realm exists for any theory of lexical meaning, including this one — but that the evidence used to characterize the profile be independent of the specific synchronic distribution the profile is then used to explain. A profile inferred from historical citations spanning centuries prior to, and methodologically distinct from, the collocate-frequency count of a contemporary corpus satisfies this requirement; a profile inferred from, and immediately reapplied to, that same contemporary count would not. Section three develops what independence of this kind requires beyond the historical-source condition alone.

Meaning as Dispositional Usage

The weighted-bundle formalism describes the structure of a lexical profile without yet saying what a profile is a profile of. This paper's answer is functionalist and, in the relevant sense, deflationary: meaning, for a given morpheme or word, simply is the pattern of usage a speech community is disposed to produce and accept — not a separate fact standing behind usage and merely evidenced by it, but the dispositional pattern itself. An agreement score α_i for feature f_i is an operationalization of how convergent that dispositional pattern is across speakers of a community, for that feature, at a given time.

Stated this way, the construct is deliberately agnostic about neural implementation. It does not claim, and does not need to claim, that a given feature corresponds to an identifiable, individuable pattern of neural activity locatable by current or foreseeable neuroimaging. Two speakers instantiate "the same" meaning-fact with respect to a feature just in case their dispositions to use and accept the relevant expression converge in the relevant respects — their inferences, their acceptability judgments, their processing speed under priming, their production choices — regardless of whatever differs at the level of neural microstructure that produces that disposition in each of them.

One clarification is important given this identification of meaning with dispositional usage, because it bears directly on the paper's own reliance on corpus-frequency evidence throughout. Identifying meaning with usage is not the same as identifying meaning with the *actual, historically realized token record* a corpus happens to contain. A corpus samples one path through a much larger space of what a speech community is disposed to produce; a slow year for natural disasters would shift *cause*'s realized collocate frequencies without any change whatsoever in what speakers of English are disposed to accept as felicitous uses of the word. The dispositional pattern is what is stable at a given synchronic moment — fixed, in the sense developed in the preceding section, prior to and independent of which slice of it any particular corpus happens to sample. Frequency counts are evidence about this disposition, in the same sense that a survey is evidence about an underlying population attitude: genuine evidence, subject to sampling variation, and not to be confused with the disposition itself. This is the sense in which the paper's repeated claim that frequency is evidence *of profile* should be read: not evidence of some further fact standing behind usage, but evidence of the dispositional pattern that usage, correctly generalized beyond any one sample, actually consists in.

The relevant empirical tests for a given α_i value follow directly from this identification: acceptability judgments, semantic priming, and comprehension or production asymmetries — precisely the methods proposed in section nine's experimental program — are not indirect proxies for some hidden meaning-fact; they are direct probes of the dispositional pattern that the meaning-fact, on this account, simply is. The corpus-frequency and counterexample evidence marshaled in section four plays the same evidential role, mediated by the diagnostic developed in the next section, which connects observable felicity patterns to the underlying α_i values those felicity patterns are generated by. It bears emphasis at the outset, before that diagnostic is developed, that the paper does not treat its own case-by-case α assignments as measurements in this stronger sense; they are qualitative estimates, ordered relative to one another by the evidence reviewed in section four, and section nine returns to what would be required to convert them into the kind of direct probe just described.

The Contradiction Test: A Diagnostic for Locating Structural and Valence Content

From an Independence Requirement to a Falsification Procedure

One might attempt to defend profile characterizations of the kind sketched in section two by appeal to their evidential source alone: a structural feature such as *cause*'s agentivity is drawn from lexicographic history and force-dynamic analysis, methodologically and temporally independent of the synchronic corpus data the profile is then used to explain. That defense, developed in the preceding section as the independence requirement governing non-circular explanation, is not wrong, but it is incomplete, because it leaves open a question a careful reader will still press: granting that the evidence for a profile is independent of the data it explains, what procedure determines which candidate feature is correct? If the answer is simply "whichever feature the lexicographic tradition happens to associate with the word," the account risks trading one form of unprincipled fit for another, since dictionary glosses are themselves informal and were not written with this diagnostic purpose in mind.

The diagnostic developed here answers that question directly, and it does so in a way that is falsifiable rather than interpretive. Call it the *contradiction test*. Given a candidate feature *f* proposed as part of a word *w*'s strongly conventionalized profile, the test asks: does the attested distribution of *w* contain a structured, non-marginal, unmarked class of uses in which *f* fails to hold? If such a class exists, *f* is disconfirmed as strongly conventionalized content; whatever standing it retains in Profile(*w*) must be weaker, and the burden shifts to explaining the observed skew toward *f*-consistent uses by some mechanism other than *f*'s being definitional. If no such class exists — if every attested and elicitable use of *w* is consistent with *f*, including uses that run counter to *w*'s typical evaluative coloring — then *f* survives as a candidate for strongly conventionalized status, on the same logical footing as any other well-supported entailment.

This is falsification, not curve-fitting, and it is worth being explicit about what keeps it from collapsing into the latter, since the risk is real. A candidate feature proposed after the fact, tailored broadly enough to cover whatever distribution it is asked to explain, would trivially survive the test without having earned survival; a diagnostic that always returns a passing structural feature for any word examined is not a diagnostic at all. The discipline the test imposes against this is that the candidate feature must be motivated independently of the counterexample search that will test it — typically by etymology, by force-dynamic or aspectual analysis of the kind developed by Talmy (2000) and Vendler (1957), or by a lexicographic definition antecedently on record — and only then checked against a distribution the feature was not built to accommodate. When "negative" is proposed as *cause*'s core content and *cause celebration* is found to be unmarked and fully felicitous, "negative" fails the test. When "agentive" is proposed instead, drawing on an independently motivated force-dynamic analysis of causation rather than one custom-fitted to *cause*'s own collocate list, and no structured class of non-agentive uses of *cause* can be located, "agentive" survives. The evidential role of counterexamples is therefore inverted relative to how the semantic-prosody literature has typically treated them. They have usually been treated as noise to be explained away, exceptions tolerated at the margin of an otherwise absorbed prosody. Under the contradiction test they are the primary evidence: a candidate feature's survival is measured precisely by the absence of exactly this kind of structured counterexample, and a feature's failure is diagnosed precisely by their presence.

A threshold for what counts as "structured" and "non-marginal" rather than idiomatic residue cannot be given in the abstract, and this paper does not attempt to legislate one in advance of the case studies. The working standard applied throughout section four is productivity across more than one lexical collocate and more than one register: a single fossilized idiom (*cause for celebration* alone, say, with no further positive uses available) would not by itself disconfirm a candidate valence feature, since idioms are routinely exceptional without unsettling the generalization they depart from. What disconfirms a candidate feature in the cases below is a class that is generative — that extends productively to new collocates a speaker has not simply memorized — and that shows no markedness or contextual strain in doing so. This is a working operationalization rather than a formal one, and section nine returns to what a fully general threshold would require.

Two further points sharpen the test's structure. First, it is a comparative test between two or more candidate features for the same distribution, not an isolated check on one hypothesis. *Cause*'s distribution must be checked against both "negative" and "agentive" to show that the latter dominates the former in coverage; a test run against only one candidate cannot show that a rival does better. Second, the test does not require that a disconfirmed feature vanish from the profile altogether. A feature that fails the strongly-conventionalized test may nonetheless survive at a weaker standing, exactly as valence does for most of the words examined below — present, statistically real, influential on discourse frequency, but not part of the entailed, non-cancelable content of the word.

Distinguishing the Test from Gradient Absorption

The contradiction test needs to be distinguished carefully from a move it might initially be mistaken for: simply allowing that a word's prosody is a matter of degree, as Hunston (2007) argues in revisiting the construct from within the semantic-prosody tradition itself. That revisiting is a genuine advance over the strong absorption claims of Sinclair and Louw, and nothing in this paper is aimed at it as though it were the same target as the earlier, more rigid formulations. A gradient-absorption view can accommodate *cause a sensation* by treating *cause*'s negative prosody as weak, partially suspended, or simply absent in that particular context, without needing to abandon the underlying claim that *cause* leans negative overall. On its own terms, this is not obviously wrong, and the contradiction test as stated so far does not refute it.

What gradience on a single evaluative dimension cannot do, however, is answer a question the contradiction test is specifically built to ask: does the pattern of where the prosody weakens or disappears track a specifiable, independently motivated non-evaluative feature, or is it simply noise around the central evaluative tendency? A gradient-absorption account has no principled way to distinguish these two possibilities, because it has only one dimension along which the word's meaning is represented; wherever the evaluative signal is weak, that is simply recorded as a region of weak signal, with no further question to ask about why weakness clusters exactly where it does. The two-axis architecture developed in section two supplies that further question by giving the exception pattern somewhere specific to go: if the uses in which a word's typical valence fails to hold form a class unified by a describable structural property — agentivity, volitional binding, gradual onset — rather than a scatter of otherwise-unrelated idiosyncratic cases, that unity is itself evidence that a second, independent feature is doing explanatory work the valence dimension alone cannot account for. The contradiction test operationalizes exactly this: it is not merely a test for the presence of exceptions, which gradient absorption already predicts and tolerates, but a test for whether those exceptions are structured by an independently motivated feature. A word that failed this stronger test — whose counterexamples to a proposed valence feature turned out to be an unstructured scatter, explicable by no candidate structural feature motivated independently of the scatter itself — would be better described by gradient absorption than by the account developed here, and the diagnostic would say so. The claim of this paper is not that gradient absorption is false as a general schema; it is that for the words examined in section four, the exceptions are structured in exactly the way gradient absorption has no mechanism to predict, and that a diagnostic capable of detecting that structure is doing more explanatory work than one that merely allows for its existence.

A Second Diagnostic: The Euphemism Signature

The contradiction test locates the strength of a feature by searching for counterexamples within the felicity space of a single word. A complementary diagnostic looks outside that space, at speaker behavior when the word itself is avoided. Strongly conventionalized valence features are predicted to generate euphemistic substitution: when speakers wish to describe an event that would trigger the word's core sense without endorsing the negative valence bound into that sense, they reach for a different word entirely rather than using the original word felicitously in a neutral or positive frame. A state that carries out a killing it wishes to present as justified does not describe itself as having *murdered* anyone; it describes itself as having *executed* them. The absence of a felicitous positive or neutral use of *murder*, together with the presence of a distinct substitute vocabulary reserved for exactly the cases that would otherwise require one, is independent behavioral confirmation that the valence in question is non-cancelable.

Weakly conventionalized valence features generate no comparable substitution pressure, because the valence itself is weak enough that there is nothing for a speaker to route around. Speakers wishing to describe *cause* bringing about a positive outcome do not reach for a different verb; they use *cause* directly, as in *cause for celebration* or *cause a sensation*, because the word's core content was never evaluatively loaded to begin with. It bears emphasis that this is not a case of a normally-negative default being overridden or suspended in these uses; on the profile assigned in section four, there was no substantial negative weight present to suspend, and the felicity of *cause a sensation* is simply what a weak valence feature predicts directly, without any intervening act of contextual defeat. The euphemism test and the contradiction test are therefore expected to converge: wherever the contradiction test finds a structured class of felicitous counterexamples, the euphemism test should find no corresponding substitution pattern, and wherever the contradiction test fails to find such a class, the euphemism test should find one. Section four checks this convergence explicitly for the paper's contrast case, and section nine proposes a systematic check of the same convergence across a larger item set than the present paper undertakes.

Case Studies

The following seven case studies apply the contradiction test, supplemented where relevant by the euphemism test, to the six words most frequently discussed in the semantic-prosody literature together with one additional word, *murder*, chosen specifically because its valence is predicted to sit near the strong end of the scale rather than in the weak-to-moderate range occupied by the other six. Presenting a case in which the diagnostic locates strong valence content alongside six cases in which it locates weak-to-moderate valence content demonstrates that the test is a genuine two-sided diagnostic rather than a procedure tuned to produce a single predetermined outcome. Each case study reports both readouts of the profile distinguished in section two: the structural-axis verdict, concerning the coverage of the word's non-evaluative content, and the valence-axis verdict, concerning the relative strength of the word's evaluative feature specifically.

The profile notations given below use qualitative standing — *strong*, *moderate*, *weak* — rather than point-value estimates. This is a deliberate departure from treating these judgments as measurements. What the contradiction and euphemism tests license, applied to published corpus evidence and to counterexample classes drawn from COCA and the BNC, is an ordering: whether a feature belongs among a word's entailed, non-cancelable content or among its weaker, statistically-visible-but-cancelable content. That ordering is defensible from the evidence marshaled here. A specific numerical value along that ordering is not, and section nine returns to what would be required to earn one.

Case 1: Cause

The verb *cause* is the most extensively documented case in the prosody literature. Stubbs (2001, p. 106–107) reports that in the Bank of English corpus, *cause* collocates predominantly with nouns denoting adverse outcomes — *problems*, *damage*, *concern*, *trouble*, *death* — with Sinclair (1991) estimating that negatively valenced direct objects account for roughly seventy percent of transitive uses. The standard prosodic interpretation treats this skew as evidence that *cause* has absorbed a negative evaluative charge through repeated exposure to harmful contexts — a valence-axis reading of the frequency data, and, on the account developed here, the wrong axis for this particular skew to be read onto.

Applying the contradiction test to the candidate feature "negative valence" as *cause*'s core content immediately produces disconfirmation. Corpus data yield a substantial, unmarked, nonironic class of uses in which *cause* combines with positive or neutral outcomes: *cause a sensation*, *cause excitement*, *cause great joy*, *cause a stir*, and the fully lexicalized *cause for celebration*. These are not rhetorically marked or requiring special contextual licensing; a speaker encountering *the discovery caused great excitement in the scientific community* processes it without any sense of semantic tension, and the pattern extends productively to new positive collocates rather than being confined to a short list of fossilized phrases. "Negative valence" therefore fails the test as a candidate for strongly conventionalized status: valence-axis verdict, weak.

Applying the same test to the candidate feature "agentive" — direct, forceful, immediate bringing-about by an agent or agent-like force — produces no disconfirming class. Every attested use of *cause*, whether the outcome described is harmful, beneficial, or neutral, involves this force-dynamic structure: some entity or event is construed as directly and forcefully responsible for bringing about a subsequent state. *Cause* survives as strongly conventionalized content on this feature — structural-axis verdict, near-total coverage — and Profile(*cause*) can accordingly be represented as {(agentive, strong), (direct, strong), (forceful, strong), (immediate, strong), (negative-valence, weak)}, with the valence weight set weak to reflect the ease, frequency, and total absence of markedness with which positive and neutral uses occur.

The remaining question is what produces the seventy-percent skew if valence is not doing the explanatory work, and the answer is where the structural-axis reading earns its keep. The mechanism proposed here is discourse-reportability asymmetry, attaching not to the valence slot but to the *agentive* feature itself: harmful, forceful impositions by an agent are more newsworthy, more consequential for the hearer's practical interests, and more likely to demand causal explanation in the first place than are beneficial ones — nobody asks what caused the sunrise, but many people ask what caused the crash. This asymmetry is a fact about which topics generate discourse, mediated through the structural feature that makes *cause* the natural word for describing direct forceful imposition regardless of outcome, not a fact about *cause*'s valence content. It predicts exactly the observed pattern: a word whose felicity conditions are evaluatively neutral, sampled disproportionately from a discourse environment that is not, because the discourse environment's asymmetry tracks *agentive imposition*, a category with no built-in polarity, rather than tracking negativity as such.

Case 2: Utterly

The intensifier *utterly* shows a comparable pattern at a smaller scale. Louw (1993) and Sinclair (1991) document its predominant function magnifying negative adjectives — *utterly disastrous*, *utterly appalling*, *utterly ridiculous* — and Partington (2004, p. 145) quantifies seventy-eight percent of *utterly*'s collocates in the Cambridge International Corpus as negatively valenced.

The candidate feature "negative valence" again fails the contradiction test. COCA and BNC data yield a substantial class of *utterly* modifying positive or neutral adjectives — *utterly beautiful*, *utterly compelling*, *utterly delightful*, *utterly charming*, *utterly captivating* — used without irony or the sense of contextual strain that a genuinely encoded negative feature would predict, and the class is productive rather than idiomatic, extending readily to adjectives a speaker has not previously encountered paired with *utterly*. The candidate feature "intensive" — denoting completeness, totality, or an extreme degree, etymologically continuous with *utter* in the sense of complete or absolute — survives the test without disconfirmation: every attested use, regardless of the polarity of the adjective modified, involves magnification to an extreme or total degree.

Utterly's valence weight, while still weak relative to its intensive core, is plausibly somewhat less weak than *cause*'s. The asymmetry documented by Pomerantz (1984) in conversation-analytic work on assessment sequences — upgrades to negative evaluations are pragmatically less costly and more common than upgrades to positive ones, since intensified praise risks sounding disproportionate or insincere in a way intensified complaint does not — suggests a genuine, if modest, processing cost attaching to *utterly*'s positive uses that *cause*'s positive uses do not carry. Profile(*utterly*) is therefore represented as {(intensive, strong), (totality, strong), (extremity, strong), (negative-valence, weak-to-moderate)}, with the comparison to *cause* flagged as provisional pending the acceptability-rating study proposed in section nine, which would compare processing costs for *utterly* plus positive versus negative adjectives directly rather than relying on the qualitative contrast drawn here.

Case 3: Commit

Commit is frequently cited for its high-frequency association with transgression: *commit a crime*, *commit murder*, *commit suicide*, *commit fraud*. Stubbs (2001, p. 109) reports that in the Bank of English, over sixty percent of transitive uses take direct objects denoting illicit, harmful, or transgressive acts.

The candidate feature "negative valence" fails the contradiction test decisively. *Commit*'s neutral and positive uses — *commit to a relationship*, *commit resources to a project*, *commit oneself to a goal*, *commit funds to research*, *commit troops to battle*, and the institutional senses *commit a patient to care* or *commit a defendant to trial* — are not marginal exceptions but, as their range across ordinary, legal, and institutional registers indicates, central to the verb's polysemy and productive across each of those registers independently. The candidate feature "volitional-binding" — drawing on the force-dynamic analysis of self-compulsion and constraint developed by Talmy (2000) — survives without disconfirmation: every attested use, transgressive or otherwise, involves an agent's deliberate, weighty, and not-easily-reversed binding of themselves or something else to a course of action.

The skew toward transgressive objects is explained, as with *cause*, by reportability attaching to the structural feature rather than to valence: binding oneself to a norm-violating act is a highly salient, socially consequential, and narratively reportable event type, in a way that binding oneself to a mundane obligation typically is not. Profile(*commit*) is represented as {(volitional, strong), (binding, strong), (consequential, strong), (negative-valence, weak)}, comparable to *cause*'s valence weight and for the same reason: the counterexample class is broad, unmarked, and central rather than peripheral.

Case 4: Set In

The phrasal verb *set in* is, in Louw's (1993, p. 163) description, "unfailingly negative," collocating with *rot*, *decay*, *despair*, and *boredom*. The candidate feature "negative valence" is disconfirmed by a documented class of neutral and positive counterexamples: *a sense of peace finally set in*, *healing set in*, *a comfortable routine set in*, *after the chaos, calm set in* (Davies, 2008–). The candidate feature "gradual-onset" — encoding inchoative beginning combined with entrenchment, persistence, and difficulty of reversal, aligning with Vendler's (1957) category of achievements transitioning into states — survives: every attested use, whether the resulting state is welcome or unwelcome, describes a state establishing itself gradually and persistently rather than abruptly.

Set in's valence weight is harder to place with equal confidence relative to *cause*'s and *commit*'s, for a reason worth registering rather than glossing over: the documented counterexamples (*peace set in*, *calm set in*) are drawn from a narrower set of established collocations than the correspondingly broad and evidently productive counterexample classes found for *cause*, *foster*, and *commit*. Whether *peace/calm/routine set in* represents a fully live, generatively available pattern or a smaller set of semi-conventionalized phrases is an empirical question this paper cannot settle without the corpus-productivity check proposed in section nine.

Provisionally, Profile(*set in*) is represented as {(gradual-onset, strong), (persistent, strong), (entrenched, strong), (negative-valence, weak-to-moderate)}, with the valence standing flagged as the least secure of the six classical cases and the productivity of its counterexample class specifically marked as needing direct corpus verification rather than the illustrative examples given here.

Case 5: Foster

Foster offers the positive-pole mirror of *cause*. Stubbs (2001, p. 110) documents its frequent collocation with desirable outcomes — *foster growth*, *foster development*, *foster understanding*, *foster cooperation* — with roughly eighty percent of its direct objects in the Bank of English positively valenced.

The candidate feature "positive valence" fails the contradiction test on evidence just as strong as *cause*'s. *Foster resentment*, *foster dependence*, *foster inequality*, *foster corruption*, *foster distrust*, *foster complacency* (Davies, 2008–) are, again, not ironic or marked uses but straightforward descriptions of the gradual nurturing of a negative outcome; *policies that foster inequality* describes exactly what it says without any sense of contradiction, and the pattern extends productively to new negative outcomes rather than being limited to a fixed list. The candidate feature "facilitative" — an agent or environment providing conditions conducive to gradual, indirect development, contrasting systematically with *cause*'s direct, forceful, immediate imposition — survives without disconfirmation across the full range of attested uses.

The positive skew is explained by the same domain-asymmetry mechanism operating in reverse, and attaching, as with *cause*, to the structural rather than the valence feature: contexts of cultivated, gradually-supported development — child-rearing, education, economic growth, social cohesion — are disproportionately positive in the events that recruit this kind of causation, without this constituting any semantic requirement that the outcome be positive. Profile(*foster*) is represented as {(facilitative, strong), (gradual, strong), (nurturing, strong), (indirect, strong), (positive-valence, weak)}, structurally the mirror image of *cause*'s bundle, as the earlier force-dynamic contrast between the two verbs predicts.

Case 6: Regime

Regime is, as Channell (2000, p. 46) documents, associated with collocates like *military*, *communist*, *Nazi*, *Soviet*, *fascist* — a pattern the standard prosodic account reads as absorbed pejoration.

The candidate feature "negative valence" fails the contradiction test, though the counterexample class here is drawn from a register shift rather than from within the political domain itself: *a new training regime*, *a regime of physical therapy*, *a regulatory regime for data privacy*, *a dietary regime*, *the regime of international law*. In these uses *regime* denotes a systematic plan or set of rules with no negative evaluation whatsoever, and the uses are entirely unmarked in medical, athletic, and administrative registers. The OED's earliest citations of the word in English, from 1792 onward, confirm that this system-control sense was present from the point of borrowing from French *régime* and was not restricted to authoritarian government at any point in its history; the association with fascist, communist, and military dictatorships intensifies specifically in twentieth-century political usage, alongside, not in place of, the persisting neutral administrative and medical senses. The candidate feature "system-control" — formal, structured, top-down administration — survives across both registers.

Regime differs from the other five cases in one important respect that the diagnostic itself brings into focus rather than obscures: its valence weight is measurably higher than the near-floor weight found for *cause*, *commit*, and *foster*, though not near the strong end of the scale occupied by *murder*, examined next. This is consistent with the counterexample evidence: the neutral uses of *regime* are real and unmarked, but they are concentrated in specific registers —

medical, athletic, regulatory — rather than distributed evenly across contexts the way *cause*'s positive uses are, and *regime*'s political-register uses show a stronger, more consistently negative default than *cause*'s default in any register. Profile(*regime*) is represented as {(system-control, strong), (institutional, moderate-to-strong), (negative-valence, moderate)}, explicitly the strongest valence standing among the six classical cases and the one for which the diachronic evidence discussed in section five is most directly relevant. This moderate standing is itself a qualitative judgment drawn from the register-concentration pattern just described, not a measured quantity, and section nine proposes the priming methodology that would be needed to test it directly.

Regime is also the case that most directly tests the general shape of the claim being made about valence's contribution to skew. For *cause*, *commit*, and *foster*, the claim is that valence contributes negligibly to observed frequency skew because valence standing is weak; there is, on this account, almost nothing for valence to contribute. The stronger and more general form of the claim is not that valence never contributes, but that whatever contribution it makes should itself scale with its own standing — negligible when that standing is weak, and potentially non-negligible as it rises toward the moderate range *regime* occupies. *Regime*'s political-register skew is accordingly not predicted to be purely a structural-axis phenomenon in the way *cause*'s is; some portion of its political-context frequency pattern plausibly reflects the word's own moderate valence weight operating directly, alongside the system-control feature's own reportability-linked skew. Distinguishing how much of *regime*'s political-register skew is attributable to each axis is a further empirical question the priming methodology proposed in section nine is suited to, since priming effects tied specifically to the valence feature should scale with its own standing in a way structural-axis effects need not.

Case 7: Murder — A Strong-Valence Contrast

The preceding six cases share a common outcome: a candidate valence feature is proposed, disconfirmed or substantially weakened by a structured counterexample class, and replaced by a structural feature that survives the test at strong standing. This pattern could, in principle, be an artifact of case selection rather than a general property of the diagnostic — a concern worth taking seriously rather than dismissing. The purpose of this seventh case is to show that the same diagnostic, applied to a word chosen specifically because its valence is predicted to be strong rather than weak, returns the opposite verdict.

Murder denotes an unlawful, morally condemned killing. Applying the contradiction test to the candidate feature "wrongful/condemned" as part of *murder*'s core profile, the search for a structured, unmarked counterexample class comes back empty. There is no register, formal or informal, legal or colloquial, in which *murdered someone kindly*, *murdered him for a good cause*, or *a justified murder* functions as an ordinary, unmarked description of a killing the speaker regards as warranted. Where English needs to describe a killing without condemning it — a lawful killing in war, a state execution, a killing in self-defense — it reaches for a distinct vocabulary: *executed*, *killed in action*, *killed in self-defense*. This is precisely the euphemism signature predicted in section three for a strongly conventionalized valence feature: rather than *murder* being used felicitously in a neutral or approving frame, speakers route around the word entirely, deploying dedicated substitute vocabulary reserved for exactly the cases that would otherwise force a contradiction.

This case therefore behaves oppositely from the other six at every diagnostic point that matters. The contradiction test finds no disconfirming counterexample class for "wrongful" as strongly conventionalized content. The euphemism test finds active substitution rather than direct felicitous override. Profile(*murder*) is accordingly represented as {(killing, strong), (intentional, strong), (negative-valence, strong)}, with the valence feature sitting alongside, not beneath, the word's descriptive content. The contrast with *cause*'s profile — where the analogous slot is weak — demonstrates that the diagnostic is genuinely two-sided: it is not a procedure that relabels every candidate valence feature as weak by construction, but one that returns strong standing when the evidence supports it and weak standing when it does not.

A Forward Prediction: Decorrelation Between Skew and Valence Standing

The seven case studies motivate a further prediction that has not, to this point, been stated explicitly, and it follows directly from treating the six classical words' skews as structural-axis rather than valence-axis phenomena. If discourse-reportability asymmetry attaches to a word's structural feature rather than to its valence feature, then the magnitude of a word's collocational skew should not track the strength of its valence standing across the six classical cases; the two quantities are generated by different mechanisms and have no reason to move together. If, by contrast, the classical prosodic account were correct and skew magnitude is directly generated by valence strength, skew and valence should correlate closely: words assigned a weaker valence standing should show a correspondingly weaker skew.

This paper does not attempt to confirm that prediction against its own case studies, and it is worth being explicit about why. The valence standings assigned in the case studies above are qualitative judgments, drawn in part from the same breadth-of-counterexample-class evidence that also informs intuitions about a word's overall profile; checking a decorrelation prediction against estimates calibrated by exactly the reasoning the prediction is meant to test would not constitute an independent confirmation, whatever its numerical appearance might suggest. What can be said at this stage is only the qualitative observation that four of the six classical cases — *cause*, *utterly*, *commit*, and *foster* — cluster together at weak-to-moderate valence standing while showing skew magnitudes that are not obviously ordered by that clustering, and that *regime* combines a distinctly stronger valence standing with a distinctly more register-concentrated skew pattern than the other five. That combination is consistent with the prediction. It is not a test of it.

A genuine test requires valence standing measured independently of the counterexample-breadth judgments used throughout section four — by the acceptability-rating and priming methodology proposed in section nine, which would derive a valence estimate for each word from processing data rather than from the analyst's own reading of its collocational range, and could then be checked against skew magnitude without the circularity risk just described. Until that methodology has been applied, the decorrelation prediction is offered as exactly that: a prediction, stated precisely enough to be tested, rather than a result.

Diachronic Evidence

The synchronic case studies in section four establish that, at the present moment, structural features rather than valence features constitute the strongly conventionalized core of each word's profile. A separate question is whether this architecture is stable across historical time or an artifact of contemporary usage. If a word's structural profile is genuinely foundational while its valence weight is a comparatively recent and shallow overlay, the historical record should show the structural feature persisting across centuries of use while the valence weight fluctuates, appears, strengthens, or weakens independently. If, conversely, the two were bundled together as a single absorbed prosody from the outset, no such dissociation should be observable.

The historical record supports dissociation for every case examined. *Cause* has been attested with the same agentive, forceful, direct force-dynamic structure across six centuries of English, over a period during which the specific events most commonly described using the verb have shifted dramatically — from disease and famine in earlier periods to industrial accidents, and more recently to phenomena such as economic stress and climate change. Across this entire span, positive and neutral uses of *cause* persist at every period for which citations survive; the proportion of negative to positive uses fluctuates with what is topical in a given era, but neither pole is ever rendered infelicitous, which a genuine absorption account — predicting that accumulated negative association should progressively crowd out or mark positive uses — would not predict.

Regime offers the clearest case of a valence weight changing independently of a stable structural core, and is for that reason the most diachronically interesting of the six. The OED's earliest English citations, from 1792 onward, apply *regime* to a range of governmental systems — *the ancient regime of France*, *the present regime*, *under the regime of Louis XIV* — without inherent negative evaluation; the system-control profile is present from the point of borrowing. The word's association with strongly negatively-evaluated government emerges primarily in twentieth-century usage, concentrated on fascist, communist, and military dictatorships, a shift plausibly tied to the geopolitical prominence of exactly those regime types in the discourse of the period rather than to any change in the word's structural content. Non-political uses — medical and scientific regimes of treatment, diet, and exercise — persist throughout the nineteenth and twentieth centuries without interruption, continuing to license the system-control profile across domains even as the political-register valence weight was, on the account developed here, rising. This is the diachronic signature the weighted-bundle architecture predicts and a single-scalar absorption account, gradient or otherwise, does not: one feature of the bundle (system-control) essentially flat across two centuries, another feature of the same bundle (negative-valence) shifting substantially within a subset of registers over the same period, with no requirement that the two move together.

The general pattern across all six classical cases is that historical stability of structural profile coexists with historical variability of collocational and evaluative pattern. This is not, contrary to a natural first reading, evidence that is merely correlational and silent on causal direction. The specific prediction the weighted-bundle account makes — and a single-dimension absorption account, however gradiently construed, does not — is that the strongly conventionalized components of a profile should show markedly greater diachronic stability than the weakly conventionalized components, because conventionalization strength is precisely a measure of how resistant a feature is to revision by ongoing usage at a given synchronic slice. An account that represents all of a word's evaluative content as variation along a single dimension has no principled basis for predicting that some parts of a word's meaning should be far more stable than others; the weighted-bundle account predicts this differential stability directly from its formal architecture, and the historical record for *cause* and *regime* is consistent with it.

The dissociation documented for *regime* also bears on the forward prediction stated at the close of section four. *Regime*'s valence standing plausibly rose over the twentieth century within the political register specifically, while its skew pattern became correspondingly more register-concentrated over the same period — a joint historical movement that the account anticipates, because it treats a rising valence weight as capable of contributing non-negligibly to skew once that weight leaves the near-floor range the other five classical cases occupy. This historical pattern is offered in the same spirit as the synchronic observation at the close of section four: consistent with the account, and suggestive, but not a substitute for the independently measured valence estimates that section nine's proposed methodology would supply. The historical record and the synchronic case studies are, on this reading, two bodies of evidence pointing the same direction rather than one confirming the other.

Language Models and the Limits of Distributional Learning

Large language models trained on corpus data provide a further source of evidence bearing on the profile-versus-absorption question, because their behavior reveals what can and cannot be learned from distributional exposure alone. Models trained predominantly on political and historical text containing *regime* nonetheless generalize appropriately to *exercise regime*, *dietary regime*, and *drug regime* in medical and athletic contexts, extending the word correctly to registers in which its valence association is largely absent. This transfer is predicted directly by a system-control profile that is register-general by nature; it would be considerably harder to explain under an account in which the model had simply memorized a context-specific negative association restricted to the political domain in which the bulk of its exposure occurred. Similarly, models readily produce both *policies that cause harm* and *discoveries that cause excitement*, and both *policies that foster growth* and *policies that foster inequality*, despite skewed training distributions for each verb, suggesting that what is extracted from the distributional signal is closer to abstract force-dynamic or facilitative structure than to a fixed evaluative polarity.

This capacity for flexible, profile-consistent generalization does not, on its own, establish that a model has represented anything resembling a discrete feature bundle in the sense developed here, and this paper does not overclaim on that point. An alternative explanation remains available: a model exposed to a sufficiently large and varied corpus may simply have encountered enough counterexamples during training to represent the full empirical distribution, including its minority classes, without extracting any underlying structural generalization at all — in which case its flexible output would reflect rich memorization of distributional variance rather than genuine feature abstraction. Distinguishing these two explanations is, in the end, an empirical question about model internals that behavioral output alone cannot settle, and the present paper's argument for the weighted-bundle account of human lexical competence does not depend on resolving it. What the language-model evidence does establish, more modestly, is that the distributional signal available in a large corpus is informationally rich enough to support profile-consistent generalization, whichever mechanism a given model uses to exploit that richness — which is itself a point against treating collocational skew as containing only frequency information about evaluative polarity and nothing more.

A further question, more consequential and more speculative, concerns what distributional learning of any scale should be expected to achieve in principle, if meaning is a matter of dispositional convergence across a speech community rather than a fact fully recoverable from any single sampled trace of that disposition, however large. That question is not pursued further here; it belongs to a separate line of argument concerning the relationship between distributional learning and the structural preconditions for reliable assertion, and the present paper's claims about human lexical competence do not depend on how that further question is resolved. The evidence reviewed in this section supports the narrower point directly relevant here: whatever a model's internal mechanism, its output is more consistent with sensitivity to structural, non-evaluative profile content than with a picture in which collocational skew encodes only evaluative polarity.

The Lexical-Encyclopedic Consequence

The weighted-bundle architecture developed in this paper carries an implication for a long-standing dispute in lexical semantics: the boundary between properly lexical meaning and merely encyclopedic world knowledge. Traditional accounts have generally required this boundary to be drawn categorically, with dictionary-style semantic content on one side and pragmatically-recruited background knowledge on the other. Cognitive-linguistic traditions extending from Langacker (1987) and Fillmore (1982) have long argued that this categorical boundary is unmotivated, that lexical meaning is inherently encyclopedic, and that frame-semantic structures spanning both domains are required even for apparently simple lexical items.

The present account is not neutral with respect to this dispute, though it arrives at a similar conclusion by a different route, and the two-axis distinction developed in section two is precisely what keeps that conclusion from collapsing into an unprincipled one. If a word's meaning is a weighted bundle in which structural features and valence features occupy the same representational architecture, differing only in their degree of conventionalization rather than in their kind, then the traditional lexical-encyclopedic boundary is not so much refuted as rendered the wrong sort of distinction to draw. A feature like *regime*'s political-register valence — arguably closer to what a categorical account would call encyclopedic knowledge about which governments tend to be evaluated negatively — sits in the same bundle, subject to the same formal apparatus and the same diagnostic procedure, as a feature like *regime*'s system-control content, which a categorical account would comfortably call lexical. What distinguishes them is not their category but their degree of conventionalization, and the structural axis and valence axis introduced in section two are, correspondingly, not two different kinds of representation but two different questions that can be asked of one flat bundle. This is the sense in which the axis distinction and the lexical-encyclopedic consequence are the same architectural claim viewed from two directions: treating valence and structural content as differing only in degree is what licenses asking of a single word's profile both how strong its evaluative feature is and how complete its structural coverage is, without thereby positing two separate representational systems that the categorical tradition would need justified independently.

This paper does not develop the full consequences of that reframing, which belong to separate theoretical work; it registers the consequence because the case studies above depend on it. Treating *cause*'s agentive structure and its valence weight within a single formal object is only a coherent move if the lexical-encyclopedic boundary is gradient rather than categorical to begin with. It is worth noting, in closing this section, that the misattribution diagnosed throughout this paper is itself a symptom of assuming the categorical boundary rather than the gradient one: a methodology that expects lexical content and evaluative coloring to belong to two different kinds of thing has, correspondingly, no natural place to locate a feature that is structural in character yet happens to correlate with an evaluatively skewed discourse domain — and will default to sorting it into whichever of the two available categories seems closest, typically the evaluative one, since evaluative coding from collocate inspection is the more readily available operation. The gradient architecture removes the forced choice, and with it the systematic pressure toward misattribution that a categorical architecture creates.

Implications and Significance

Lexicography

Current dictionary practice typically marks a word as carrying a usage label — *derogatory*, *informal*, *literary* — as a binary annotation attached to a sense, with no indication of how strong or how cancelable the labeled association is, and with no vocabulary at all for distinguishing a genuinely lexical structural regularity from an evaluative overlay riding on top of it. A weighted-bundle entry would instead record, alongside each sense, the specific features contributing to that sense's felicity conditions and an approximate standing for each, distinguishing a weakly conventionalized association like *cause*'s residual negative valence from a strongly conventionalized one like *murder*'s, and distinguishing both from the structural features — *agentive*, *volitional-binding*, *system-control* — that a purely evaluative labeling scheme has no slot for at all. This is a more demanding standard of lexicographic annotation, but it is also a more falsifiable one, since it generates specific predictions — about which uses should sound felicitous, which should require euphemistic substitution, and which should sound merely unusual — that a binary label does not. The significance of this extends beyond tidiness of presentation: a dictionary built on the two-axis architecture would be making an empirical claim, checkable against acceptability judgments, every time it assigned a standing to a feature, in a way a derogatory-or-not label never was.

Natural Language Processing

The two diagnostics developed in section three offer a concrete evaluation procedure independent of raw collocational frequency, and the forward prediction developed in section four sharpens that procedure into something a system's internal representations, not merely its outputs, could in principle be tested against. A system intended to model connotative meaning could be tested directly against the contradiction test's predictions: does the system correctly rate *cause celebration* as felicitous and unmarked while correctly rating a hypothetical *murdered kindly* as infelicitous, rather than treating both as equally low-frequency and therefore equally anomalous constructions? Frequency-based training signals alone conflate these two cases, since both are statistically rare in a training corpus; a system that has captured the underlying structural-versus-valence distinction should not. Beyond output-level testing, the structural-axis and valence-axis distinction suggests a more targeted probing methodology than has typically been applied to connotation in language models: if a model has learned something resembling separate structural and valence readouts of a word's representation, targeted interventions on one axis — for instance, ablating or amplifying whatever internal direction correlates most strongly with evaluative polarity judgments — should leave structural-axis behavior, such as felicitous extension to novel agentive contexts, comparatively undisturbed. A model that has instead learned a single undifferentiated evaluative-cum-structural representation for words like *cause* would not be expected to show this dissociation. This is a more demanding and more informative test than accuracy on a held-out sentiment-classification benchmark, and it follows directly from treating the two axes as genuinely separable readouts of a shared representation rather than as a single scalar, gradient or otherwise.

Language Pedagogy

The reframing has a direct pedagogical benefit. Presenting *cause* and *foster* to learners as near-synonyms distinguished by fixed polarity — *cause* for bad things, *foster* for good things — produces confusion the first time a learner encounters *cause excitement* or *foster dependence*. A profile-based explanation — *cause* implies direct, forceful bringing-about; *foster* implies gradual, indirect enabling; English discourse happens to frame harmful events as forcefully imposed and beneficial developments as gradually nurtured more often than the reverse, but the words themselves are evaluatively neutral at their structural core — equips learners for both the prototypical and the creative uses they will encounter, and generalizes to the diagnostic itself as a tool learners can apply to new items: does a proposed reading survive counterexample, or does it require euphemistic substitution. A curriculum built around this distinction would teach the two-axis question directly — what does this word describe, independent of what English discourse tends to say about it — rather than teaching a list of polarity associations that learners must later unlearn piecemeal whenever a counter-skewed use is encountered.

Beyond Polarity-Contrastive Vocabulary

The most consequential implication of treating reportability asymmetry as attaching to structural features rather than to valence is one this paper has so far only gestured toward: the mechanism predicts comparable collocational skew for words that carry no evaluative polarity at all, wherever their structural feature happens to intersect unevenly with what is discourse-salient. Consider a neutral technical verb whose structural content picks out some category of event that is disproportionately newsworthy for reasons entirely unconnected to any evaluative coloring — a verb of quantification, say, whose objects skew toward large or unusual magnitudes simply because large or unusual magnitudes are what get remarked upon, with no plausible positive or negative valence available to misattribute the resulting skew to in the first place. If such cases show collocational skew of a magnitude comparable to *cause*'s or *foster*'s, driven by the same reportability logic but with no polarity story even available as a rival explanation, that would constitute independent confirmation of the mechanism uncontaminated by the interpretive ambiguity that has attended every polarity-contrastive case examined in this paper. This is offered here as a further falsifiable prediction rather than a demonstrated result — no such case has been worked through with the same care as the seven examined above — and it is proposed as a natural next step precisely because it would test the account's central mechanism in a setting where the valence-axis misattribution this paper diagnoses could not, even in principle, arise. Section nine lists this among the empirical priorities the account's central mechanism most needs.

Limitations and Future Directions

This paper is a theoretical reanalysis grounded in published corpus findings, lexicographic evidence, and the two diagnostics developed here; it does not report original quantitative corpus queries, significance testing, or systematic sampling beyond the seven words examined. Several limitations follow directly from this scope and are registered rather than minimized.

First, and most centrally, the qualitative standings assigned throughout section four — weak, moderate, strong — are estimates calibrated against the strength and breadth of the documented counterexample classes, not measurements derived from a dedicated study. The relative ordering among the six classical cases — *regime*'s valence standing clearly exceeding that of *cause*, *foster*, and *commit* — is better supported by the available evidence than any finer-grained ranking would be, and this paper has deliberately avoided assigning point-value estimates for exactly that reason. The forward prediction stated at the close of section four is, correspondingly, a prediction and not a confirmed result: checking it properly requires valence standings derived independently of the counterexample-breadth judgments used to motivate the case studies in the first place, which only the acceptability and priming methodology outlined below can supply. Readers should not treat the qualitative consistency noted in sections four and five as evidence beyond what it is — an observation that the case studies do not obviously contradict the prediction, not a test that the prediction has survived.

Second, the seven-word case set, while now spanning from weakly to strongly conventionalized valence content, remains a small and non-randomly-sampled set. Demonstrating that the diagnostic returns coherent verdicts at both ends of the scale for seven well-chosen words is not equivalent to demonstrating that the profile-based account generalizes across the lexicon; the appropriate scope for this paper's claim is that it establishes and motivates a diagnostic procedure on a representative range of cases, with systematic application across a larger, more systematically sampled item set left as a distinct empirical undertaking.

Third, the contradiction test's discipline against unprincipled feature-fitting, discussed in section three, depends on candidate features being motivated independently of the distribution they are tested against, and on a workable threshold for what counts as a structured, non-marginal counterexample class rather than idiomatic residue. This paper has aimed to satisfy the first condition throughout, drawing candidate structural features from established force-dynamic and aspectual analyses predating and independent of this paper's own case selection, and has offered a working standard — productivity across more than one collocate and more than one register — for the second. Neither condition, however, is formalized beyond the case-by-case judgment applied in section four, and a systematic account of what independently counts as an admissible candidate feature, as opposed to a post hoc fit dressed in etymological language, together with a quantitative threshold for structured productivity, is not fully worked out here. This is arguably the most consequential item on this list, since the contradiction test's evidential force depends on it, and a fuller treatment belongs to future methodological work rather than to the case-by-case application undertaken in this paper.

A program of six studies would substantially strengthen the empirical foundation laid out here, and is proposed rather than executed within the scope of the present paper. An acceptability-judgment study would have participants rate the naturalness of profile-matching but low-frequency expressions (*foster resentment*, *cause celebration*, *utterly beautiful*) against high-frequency expressions and profile-mismatching controls, predicting comparable ratings for the former pair against frequency-based accounts predicting a substantial gap; the same design would test the finer-grained prediction that partial profile matches receive intermediate ratings, as when *foster apathy*, involving gradual development, should outrate *foster rage*, whose suddenness mismatches the facilitative profile's gradual-onset component. This study is also the most direct route to converting the qualitative standings used throughout this paper into independently measured values, and should be treated as a precondition for any future paper that wishes to state the forward prediction of section four as a confirmed rather than a projected result. A semantic-priming study would test whether activating a profile feature facilitates lexical decision on the target word regardless of the evaluative polarity of the priming context, predicting facilitation for both a vignette describing forceful positive imposition and one describing forceful harm when priming *cause*, contrary to an absorption account predicting interference for the mismatched-polarity case; the same study is the natural venue for disentangling *regime*'s dual structural and valence contributions to skew, flagged as unresolved in section four, by isolating priming effects specific to the valence feature from those specific to the system-control feature. A cross-linguistic study would test whether profile-pattern relationships replicate for semantically parallel causative and aspectual predicates in typologically diverse languages, predicting that more agentive causative constructions show negative skews comparable to *cause* regardless of the specific language's discourse conventions. A computational-modeling study would test whether the profile-consistent generalization documented in section six reflects genuine feature abstraction or rich distributional memorization, using targeted probing methods on model internals — along the lines sketched in section eight's discussion of axis-specific ablation — rather than output behavior alone. A scale-invariance study would partition a large, freely available corpus into progressively smaller subsets and test whether profile-relevant collocational rankings remain stable across scale, predicting that this stability should hold more strongly for strongly conventionalized features than for weakly conventionalized ones, since genuinely fixed lexical content should stabilize at small sample sizes while variable, register-dependent valence weights should remain comparatively unstable even at scale. A sixth study, motivated directly by section eight's discussion of vocabulary beyond polarity contrast, would identify a set of evaluatively neutral technical or quantificational verbs with no plausible valence dimension at all, and test whether their collocational distributions show reportability-driven skew of comparable magnitude to the polarity-contrastive cases examined here — a test capable, in principle, of confirming or disconfirming the structural-axis mechanism in a setting where valence-axis misattribution could not arise even as a rival interpretation.

None of these six studies has been run as part of the present paper; each is offered as a specific, falsifiable theoretical commitment, on the view that a framework which states its predictions in advance of the resources needed to test them has done more useful theoretical work than one that withholds prediction until testing is feasible.

Finally, the account of meaning as dispositional usage developed in section two invites a clarification this paper should preempt rather than leave implicit. Identifying meaning with a speech community's disposition to use and accept an expression, rather than with the specific historical record of tokens a corpus happens to contain, is not a retreat into some further fact standing behind usage; it is a claim about what usage, properly generalized beyond any one sample, actually consists in. The qualitative standings reported throughout this paper should be read in that light: not as claims about a hidden meaning-fact for which usage is mere evidence, nor as claims about neural substrate, and not yet as the measured quantities the formal notation of section two anticipates — but as ordered judgments about the degree of cross-speaker convergence a given feature plausibly exhibits, measurable in principle by the acceptability-judgment and priming methods proposed above, and only imperfectly sampled by any single corpus, including the corpora this paper has itself relied on throughout.

Conclusion

The central finding of this paper is that the collocational skews documented across three decades of semantic-prosody research are real, replicable, and correctly measured, but have been persistently misattributed. Where the field has typically read a skew toward negative contexts as evidence that a word has absorbed negative evaluative content through repeated exposure — a reading that survives even in gradient formulations of absorption, which allow the association to vary in strength but still locate it on a single evaluative dimension — this paper has argued that the skew more often reflects a discourse-level asymmetry attaching to a structural, evaluatively neutral feature already present in the word's strongly conventionalized profile. This misattribution is not an occasional analytic slip but a predictable consequence of a methodology and a taxonomy that, as documented in section two, offers evaluative polarity as the only available category for a pattern that does not fit neatly within a single semantic field, and that gradience within that category does not by itself supply the missing alternative. The contradiction test developed here provides a falsifiable procedure for adjudicating, case by case, whether a candidate feature belongs to that strongly conventionalized structural core or to a separate, weaker layer within the same profile bundle; the euphemism test provides an independent behavioral check on the same distinction; and the forward prediction stated at the close of section four offers a further test of the misattribution diagnosis, one this paper states precisely without claiming to have confirmed it. Applied across *cause*, *utterly*, *commit*, *set in*, *foster*, and *regime*, the diagnostic consistently locates valence at weak-to-moderate standing beneath a strongly conventionalized structural core; applied to *murder*, chosen specifically to test the diagnostic's other pole, it locates valence at strong standing, fused with rather than separable from the word's descriptive content.

Treating valence as one weighted feature among others within a single formal profile, rather than as a separate absorbed layer standing outside the lexicon proper, does more than resolve the circularity that has attended earlier attempts to argue for profile-based explanations of these phenomena. It supplies a uniform representational format — $\text{Profile}(w) = \{(f_1, \alpha_1), \ldots, (f_n, \alpha_n)\}$ — capable of stating precisely what a given word's meaning consists in, at a given synchronic moment, in terms that are simultaneously historically grounded, behaviorally testable, and formally explicit, while distinguishing two independent questions — how complete is the structural coverage, and how strong is the evaluative weight — that a single collocational-frequency statistic has too often been asked to answer at once, and that a single evaluative dimension, however gradiently construed, still cannot ask separately. The account developed in this paper is offered as a template for that kind of precision applied to connotative meaning specifically, and as a diagnosis of a specific, nameable error — reading structural-axis evidence onto the valence axis by taxonomic default — that a purely case-by-case correction would have left unexplained as a recurring pattern across six independently studied words.

References

Ben Ghozlen, B. (2022). A critical overview of semantic prosody. *Legon Journal of the Humanities*, *33*(2), 59–82.

Bybee, J. (2010). *Language, usage and cognition*. Cambridge University Press.

Channell, J. (2000). Corpus-based analysis of evaluative lexis. In S. Hunston & G. Thompson (Eds.), *Evaluation in text: Authorial stance and the construction of discourse* (pp. 38–55). Oxford University Press.

Davies, M. (2008–). *The Corpus of Contemporary American English (COCA)*.

Fillmore, C. J. (1982). Frame semantics. In *Linguistics in the morning calm* (pp. 111–137). Hanshin.

Hunston, S. (2007). Semantic prosody revisited. *International Journal of Corpus Linguistics*, *12*(2), 249–268.

Langacker, R. W. (1987). *Foundations of cognitive grammar, Vol. 1: Theoretical prerequisites*. Stanford University Press.

Louw, B. (1993). Irony in the text or insincerity in the writer? The diagnostic potential of semantic prosodies. In M. Baker, G. Francis, & E. Tognini-Bonelli (Eds.), *Text and technology* (pp. 157–176). John Benjamins.

McEnery, T., & Hardie, A. (2012). *Corpus linguistics: Method, theory and practice*. Cambridge University Press.

Partington, A. (2004). "Utterly content in each other's company": Semantic prosody and semantic preference. *International Journal of Corpus Linguistics*, *9*(1), 131–156.

Pomerantz, A. (1984). Agreeing and disagreeing with assessments: Some features of preferred/dispreferred turn shapes. In J. M. Atkinson & J. Heritage (Eds.), *Structures of social action* (pp. 57–101). Cambridge University Press.

Sinclair, J. (1991). *Corpus, concordance, collocation*. Oxford University Press.

Stubbs, M. (2001). *Words and phrases: Corpus studies of lexical semantics*. Blackwell.

Talmy, L. (2000). *Toward a cognitive semantics, Vol. 1: Concept structuring systems*. MIT Press.

Traugott, E. C., & Dasher, R. B. (2002). *Regularity in semantic change*. Cambridge University Press.

Vendler, Z. (1957). Verbs and time. *The Philosophical Review*, *66*(2), 143–160.