Translation Literacy in Asylum Interpreting: A Pilot Study of TRACE and Structured Reflection
Mohammad Shehadeh
PREPRINT
August 2026 - v1

Abstract
In UK asylum proceedings, interpreting quality can shape how an applicant's account is understood, yet concerns remain about the consistency of interpreter preparation and quality assurance. This pilot study examines translation literacy, defined here as the ability to reflect explicitly on translation choices, alternatives, evidence and consequences. Twenty-one Arabic-English bilingual adults completed pre- and post-framework translation-evaluation tasks after exposure to TRACE, a mnemonic self-examination framework, or an active structured control checklist. Eighty-four responses were scored on a six-dimension rubric by the original marker and three independent blinded raters, and analyses used the mean of the four ratings. TRACE did not produce a statistically significant advantage over the control framework on combined change scores (TRACE +1.48/24; control +2.33/24; U=47.0, p=.597). Exploratory patterns included a larger TRACE gain in Alternatives and a larger descriptive increase in breadth of dimensions used, although neither supported a TRACE-specific effect. The findings support treating TRACE as a hypothesis-generating approach to structured translation reflection. Larger studies should use balanced tasks, a minimal or no-framework comparator and delayed retention testing.

Keywords: interpretation, translation, asylum, translation studies, metacognition

Introduction
In the asylum context, paragraph 339ND of the Immigration Rules requires the Secretary of State to provide an interpreter where necessary and to select an interpreter who can ensure appropriate communication during the interview. It does not prescribe a named professional qualification. In 2025, Baroness Coussins proposed Amendment 79, which would have replaced "an interpreter" with "a qualified, professional interpreter". The amendment was withdrawn after debate. The proposal nevertheless illustrates continuing concern about how interpreter competence is specified in asylum procedures (Home Office, n.d.; UK Parliament, 2025).

Background and Literature Review
Interpreter quality matters in asylum proceedings because linguistic choices can shape how an applicant's account is understood. Williamson (2024) describes UK public-service interpreting as a highly deregulated sector and identifies asylum interpreting as a setting with substantial linguistic, institutional and training demands. His analysis provides a useful context for considering reflective competence alongside formal access to interpreting services.
Inspection evidence also shows uneven preparation and language matching. The ICIBI's 2020 inspection reported a "rare and difficult" route for languages in which recognised qualifications were less readily available. The same inspection found no Home Office training on languages or dialects for relevant booking and asylum staff. It also reported that Home Office interpreters received guidance through the Code of Conduct without a dedicated training programme (ICIBI, 2020).

Hu, Zheng and Wang (2021) tested repeated use of a metacognitive self-regulation inventory in English-Chinese translator self-training. Their pre-post study linked the inventory to greater metacognitive awareness and improvements in translation quality, including clarity, vocabulary, morphosyntax, genre conventions, purpose and target audience. This provides a relevant precedent for investigating structured self-evaluation in translation. It does not establish that the same effects will occur in asylum interpreting, which involves different tasks and pressures.

More recent inspection evidence identified practical problems in substantive asylum interviews, including concerns about some interpreters' English proficiency, interpreting quality and access to the appropriate language or dialect (ICIBI, 2024).

Dialect matching was a recurring concern. The ICIBI recorded evidence that Arabic speakers could be assigned interpreters from a different regional variety, including a Moroccan applicant who had difficulty obtaining a Moroccan Arabic interpreter. The report described similar problems for less commonly spoken languages and specific dialects or regions (ICIBI, 2024).

The interpreting context can also be demanding. Morrison, Given-Wilson and Memon (2024) studied 28 Arabic-English interpreters in a mock asylum interview and found that emotionally evocative responses were interpreted 4% to 8% less accurately than neutral responses. The differences were statistically significant and had medium to large effects. This finding provides direct evidence that characteristics of asylum-style content can affect interpreting accuracy.

Geiling et al. (2022b) analysed 164 interpreters working in refugee care in Germany. Participants showed elevated psychological distress, and work-related exhaustion was associated with dissatisfaction with payment and greater exposure to traumatic content. In a separate study of 158 interpreters, Geiling et al. (2022a) found significant differences between work settings in psychological distress ($p=.01$) and work-related exhaustion ($p=.04$). Counselling showed the highest median psychological distress and, descriptively, the highest work-related exhaustion. These studies provide evidence about interpreter wellbeing and working conditions. They do not demonstrate that distress causes interpreting errors.

The ICIBI's 2024 asylum casework inspection did not make a standalone recommendation on interpreter accreditation. Its Recommendation 6 concerned routine quality assurance of asylum interviews and decisions. This is relevant to interpreting because interview quality assurance can expose problems in the communication record, although the recommendation does not itself set qualification or verbatim-recording standards for interpreters (ICIBI, 2024).

Against this background, TRACE was developed as a self-examination framework for interpreters and translators. It focuses on reflective decisions within the translation task, including uncertainty, semantic range, alternatives and consequences. It is not designed to address structural conditions such as dialect availability, workload or interpreter procurement. Its relevance to asylum interpreting is therefore exploratory.

The DPSI and the scope of TRACE

The Chartered Institute of Linguists Qualifications Level 6 Diploma in Public Service Interpreting (DPSI) is an Ofqual-regulated, nationally recognised qualification that CIOL describes as a benchmark for public-service interpreting. It consists of five units covering consecutive and simultaneous interpreting, sight translation into and from English, and written translation into and from English (Chartered Institute of Linguists [CIOL], n.d.).
The DPSI assesses performance across established public-service interpreting and translation tasks. TRACE addresses a narrower question about whether explicit prompts can elicit reasoning about semantic range, alternatives, consequences and evaluation. This pilot did not recruit or compare DPSI candidates as a defined group. Any use of TRACE alongside DPSI preparation would therefore require separate evaluation.

The TRACE framework

TRACE is a self-examination framework for interpreters and translators that focuses on reflective translation decision-making. It draws on Hu et al. (2021), whose metacognitive inventory emphasises planning, monitoring and evaluation, and on PACTE's empirical work on translation competence and decision-taking (PACTE, 2005). TRACE contains five prompts.

- T: Text and transmission
- R: Range of semantic meaning
- A: Alternatives
- C: Consequences
- E: Evaluation

Text and transmission

The first stage identifies why a source is being translated, and what kind of source it is. Although less relevant for the asylum interview as there may not be a written text, questions such as:

- What is the dialect being spoken?
- How is this being spoken?

Can be asked, and more generally:

- What genre is this text?
- Who is the intended audience?
- Who wrote the original?

These questions create a baseline assessment for the interpreter to assess whether or not they are the appropriate mediator and what action needs to be taken to translate appropriately.

Range of semantic meaning

The second stage asks what the significant words might mean.

Words do not map neatly between languages, and being aware of this is key. An Arabic term may contain associations that no single English word preserves. Baker (2018), for example, discusses non-equivalence at word level and the strategies translators use when no direct target-language equivalent captures the source item. Conversely, an English choice may introduce associations that are absent from the original.

Alternatives and Consequences

The third stage asks interpreters to generate or compare other possible translations.

Alternative generation is valuable because it challenges the idea that the supplied translation is inevitable and prevents criticism from remaining vague. Borg (2019) describes alternative translation solutions as a recurrent part of the translation process and examines their role in translatorial decision-making. Saying that a translation "sounds wrong" is therefore less analytically useful than proposing another expression and explaining why it might work better.

Similarly, the fourth stage asks why the choice matters.

Does it alter the tone? Does it make a character sound more sincere, hostile, confident or uncertain? Does it remove a joke? Does it change the apparent relationship between speakers? Does it make a religious action sound reverent, absurd or cynical? Does it turn an animate creature into something mechanical?

Evaluation

The final stage brings the previous questions together.

Is the translation effective for its intended purpose? What does it preserve? What does it alter? Is there a preferable alternative, or are several versions defensible? How confident can we be?

Evaluation should be provisional and evidenced. Interpreters should avoid assuming that every difference is a failure. A translation may be successful overall while containing one debatable choice. Equally, a grammatically accurate rendering may still produce an inappropriate tone.

The final judgement may be that the original translation should be retained. TRACE does not require change for its own sake. A translator may decide that a particular wording is the best available option after considering the alternatives.

Method

Recruitment Ethics

Participants were recruited from Prolific and provided with informed consent before participation. Participation was voluntary, responses were anonymised and participants could withdraw in accordance with study procedures.

Participants were pre-screened using the screening tools provided by Prolific. The criteria were fluency in both Arabic and English. Further to this, participants also had to be over 18 to comply with Prolific's Terms of Service.

Survey Structure

Texts

The survey used four task types across assessment and guided practice, comprising six constructed passages generated with ChatGPT 5.6 Sol:

- Two Arabic social-media posts (versions A and B), each paired with a supplied English translation for critique
- Two English diary entries (versions A and B), each translated into Arabic and then self-reviewed
- One Arabic storybook passage with a supplied English translation, used only during guided practice
- One English historical-style passage, translated into Arabic and self-reviewed only during guided practice

The study used constructed source texts to keep the assessment structure controlled and to avoid legally sensitive, traumatic or personally identifying asylum material. The A/B manipulation counterbalanced which social-media post and diary entry appeared before and after the intervention; genre itself did not vary. The two guided-practice passages were identical for all participants. Constructed tasks reduce ecological validity because they do not reproduce the full linguistic, emotional and procedural complexity of real asylum interviews.

Consent and Screening

At the beginning of the survey, participants gave informed consent and indicated whether their anonymised data could be shared and quoted publicly. They were screened again for Arabic and English fluency. The survey also collected age, education, Arabic variety, translation experience, translation frequency and self-rated language and translation ability. These variables were used to describe the sample and to inspect potential baseline imbalances between conditions.

Groups

Participants were assigned using permuted blocks of four. Within each block, the four condition and passage combinations, TRACE+A, TRACE+B, control+A and control+B, were shuffled and issued sequentially to new participant IDs. Each complete block therefore targeted a 1:1 TRACE/control allocation while balancing passage version within condition. Returning participants retained their original assignment. The control condition used the structured checklist shown below, while the TRACE condition used the TRACE prompts.

(1) Read
  - Read and summarise the text's main meaning in your own words

(2) Clarify
  - Is the translation clear and easy to understand? Explain briefly.

(3) Language
  - Check vocabulary, grammar and naturalness. Note anything that may need attention.

(4) Problems
  - Identify any possible problem or uncertainty you notice.

(5) Judgement
  - Give your overall judgement of the translation and explain it

The shared Google Apps Script serialised assignments with a lock and recorded them in an Assignments sheet. A browser-local fallback used the same four-combination scheme. Because fallback assignments occurred independently within a browser, they did not guarantee balance across browsers.

Questions

Before exposure to either framework, participants completed two unguided baseline tasks. First, they critiqued a supplied English translation of an Arabic social-media post. Second, they translated an English diary entry into Arabic, reviewed their own translation, explained their choices and rated their confidence. Participants assigned version A saw Social A and Diary A at baseline; participants assigned version B saw Social B and Diary B.
Participants were then introduced to their allocated framework: either the generic control checklist or TRACE. Guided practice used the same two passages for everyone regardless of A/B version: an Arabic storybook passage with a supplied English translation for critique, and an English historical-style passage for translation into Arabic and self-review. The passages were shared across participants, but the reflective prompts followed each participant's assigned TRACE or control framework.
For the post-framework assessment, the scaffold was removed. Participants again completed an unguided social-media critique and an English-to-Arabic diary translation with self-review. Version A participants now received Social B and Diary B, while version B participants received Social A and Diary A. The before and after questions were otherwise matched, allowing the study to examine whether framework-related reasoning appeared after the prompts themselves had disappeared.

Framework condition and passage version were therefore assigned jointly through the block-of-four procedure. Despite the intended balance, the realised analysed sample contained five participants in Set A and sixteen in Set B. The retained study records do not establish why this imbalance appeared. Passage-set analyses are consequently descriptive, and the realised allocation is treated as an implementation limitation.

Rubric and Scoring

Each response was scored on six dimensions, with every dimension rated from 0 to 2. The rubric assessed specificity, semantic range, alternatives, evidence, consequences and evaluation. A score of 0 indicated that the feature was absent. A score of 1 indicated that it was present but limited, vague or supported by only one example. A score of 2 indicated a clear and justified demonstration, usually with multiple examples. Table 1 illustrates the scoring structure, and Table 2 provides examples from participant responses.

Pre critique , Specificity	Pre critique , Range	Pre critique , Alternatives	Pre critique , Evidence	Pre critique , Consequences	Pre critique , Evaluation	Total /12
0	2	1	1	2	0	6

Table 1 - An example of scoring an answer. Each column is one of the dimensions. This answer would score a total of 6/12. The pre- indicates that this answer was provided before a framework was given.

Dimension	Score	Descriptor	Example	Participant / Group / Task
Specificity	0	No specific feature named	“Nothing. My initial translation was accurate.”	CONTROL / Post self
Specificity	1	One feature named	“‘laugh’ should have been used, not ‘smile’. The translation is acceptable but not quite accurate.”	CONTROL / Pre critique

Dimension	Score	Descriptor	Example	Participant / Group / Task
Specificity	2	Multiple specific features named with precision	"It's okay, but feels a bit stiff for a casual post. 'Smile' should be 'laugh', 'cheerful' misses the joking aspect, and 'explained their opinion' is too formal for 'شو بباله'."	TRACE / Pre critique
Range	0	No awareness of semantic range	"Nothing I just translated the meaning"	CONTROL / Pre self
Range	1	General tension named, but not explored	"i focused on meaning and naturalness of the text more than anything else."	CONTROL / Post self
Range	2	Semantic possibilities mapped across languages	"it's good and captures the overall meaning well. It translates the informal colloquial phrase 'جروب العماره' as 'building group' and 'على كيفه' as 'differently,' which fits standard English. However, 'رضا' literally means 'satisfaction' or 'approval,' so using 'agreement' changes the nuance slightly."	TRACE / Post critique
Alternatives	0	No alternative proposed	"it's a literal translation in either Jordanian or Lebanese dialect, clear and understandable more on the common language level than a professional level."	CONTROL / Pre critique
Alternatives	1	One alternative for one word	"'laugh' should have been used, not 'smile'."	CONTROL / Pre critique
Alternatives	2	Multiple alternatives, each justified	"Saying 'the building's group chat' would be clearer. Also, 'everyone understands the situation differently' is fine, but 'everyone is interpreting the situation in their own way' is closer to the original Arabic."	CONTROL / Post critique
Evidence	0	Claim without textual support	"Very, very real and accurate translation."	CONTROL / Post critique
Evidence	1	Choices named but not tied to the text	"i tried to keep the meaning as much as possible without adding any words, i also added vowels in some words to help with pronunciation"	CONTROL / Pre self

Dimension	Score	Descriptor	Example	Participant / Group / Task
Evidence	2	Claims anchored in source and target words	“‘ضحكة’ is closer to ‘laugh’ than ‘smile,’ and ‘يمزح’ means that everyone has been joking, not necessarily that they have been cheerful. This distinction matters because the joking may be a way of hiding discomfort after the news.”	TRACE / Pre critique
Consequences	0	No consequence discussed	“Very, very real and accurate translation.”	CONTROL / Post critique
Consequences	1	Consequence named but not explained	“It's a 6/10. Because it's translated literally so the real meaning of the quote is not met in the English version.”	CONTROL / Pre critique
Consequences	2	Specific consequences tied to choices	“‘Smile’ should be ‘laugh’, ‘cheerful’ misses the joking aspect, and ‘explained their opinion’ is too formal for ‘شو بباله’.”	TRACE / Pre critique
Evaluation	0	Bare assertion	“Very, very real and accurate translation.”	CONTROL / Post critique
Evaluation	1	Judgement with a single reason	“It's a 6/10. Because it's translated literally so the real meaning of the quote is not met in the English version.”	CONTROL / Pre critique
Evaluation	2	Balanced, multi-criteria judgement	“The translation communicates the general contrast between outward behaviour and hidden feelings, but it changes several important details... A closer translation would be...”	TRACE / Pre critique

Table 2 - Examples of scored answers across each dimension of the rubric, showing what scores of 0, 1 and 2 look like in participant responses.

Independent marking and score aggregation

To reduce the effect of one marker’s judgement on the final results, three independent Arabic-English reviewers were recruited through Prolific after the initial marking. Each reviewer was given all 84 responses in a blinded and randomised order. They were not shown participant IDs, whether a response came from TRACE or the control group, whether it was produced before or after the framework, the passage set, AI flags, or the scores given by the original marker. Each reviewer used the same 0-2 six-dimension rubric.

This produced four ratings for every response: the author’s original rating and three independent ratings. For each dimension, the arithmetic mean of the four ratings was calculated. The six dimension means were then added to produce the final score out of 12 used in the analyses below. Using an average was chosen before comparing the new TRACE and control results so that the final score was not selected according to which aggregation method produced the strongest effect.

Inter-rater reliability was assessed before the consensus scores were interpreted. A two-way random-effects absolute-agreement intraclass correlation was calculated for total scores. Reliability was ICC(A,1)=.303 for a single rater and ICC(A,4)=.635 for the mean of four raters. Exact agreement on the individual 0-2 categories was assessed with Fleiss' kappa, which ranged from -.002 to .139 across the six dimensions. The category-level agreement was therefore low, while averaging the four ratings produced a more reliable total score than relying on one rater. This pattern is consistent with evidence that human assessment of interpreting involves complex rater judgement and interaction with analytic rating scales (Han et al., 2024).

Rater score distributions

Because the four raters used the 0–2 rubric somewhat differently, the distribution of their total response scores is also shown graphically. This helps to explain why aggregation across raters was preferable to relying on a single marker.

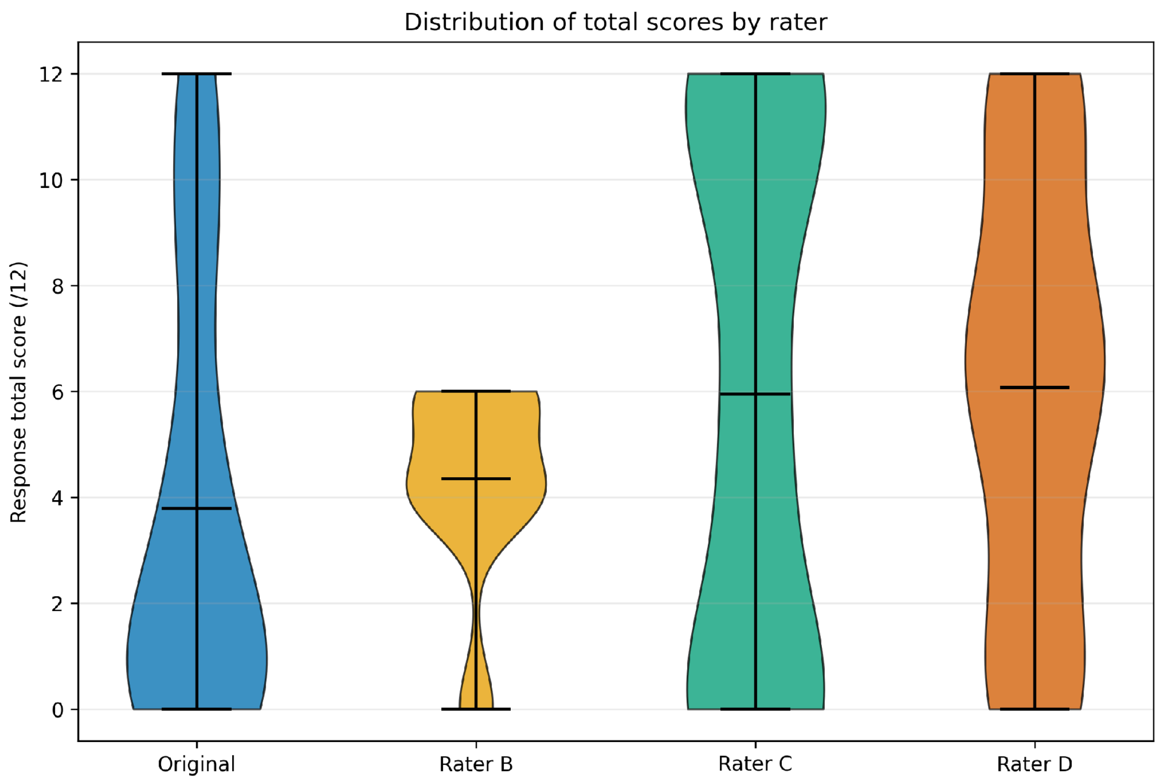

Figure 1 - Distribution of total response scores (/12) by rater. Raters are anonymised as Original and Raters B-D. The violin plot shows that the raters used the scoring scale differently. Rater B used a comparatively compressed range of scores, while Raters C and D used the upper end of the scale more readily. The original ratings were also relatively strict. This variation supports the decision to use the four-rater mean for the primary scores.

As the study had a small sample (n=21) and bounded outcome scores, the inferential analysis was treated as exploratory. Wilcoxon signed-rank tests were used for pre- to post-framework comparisons within each group, and Mann-Whitney U tests compared change scores between TRACE and control participants. Tests were two-sided.

P-values were not adjusted for multiple comparisons, so isolated values below .05 are interpreted cautiously and do not constitute definitive evidence of an intervention effect.
Two further exploratory analyses were used to examine patterns that could be hidden by mean score change. Participants were classified by the direction of their combined /24 change (improved, unchanged or declined), and the magnitude of each participant's change was inspected graphically. A breadth score was also calculated as the number of the six rubric dimensions receiving a score above zero in each response; self-critique and critique breadth were then averaged within each phase to give a phase-level breadth score out of six. This was intended to test whether participants spontaneously used a wider range of translation-literacy behaviours after the framework was removed.

Results

Group	n	Pre-self	Post-self	Self change	Pre-critique	Post-critique	Critique change
ALL	21	3.83 (2.75) [2.75]	4.39 (2.78) [3.25]	+0.56 (3.14) [0.25]	5.30 (2.59) [5.00]	6.62 (2.39) [7.25]	+1.32 (2.59) [1.25]
TRACE	11	4.52 (3.45) [3.00]	5.14 (3.02) [3.25]	+0.61 (2.87) [0.50]	6.07 (2.92) [5.50]	6.93 (2.74) [8.00]	+0.86 (2.13) [1.00]
Control	10	3.08 (1.51) [2.75]	3.58 (2.38) [2.75]	+0.50 (3.57) [−0.38]	4.45 (1.98) [4.50]	6.28 (2.04) [6.88]	+1.83 (3.05) [2.50]

Table 3 - Changes in scores across self-critique and critique tasks for all participants, using the mean score across four raters. Results are shown as mean (standard deviation) [median].

Table 3 presents the changes in scores across the two main groups of tasks: critiquing ones' own work (shown as -self) and critiquing others' work (shown as -critique). Scores now use the mean across four raters. The total sample size is n=21.
For self-critique, the mean score across all participants increased from 3.83 to 4.39, giving a mean change of +0.56. Across the TRACE group, the mean increased from 4.52 to 5.14 (+0.61), compared to the control group which increased from 3.08 to 3.58 (+0.50). Therefore, the difference in mean self-critique change between the two groups was small once the four ratings were averaged.
For critiquing another translation, scores across all participants increased from a mean of 5.30 to 6.62, giving a mean change of +1.32. The TRACE group increased by +0.86, from 6.07 to 6.93, whereas the control group increased by +1.83, from 4.45 to 6.28. This reverses the descriptive pattern found in the initial single-rater scoring, as the control group shows the larger mean increase on critique tasks when the four ratings are averaged.
Exploratory Mann-Whitney U tests did not find a statistically significant difference in change between TRACE and control participants for self-critique (U=63.0, p=.597) or critique of another translation (U=42.0, p=.378). Within-group Wilcoxon tests were also non-significant for both task types in both groups (all p>.10).

These findings therefore do not provide evidence of a statistically significant TRACE-control difference at task level in the full sample. The descriptive increases remain useful for examining how scores moved, but they should not be interpreted as showing that TRACE outperformed the control framework.

Group	n	Pre-self	Post-self	Self change	Pre-critique	Post-critique	Critique change
ALL	18	3.74 (2.63) [2.88]	4.36 (2.89) [3.00]	+0.63 (2.84) [0.00]	4.82 (2.40) [4.50]	6.36 (2.49) [7.00]	+1.54 (2.68) [2.00]
TRACE	9	4.33 (3.38) [3.00]	5.58 (3.18) [6.75]	+1.25 (2.18) [0.50]	5.42 (2.76) [4.50]	6.61 (2.95) [7.50]	+1.19 (2.13) [1.25]
Control	9	3.14 (1.59) [2.75]	3.14 (2.05) [2.25]	+0.00 (3.40) [−0.50]	4.22 (1.96) [4.00]	6.11 (2.09) [6.50]	+1.89 (3.23) [3.25]

Table 4 - Changes in scores across self-critique and critique tasks with participants flagged for possible AI use excluded, using the mean score across four raters. Results are shown as mean (standard deviation) [median].

Table 4 presents the same task-level results with the three participants flagged for possible AI use excluded, leaving n=18 (TRACE n=9; control n=9).
After AI exclusions, self-critique increased by +0.63 across all participants. TRACE increased by +1.25, while the control-group mean was unchanged (+0.00). For critique tasks, TRACE increased by +1.19 and control increased by +1.89.
Neither between-group difference was statistically significant: self-critique U=55.0, p=.215, and critique U=32.5, p=.507. Within TRACE, neither self-critique (p=.122) nor critique (p=.137) reached the .05 threshold. Within control, neither self-critique (p=.722) nor critique (p=.123) was statistically significant.
No within-group or between-group test in this corrected AI-excluded task-level sensitivity analysis reached the .05 threshold.
The corrected sensitivity analysis therefore does not produce a clear TRACE-specific advantage once the three AI-flagged participants are excluded and the four-rater mean is used.

Group	Dimension	Pre	Post	Change
ALL	Specificity	0.88	1.03	+0.15
ALL	Range	0.85	0.98	+0.13

Group	Dimension	Pre	Post	Change
ALL	Alternatives	0.61	0.78	+0.17
ALL	Evidence	0.64	0.82	+0.18
ALL	Consequences	0.79	0.95	+0.17
ALL	Evaluation	0.81	0.95	+0.14
TRACE	Specificity	1.03	1.14	+0.10
TRACE	Range	0.97	1.05	+0.08
TRACE	Alternatives	0.73	0.92	+0.19
TRACE	Evidence	0.82	0.92	+0.10
TRACE	Consequences	0.86	1.01	+0.15
TRACE	Evaluation	0.89	1.00	+0.11
Control	Specificity	0.70	0.91	+0.21
Control	Range	0.71	0.90	+0.19
Control	Alternatives	0.49	0.63	+0.14
Control	Evidence	0.44	0.71	+0.28
Control	Consequences	0.70	0.89	+0.19
Control	Evaluation	0.73	0.89	+0.16

Table 5 - Mean scores across each dimension of the rubric before and after exposure to a framework, using the mean across four raters and shown across all participants and both groups.

Table 5 shows the mean scores across all six rubric dimensions after averaging the four raters. Across all participants, every dimension increased, with mean changes ranging from +0.13 for range to +0.18 for evidence. The average change across the six dimensions was approximately +0.16.

Across the TRACE group, every dimension increased descriptively. The largest increase was in Alternatives (+0.19), followed by Consequences (+0.15). The average change across the six dimensions was approximately +0.12. Evaluation also increased (+0.11), reversing the direction seen in the original single-rater scoring.

Across the control group, the average dimension change was larger at approximately +0.19. The largest increase was in evidence (+0.28), followed by specificity (+0.21). Therefore, the four-rater dimension scores do not show a general pattern of TRACE outperforming control; the clearest TRACE-specific descriptive feature remains the relatively larger focus on alternatives compared with its other dimensions.

Group	Dimension	Pre	Post	Change
TRACE	Specificity	0.94	1.15	+0.21
TRACE	Range	0.89	1.03	+0.14
TRACE	Alternatives	0.64	0.94	+0.31
TRACE	Evidence	0.75	0.92	+0.17
TRACE	Consequences	0.79	1.04	+0.25
TRACE	Evaluation	0.86	1.01	+0.15

Table 6 - Mean scores across each dimension of the rubric for the TRACE group, using the mean across four raters and with participants flagged for possible AI use excluded.

Group	Dimension	Pre	Post	Change
Control	Specificity	0.69	0.88	+0.18
Control	Range	0.71	0.85	+0.14
Control	Alternatives	0.50	0.61	+0.11
Control	Evidence	0.42	0.67	+0.25
Control	Consequences	0.67	0.81	+0.14
Control	Evaluation	0.69	0.82	+0.13

Table 7 - Mean scores across each dimension of the rubric for the control group, using the mean across four raters and with participants flagged for possible AI use excluded.

Tables 6 and 7 show the dimension results after the three AI-flagged participants were removed. TRACE improved across all six dimensions, with the largest changes in alternatives (+0.31) and consequences (+0.25). Control also improved across all six dimensions, with its largest change in evidence (+0.25). The average dimension change was approximately +0.20 for TRACE and +0.16 for control. These results are descriptive; the sensitivity analysis does not establish a TRACE-specific dimension effect.

Immediate unguided retention and transfer of TRACE dimensions

A secondary aim of TRACE is to make its components memorable once the framework is taken away. The post-framework tasks in this study were completed without either group being shown its framework again. The dimension scores can therefore also be read as a measure of unguided transfer and whether the kinds of reasoning prompted during training continued to appear when participants were asked to evaluate translations independently. This should not be interpreted as long-term retention, as there was no delayed follow-up assessment.

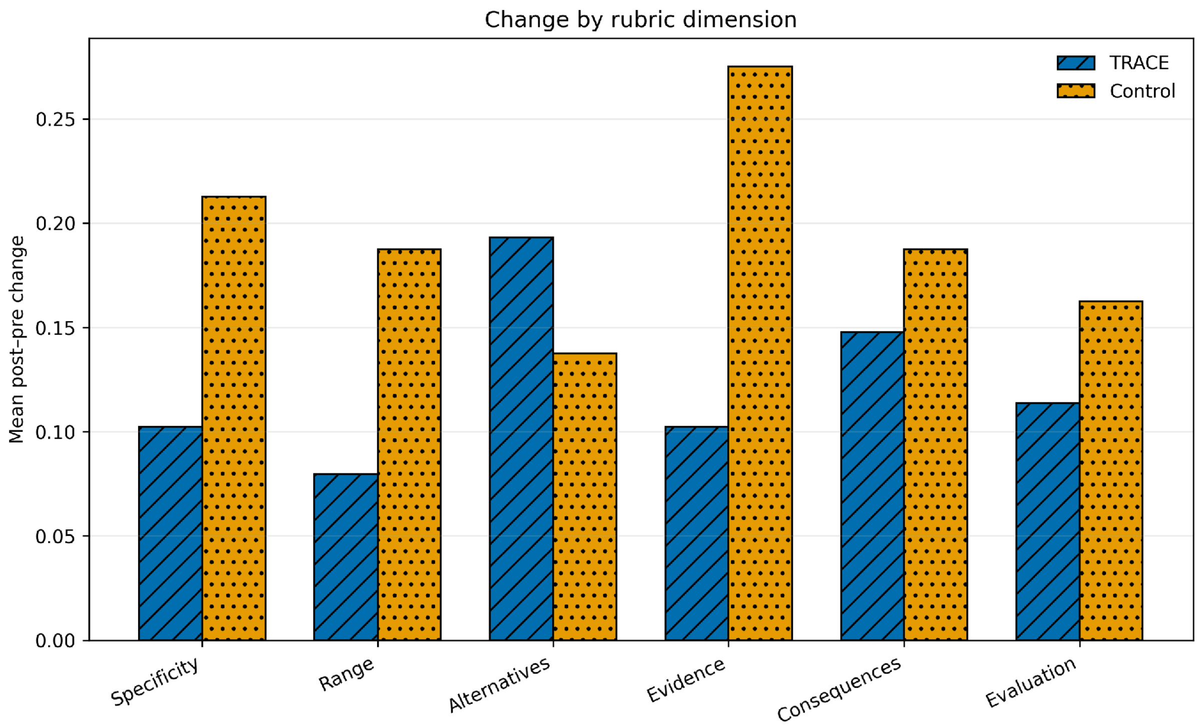

Figure 2 - Mean post- to pre-framework change by rubric dimension. Positive values indicate improvement. TRACE and control are distinguished by both colour and hatch pattern.

Figure 2 shows the dimensional change pattern directly. All six dimensions moved in a positive direction in TRACE, and Alternatives was the only dimension for which TRACE showed a larger mean increase than control. Control showed larger mean gains on Specificity, Range, Evidence, Consequences and Evaluation.

Across TRACE participants, all six dimensions increased from the unguided pre-framework tasks to the unguided post-framework tasks. The largest descriptive increase was in alternatives, from 0.73 to 0.92 (+0.19), followed by consequences (+0.15). This is relevant to the mnemonic aim of TRACE because alternatives is explicitly represented by the A in TRACE, and participants continued to provide more alternative translations after the framework itself had been removed.

However, this pattern does not provide evidence that the TRACE mnemonic produced stronger immediate retention than the control framework. Alternatives was the only dimension in which TRACE showed a larger mean increase than control (+0.19 compared with +0.14), but the difference in change scores was not statistically significant (U=56.5, p=.943). TRACE also had a higher post-framework alternatives score than control (0.92 compared with 0.63), but this difference was not statistically significant (U=78.0, p=.110), and TRACE had already started higher at baseline (0.73 compared with 0.49). No dimension showed a statistically significant TRACE-versus-control difference in change.

Breadth of retained translation-literacy reasoning

Breadth measures how many of the six translation-literacy dimensions appeared at all in a response. Each dimension with a mean four-rater score above zero contributed one point, giving a score from 0 to 6. A breadth score of 5/6 therefore means that five of the six dimensions appeared to some extent. Pre- and post-breadth were calculated for each participant and averaged across the self-critique and critique tasks within each phase.

Analysis	Group	n	Pre breadth	Post breadth	Change
All responses	ALL	21	5.38 (0.96) [6.00]	5.83 (0.33) [6.00]	+0.45 (1.08) [0.00]
All responses	TRACE	11	5.36 (1.03) [6.00]	5.95 (0.15) [6.00]	+0.59 (1.07) [0.00]
All responses	CONTROL	10	5.40 (0.94) [5.75]	5.70 (0.42) [6.00]	+0.30 (1.14) [0.00]
AI-excluded	ALL	18	5.33 (1.01) [6.00]	5.81 (0.35) [6.00]	+0.47 (1.16) [0.00]
AI-excluded	TRACE	9	5.33 (1.12) [6.00]	5.94 (0.17) [6.00]	+0.61 (1.17) [0.00]
AI-excluded	CONTROL	9	5.33 (0.97) [5.50]	5.67 (0.43) [6.00]	+0.33 (1.20) [0.00]

Table 7a - Breadth of translation-literacy reasoning before and after the framework was removed. Breadth is the number of rubric dimensions present in a response (0-6), averaged across the self-critique and critique tasks within each phase. Results are shown as mean (standard deviation) [median].

In the full sample, TRACE breadth increased from 5.36 to 5.95 (+0.59), compared with 5.40 to 5.70 (+0.30) in control. The TRACE within-group change approached but did not reach the .05 threshold ($p=0.078$), while the control change was non-significant ($p=0.457$). The difference in breadth change between groups was also non-significant ($p=0.410$).

After the three AI-flagged participants were excluded, TRACE showed a mean breadth increase of +0.61, compared with +0.33 in control. The between-group difference remained non-significant ($U=48.0$, $p=.516$). Baseline breadth was already close to the six-dimension ceiling and median change was zero in both groups.

Group	n	Pre /24	Post /24	Change
ALL	21	9.13 (4.66) [8.25]	11.01 (4.38) [10.00]	+1.88 (4.64) [1.75]
TRACE	11	10.59 (5.57) [11.50]	12.07 (4.86) [11.25]	+1.48 (4.03) [0.75]

Group	n	Pre /24	Post /24	Change
Control	10	7.53 (2.88) [7.13]	9.85 (3.70) [8.75]	+2.33 (5.43) [3.25]

Table 8 - Combined pre- and post-framework scores out of 24 across all participants and both groups, using the mean score across four raters. Results are shown as mean (standard deviation) [median].

Group	n	Pre /24	Post /24	Change
ALL	18	8.56 (4.61) [6.88]	10.72 (4.63) [9.50]	+2.17 (4.58) [1.88]
TRACE	9	9.75 (5.75) [7.00]	12.19 (5.42) [10.00]	+2.44 (3.66) [1.75]
Control	9	7.36 (3.00) [6.00]	9.25 (3.36) [8.75]	+1.89 (5.57) [2.50]

Table 9 - Combined pre- and post-framework scores out of 24 with participants flagged for possible AI use excluded, using the mean score across four raters. Results are shown as mean (standard deviation) [median].

Tables 8 and 9 show the combined scores out of 24 across both groups, with Table 8 showing all participants and Table 9 showing AI-excluded participants. These combined scores use the four-rater mean for both the self-critique and critique task.

Across all participants, the combined mean score increased from 9.13/24 to 11.01/24, a mean increase of +1.88. TRACE began higher at 10.59/24 and increased to 12.07/24 (+1.48), while control increased from 7.53/24 to 9.85/24 (+2.33). A Mann-Whitney U test found no statistically significant difference between the TRACE and control change scores ($U=47.0$, $p=.597$).

After AI exclusions, the combined mean increase was +2.17 across all participants. TRACE increased by +2.44 and control by +1.89. The between-group difference was not statistically significant ($U=42.0$, $p=.930$). Within TRACE, the combined pre-post change approached but did not reach the .05 threshold ($p=.066$), while the control change was also non-significant ($p=.343$).

Taken together, the inferential tests do not show evidence that TRACE produced a greater overall improvement than the control framework in this sample. This differs from the pattern seen using the author’s original ratings alone and shows why averaging across independent raters materially changes the interpretation of the study.

The absence of a significant group difference does not establish equivalence between the frameworks, particularly with $n=21$. The results are exploratory. Both groups changed after exposure to structured evaluation guidance, and this pilot sample does not establish a TRACE-specific effect.

Participant-level direction and magnitude of change

Participant-level change was inspected to determine whether the group means were being driven by a small number of extreme responders. In the full sample, 7/11 TRACE participants (63.6%) improved on the combined score, compared with 7/10 control participants (70.0%). Four TRACE participants and three control participants declined, and no participant was exactly unchanged. After the corrected AI exclusions, 7/9 TRACE participants (77.8%) and 6/9 control participants (66.7%) improved.

The paired pre- to post-framework plot shows that the control-group increase was not produced by a single extreme participant: several control participants moved upward by four points or more. TRACE also contained strong positive responders, including changes of +10.00 and +5.75, but also contained several small gains and four declines in the full sample. The figure therefore supports the same overall interpretation as the group means: improvement was common in both groups, while there was no clear participant-level advantage for TRACE.

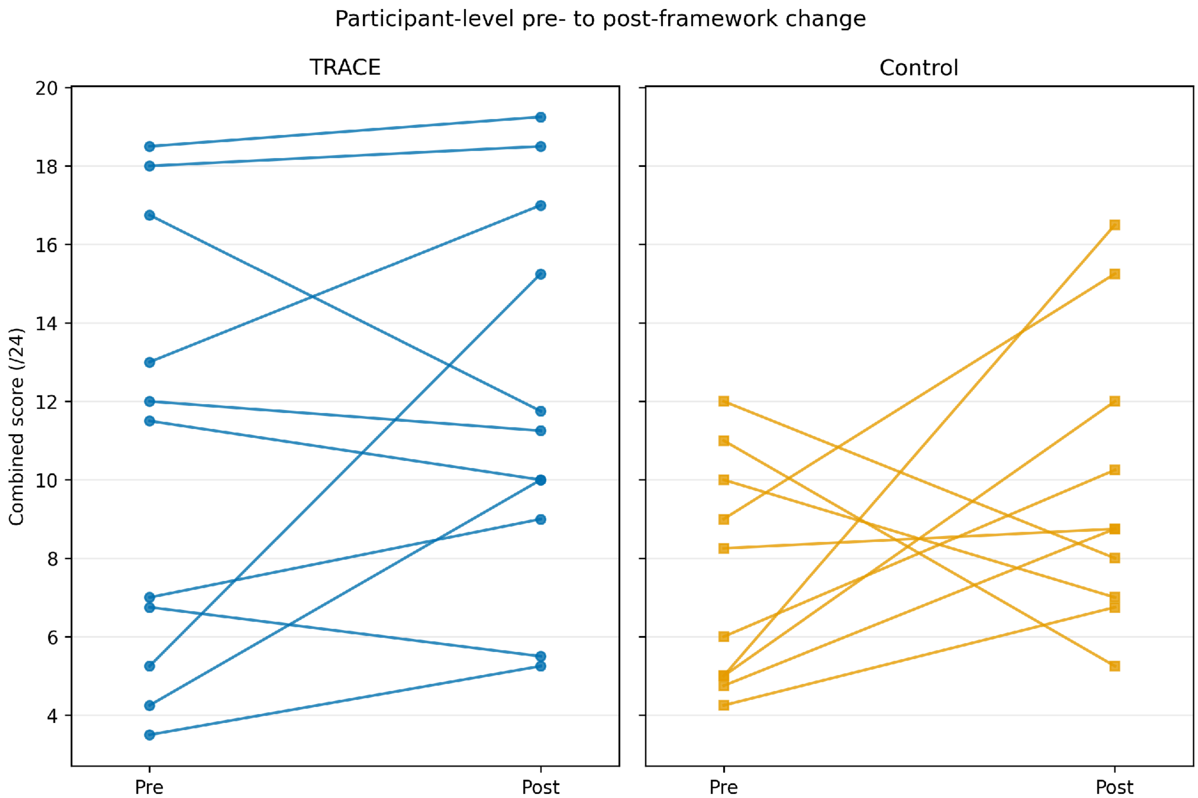

Figure 3 - Participant-level pre- to post-framework change in the combined /24 score. Each line represents one participant; TRACE uses circular markers and control square markers.

Set	n	Pre /24	Post /24	Change
A	5	5.00 (1.31) [5.00]	10.30 (3.70) [10.00]	+5.30 (3.49) [5.75]
B	16	10.42 (4.59) [10.50]	11.23 (4.67) [10.13]	+0.81 (4.51) [0.50]

Table 10 - Combined scores out of 24 across task sets A and B, showing the difference between pre- and post-framework performance.

Table 10 shows combined scores across task sets A and B. Set A began much lower at 5.00/24 and increased to 10.30/24, a mean change of +5.30. Set B began at 10.42/24 and increased to 11.23/24, a smaller mean change of

+0.81. The difference in starting scores again suggests that the two sets were not equivalent in difficulty and leaves substantially more room for improvement in Set A.

Group × Set	n	Pre /24	Post /24	Change
TRACE A	4	5.00 (1.51) [4.75]	9.88 (4.13) [9.50]	+4.88 (3.88) [3.88]
TRACE B	7	13.79 (4.23) [13.00]	13.32 (5.07) [11.75]	−0.46 (2.73) [−0.75]
Control A	1	5.00 [5.00]	12.00 [12.00]	+7.00 [7.00]
Control B	9	7.81 (2.90) [8.25]	9.61 (3.84) [8.75]	+1.81 (5.48) [2.50]

Table 11 - Combined scores out of 24 across task sets A and B, split further by the TRACE and control groups.

Table 11 separates the set results further by TRACE and control group. Set A remains heavily TRACE-weighted, with four TRACE participants and only one control participant, while Set B contains seven TRACE and nine control participants. These small and uneven cells mean that the group-by-set results are descriptive only.
Within TRACE, Set A increased from 5.00/24 to 9.88/24 (+4.88), whereas TRACE Set B decreased slightly from 13.79/24 to 13.32/24 (−0.46). The single Control A participant increased from 5.00 to 12.00 (+7.00), while Control B increased by +1.81 on average.
The very high starting score of TRACE Set B, combined with the small Set A sample, creates an obvious ceiling and task-difficulty problem. These results cannot separate a framework effect from differences between task sets and should not be used as a causal comparison.

Metadata Analysis

Education	ALL n (%)	TRACE n (%)	Control n (%)
Bachelor’s degree	13 (61.9%)	6 (54.5%)	7 (70.0%)
Master’s degree	5 (23.8%)	4 (36.4%)	1 (10.0%)
Doctoral degree	3 (14.3%)	1 (9.1%)	2 (20.0%)
Total	21 (100%)	11 (100%)	10 (100%)

Table 12 - Highest level of education reported by participants, shown across the full sample and both groups.

Table 12 shows the highest level of education across participants. There is no clear descriptive difference in education across groups, with most participants in both groups having a bachelor’s degree.

Translation frequency	ALL n (%)	TRACE n (%)	Control n (%)
Daily or almost daily	4 (19.0%)	3 (27.3%)	1 (10.0%)
A few times a week	4 (19.0%)	3 (27.3%)	1 (10.0%)
A few times a month	6 (28.6%)	3 (27.3%)	3 (30.0%)
A few times a year	6 (28.6%)	1 (9.1%)	5 (50.0%)
Never or almost never	1 (4.8%)	1 (9.1%)	0 (0.0%)
Total	21 (100%)	11 (100%)	10 (100%)

Table 13 - Self-reported translation frequency across the full sample and both groups.

Group	n	Translation-frequency score
ALL	21	2.19 (1.21) [2]
TRACE	11	2.55 (1.29) [3]
Control	10	1.80 (1.03) [1.5]

Table 14 - Translation frequency converted to a score from 0-4, where 0 represents never or almost never and 4 represents daily or almost daily. Results are shown as mean (standard deviation) [median].

Tables 13 and 14 describe self-reported translation frequency. Daily translation was reported by 27.3% of TRACE participants and 10.0% of control participants. On the 0-4 frequency score, TRACE had a mean of 2.55 and median of 3, while control had a mean of 1.80 and median of 1.5. A Mann-Whitney U test did not show a statistically significant between-group difference (U=76.0, p=.137), so the imbalance is treated descriptively.

Measure	ALL	TRACE	Control	TRACE − Control
Arabic → English translation	8.52 (1.17) [9]	8.82 (0.98) [9]	8.20 (1.32) [8]	+0.62
English → Arabic translation	8.95 (0.92) [9]	9.09 (0.83) [9]	8.80 (1.03) [9]	+0.29
English reading	9.14 (1.35) [10]	9.55 (0.69) [10]	8.70 (1.77) [9.5]	+0.85

Measure	ALL	TRACE	Control	TRACE − Control
English writing	8.76 (1.26) [9]	8.91 (0.94) [9]	8.60 (1.58) [9]	+0.31
MSA reading	9.90 (0.30) [10]	10.00 (0.00) [10]	9.80 (0.42) [10]	+0.20
MSA writing	9.90 (0.30) [10]	10.00 (0.00) [10]	9.80 (0.42) [10]	+0.20

Table 15 - Self-reported language and translation ability out of 10 across the full sample and both groups. Results are shown as mean (standard deviation) [median].

Table 15 shows self-reported ability levels, out of 10, across all participants. Across all skills participants were asked to rate, the TRACE group consistently scored themselves higher than the control group. The greatest difference was in English reading (+0.85), followed by Arabic to English translation (+0.62). The smallest differences were MSA reading and writing (+0.20 each). However, scores across both groups remained high overall.

Data Quality Statement

Three participants were flagged for authorship uncertainty in at least one response. Flagging was based on a marked within-participant discontinuity in style and level of elaboration compared with that participant's other answers. Polished writing alone was not used as evidence of AI assistance. These qualitative cues are not diagnostic, and no automated detector was treated as ground truth. The primary analysis retains all submitted participants. A sensitivity analysis excludes any participant with a flagged response to test whether the overall pattern depends on uncertain-authorship responses. The flagging process relied on researcher judgement and is not a validated detection procedure. The four-rater reliability analysis also showed limited exact agreement on individual 0-2 rubric categories. Reliability improved when ratings were averaged, so the rubric and consensus scores remain exploratory measures.

Discussion

The central result is the absence of evidence that TRACE outperformed the active control. Mean combined change was +1.48/24 in TRACE and +2.33/24 in control, and the between-group difference was not statistically significant ($U=47.0$, $p=.597$). With a sample of 21 participants, this result is imprecise and cannot establish equivalence between the frameworks. It does show that the present data do not support a claim of TRACE superiority. The higher TRACE baseline also limits interpretation of post-framework scores because participants with high starting scores had less available headroom.

Both conditions received structured reflective guidance. The study therefore tests TRACE against an active checklist and does not estimate the effect of structured reflection relative to no intervention. Mean scores rose in both groups, but several explanations remain plausible, including framework exposure, repeated-task familiarity, regression to the mean and differences between passage versions. Hu, Zheng and Wang (2021) provide evidence that repeated metacognitive self-regulation can influence translator behaviour and translation quality in a student-training context. The present pilot extends that question to Arabic-English bilingual participants, although its design does not isolate a general effect of structured reflection.

The dimension results identify a narrower question for future testing. Control showed larger mean gains on five of the six dimensions. Alternatives was the exception and was TRACE's largest dimension gain, increasing by

+0.19 compared with +0.14 in control. The difference was not statistically significant (p=.943). Alternative generation is explicitly represented by the A in TRACE, and Borg (2019) identifies alternative translation solutions as a recurrent component of translatorial decision-making. The alignment between the framework and this descriptive pattern makes Alternatives a reasonable pre-specified outcome for a larger study, while the current result remains exploratory.

Breadth provides a second hypothesis about mnemonic carry-over. TRACE breadth increased by +0.59 dimensions and control breadth by +0.30 in the full sample. The between-group difference was non-significant (p=.410). Baseline breadth was already close to the six-dimension maximum, which leaves little room for this measure to detect improvement. A future retention measure should therefore be more demanding than simple presence or absence of a dimension and should be repeated after a delay without displaying the mnemonic.

Baseline and passage-set imbalances complicate the group comparison. TRACE participants began higher on the combined score and on every rubric dimension. The realised passage allocation was also highly uneven, with Set A containing five participants and Set B sixteen, despite the intended block allocation. Set A began at 5.00/24 and Set B at 10.42/24, indicating substantial empirical differences in task difficulty or sample composition. These features can affect observed change through ceiling and headroom effects. The current data cannot separate those influences from any framework effect.

Participant-level results reinforce the need to look beyond group means. Seven of eleven TRACE participants and seven of ten control participants improved in the full sample. Other participants declined, and the change scores had large standard deviations. The apparent response to structured reflection was therefore heterogeneous. A larger sample could test whether baseline translation ability, translation frequency, professional experience or other participant characteristics predict change. The present sample is too small for credible moderator analysis.

Measurement is another important finding of the pilot. Exact 0-2 agreement was low across dimensions, and the reliability of the four-rater mean reached ICC(A,4)=.635 compared with ICC(A,1)=.303 for one rater. Human interpreting assessment is known to involve complex rater processes even when analytic scales are available (Han et al., 2024). The present rubric adds further judgement because concepts such as semantic range, evidence and consequences can be demonstrated in several ways. Future development should refine behavioural anchors, provide calibration examples before scoring, assess rater drift and evaluate validity independently of an intervention study. Averaging raters reduces individual scoring noise, but it does not by itself validate the construct being measured.

The asylum relevance of TRACE should be interpreted at the level of reflective reasoning. Public-service interpreters make decisions involving dialect, semantic range, alternative renderings and the consequences of wording choices. Morrison et al. (2024) shows that emotionally evocative content can reduce interpreting accuracy in a controlled mock asylum interview, while ICIBI reports identify practical concerns about quality and dialect matching. The present study did not reproduce those conditions. Its written, constructed tasks were completed without the time pressure, turn-taking, memory load, emotional intensity and institutional consequences of a live asylum interview. Evidence of transfer to professional asylum interpreting therefore requires a different study design.

Formal qualifications and reflective training address different parts of interpreter competence. The DPSI assesses professional public-service interpreting and translation tasks, while TRACE prompts explicit reflection on translation choices. The present results support further investigation of TRACE as a possible reflective training activity. They do not support its use as a competency standard, accreditation substitute or evidence of readiness for asylum interpreting. Testing practising interpreters and DPSI-qualified participants would help establish whether TRACE adds anything to existing professional preparation.

The sensitivity analysis provides a limited robustness check. Excluding the three participants whose authorship was considered uncertain did not reveal a TRACE-specific advantage, and the corrected AI-excluded comparison remained non-significant. This stability reduces concern that the main conclusion depends entirely on those responses. The flagging process was qualitative and was not a validated method for detecting AI assistance, so it cannot establish which participants used AI. The full-sample analysis remains the primary analysis.

Limitations and next steps

The sample size of n=21 limits statistical power and produces imprecise estimates of group differences. It also restricts meaningful analysis of participant characteristics that may influence response to the frameworks. Participants were Arabic-English bilingual adults recruited through Prolific and were not restricted to practising asylum interpreters or DPSI holders. Generalisation to professional public-service interpreting is therefore limited.

The active control creates an important design constraint. It allows a comparison between TRACE and another structured checklist, but there was no minimal or no-framework condition. The study cannot determine whether either framework caused the observed pre-post change. Repeated exposure to similar tasks, familiarity with the evaluation format and regression to the mean remain possible explanations. A future three-arm design should include TRACE, an active structured control and a comparator that repeats the assessment without an intervening reflective scaffold.

The realised allocation also departed substantially from the intended permuted block design. Only five participants were analysed in Set A and sixteen in Set B, and Set A was heavily TRACE-weighted. The retained records do not identify the cause of this implementation failure. The sets were not empirically equivalent at baseline, with Set A beginning much lower than Set B. Passage difficulty, allocation imbalance and ceiling effects are therefore entangled with the framework comparison. Future work should verify allocation logs during data collection and pilot-test passages for comparable difficulty before recruitment.

Ecological validity is limited by the use of constructed written passages. These tasks improved experimental control and avoided the ethical difficulties of using sensitive asylum material, yet they omitted central features of live interpreting such as spoken delivery, turn-taking, note-taking, memory demands, time pressure and interaction with an interviewer and applicant. They also lacked the emotional and procedural stakes of an actual asylum decision. Morrison et al. (2024) demonstrates that emotional content can affect interpreting accuracy, which makes this difference particularly important for the intended application.

The outcome measure also requires further validation. Exact category agreement was low and the four-rater average achieved only moderate reliability. The original marker was one of the four raters, and averaging ratings cannot remove limitations in the rubric's construct definition. The breadth measure had an additional ceiling problem because most participants already displayed nearly all six dimensions at baseline. Multiple exploratory tests were conducted without adjustment, increasing the chance of isolated false-positive findings. The statistical analyses should therefore be read as hypothesis-generating.

Future evaluation should use a larger sample with pre-specified primary outcomes and an analysis plan, verify the allocation procedure as recruitment proceeds, and use balanced passages that have been pilot-tested for difficulty. The study should include practising interpreters and could stratify or report professional qualification and asylum-interpreting experience. A refined rubric should be tested independently with trained raters before it is used as an efficacy outcome. A delayed unguided assessment after several days or a week would test retention, while a spoken or simulated asylum-interview task would test transfer under more realistic conditions. Together, these changes would allow separate estimates of the effect of structured reflection, any additional TRACE-specific effect and the durability of mnemonic recall.

Conclusion

This pilot found no statistically significant advantage for TRACE over an active structured control. It also cannot attribute the overall pre-post increase to structured reflection because no minimal comparator was included. The study nevertheless demonstrates a feasible way to operationalise several forms of translation reasoning and identifies Alternatives and breadth as outcomes worth testing prospectively. Further evaluation requires a larger and better-balanced design, stronger measurement, delayed retention testing and tasks that more closely represent professional interpreting. Evidence from such studies would be needed before TRACE could be considered for use in high-stakes settings such as asylum proceedings.

References

Baker, M. (2018). In Other Words: A Coursebook on Translation (3rd ed.). Routledge.

Borg, C. (2019). Written alternative translation solutions in the translation process: A case study. Translation, Cognition & Behavior, 2(1), 55–78. https://doi.org/10.1075/tcb.00019.bor

Chartered Institute of Linguists. (n.d.). DPSI overview. https://www.ciol.org.uk/dpsi-overview

Geiling, A., Böttche, M., Knaevelsrud, C., & Stammel, N. (2022a). A comparison of interpreters' wellbeing and work-related characteristics in the care of refugees across different work settings. BMC Public Health, 22, 1635. https://doi.org/10.1186/s12889-022-14034-7

Geiling, A., Knaevelsrud, C., Böttche, M., & Stammel, N. (2022b). Psychological distress, exhaustion, and work-related correlates among interpreters working in refugee care: Results of a nationwide online survey in Germany. European Journal of Psychotraumatology, 13(1), 2046954. https://doi.org/10.1080/20008198.2022.2046954

Han, C., Zheng, B., Xie, M., & Chen, S. (2024). Raters' scoring process in assessment of interpreting: An empirical study based on eye tracking and retrospective verbalisation. The Interpreter and Translator Trainer, 18(3), 400-422. https://doi.org/10.1080/1750399X.2024.2326400

Home Office. (n.d.). Immigration Rules part 11: Asylum (paragraph 339ND). GOV.UK. https://www.gov.uk/guidance/immigration-rules/immigration-rules-part-11-asylum

Hu, Z., Zheng, B., & Wang, X. (2021). The impact of a metacognitive self-regulation inventory in translator self-training: A pre-post study with English-Chinese translation students. The Interpreter and Translator Trainer, 15(4), 430–449. https://doi.org/10.1080/1750399X.2020.1844436

Independent Chief Inspector of Borders and Immigration. (2020). An inspection of the Home Office's use of language services in the asylum process. GOV.UK. https://www.gov.uk/government/publications/an-inspection-of-the-home-offices-use-of-language-services-in-the-asylum-process

Independent Chief Inspector of Borders and Immigration. (2024). An inspection of asylum casework (June–October 2023). GOV.UK. https://www.gov.uk/government/publications/an-inspection-of-asylum-casework-june-october-2023

Morrison, L., Given-Wilson, Z., & Memon, A. (2024). The impact of emotionally evocative information on interpreting accuracy in a mock asylum interview. Applied Cognitive Psychology, 38(2), e4185. https://doi.org/10.1002/acp.4185

PACTE. (2005). Investigating translation competence: Conceptual and methodological issues. Meta, 50(2), 609–619. https://doi.org/10.7202/011004ar

UK Parliament. (2025). Amendment 79: Provision of interpreters and translators under the Immigration Rules. https://bills.parliament.uk/bills/3929/stages/20214/amendments/10028010

Williamson, A. J. (2024). Defining the role of the UK asylum interpreter: Expectations, realities and training needs [MPhil thesis, University of Glasgow]. Enlighten Theses. https://doi.org/10.5525/gla.thesis.84082